




reference: optimum, overlapped

no overlap

no foresight (best online policy)

policy as deployed (deployed)

GPU work at batch 1 (measured phases)

host launch and bookkeeping

residual (admission contention, unmodelled latency)