


25% of experts per layer on the GPU (exact optimum)



W50: foresight closing half the gap

