

9 MoE models, own text, $C = E/8, E/4, 3E/8$ per layer

host reads per token, relative to Belady

