

Where the Seconds Go: Bounding Mixture-of-Experts Decode from Host Memory, and Measuring What Foresight Is Worth

Harshith Kantamneni

Independent Researcher

harshithkantamneni@users.noreply.github.com

Abstract—A Mixture-of-Experts model too large for one consumer GPU keeps most of its experts in host memory, and decoding one token at a time then waits on host reads. We ask how fast such decoding could be on a given machine, where the remaining time goes, and what knowing the future routing is worth. Our bound is the host reads of Belady’s MIN with bypass, the fewest that any cache holding C experts per layer can make, at the machine’s probed read rate. Our llama.cpp expert cache, faster than two other systems at most configurations, reaches 31–54% of the bound’s speed on machines with consumer processors at the smallest budget we test. We call its time beyond the bound the gap. Oracles built into the engine close about half of the gap; we attribute the rest without measuring it part by part. In our engine, at that budget and on machines whose link reads at least half as fast as the CPU, caching the experts MIN’s greedy schedule keeps pays only if each admitted expert is read once, not twice. Together the two changes close 35% of the gap in a test registered before its machines were rented. There, policies without foresight, built from the engine’s own mechanisms, recover at most 6% of what an oracle that reads ahead gains. In host reads, half of foresight’s value needs the routing of about $0.65C$ distinct experts per layer ahead.

Index Terms—Caching, foresight, mixture of experts, offloading, performance bounds

I. INTRODUCTION

Open-weight Mixture-of-Experts (MoE) models now carry 30–120B parameters and activate a few billion per token; a desktop GPU has 24–32 GB. Running such a model for one user therefore means keeping most experts in host DRAM. Each decoded token reads its routed experts from GPU memory if they are resident, and otherwise from host memory: either the CPU runs the expert, or the expert is copied over PCIe into a GPU slot [1], [2].

Many systems build on this: GPU expert caches [3]–[5], CPU execution of misses [6]–[8], a per-miss choice between the two [9], [10] and expert prediction [11]–[13]. Each reports a speed-up over a baseline of its choosing. None says how fast the same machine could decode the model, or which part of the remaining time a better design could remove. Without that, a builder cannot tell whether to work on what is cached, how it is read, when it is read, or how far ahead the routing must be predicted.

The appendices and the tables and figures numbered S1, S2, ... are in the supplementary material (ieee-supplement.pdf).

We answer these questions for batch-1 decode of gpt-oss-120b [14] and Qwen3-30B-A3B [15] on rented RTX 5090 machines. We derive a bound, build oracles into a real engine that remove one part of the gap at a time, and run them on many machines, with every prediction committed before its machine started. Table I lists the answers and the evidence behind each.

The headroom is large (Section III). MIN’s host reads at the machine’s probed rate bound the time per token of any system with C slots per layer, and that read time is nearly reachable in a microbenchmark. Our cache, which leads FreeToken and llama.cpp at most configurations, runs at 31–54% of the bound’s speed at the smallest gpt-oss budget on machines with consumer processors; published systems, scored the same way, reach a median 13.6%.

What is cached and how it is read pay only together (Section IV). In our engine, at the smallest gpt-oss budget, caching MIN’s set on the deployed read path, where the CPU serves a miss and a background copy fills its slot, closes almost none of the gap; fetching the deployed set into its slots in the step closes none either. Doing both closes 33% of the gap on 15 panel machines, and 35% on 5 machines rented later for a registered test. A registered control that issued the copies a step earlier failed its own check and ran no faster. What the oracles leave we attribute to the GPU’s own work in series with the reads and to the oracles’ extra reads (Table IV).

Without foresight, little of it is reachable (Section VI). Four variants of the deployed cache built from the engine’s admission margin and its layer-ahead prediction recover at most 6% of the read-ahead oracle’s gain on those machines: they read about as many experts as the deployed cache.

The machine decides how foresight should be spent (Sections V and VII). Each admitted expert should be read once, with as few admissions as optimality allows, on the read path the link-to-CPU ratio favours. A trend in that ratio, fitted on RTX 5090 machines, predicted the oracles’ speed on five RTX 4090 machines at ratios 0.37–0.77 within 7%; new RTX 5090 machines with fast links scatter more around it.

Foresight has a price (Section VIII). Half of its value needs the routing of about $0.65C$ distinct experts per layer ahead, on 9 models; exact routing of the next token recovers at most a third of it on our workload.

Section IX turns these results into rules for system builders;




Fig. 1. One decode step of the deployed cache, schematically. Blue: the GPU’s own kernels. Grey: host reads of the misses, which the step waits for. Orange: the background copy that fills an admitted expert’s slot, a second read of the same expert.

Appendices D and N give every prediction, including the failures, and the artifact.

II. SETTING

a) Models and workload: gpt-oss-120b has 36 MoE layers of 128 experts, top-4, in MXFP4 (13.25 MB per expert). Qwen3-30B-A3B has 48 layers of 128 experts, top-8, in BF16 (9.44 MB per expert). Prompts are AIME-25 problems, and the cache carries over between problems, as in one user’s session. The system comparison decodes 256 tokens of each of the 30 problems through each system’s server. The engine experiments replay the model’s own greedy output, teacher-forced, so that every configuration sees the same routing. A *budget* is the share C/E of each layer’s experts that fit in GPU slots: gpt-oss 11% and 25% are $C=14$ and $C=32$ of 128. The two smallest budgets of each model are *host-bound*: at the GPU’s datasheet rate, host reads, not the GPU, limit their speed in Eq. (1), except at gpt-oss 25% on the few machines with the fastest host memory. At the GPU’s measured batch-1 rate the GPU term also binds at gpt-oss 25% on both hosts of Table III, and at Qwen3 25% on host B. Most results are at gpt-oss 11% and 25%.

b) The engine: Our cache is a 3,500-line patch to llama.cpp [16]. Each MoE layer has C GPU slots. A *miss* is read from host memory in one of two ways: CPU helper threads run it, or it is *fetched* (copied into a slot over PCIe and run on the GPU in the same step). A per-machine table, computed from the machine’s bandwidth probe, sets how many of a layer’s misses are fetched. The *deployed* policy admits a miss into a slot when its decayed request count beats the lowest resident’s by a margin $\kappa=1$. The CPU serves the admitted miss and a background copy then fills its slot, so the expert is read from host memory twice (Fig. 1).

c) Terms: *MIN* is Belady’s optimum with bypass [17], [18]: it *admits* a miss only if some resident is needed later than the miss, which gives the most hits possible with C slots [19], [20]. Many schedules reach that optimum. The *greedy* one admits whenever it may; the *fewest-admission* one admits as rarely as any of them. An admitted expert is read *once* (1R), fetched into its slot in the step that needs it, or *twice* (2R), served by the CPU in that step and copied in the background;

the engine publishes a background copy at the start of the second step after the one that admitted it. A copy issued a few steps before the expert’s use is read *ahead*. *Oracles* read the future routing from a trace recorded on the same machine. Configurations are named by the set admitted and the number of reads: MIN-1R is MIN’s greedy set read once, MIN-2R the same set read twice, Dep-1R the deployed set read once, and Few-1R and Few-2R MIN’s fewest-admission set. Table II lists the configurations, machine sets and yardsticks (the window policies of Section VIII are in Appendix G), and Appendix C maps the codes of the appendix tables to these names.

d) Machines: Every machine is a rented host with one RTX 5090 (an RTX 4090 in Section VII). Each runs our bandwidth probe before any model run. It reads host memory from CPU threads, at the number of helper threads the engine uses (B_c), and over PCIe into the GPU, at the rate the engine’s copies use (B_p); B_{host} is its highest reading by any method. The *link-to-CPU ratio* is B_p/B_c . We call a link *fast* when the ratio is at least 0.5 and *slow* otherwise; the term compares the link with the CPU, not with other links, and we drew the 0.5 line after the panel’s data, before the registered tests of Sections IV, VI and VII. Ryzen, Core and Threadripper parts, high-end desktop ones included, are *consumer* processors; EPYC, Xeon and engineering-sample parts are *server* processors.

e) Statistics and registration: Time per token is averaged over problems, and speed ratios are ratios of mean times. Machines differ far more than problems (two panel machines with the same CPU model differ by up to 29% in the deployed cache’s time), so the machine is the unit. Claims are about orderings that hold on every machine of a stated class, and summaries over machines carry 95% t -intervals over machines, since with a handful of machines a percentile bootstrap over machines is too narrow [21]; comparisons on one machine carry 95% intervals over problems (a paired bootstrap). A *launch* is one rental; in the later experiments each configuration ran in two or three processes, *rounds*, per launch (at 25% in the registered tests, one). A machine *ran stably* when its outputs matched the reference and, where it ran more than one process, their times agreed within 2%; two experiments ran each configuration in a single process, so only the first condition applies to them (Appendix D). Intervals over problems leave out rounds and rentals: a speed ratio moved by a median 0.2% (at most 1.8%) between rounds and 0.5% (at most 5.3%, where a relaunch’s probe changed the fetch table) between rentals. Widened by its round range, each of the 89 per-machine registered clauses that held with an interval still does; widened by the largest rental difference, 49 do, and 29 of the other 40 are bands of ± 0.06 or narrower (Appendix D). Every experiment’s predictions were committed to a public branch before its machine started; Appendix D scores all of them, and Table I marks the thresholds that earlier data made easy to meet.

TABLE I

OUR CLAIMS, THE THRESHOLD REGISTERED FOR EACH BEFORE ITS MACHINES STARTED, AND THE RESULT (APPENDIX D SCORES EVERY CLAUSE). *Derived*: FOLLOWS FROM THE DEFINITIONS; *post hoc*: FOUND IN THE DATA IT DESCRIBES; *loose*: A THRESHOLD THE EARLIER DATA MADE EASY TO MEET; *held*: THE POINT ESTIMATE MEETS THE THRESHOLD (APPENDIX D SAYS WHICH CLAUSES ALSO HOLD OVER THEIR WHOLE INTERVAL). *Mach.*: MACHINES THAT PASSED EVERY CHECK; “ $a(b) + c$ ”: a OF THE b PANEL MACHINES AND c NEW ONES.

Claim	Registered threshold	Result	Mach.	§
MIN’s reads at the probed rate bound any exact-routing cache with C slots per layer	derived	–	–	III
The bound’s read time is nearly reachable	$\geq 50\%$ of it per layer (loose)	86–96%, held; 5 of the job’s 33 clauses failed	3	III
Our cache is faster than FreeToken and llama.cpp at most configurations	on B: faster than FreeToken (interval > 1), $\geq 1.8 \times$ llama.cpp ($\geq 2 \times$ at Qwen3 12.5%); on S: within ± 0.06 of B’s ratios	held on B; on S, five of six ratios fell outside ± 0.06 of B’s; faster at 11 of 12 in all	B, S	III
MIN’s set and one read pay only together (gpt-oss 11%, fast links, greedy schedule)	interaction > 0 on every machine; MIN-1R’s mean share of the gap 0.25–0.55 (loose)	positive on 9 of 10 panel machines (failed), 5 of 5 new ones (held; fast-link scope set in between); share 35%	10(15) + 5	IV
Issuing the copies a step earlier does not rescue the two-read path	misses $\leq 0.95 \times$ Few-2R’s; slowdown ≤ 0.01 ; per machine and budget	at 11%: misses 0.97–0.98 $\times$, 0.03–0.05 slower; failed on 5 and 7 of 8 machine-budgets; the second read costs more than the late landing (held)	4	IV
Policies without foresight recover little of the oracle’s gain	speed $\leq 1.10 \times$; capture $\leq 35\%$ (loose), per machine	best 1.007–1.024 $\times$, capture at most 6% (held)	5	VI
MIN’s fewest-admission set read once gains on stable machines	> 1 on every machine	1.24 $\times$ (geometric mean) on all 13 stable machines, held on the 7 registered; lost on the two unsteady ones (one registered, failed; one excluded by its job’s validity rule)	13	V
The link-to-CPU ratio predicts the oracles’ speed on a second card	within 0.10 in log of the RTX 5090 trend (trend fitted post hoc)	within 7% at ratios 0.37–0.77 (held); new fast-link RTX 5090s off by up to 18% (post hoc)	5	VII
Our cache reaches 31–54% of the bound on consumer machines (gpt-oss 11%)	post hoc; consumer scope set after the data	server machines: 12–46%	25	III
Published systems reach a median 13.6% of their bound	post hoc	their papers’ hardware at datasheet rates	–	III
Half of foresight’s value needs $\approx 0.65 C$ distinct experts ahead	post hoc	on 9 models’ traces	–	VIII

TABLE II

THE NAMES THIS PAPER USES: CONFIGURATIONS (LEFT), MACHINE SETS AND YARDSTICKS (RIGHT).

Admitted set	read twice	read once	read ahead	<i>Machine sets</i> (RTX 5090 unless noted)	
<i>No foresight</i>				panel	the 15 machines that ran every state of Table IV before the new machines; 10 ran its main job
deployed: decayed count beats the lowest resident’s by $\kappa=1$	deployed cache	Dep-1R	–	new machines	5 rented for the registered test of Section IV (ratio ≥ 0.5)
the same, larger margin ($\kappa=3$ at 11%, 2 at 25%)	Margin	–	–	O3–O5	three machines that ran every budget (Section V)
the same, plus a layer-ahead copy ^a	LA	LA-1R, LA-1R-margin ^b	–	hosts B, S	the system comparison (Table III)
every miss	–	admit every miss	–	RTX 4090s	5 machines of the second-card tests (Section VII)
<i>Oracles</i> (the future routing, from a trace recorded on the same machine)	trace recorded	MIN-1R	read-ahead oracle ^c	<i>Yardsticks</i> (T : time per token on one machine)	
MIN with bypass, greedy schedule	MIN-2R	MIN-1R	–	speed	T_{deployed}/T ($\times$), the default
MIN with bypass, fewest admissions	Few-2R ^d	Few-1R	–	share of the bound	T_{bound}/T , the bound being Eq. (1)
Belady’s rule, no bypass	Belady-2R	–	Belady-1R	share of the gap	$(T_{\text{deployed}} - T)/(T_{\text{deployed}} - T_{\text{bound}})$
				capture	speed gain over the read-ahead oracle’s gain
				value of	share of MIN’s host-read saving over admitting
				foresight	every miss

Read twice: the CPU serves the miss, then a background copy fills the slot; *read once*: fetched into the slot in the step. ^a A copy, one layer ahead, of the non-resident expert that the next layer’s router, applied to this layer’s input, ranks highest; an admission it predicted is read once. ^b LA-1R with Margin’s margin. ^c MIN’s admissions copied three steps ahead where the next use allows; its early copies evict sooner. ^d Few-2R-early: each copy issued one step earlier, so that it lands for the next use.

III. HOW FAST COULD IT BE?

Consider any system that executes the model’s exact routing with at most C experts of each layer resident on the GPU. Per token, each of the Lk routed experts is served by a resident copy or read from host memory. Let R^* be MIN’s host reads

per token, summed over layers, on the routing trace the engine runs (Appendix N); no policy, even one that knows the future, reads fewer. If m experts per token run on the CPU, host memory supplies at least $\max(R^*, m)S$ bytes for experts of S bytes, and the GPU reads the dense weights D plus $(Lk - m)S$.

With everything overlapped,

$$\bar{T} \geq \min_{0 \leq m \leq Lk} \max \left(\frac{D + (Lk - m)S}{B_{\text{gpu}}}, \frac{\max(R^*, m)S}{B_{\text{host}}} \right). \quad (1)$$

At a host-bound budget this is MIN's read time, R^*S/B_{host} . The bound assumes exact routing, whole experts, at most C residents per layer and one token per step. It drops all latencies, lets the two read paths share B_{host} , and is relative to the probe: a reader faster than the probe could beat it (Section XI).

a) The demand bound: reads that wait for the router:

A system that reads an expert only after its layer's router has chosen it cannot overlap that read with the GPU's work: layer l 's misses are known only after its attention and router, and layer $l+1$ starts only after layer l 's experts. Its time is at least the GPU's non-expert work plus the reads, in series:

$$\bar{T}_{\text{demand}} \geq T_{\text{GPU}} + \frac{R^*S}{B_{\text{host}}}, \quad (2)$$

where T_{GPU} is the GPU's attention, dense layers, norms and routers per token: 2.9 ms for gpt-oss, the smallest of the 14 Nsight profiles of our engine we use; a system with faster non-expert kernels has a lower bound. Every system we compare reads on demand; only reads issued before routing can go below Eq. (2).

b) The ordered bound: reads foresight cannot free:

Foresight does not free every read. An expert run on the CPU needs its layer's attention output, and the next layer's attention needs its result, so CPU reads sit in series with the GPU's non-expert work; only copies over the link can be issued ahead. This assumes that the CPU reads an expert's weights from memory when it runs it; a read-ahead into the CPU's cache would relax it. With a share f of MIN's reads served by the CPU at the probe's highest CPU reading $\hat{B}_c$ and the rest copied at its highest link reading $\hat{B}_p$,

$$\bar{T} \geq \min_{0 \leq f \leq 1} \max \left(\frac{(1-f)R^*S}{\hat{B}_p}, T_{\text{GPU}} + \frac{fR^*S}{\hat{B}_c} \right), \quad (3)$$

and Eq. (1) holds as well, so the bound is the larger of the two. On the median consumer machine at gpt-oss 11%, Eq. (3) equals Eq. (1) (ratio 1.00); where the link is slow, it is up to 1.56 times as high.

c) The read time is nearly reachable: A microbenchmark replays MIN's per-layer reads of gpt-oss on 3 machines without the model: CPU threads read part of each layer's experts while the copy engine copies the rest. It reaches 86–96% of the bound's read time (Table S23; registered: at least 50%, a loose threshold). The GPU term of Eq. (1) is looser, but at datasheet rates it binds only at the largest budgets (Table S22).

d) Where systems stand: Table III places three systems against the bounds on two machines. Our cache leads Free-Token [9] at 11 of 12 configurations and stock llama.cpp at all of them (Appendix J). Across 25 machines with consumer processors at gpt-oss 11%, it runs at 31–54% of the speed of Eq. (1), the bound we use throughout (40–54% of Eq. (3)); the 5 machines added in Section IV fall at 41–48%. The top of that range is probably too high: on that machine the engine's

reads imply, by the closed-form account of Appendix H, a rate above the probe's best reading (Table S21).

e) Published systems (post hoc): Scored against the same bound on their papers' own hardware at datasheet rates, 20 published batch-1 measurements reach a median 13.6%; ours, scored the same way, a median 27% (Appendix M, which gives the caveats).

IV. WHERE THE SECONDS GO

The gap is the deployed cache's time beyond Eq. (1). Oracles in the engine remove three parts of it: what is cached (MIN's set instead of the deployed one), how it is read (each admitted expert fetched once in the step, instead of served by the CPU and copied in the background), and when (the read-ahead oracle issues its copies a few steps early). We change the first two alone and together, on the 15 panel machines that ran every state and, under a test registered before they started, on 5 machines rented for it (Fig. 2 and Table IV).

a) On this engine, what is cached and how it is read pay only together: At gpt-oss 11%, MIN-2R, MIN's set on the deployed read path, closes 0% of the gap, and Dep-1R, the deployed set read once, closes 0%. MIN-1R, both changes, closes 33% on the panel and 35% on the new machines. MIN's set admits more experts than the deployed one, and on the deployed path each extra admission costs a second read. Its copies also land two steps after the admission, so MIN-2R misses more than MIN-1R: the factor is the read path, which sets both the number of reads and when a copy lands. A registered control tried to separate the two (Table V): **Few-2R-early** issues each copy of MIN's fewest-admission set one step earlier, so that it lands for the next use. It failed its own check. Its misses fell only to 0.97–0.98 of **Few-2R**'s (registered: at most 0.95), because the fetch table still fills slots the schedule did not choose, and it admitted 11.7–11.8 experts per token against **Few-2R**'s 6.2–7.4, reading 4–7% more from the host. On four machines at 11% it ran $0.98\text{--}1.00\times$ the deployed cache, **Few-2R** $1.03\times$ and **Few-1R** $1.21\text{--}1.41\times$. So the control cannot say how much of the two-read arms' loss is late landing; it shows only that issuing the copies earlier, in this engine, does not recover it. We registered a positive interaction on every machine: it was positive on 9 of the 10 machines of the panel's main job (failed) and on all 5 new machines (held). The claim is about MIN's greedy set: its fewest-admission set read twice gains a little alone (**Few-2R**, $1.03\times$ on the control's machines) and more on slow links (Section VII). At 25%, where admissions are fewer, MIN's set alone already closes 21%.

b) What is left: The read-ahead oracle closes another 15% at 11%, and 52% of the gap remains (Table IV). On fast links, two named parts account for it (at 25% they slightly overshoot). The oracle still reads its remaining misses after their router runs, so the GPU's non-expert time stays in series with them, and it reads more than MIN, because its early copies evict sooner. A system that issued every read before routing and read no more than MIN would remove both, but only copies over the link can be issued early (Eq. (3)). On slow links the

TABLE III

DECODE SPEED AT EQUAL GPU EXPERT MEMORY ON TWO RTX 5090 HOSTS (30 AIME-25 PROBLEMS, THE FIRST 256 GREEDY TOKENS, ONE SESSION).

Model	Experts on GPU	llama.cpp (tok/s)	FreeToken (tok/s)	Ours (tok/s)	Ours ÷ FreeToken (95% CI)	Ours ÷ llama.cpp	Bound, datasheet tok/s	GPU ours	Bound, measured tok/s	GPU ours
Host B: Ryzen 9 9950X3D, card memory 17001 MHz; CPU 72, link 53, together 78, highest 88 GB/s (the bound's rate)										
gpt-oss	11%	34.9	54.0 ^h	69.9	1.29 [1.28, 1.31]	2.00×	172	41%	172	41%
gpt-oss	25%	40.0	85.6 ^h	109.2	1.27 [1.25, 1.29]	2.73×	430	25%	279	39%
gpt-oss	40%	47.7	132.2 ^o	152.5	1.15 [1.13, 1.17]	3.20×	519	29%	279	55%
Qwen3	12.5%	19.5	38.7 ^h	40.0	1.03 [1.02, 1.04]	2.05×	96	42%	96	42%
Qwen3	25%	22.1	54.7 ^h	63.1	1.15 [1.13, 1.17] [†]	2.86×	214	30%	192	33%
Qwen3	43.75%	28.4	103.1 ^o	108.2	1.05 [1.04, 1.06]	3.81×	308	35%	192	56%
Host S: Ryzen 9 9950X, stock clock, card 14001 MHz; CPU 56, link 53, together 66, highest 71 GB/s; FreeToken backends as on B										
gpt-oss	11%	28.9	48.0 ^h	57.9	1.21 [1.19, 1.22]	2.00×	140	41%	140	41%
gpt-oss	25%	33.6	78.7 ^h	94.1	1.20 [1.16, 1.22]	2.80×	351	27%	275	34%
gpt-oss	40%	40.1	122.2 ^o	133.6	1.09 [1.08, 1.11]	3.34×	514	26%	275	49%
Qwen3	12.5%	16.1	32.7 ^h	33.5	1.03 [1.02, 1.04]	2.08×	78	43%	78	43%
Qwen3	25%	18.5	51.4 ^h	53.9	1.05 [1.04, 1.06]	2.91×	174	31%	174	31%
Qwen3	43.75%	23.8	98.3^o	95.8	0.97 [0.96, 0.99]	4.02×	305	31%	189	51%

Ratios are of mean speeds; the ratio over FreeToken is paired by problem within one launch, with a 95% bootstrap interval over problems. Host B's rows come from three rentals of the same machine (the same GPU), so its ratios over llama.cpp compare launches. Our cache and FreeToken ran twice per cell and the second launch is shown (on host B both launches ran ours first); llama.cpp and the [†] cell ran once. Each system decodes its own greedy text. Our cache uses each machine's fetch table; FreeToken uses its faster backend per budget (^h hybrid, ^o offload), picked on host B on a separate launch. gpt-oss: gpt-oss-120b; Qwen3: Qwen3-30B-A3B. *Bound*: Eq. (1) on that machine (MIN's reads, the host's highest probed rate) with the GPU at its datasheet rate, and at the rate all-in-VRAM decode measures at batch size 1 (52% of datasheet for gpt-oss, 61% for Qwen3; Table S22 varies the other assumptions; Appendix J gives the pooled bounds and all-in-VRAM speeds). [†] Selection launch only, which favours FreeToken.




Fig. 2. Where the seconds go, measured: each state's time per token relative to the deployed cache on the same machine (lower is faster), as oracles change what is cached and how it is read (alone, then together) and when it is read, with the two bounds for reference. One line per machine. Blue and orange: the panel, split at a link-to-CPU ratio of 0.5; on the two slow-link panel machines (orange), reading once in the step saves little or costs time. Green: the machines of the registered test. Thick dark lines: medians over the fast-link panel machines, changing MIN's set first (solid) or the number of reads first (dashed).

two parts fall short: a residual of 35–47% of the gap at 11% stays unexplained on the two slow-link panel machines, and 31–36% on the three slow-link RTX 4090s of Section VII (27–31% with that card's own non-expert time). Of the gap to Eq. (3), the tighter bound there, the read-ahead oracle still leaves 83–99% on the two slow-link panel machines and 73–81% on the three RTX 4090s (median 45% on fast links): the oracles fetch over the link, the slow path on these machines.

V. HOW FORESIGHT MUST BE SPENT

This section measures each way of spending foresight as a speed relative to the deployed cache on the same machine. A

running example: on a Ryzen 9 9950X at gpt-oss 11%, Eq. (1) gives 10.0 ms per token and Eq. (2) 13.0. The deployed cache takes 20.6 ms, MIN-1R 15.2, and the read-ahead oracle 13.8.

a) *Read each admitted expert once*: MIN-1R is faster than the deployed cache at every budget on the three factorial machines (Fig. 3). The usual ways of spending foresight waste reads. Prefetching without bypass, Belady's original rule, admits experts MIN would bypass. Serving a miss on the CPU before copying it reads every admission twice. At the smallest budget of each model these gain at most 15% or lose, a trade-off known for disks [22] and CPU caches [23]. (Registered as

TABLE IV

THE DEPLOYED CACHE’S GAP TO EQ. (1), SPLIT BY ORACLES IN THE ENGINE (FIG. 2): MEAN SHARE OF THE GAP, IN PERCENT, WITH 95% t -INTERVALS OVER MACHINES.

Part of the gap	gpt-oss 11%		gpt-oss 25%	
	Panel (15)	New (5)	Panel (15)	New (4)
<i>Measured</i>				
MIN’s set alone (MIN-2R)	0 [−2, 3]	3 [−1, 8]	21 [19, 22]	20 [18, 21]
One read alone (Dep-1R)	0 [−2, 2]	0 [−2, 1]	−3 [−7, 1]	−2 [−3, −1]
Both (MIN-1R)	33 [21, 44]	35 [22, 48]	31 [23, 39]	27 [23, 32]
Then the read-ahead oracle	15 [11, 19]	19 [14, 24]	21 [18, 23]	24 [21, 27]
Left after all three	52 [43, 62]	46 [38, 55]	48 [41, 55]	48 [44, 53]
<i>Attributed: of what is left</i>				
the GPU’s non-expert time (T_{GPU})	27 [25, 30]	28 [20, 36]	34 [31, 37]	38 [33, 43]
the oracle’s reads beyond MIN’s, at B_{host}	21 [20, 23]	18 [15, 20]	19 [18, 21]	16 [15, 17]
residual	4 [−6, 13]	1 [−3, 4]	−5 [−12, 2]	−5 [−7, −3]

Panel: the 15 machines of jobs 096–101 that ran every state. *New*: the 5 machines of job 109, rented for this test, whose population (desktop-class, link-to-CPU ratio at least 0.5) and predictions were registered before any of them started; the job’s deadline cut one machine’s 25% round. The first two rows each change one thing; the third changes both. *Both*, *then the read-ahead oracle* and *left* sum to the gap. The last three rows are attributed, not measured: they split *left* with T_{GPU} as in Eq. (2) (2.9 ms, the smallest profile) and the read-ahead oracle’s reads beyond R^* at B_{host} ; the residual is what they leave. Probe sensitivity: with every machine’s B_{host} 10% (22%) higher, the shares on fast links move by at most 5 (9) points for *Both* and 6 (11) for *left* at 11%.




Fig. 3. What foresight is worth, by how it is spent: speed of each configuration of Table II relative to the deployed cache on the same machine, at the six budgets († host-bound). Three machines at every budget (O3–O5: two Ryzen 9 9950X and a 9950X3D); the 10 machines of the panel’s main job (grey ticks) at the two host-bound gpt-oss budgets, and 3 of them at Qwen3 12.5%. MIN-1R gains at every budget on these machines; reading twice or admitting what MIN bypasses gains at some budgets and loses at others.

gain bands per budget on two machines; held at most budgets.)

b) Admit as rarely as optimality allows: The greedy schedule admits whenever MIN allows and may evict an expert again before its next use. The fewest-admission schedule makes 0.60–0.73 of the greedy one’s copies for the same hits in replay. Read once in the step at gpt-oss 11% (Few-1R), it beat the deployed cache on all 13 machines that ran it stably, by a geometric mean of $1.24\times$ (registered on 7 of them; two are RTX 4090s and two have server processors). Counting the machines that ran unsteadily as well, it lost on two, the two with the slowest links (ratios 0.14 and 0.21; Appendix I): the stability rule leaves out exactly those.

VI. WITHOUT FORESIGHT, THE ENGINE RECOVERS LITTLE

Two of the engine’s mechanisms act on the same levers without foresight, in four variants (Table II). *Margin* admits

less: it cuts the deployed cache’s host reads by 4–6% at 11%. *LA* copies, one layer ahead, the expert that the next layer’s router predicts, so that a correctly predicted admission is read once, by the GPU; *LA-1R* and *LA-1R-margin* also fetch every admission in the step. We registered that none of the four would run faster than $1.10\times$ the deployed cache or capture more than 35% of the read-ahead oracle’s gain on any machine; the second threshold was loose. At 11% the best of them runs 1.007 – $1.024\times$ the deployed cache, at most 6% of the read-ahead oracle’s gain. The counters say why: the four read 1.57 – $1.76 R^*$ host experts per token, about as many as the deployed cache. A prediction one layer ahead moves reads earlier but does not avoid them. On the slow-link RTX 4090s of Section VII, the variants built on the layer-ahead copy lose time (0.64 – $0.72\times$ at 11%), as registered, and *Margin* is the best, at 1.033 – $1.041\times$: up to 26% of an oracle gain that is

TABLE V

WHEN THE COPY LANDS, AGAINST HOW OFTEN THE EXPERT IS READ: MIN’S FEWEST-ADMISSION SET READ ONCE (FEW-1R), READ TWICE WITH THE COPY PUBLISHED TWO STEPS AFTER THE ADMISSION (FEW-2R), AND READ TWICE WITH EACH COPY ISSUED A STEP EARLIER SO THAT IT LANDS FOR THE NEXT USE (FEW-2R-EARLY).

Machine	Ratio	Speed vs deployed			Misses per token		
		1R	2R	early	1R	2R	early
<i>RTX 5090</i>							
Ryzen 9 5950X	0.74	1.41	1.03	0.98	40.7	55.6	54.5
Ryzen 9 7945HX	0.88	1.25	1.03	1.00	40.7	55.7	54.6
<i>RTX 4090</i>							
Ryzen 7 7800X3D	0.57	1.21	1.03	1.00	40.9	54.7	53.0
Core i5-12400	0.77	1.35	1.03	0.98	40.8	55.7	54.5

Speed relative to the deployed cache on the same machine (every 95% interval over problems within ± 0.01) and host misses per token, at gpt-oss 11%; registered before the machines were rented.

smaller there.

VII. THE LINK-TO-CPU RATIO ON A SECOND CARD

On the RTX 5090 machines, the gain of every configuration that moves reads from the CPU to the link follows the link-to-CPU ratio: for MIN-1R, Spearman 0.89 [0.63, 0.97] over the 19 machines of jobs 093–104 (Fig. 4). We froze that trend on the 15 panel machines, one line in the log of the ratio for each of MIN-1R, MIN-2R, Dep-1R and the read-ahead oracle, and registered that their speed on RTX 4090 machines, a slower card behind a PCIe 4.0 link, would fall within 0.10 of it in log. The first test ran on three machines, all at ratios 0.37–0.40; a second, with the same predictions, added two at 0.57–0.77. On all five, MIN-1R and the read-ahead oracle came within 7% of the prediction at 11% and 5% at 25% (Table VI), and every machine’s 95% interval over problems lies inside the band; the interaction of Section IV is positive on the two with fast links, as registered. (The first test was relaunched with the RTX 4090’s own reference loss; Appendix D gives the first rounds.) Predicting no change would have missed by up to 46% at 11% and 69% at 25%. Applied after the fact to the new RTX 5090s of Section IV, all fast-linked, the same trend misses five of their 18 cells by more than the band, by up to 18%: on fast links the ratio sets the direction of the gain more reliably than its size. Where the link reads at under a third of the CPU rate, Few-2R pays instead of MIN-1R (1.11–1.22 $\times$ over the two budgets on RTX 5090s at ratios 0.28–0.32).

VIII. HOW MUCH FORESIGHT IS NEEDED

The oracles know the whole future; a predictor would know part of it, imperfectly. We give the cache a *window*, the routing of the next W tokens, exact or with each future expert right only with probability r . Throughout this section, the value of a window is its share of MIN’s gain over admitting every miss, the online policy that reads least at these budgets (Appendix G gives the policies).

a) Next-token prediction buys little: Exact routing of the next token recovers at most 0.31 of MIN’s gain at any host-bound budget of our two models; half of the gain at gpt-oss 11% needs 4 tokens (Fig. 5). In the engine, an exact 16-token window recovers 0.78–0.93 of the gain in time at 11% (registered at 0.80 or more; the lowest machine missed). Against the deployed cache, the same window copied in the step runs 0.89–1.32 $\times$ at 11% on 4 machines, below 1 on the 2 with the lowest link-to-CPU ratios. Accuracy matters as much as horizon: an 8-token window that is right half the time recovers 0.29 of the gain at 11%, about what 2 exact tokens buy (0.31).

b) The horizon is measured in distinct experts (exploratory): On 9 models’ own sampled text, the window that closes half of the read gap ranges from 0.6 to 34 tokens. Paging theory measures lookahead in distinct pages [24], [25], and so does the rule that fits: within that window a layer requests about 0.65 C distinct experts (0.66–0.81 C on the AIME routing the engine runs). Given a model’s routing trace, this one-constant rule predicts its horizon as well as a two-constant power law in tokens (Appendix F).

c) Nothing realisable reaches it: Forecasters trained on routing histories, linear or recurrent, close at most 15% of the read gap on gpt-oss: their precision falls with the horizon. Speculative batches save an online cache about a tenth of its reads at 90% draft acceptance and cost reads at 80% [26]–[28] (Appendix F). Layer-ahead predictors [11], [12] reach 84% recall in our engine but see less than one token ahead (Section VI).

IX. WHAT THIS MEANS FOR SYSTEM BUILDERS

These rules come from one engine (ours), batch-1 decode and one prompt set, and rules 2–4 from one model at two budgets; elsewhere they are hypotheses to test, with Eq. (1) as the yardstick.

- 1) **Report the distance to the machine’s bound.** A speed-up over a chosen baseline hides how much is left. Equation (1) needs only the routing trace and one probe run.
- 2) **Change what is cached and how it is read together.** Serving a miss on the CPU and then copying it into a slot reads it twice. At the smallest budget, a better set read twice, or the same set read once, closes almost none of the gap; where the link reads at least half as fast as the CPU, a better set read once closes about a third of it.
- 3) **Do not expect a policy without foresight to get there.** Admitting less and predicting one layer ahead recover at most 6% of what foresight is worth where it is worth most, and the layer-ahead copy loses time on slow links; the reads they leave are the ones only foresight avoids.
- 4) **Choose the read path by the link-to-CPU ratio.** Where the link is fast, copy misses into slots in the step; where it is slow, serve them on the CPU and copy in the background. The ratio predicts the direction of the gain more reliably than its size (Section VII).

TABLE VI

THE REGISTERED TESTS ON MACHINES RENTED FOR THEM: SPEED RELATIVE TO THE DEPLOYED CACHE ON THE SAME MACHINE (GEOMETRIC MEAN OF ITS ROUNDS) OF MIN-1R, THE READ-AHEAD ORACLE (*ahead*), AND THE BEST OF THE FOUR VARIANTS WITHOUT FORESIGHT (*none*: MARGIN, LA, LA-1R, LA-1R-MARGIN; SECTION VI), AND THAT VARIANT’S CAPTURE, ITS SHARE OF THE READ-AHEAD ORACLE’S GAIN.

Machine	Ratio	gpt-oss 11%				gpt-oss 25%			
		MIN-1R	Ahead	None	Capture	MIN-1R	Ahead	None	Capture
RTX 5090, job 109									
Core Ultra 9 285K	0.54	1.22 $\pm$.01	1.46 $\pm$.01	1.02 $\pm$.00	4 $\pm$ 0%	1.24 $\pm$.01	1.64 $\pm$.02	1.03 $\pm$.00	4 $\pm$ 0%
Core i9-13900KF	0.59	1.25 $\pm$.01	1.46 $\pm$.01	1.01 $\pm$.00	2 $\pm$ 0%	1.26 $\pm$.01	1.61 $\pm$.02	1.01 $\pm$.01	2 $\pm$ 1%
Ryzen 9 5950X	0.84	1.37 $\pm$.01	1.50 $\pm$.01	1.01 $\pm$.00	2 $\pm$ 0%	—	—	—	—
Ryzen 7 9800X3D	0.88	1.15 $\pm$.01	1.34 $\pm$.01	1.01 $\pm$.01	4 $\pm$ 2%	1.22 $\pm$.01	1.55 $\pm$.02	1.01 $\pm$.00	2 $\pm$ 1%
Ryzen 9 7945HX	0.89	1.24 $\pm$.01	1.42 $\pm$.01	1.02 $\pm$.01	6 $\pm$ 1%	1.28 $\pm$.01	1.63 $\pm$.02	0.99 $\pm$.01	−1 $\pm$ 1%
RTX 4090, job 111									
Core i9-14900KF	0.37	0.96 $\pm$.00 [0.99]	1.13 $\pm$.01 [1.11]	1.03 $\pm$.00	26 $\pm$ 3%	1.02 $\pm$.00 [1.07]	1.30 $\pm$.01 [1.29]	1.04 $\pm$.00	13 $\pm$ 1%
Core Ultra 9 285K	0.39	1.01 $\pm$.00 [1.01]	1.20 $\pm$.01 [1.13]	1.04 $\pm$.00	20 $\pm$ 2%	1.06 $\pm$.01 [1.09]	1.36 $\pm$.01 [1.31]	1.04 $\pm$.00	12 $\pm$ 1%
Core Ultra 9 285K	0.40	1.01 $\pm$.00 [1.01]	1.20 $\pm$.01 [1.13]	1.04 $\pm$.00	20 $\pm$ 2%	1.06 $\pm$.01 [1.09]	1.37 $\pm$.01 [1.32]	1.05 $\pm$.00	12 $\pm$ 1%
RTX 4090, job 112 (higher ratios)									
Ryzen 7 7800X3D	0.57	1.13 $\pm$.01 [1.14]	1.30 $\pm$.01 [1.26]	1.01 $\pm$.01	4 $\pm$ 2%	1.17 $\pm$.01 [1.21]	1.48 $\pm$.01 [1.47]	1.02 $\pm$.00	5 $\pm$ 1%
Core i5-12400	0.77	1.31 $\pm$.01 [1.25]	1.46 $\pm$.01 [1.37]	1.01 $\pm$.00	3 $\pm$ 1%	1.35 $\pm$.01 [1.31]	1.69 $\pm$.02 [1.61]	1.02 $\pm$.00	3 $\pm$ 0%

Ratio: link-to-CPU. RTX 5090: job 109, whose population (desktop-class, link-to-CPU ratio at least 0.5) was fixed before any machine started. RTX 4090: job 111 (job 110’s predictions, committed before any RTX 4090 ran) and job 112 (the same predictions, on machines at higher ratios); in brackets, what the RTX 5090 machines’ trend in the link-to-CPU ratio predicted for that machine. Subscripts: half-width of the 95% interval over problems (paired bootstrap within the machine).




Fig. 4. The machine decides which way of spending foresight pays: speed relative to the deployed cache on the same machine against the machine’s link-to-CPU ratio, at the two host-bound gpt-oss budgets. Every marker is a machine that passed our checks. Stars: Few-1R. Circles and triangles: MIN-1R and MIN-2R, both on the greedy schedule. Crosses: LA. Purple markers: the RTX 4090 machines of the registered second-card test (Section VII), not used in the fits. Lines: least-squares fits in the log of the ratio, drawn for the trend only.

- 5) **Issue reads before routing, and predict far enough ahead to save reads.** Even with full foresight, the GPU’s non-expert time stays in series with every read that waits for the router; only a prediction lets reads start earlier, and only over the link (Eq. (3)). To save reads, predictors should target the next $\approx 0.65 C$ distinct experts of each layer, not the next token; accuracy counts as much as horizon.

X. RELATED WORK

a) *Offloading systems*: Expert caches copy misses over PCIe [3]–[5], [29]; CPU-execution systems run non-resident experts on the host [6]–[8], [30]; FreeToken [9] and DALI [10] split each layer’s misses between the two. Predictors guess upcoming experts from the previous layer [11], [12],

several tokens ahead [13], [31], from draft tokens [26]–[28], or decouple the router [32]. Each reports speed-ups over its own baseline; none bounds the machine. Appendix K covers these and many more, with batched and flash-tier systems.

b) *Bounds and caching theory*: WiSP [33], Budgeting Bytes [34], MoE-CAP [35] and the roofline [36] bound decode by bytes over bandwidth; Zhang [37] and Liang et al. [38] compare caches with MIN in hits. We add the time domain. MIN with bypass maximises hits [19], [20], not speed [18]; our fewest-admission schedule extends offline caching over reuse intervals [39].

XI. LIMITATIONS

- **Hardware, engine, sample and workload.** One engine, ours, for every engine result; one GPU model for most




Fig. 5. Baseline here: *admit every miss*, not the deployed cache. The share of MIN’s gain over admitting every miss recovered, in host reads on the routing the engine runs, by exact routing of the next W tokens (solid), by a window whose experts are each right with probability 0.5 (dashed) and by a linear forecaster fitted on other text (dotted). Markers: the same share in time per token (mean and range over the panel’s main-job machines: 10 for gpt-oss, 3 for Qwen3 12.5%). Grey line: half of the gain.

results and a second (RTX 4090) for two registered tests, all rented, mostly in desktop hosts; clocks could not be locked; batch-1 decode of two models on AIME prompts, also used while developing the system, teacher-forced on each model’s own greedy text. The first RTX 4090 test deviated from its registration, and the last experiment replaced one machine early (Appendix D); where the crossovers between read paths lie rests on few slow-link machines. The machines are a convenience sample of the cheapest offers that passed our checks, and rented machines are often shared or unsteady (of 8 launches on server processors only three passed every check).

- **The bounds.** Exact routing and per-layer slots; pooling slots across layers lowers MIN’s reads by 4.5–18.6%. The bounds drop latencies and are relative to one run of our probe, mostly taken while the model downloaded: the engine’s reads imply a rate more than 5% above the probe’s highest reading on 12 of 84 launch-budgets (Table S21). A second probe on the idle machine came within 4.9% of the first on 4 machines, but between two rentals of one machine the ratio moved by up to 12%. T_{GPU} is our engine’s kernel time on an RTX 5090.
- **The decomposition and the oracles.** The parts are measured at two budgets of one model; what the oracles leave is split by attribution, not measurement. Oracles read a recorded trace, and on the two-read path the engine does not keep MIN’s set (Section IV). The engine’s only predictor looks one layer ahead; the forecasters of Section VIII were evaluated offline.
- **Outputs and more.** Parity is measured by teacher-forced KL divergence against an all-VRAM reference (mean 0.0019 and 0.0005 nats), not by task accuracy; Appendix L gives the remaining caveats.

XII. CONCLUSION

MIN’s host reads at a machine’s probed rate bound batch-1 MoE decode from host memory, and our cache reaches 31–54% of that bound at the smallest gpt-oss budget. In our engine, what is cached pays only when each admitted expert is read once, and policies without foresight recover little of the rest.

REFERENCES

- [1] A. Ivanov, N. Dryden, T. Ben-Nun, S. Li, and T. Hoefler, “Data movement is all you need: A case study on optimizing transformers,” in *Proceedings of Machine Learning and Systems (MLSys)*, 2021.
- [2] T. Hoefler, D. Alistarh, T. Ben-Nun, N. Dryden, and A. Peste, “Sparsity in deep learning: Pruning and growth for efficient inference and training in neural networks,” *Journal of Machine Learning Research*, vol. 22, no. 241, 2021.
- [3] A. Eliseev and D. Mazur, “Fast inference of mixture-of-experts language models with offloading,” *arXiv preprint arXiv:2312.17238*, 2023.
- [4] L. Xue, Y. Fu, Z. Lu, L. Mai, and M. Marina, “MoE-Infinity: Efficient MoE inference on personal machines with sparsity-aware expert cache,” *arXiv preprint arXiv:2401.14361*, 2024.
- [5] X. Song, Z. Zhong, R. Chen, and H. Chen, “ProMoE: Fast MoE-based LLM serving using proactive caching,” *arXiv preprint arXiv:2410.22134*, 2024.
- [6] K. Kamahori, T. Tang, Y. Gu, K. Zhu, and B. Kasikci, “Fiddler: CPU-GPU orchestration for fast inference of mixture-of-experts models,” *arXiv preprint arXiv:2402.07033*, 2024.
- [7] H. Chen, W. Xie, B. Zhang, J. Tang, J. Wang, J. Dong, S. Chen, Z. Yuan, C. Lin, C. Qiu, Y. Zhu, Q. Ou, J. Liao, X. Chen, Z. Ai, Y. Wu, and M. Zhang, “KTransformers: Unleashing the full potential of CPU/GPU hybrid inference for MoE models,” in *Proceedings of the 31st ACM Symposium on Operating Systems Principles (SOSP)*, 2025.
- [8] S. Zhong, Y. Sun, L. Liang, R. Wang, R. Huang, and M. Li, “HybriMoE: Hybrid CPU-GPU scheduling and cache management for efficient MoE inference,” in *Design Automation Conference (DAC)*, 2025, arXiv:2504.05897.
- [9] S. Yang, X. Fan, M. Pan, H. Xi, Z. Wang, S. Sun, K. Keutzer, S. Han, M. Zaharia, C. Xu, and I. Stoica, “FreeToken: Efficient edge-native MoE serving with bandwidth-adaptive execution,” *arXiv preprint arXiv:2608.16157*, 2026.
- [10] Z. Zhu, G. Li, P. Wang, Z. Mo, M. Pei, Z. Song, X. Liang, and J. Cheng, “DALI: A workload-aware offloading framework for efficient MoE inference on local PCs,” *arXiv preprint arXiv:2602.03495*, 2026.
- [11] R. Hwang, J. Wei, S. Cao, C. Hwang, X. Tang, T. Cao, and M. Yang, “Pre-gated MoE: An algorithm-system co-design for fast and scalable mixture-of-expert inference,” in *Proceedings of the 51st Annual International Symposium on Computer Architecture (ISCA)*, 2024, arXiv:2308.12066.
- [12] S. Zhu, S. Bohl, R. Oester, and G. Alonso, “Pre-attention expert prediction and prefetching for mixture-of-experts large language models,” *arXiv preprint arXiv:2511.10676*, 2025.
- [13] Z. Wang, Y. Wang, L. Gong, C. Tang, W. Lou, T. Wang, C. Wang, and X. Zhou, “SeqMoE: Toward full-load performance via predictive and graph-compatible MoE offloading,” *arXiv preprint arXiv:2609.12978*, 2026.
- [14] OpenAI, “gpt-oss-120b & gpt-oss-20b model card,” *arXiv preprint arXiv:2508.10925*, 2025.
- [15] A. Yang *et al.*, “Qwen3 technical report,” *arXiv preprint arXiv:2505.09388*, 2025.
- [16] G. Gerganov and contributors, “llama.cpp: LLM inference in C/C++,” <https://github.com/ggml-org/llama.cpp>, 2023–2026.

- [17] L. A. Belady, "A study of replacement algorithms for a virtual-storage computer," *IBM Systems Journal*, vol. 5, no. 2, pp. 78–101, 1966.
- [18] L. Zhang, R. Karimi, I. Ahmad, and Y. Vigfusson, "Optimal data placement for heterogeneous cache, memory, and storage systems," *Proceedings of the ACM on Measurement and Analysis of Computing Systems (POMACS)*, vol. 4, no. 1, pp. 6:1–6:27, 2020.
- [19] M. C. Carlisle and E. L. Lloyd, "On the k -coloring of intervals," *Discrete Applied Mathematics*, vol. 59, 1995.
- [20] A. Jain and C. Lin, "Back to the future: Leveraging Belady's algorithm for improved cache replacement," in *Proceedings of the 43rd International Symposium on Computer Architecture (ISCA)*, 2016.
- [21] T. Kalibera and R. Jones, "Rigorous benchmarking in reasonable time," in *International Symposium on Memory Management (ISMM)*, 2013.
- [22] P. Cao, E. W. Felten, A. R. Karlin, and K. Li, "A study of integrated prefetching and caching strategies," in *Proceedings of ACM SIGMETRICS/RICS*, 1995.
- [23] A. Jain and C. Lin, "Rethinking Belady's algorithm to accommodate prefetching," in *International Symposium on Computer Architecture (ISCA)*, 2018.
- [24] S. Albers, "On the influence of lookahead in competitive paging algorithms," *Algorithmica*, vol. 18, no. 3, pp. 283–305, 1997.
- [25] D. Breslauer, "On competitive on-line paging with lookahead," BRICS, Tech. Rep., 1995.
- [26] L. Chen, Z. Wen, T. Wu, X. Zhang, and C. Wu, "SP-MoE: Speculative decoding and prefetching for accelerating MoE-based model inference," *arXiv preprint arXiv:2510.10302*, 2025.
- [27] S. Li, J. Lin, D. Ge, and Y. Ye, "MoE-SpAc: Efficient MoE inference based on speculative activation utility in heterogeneous edge scenarios," *arXiv preprint arXiv:2603.09983*, 2026.
- [28] A. Saxena, P.-A. Tsai, H. Taneja, A. Jaleel, and M. Qureshi, "Cascade: Utility-driven speculative decoding for mixture-of-experts," in *Proceedings of Machine Learning and Systems (MLSys)*, 2026, arXiv:2506.20675 (2025).
- [29] P. Tang, J. Liu, X. Hou, Y. Pu, J. Wang, P.-A. Heng, C. Li, and M. Guo, "HOBbit: A mixed precision expert offloading system for fast MoE inference," *arXiv preprint arXiv:2411.01433*, 2024.
- [30] Y. Zhang, S. Aggarwal, and T. Mitra, "DAOP: Data-aware offloading and predictive pre-calculation for efficient MoE inference," in *Design, Automation & Test in Europe Conference (DATE)*, 2025, pp. 1–7, arXiv:2501.10375.
- [31] S. Yadav and B. Asgari, "Mira: Memory-efficient MoE inference using adaptive caching and predictive expert staging," *arXiv preprint arXiv:2609.38090*, 2026.
- [32] R. Cai, Y. Ro, G.-W. Kim, P. Wang, B. Ehteshami Bejnordi, A. Akella, and Z. Wang, "Read-ME: Refactorizing LLMs as router-decoupled mixture of experts with system co-design," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 37, 2024, pp. 116 126–116 148, arXiv:2410.19123.
- [33] J. Zhang, L. Wu, M. Darbari, and L. Hong, "WiSP: A working-set view of mixture-of-experts serving on extremely low-resource hardware," *arXiv preprint arXiv:2606.21868*, 2026.
- [34] H. Zhang, "Budgeting bytes: A windowed storage roofline and dual-budget architecture ablations for storage-bound LLM decoding," *arXiv preprint arXiv:2609.04238*, 2026.
- [35] Y. Jiang *et al.*, "MoE-CAP: Cost-accuracy-performance benchmarking for mixture-of-experts systems," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2025, arXiv:2412.07067.
- [36] S. Williams, A. Waterman, and D. Patterson, "Roofline: An insightful visual performance model for multicore architectures," *Communications of the ACM*, vol. 52, no. 4, pp. 65–76, 2009.
- [37] Y. Zhang, "Reproducible evaluation of MoE expert caching: Replay semantics, workload contamination, and operating regimes," *arXiv preprint arXiv:2608.07911*, 2026.
- [38] J. Liang, S. Wang, M. Tian, Y. Li, D. Tang, and Z. Wei, "Not all models suit expert offloading: On local routing consistency of mixture-of-expert models," in *International Conference on Learning Representations (ICLR)*, 2026, arXiv:2505.16056.
- [39] D. S. Berger, N. Beckmann, and M. Harchol-Balter, "Practical bounds on optimal caching with variable object sizes," *Proceedings of the ACM on Measurement and Analysis of Computing Systems (SIGMETRICS)*, vol. 2, no. 2, 2018.