

Supplementary Material for “Where the Seconds Go: Bounding Mixture-of-Experts Decode from Host Memory, and Measuring What Foresight Is Worth”

Harshith Kantamneni
Independent Researcher
harshithkantamneni@users.noreply.github.com

Sections, tables, figures and equations numbered without an S are in the main paper; those numbered S1, S2, ... are here. The scorecard, which lists every registered clause, is a separate document (`supplement.pdf` in the artifact).

APPENDIX A MAP OF THE APPENDICES

- Appendix B: the scientific benchmarking checklist, rule by rule.
- Appendix C: the configuration codes and host names used below, and the terms.
- Appendix D: every registered prediction and how it scored, behind Table I’s Registered threshold and Result columns.
- Appendix E: the oracle jobs behind Section V’s factorial (Fig. 3).
- Appendices F and G: the trace study and the window policies behind Section VIII.
- Appendix H: the closed-form account of the time used in Section III, and its registered tests.
- Appendix I: MIN’s schedules machine by machine (Section V) and the probe’s sensitivity (Section III).
- Appendix J: the expert cache, the read-time microbenchmark and the system comparison (Section III).
- Appendix K: related work in full (Section X); Appendix L: the limitations in detail (Section XI).
- Appendix M: the published measurements scored against their bounds (Section III).
- Appendix N: the artifact.
- The scorecard (`supplement.pdf`), a separate document: every clause scored, and three earlier analyses the main text no longer uses: a calibrated time model, the comparison on other GPUs, and an accounting of the gap as a model.

APPENDIX B SCIENTIFIC BENCHMARKING CHECKLIST

How this paper follows the twelve rules of Hoefler and Belli [S1].

Rule	Where and how
1. Report the base case of a speed-up and its absolute performance.	Every ratio in Table III is next to both absolute speeds; llama.cpp is configured with per-layer expert offload at equal expert memory, the setting its users run.
2. Explain subsets of benchmarks or resources.	Two models chosen for size and precision (MXFP4, BF16), a third that is FreeToken’s own headline model; budgets span what fits a 32 GB card; the 43.75% budget is omitted where it does not fit. All 30 AIME-25 problems are used.
3. Arithmetic mean for costs, harmonic mean for rates.	The accounting, the panel and the value of foresight average time per token over problems (a cost) and quote speeds as the harmonic mean this implies; Table III averages per-problem rates arithmetically, which departs from this rule; every problem decodes the same number of tokens, so the harmonic mean is the right summary, and ratios of total time agree with the reported ratios of mean rates within 0.003.
4. Avoid summarising ratios.	Ratios of means are never averaged across budgets or models; ranges list the per-budget values, and the medians reported across models (Table S12) or across published rows (Appendix M) are medians of per-cell quantities, labelled as such.
5. Report whether measurements are deterministic; give confidence intervals.	Decode is greedy and deterministic per engine; timing is not. Every comparison has a 95% interval over problems; the panel (job 099) adds intervals over machines, the unit that varies most.

Rule	Where and how
6. Do not assume normality.	Intervals over problems are bootstrap percentile intervals (10,000 resamples, paired by problem), with no distributional assumption. Intervals over machines (3 to 15 of them) are <i>t</i> -intervals, which assume the per-machine values are roughly normal; the per-machine values are given beside them (Table VI, the scorecard).
7. Compare nondeterministic data soundly.	Paired bootstrap of the ratio of means within a machine; across the panel's machines, a two-stage bootstrap (machines, then problems) [S2]; the variant a system is scored with is picked on a separate launch (except the Qwen3 25% row of Table III).
8. Check that central tendency is the right summary.	Mean speed over 256 tokens is FreeToken's own metric; per-token p50 and p99 latencies are recorded for every problem in the released rows.
9. Document all factors and the setup.	Machines, probes, power limits, memory clocks (Section II), software commits (llama.cpp 4da6337, FreeToken 0d652e7), flags and job scripts are released; Appendix N.
10. Report measurement, synchronisation and summarisation.	One client times every system of Table III from streamed token arrivals; the oracle comparisons of Section V are timed in process by llama-ec-bench over teacher-forced text, every policy by the same timer (Section II).
11. Show upper performance bounds.	The bound (Section III) is a column of Table III; Table S24 tightens it four ways; Section III applies it to the published literature and to our cache with the same procedure.
12. Plot enough to interpret; connect points only for trends.	Figure 3 plots every machine's ratio with its interval; Fig. 5 connects measured windows for the trend only and marks engine measurements separately; Fig. S3 is a bar chart and Figs. S4 and 4 are scatter plots; the scorecard holds the rest.

APPENDIX C NAMES AND TERMS

Table S1 maps the configuration codes used in the appendix tables, the scorecard and the scripts to the names of the main text; Table S2 defines the other terms.

APPENDIX D PREDICTION LOG

Each job's predictions are in the header of its script on the public `gpu` branch, committed before the machine started. Each machine cloned the published branch at boot and ran the header as committed, so git history orders every prediction before its measurement. Of the 101 rentals of jobs 093–114, every commit preceded its rental by at least 2.7s, and every machine started its job at least 26s after the push. For seven of them the push and the rental's creation fall within a second of each other, and for four GitHub records the push up to 0.8s after the rental was created (`scripts/reg_timing.py`). The results of jobs up to 114 do not record which commit of the branch a machine ran; its boot script now writes it into the results. Table S4 summarises every job's clauses, and Table S5 gives the failures. Outcome notes with the statistics are in `prereg/`; the batched-verification variants of Section VIII (rejected drafts free, routed like the true token, or routed at random) are in `prereg/batchk_outcome.md`. Jobs before 073, which built and debugged the engine, also registered predictions in their headers. Their outcome notes are in `prereg/`, and they are not scored here; job 072's, for one, failed or did not apply on two of its three.

Scoring: Every prediction of jobs 073–114 is split into its separable clauses, and each clause is scored under one rule: *held* when its point estimate is on the predicted side and the 95% paired-bootstrap interval excludes the threshold (for a band, lies inside it); *held (point)* when the point estimate is on the predicted side but no interval exists or the interval crosses the threshold; *failed* when the point estimate is on the wrong side; *untested* when the clause could not be evaluated; and *void* when its job was voided. Jobs 073–098 were scored by hand from the outcome notes, into `prereg/scorecard_clauses.json`, which generates the scorecard. The hand scoring makes one exception: deterministic counts such as reads have no sampling error, so it holds them outright, not held (point). From job 099 on, a script scores each job's clauses, without that exception.

Two tallies: The hand-scored tally covers the 565 clauses of jobs 073–098: 218 held, 193 held on the point estimate, 142 failed, 8 were untested and 4 void. By era, 7 of 19 scored clauses held in jobs 073–075, when the system was being built; 19 of 23 in jobs 076–081, 11 of them sign-only predictions made after a same-machine pilot; 24 of 38 in jobs 082–087; and 168 of 473 in jobs 088–098. The last row of Table S4 is the overall tally, for jobs 073–114. The hand-scored tally is a subset of it:

TABLE S1

The configuration codes of the appendix tables, the scorecard and the scripts, and their names in the main text (Table II).

Code	Main-text name, and what it is
base	the deployed cache
foa	Dep-1R: the deployed admissions, fetched into the slot in the step
aa	admit every miss
dk	Margin: the deployed cache with a larger admission margin (admitting less)
lrn	the deployed cache ordered by a learned reuse predictor
pf	LA: the deployed cache with the engine's layer-ahead copy
R1	LA-1R: Dep-1R with the layer-ahead copy (job 109)
R2	LA-1R-margin: LA-1R admitting less (the margin of dk)
fetch	MIN-1R: the greedy schedule, admissions fetched in the step
bypass	MIN-2R: the greedy schedule, served by the CPU, then copied
fetchplan	Few-1R: MIN's fewest-admission schedule, fetched in the step
bypassplan	Few-2R: the same schedule, served by the CPU, then copied
bypassplanS	Few-2R-early: Few-2R with each copy issued one step earlier (job 112)
base0, bypass0, bypassplan0	the deployed cache, MIN-2R and Few-2R with the in-step fetches off (<code>fetch=0</code>): every miss served by the CPU (job 113)
lead2	ahead only: MIN's admissions copied two steps ahead of use, other misses served by the CPU
both2, both3p	read-ahead oracle: MIN-1R plus copies two (three, paced) steps ahead
hitopt, hitoptp	Belady-2R (paced copies)
nb2	Belady-1R, copied two steps ahead
w4, w16	window $W=4, 16$: admit every miss unless the window shows the victim needed sooner
w1, w2	admit every miss, with exact foresight of the next 1 (2) steps
w8r5, allr5	admit every miss, with foresight of 8 steps (the rest of the problem), each future expert seen with probability 0.5
b4, b16	deployed + window $W=4, 16$
b8r5	deployed + window $W=8$, each future expert seen with probability 0.5
learned	Dep-1R with the learned admission and victim order (job 097; not the same as lrn)

TABLE S2

Terms used throughout.

cell	a model at a budget: the share of each layer's experts resident on the GPU
E, k, C, L	experts per layer, experts per token, GPU slots per layer, MoE layers
miss	a requested expert not resident: read from host memory by the CPU or by a copy
fetch	copy a miss into a slot and run it there in the same step (one read)
admission	making an expert resident; <i>serve-then-copy</i> : the CPU serves the miss, then a background copy reads it again
prefetch, lead d	copy an expert into a slot at least d steps before its first use (one read)
paced	copies in 16 MB pieces, one in flight
MIN, R^*	Belady's optimum with bypass on the routing trace; its host reads per token, the fewest of any policy
R_{eng}^*	MIN's host reads per token in the engine's oracle: its window ends at the problem's end, and an expert served in a step stays for that step
1R, 2R	an admitted expert read once (fetched into its slot in the step) / twice (served by the CPU, then copied in the background)
fast / slow link	a link-to-CPU ratio B_p/B_c of at least 0.5 / below 0.5
online / oracle	a policy that sees only past routing / one that reads future routing from a recorded trace
$B_c, B_p, B_{cp}, B_{\text{host}}$	probed host read rates: CPU, link, both, highest by any method
bound	Eq. (1): the speed bound of a machine for a cell
host-/GPU-bound	the bound binds on host reads / on GPU work
$v(F), v(O, F)$	the modelled time with foresight alone / with foresight and overlap (the scorecard)
W_{50}	tokens of perfect foresight that close half of an online policy's excess reads
hosts B, S	the two hosts of the system table (Table III)
host A	the host of the ablation (Fig. S3)
O1–O6	the hosts of the oracle jobs 093–097 (Table S9, Appendix E)
Pa–Pj	the hosts of the panel, job 099 (Appendix D); later jobs return to Pe, Pf and Pg

the table's rows for jobs 073–098 are the hand scores, and the script-scored jobs 099–114 make up the rest. In the overall row, a deterministic count is therefore held in jobs 073–098 and held (point) after.

Scoring checks: A script (`scripts/scorecard_auto.py`) re-scores every hand-scored clause whose measured value is a number and whose threshold parses, 469 of 565. It agrees with the hand score on 432. Of the 37 differences, 36 are the deterministic counts (reads, divergences, table entries) that the hand scoring holds outright and the script scores held (point), and 1 is a boundary value the hand scoring read unrounded (1.249 against a band starting at 1.25). Of the clauses the script scores held on the point estimate, 213 have no interval and 7 have one that crosses the threshold. The re-scoring is a check; neither tally uses it.

TABLE S3

Where to check each pre-specified claim of Table I: the data behind it, and the jobs, predictions (P, Q, T, H numbers in each job’s header) and clause identifiers in the scorecard’s tables (`supplement.pdf`; x: the machine’s letter).

Claim of Table I	Data	Jobs, predictions and clauses
The bound’s read time is nearly reachable	Table S25	job 102, P2 (layer mode at least 0.50): 102x-P2floor-C14/C32
Faster than FreeToken and llama.cpp	Table III, Appendix J	jobs 081 (host B) and 089 (host S), hand-scored (Table S4)
On the deployed read path, MIN’s set pays only together with one read	Table IV, Fig. 2	job 099, P5a: 099Px-P5a-g11/g25; job 109, P1: 109x-P1-int-g11/g25
Control: issuing the copies a step earlier	Table S6	job 112, T1-T4: 112x-T1 to 112x-T4
With the in-step fetches off, MIN’s set pays alone, and more with one read	Table V, Table IV	job 114, H1-H6: 114x-H1a-g11/g25, 114x-H1b-*, 114x-H1c, 114x-H2 to 114x-H6
Without the fetches, the fewest-admission set pays read twice	Table S7	job 113, H1-H6: 113x-H1-g11/g25, 113x-H2 to 113x-H6
Policies without foresight recover little	Table VI	job 109, P7: 109x-P7-best, 109x-P7-cap; jobs 111 and 112, Q3
The fewest-admission set read once gains	Fig. 4, Appendix I	job 104, P3 (machine a); job 105, P4; job 106, P6 and P7; job 107, P5: 104a-P3-g11, 105x-P4-g11, 106x-P6-*, 106x-P7-*, 107x-P5-fetchplan
The ratio predicts the oracles on a second card	Table VI, Fig. 4	jobs 111 and 112, Q1 (and Q6 for the interaction): 111x-Q1-*, 112x-Q1-*
The closed-form account (not in Table I)	Appendix H	jobs 105–108, P1: 105x-P1-* to 108x-P1-*

Rounds and rentals: A machine’s intervals are over problems. To see what they leave out, `scripts/noise.py` takes every valid machine of jobs 109, 111, 112, 113 and 114 that ran a configuration in two rounds (129 machine-budget-configurations on 17 machines; 25% ran one round in these jobs) and every machine measured with the same configurations in two launches (69 comparisons from 10 pairs of launches on 9 machines: job 099’s panel and its relaunched in jobs 100, 101 and 103, job 109 and job 112, job 110 and job 111). The speed ratio moved by a median 0.2% (at most 1.8%) between rounds, against a median half-width of 0.6% for the interval over problems, and by a median 0.5% (at most 5.3%; above 5% in 2 comparisons) between launches. Each launch derives its fetch table from its own probe; the tables differed in 4 of the 10 pairs, which give the largest differences; over the 39 comparisons with the same table the median is 0.3% and the largest 2.7%. Each machine’s interval for each configuration was then widened, on the log scale and on both sides, by that configuration’s full round-to-round range on that machine (zero at 25%), and the derived intervals (capture, job 112’s differences, job 114’s interaction in ms) accordingly. All 112 per-machine clauses of these jobs that held with an interval still hold; the 4 pooled clauses, whose intervals are over machines, are not widened. With every interval widened further by the largest rental difference, 61 still hold and 51 hold on the point estimate only. Of those, 33 are the bands of ± 0.06 or narrower around configurations that barely change the speed (job 109’s P2 and P10, job 110’s Q1 for Dep-1R and MIN-2R, job 113’s H3 on one machine, job 114’s H4 and H6); the other 18 are the second-card trend checks for MIN-1R and the read-ahead oracle and the no-foresight bound (Q1 and Q3) on the RTX 4090s, P8 on one machine, T3 on one, job 113’s H5 and H6 on one machine each, and job 114’s H3 at 11% on two machines and its interaction on all three. No clause changes to failed: failure is decided by the point estimate.

Failed predictions: Most bands that failed in jobs 073–098 are the sizes of effects whose direction held. In jobs 088–098, 90 of 301 bands failed, as did 23 threshold, 3 equality and 7 sign clauses; 42 of 82 sign clauses held with an interval, and 75 held at least on the point estimate. Table S5 gives the failures job by job: those that matter for jobs 073–098 (the scorecard lists all 142), and the families of failures of jobs 099–114.

TABLE S4

Every registered job from 073 on and how its clauses scored: *held* with a 95% interval on the predicted side (or outright, for deterministic counts), held on the point estimate only, failed, and untested or void. Jobs 073–098 were scored by hand, the rest by script; the scorecard lists every clause. *Cond.*: a job’s validity condition (at least k valid machines), met or not; it is operational, not a prediction, and is not counted in the other columns.

Job	What it tested	Held	Point	Failed	Untested	Cond.
073	ours against FreeToken, same host	3	0	4	0	–
074	the 40% budget fixed	0	2	4	0	–
074b	server overhead	2	1	0	1	–
075	RTX PRO 6000 at 40–60%	2	0	1	0	–
076	the system table on an RTX 5090	2	1	0	0	–
077	competitors, gpt-oss	3	0	0	0	–
078	Qwen3, four systems	3	0	0	0	–
079	Qwen3 profile	1	3	0	0	–
080	the fetch split	4	0	0	0	–
081	the system table with the calibrated fetch split	6	0	0	0	–
082	held-out prompts and steady state (review run 1)	4	0	3	1	–
083	gpt-oss, long outputs on hard held-out problems	5	0	0	0	–
084	all-in-VRAM column and traces (voided: throttled card)	0	0	0	4	–
084b	all-in-VRAM column and traces, rerun	2	2	1	0	–
085	a slower PCIe link, gpt-oss	5	0	1	0	–
085b	a slower PCIe link, Qwen3	5	0	1	0	–
086	Qwen3.6	1	1	2	0	–
087	ablation with a static cache	2	1	2	0	–
088	slow link, fair comparison	5	3	3	0	–
089	the system table on a second host	9	7	12	0	–
090	output parity against all-in-VRAM (RTX PRO 6000)	10	0	0	0	–
091	the grid on an RTX 4090	11	3	5	0	–
092	the grid on an RTX 3090	6	1	7	1	–
093	first foresight oracles	9	2	21	4	–
094	oracles that minimise bytes	12	1	24	1	–
095	single-read oracles	44	0	12	0	–
096	the oracle factorial	29	97	12	0	–
097	the learned admission order	11	53	20	0	–
098	FreeToken tuned	22	15	7	0	–
099	the host panel	190	42	39	2	–
100	windows of foresight	41	122	27	26	–
101	the crossover by link ratio	6	19	0	0	–
102	the read-schedule microbenchmark	0	28	5	0	–
103	fewest-admission schedule (plan not run)	0	4	0	41	–
104	fewest-admission schedule	1	22	24	0	–
105	Eq. (S1), plain form; layer-ahead copy; fewest-admission set	7	20	4	0	–
106	Eq. (S1) with overlap; admitting less; slow links by round	21	112	48	0	–
107	Eq. (S1) on new machines; admitting less; fewest-admission loads	11	22	12	0	not met
108	Eq. (S1) on new desktop-class machines	0	12	6	2	met
109	the 2×2 and policies without foresight on new machines	47	43	0	0	met
110	the RTX 5090 trend on RTX 4090s (stopped by V1)	0	0	0	1	not met
111	job 110 relaunched, V1 corrected	24	13	0	0	met
112	higher-ratio RTX 4090s; timing control; second probe	22	23	12	0	met
113	the two-read path with the in-step fetches off	6	6	2	0	met
114	the 2×2 with the in-step fetches off	17	14	2	0	met
all		611	695	323	84	6 met, 2 not

TABLE S5

Failed predictions, by job. *Predicted*: the registered band, threshold or sign. Jobs 073–098 were scored by hand, and the table gives the failures that matter; from job 099 on, each job’s line gives its tally as its script scored it, without its validity condition, which Table S4 counts apart (job 107’s was not met). “—”: the reason, where known, is in the job’s outcome note. The scorecard lists every clause.

Job	What was predicted	Predicted	Measured	Why it failed, or what it means
<i>Scored by hand, jobs 073–098</i>				
085b	A slower PCIe link, Qwen3. Our lead over FreeToken larger than in Table III, at 2 or more of 3 budgets (the job’s own restatement of its fourth prediction)	at least 2 of 3	1 of 3	Downgraded from the running log; the pooled 4 of 6 held
088	A slow link, fair comparison. Our lead over FreeToken	1.10–1.50×	1.51–1.58×	Above the band; the lead itself held
089	The system table (Table III) rerun on a second machine. Our ratio over FreeToken, against host B’s Our speed at the host-bound cells, against host B’s LRU slower than decayed frequency, at the two GPU-heavy cells LRU still ahead of FreeToken at gpt-oss 11%	within ± 0.06 within $\pm 12\%$ 10–30% slower ahead	0.06–0.11 below, at five of six cells 14–17% lower 6–8% slower behind	— — — A sign failure
091, 092	The grid on an RTX 4090 and on an RTX 3090 (the scorecard). Our lead over FreeToken; our share of the limit Ours at least 1.8× llama.cpp, on the slow machines The 4090 machine’s FETCH tables Ours at least 1.3× llama.cpp on Mixtral The two launch orders agree	10–35% (gpt-oss 11%), 0–15% (Qwen3); 25–45% at least 1.8× the registered entries at least 1.3× within 3%	above the bands 1.56–1.78× one more expert copied, in two tables a tie 7%; a second difference, 11% at gpt-oss 11% on the 3090 machine, was reported, not scored	FreeToken’s carried-over backend fell far below its band — An equality failure —
093	The first oracle job (Section V). The hit-optimal oracle faster than the online policy, by 20–80% Its speed at the four host-bound cells Its hit rate and admissions, against the optimum’s The cost of no overlap The cost of all-CPU, Qwen3 The hit-optimal oracle’s gain at the two GPU-bound cells	faster; 20–80% 45–65% of the limit within 3 points; within 30% 5–25% 0–10% below 15%	slower at gpt-oss 11% and Qwen3 12.5%, the two lowest budgets 43.7–44.8% outside, at three and at two of the four cells 0% 15% 48 and 56%	A sign failure as well as a band failure Just below the band — — —
094	Oracles that minimise bytes: the bands The bytes-optimal oracle beats the hit-optimal one, at the two lowest budgets Its 16-token window keeps its gain	missed as Appendix E describes. by at least 15% at least 70% of it	by less less, at three cells	— —
095	Single-read oracles: 44 of 56 clauses held, and the 12 that failed The model on MIN-1R’s counters MIN-1R’s reads, against the simulation’s The read-ahead oracle at the host-bound cells The model’s over-prediction of the read-ahead oracle The paced variant, against the unpaced one	within 10% within 10% 55–85% of the limit 15–50% within 10%	are of size. 13–17% off, at four cells 17% below, at the two GPU-bound cells 44 and 51% at the two 25% cells 7 and 13% at the two lowest budgets 12–13% faster at the two 25% cells	Size The registered simulator gave a resident carried over from the previous sequence no next use, and so evicted it first; corrected, the simulation is within 4% of the engine’s reads at every cell (Appendix E) Size Size
096	The oracle factorial on O4 and O5: the clauses held at most cells of both machines, and Appendix E lists the failures.			

continued

Job	What was predicted	Predicted	Measured	Why it failed, or what it means
098	MIN-1R's gain, O4 against O5	differs by at most 0.15	by more, at five cells	The largest failure of the job: the difference between the machines
	Belady's single-read prefetch, Qwen3 12.5% on O4	at most 1.05× the online policy	above	—
	FreeToken tuned (prereg/freetoken_outcome_098.md): its best setting stayed within the predicted 10% of its default, and our ratios within the predicted bands.			
	FreeToken's best backend, gpt-oss 25% Stock llama.cpp against FreeToken's CPU backend, CPU only (gpt-oss; Qwen3)	hybrid at least 2×; at least 1.2×	offload about equal	An equality failure
	The fixed fetch caps; the thread count	within 8% of the calibrated split; 8 threads within 3% of the default or slower	outside	As the job's outcome note describes
<i>Scored by script, jobs 099–114</i>				
099	The host panel (scripts/panel_099.py; prereg/scorecard_099.json, prereg/panel_outcome_099.md): 273 clauses; 190 held, 42 held (point), 39 failed, 2 untested. Bands set from four earlier machines, none of them link-starved	the registered bands	most failures, on the two machines whose link reads at less than half their CPU's rate	The bands did not allow for dependence on the machine
100	The planning time of the noisy whole-future window	at most 150 μ s per step	above	—
	Windows of foresight (scripts/job100.py; prereg/window_outcome_100.md): 216 clauses; 41 held, 122 held (point), 27 failed, 26 untested.			
	The deployed-path windows miss within 10% as often as a replay that admits at once	within 10%	more misses; most of the job's failures	Their copies land late
	The probe-only model, on the half-width-link machine	within 12%	outside	—
	Four other bands or orderings	as registered	at their edges	—
101	The crossover by link ratio (scripts/job101.py; prereg/crossover_outcome_101.md): 25 clauses; 6 held, 19 held (point), 0 failed, 0 untested.			
102	The read-schedule microbenchmark on three machines (scripts/job102.py; prereg/readsched_outcome_102.md): 33 clauses; 0 held, 28 held (point; none has an interval), 5 failed, 0 untested. Waiting per layer costs time against the analytic schedule	at least 0.03 of the host term	nothing measurable on two of the three machines	Four of the five failures
103	The order of the two budgets, Intel machine	as registered	the other way	—
	MIN's fewest-admission schedule against the greedy one; the plan did not run (scripts/job103.py): 45 clauses; 4 held (point; the relaunch of Pf), 41 untested.			
	Job 103 repeated, with the plan run (prereg/minadm_outcome_104.md): 47 clauses; 1 held, 22 held (point), 24 failed.			
104	Few-1R's copies and misses at gpt-oss 25%, against MIN-1R's	copies at most 0.70×; misses within 3%	0.73×; 5.3% more	—
	The engine's misses, against the machine's replay	within 2%	more	—
	Few-2R's reads, against MIN-2R's	at most 0.90×	more	—
	The 2×2's interaction, with the plan's set	shrinks	grew	—
105	Few-2R against MIN-2R	at most 0.01 behind	0.047 behind, on Pg at gpt-oss 25%	Held at the other 5 of 6 cells
	The plain form of Eq. (S1) with G profiled on each machine, LA and Few-1R (scripts/job105.py; prereg/sumlaw_outcome_105.md): 31 clauses; 7 held, 20 held (point), 4 failed. Every failure is the relation at gpt-oss 25%.			
	The relation at gpt-oss 25%, per machine	within 6%	over-predicts the deployed cache by 6.9–9.3% on 3 of four machines (3.8% on Pf)	At 11% it is within –1.5 to 4.7%
106	Its pooled median error	at most 4%	4.2%	The 25% cells take it past the band
	The relation of Eq. (S1) with its overlap term on both models at two budgets, the two online admission rules, and MIN's sets launch by launch, on five machines (scripts/job106.py; prereg/onlineadmit_outcome_106.md): 181 clauses; 21 held, 112 held (point), 48 failed: 36 on the two unstable machines, 11 on the 3 stable ones and 1 pooled.			
	Clauses on the two server machines new to the job	as registered	36 failed	One computed wrong outputs; the other varied by more than the registered 2% between rounds (Appendix I)
	Margin's speed-up at gpt-oss	at least 1.01	below, on two stable machines	—
	The learned order's speed-up	at least 1.00	below, on the fast-memory machine	—

continued

Job	What was predicted	Predicted	Measured	Why it failed, or what it means
	The learned order's host time	at most 400 μ s (gpt-oss), 700 μ s (Qwen3) per step	above, on the fastest-link machine	—
	The relation's prediction of both rules' ratios, Pf at 25%	within 0.03	off by 0.030 and 0.045	—
	The relation's median error over all machines	at most 4%	above	—
107	The relation, Margin and MIN's fewest-admission set loaded both ways, on machines never rented before (<code>scripts/job107.py; prereg/newhosts_outcome_107.md</code>): on its two valid machines, 45 clauses; 11 held, 22 held (point), 12 failed.			
	The relation, per cell	within 6%	beyond, at 7 of 8 cells	—
	The relation's pooled median error; the overlap form against the plain one	at most 4%; overlap form better	above; not better	—
	Margin's speed-up	at least 0.995	below, at one cell	—
	The relation's prediction of Margin's speed, on the server	within 0.03	outside, at two cells	—
	Valid machines	at least three	two	Inconclusive by its registered rule
108	The relation on machines never rented before, gpt-oss only (<code>scripts/job108.py; prereg/smallhosts_outcome_108.md</code>): 20 clauses; 12 held (point), 6 failed, and the rest could not be tested.			
	The relation on the EPYC 7543, both budgets	within 8%	outside	—
	Its implied read rate, both budgets	0.85–1.25 of B_{host}	outside	—
	Pooled: the cells within 6%; the median error	at most one cell beyond 6%; at most 4%	2 of 6 within 6%; 6.7%	—
109	The 2 $\times$ 2 of Section IV, the read-ahead oracle and four variants without foresight, on desktop-class machines never rented before (<code>scripts/job109.py</code> ; Table VI): 90 clauses; 47 held with a 95% interval over problems, 43 held (point), 0 failed. Of its 11 predictions 11 held; one clause of P8, for ratios of at least 0.9, had no machine to test it.			
110	The RTX 5090 trend of Section VII, frozen in its header, on RTX 4090 machines: void as registered.			
	Valid machines	at least 3	none	The output check used the RTX 5090's reference loss
111	Job 110 relaunched with the RTX 4090's own reference: job 110's clauses on three valid machines; 24 held with an interval, 13 held (point), 0 failed. Its sixth prediction, for machines with a ratio of at least 0.5, had none to test.			
112	The second card at higher ratios, the timing control, a second probe and two relaunches of job 109's machines (<code>scripts/job112.py</code> ; Tables S6 and VI): 57 clauses; 22 held with an interval, 23 held (point), 12 failed.			
	Few-2R-early's misses	at most 0.95 of Few-2R's	0.97–0.98 at 11%, 0.95 at 25%; failed on 5 of 8 machine-budgets	The replay of the publication rule predicted 0.91; the engine's fetch table keeps MIN's set from forming
	Few-2R-early no slower than Few-2R	at least –0.01	–0.05–0.03 at 11%, –0.03–0.01 at 25%; failed on 7 of 8	The early copies admitted and read more; the control cannot separate landing from reading
113	The two-read path with the in-step fetches off (<code>scripts/job113.py</code> ; Table S7): 14 clauses; 6 held with an interval, 6 held (point), 2 failed.			
	The deployed cache with the fetches off	0.80 to 1.00 of the deployed cache	0.98–1.07 at 11%; failed on the Core Ultra 9 285K	The probe-set fetch table slows the deployed cache on that machine
	One read over the fetch-free two-read arm	at least 0.05	0.02–0.16 at 11%; failed on the Core Ultra 9 285K	Once the set is held, the second read costs little there
114	The 2 $\times$ 2 with the in-step fetches off (<code>scripts/job114.py</code> ; Table V): 33 clauses; 17 held with an interval, 14 held (point), 2 failed.			
	Dep-1R's in-step fetches without the table	at most 0.5 of those with it	1.00 on the Ryzen 9 5950X with an all-zero table	Its probe gave an all-zero fetch table; there were no table fetches to remove
	MIN-2R without the fetches over MIN-2R with them	at least 1.03	1.00 on the same machine	Likewise: on that machine the two arms are the same configuration

Changes against the running log: Against the running log kept during the study, two clauses are downgraded. A llama.cpp clause of job 074b that the log counted as held was never run, and job 085b's own restatement of its fourth prediction, “at least 2 of 3”, failed at 1 of 3, though the pooled 4 of 6 held. Job 073's “identical outputs” is scored on fully identical outputs (0–3%) rather than on agreement of the first 160 characters (27–47%). Ten clauses the log counted as held, among them job 087's first, hold on the point estimate only under the interval rule. The file `prereg/scorecard_outcome.md` lists every change.

Reruns and rentals without results: Job 084 ran on a throttled card and was rerun, with its predictions unchanged, as 084b/084c. Job 085's Qwen3 half ran as 085b on a second machine with a slow link after a hung conversion. Rentals that produced no results were destroyed and are in the rental ledger (`gpu/vast_ledger.json`): two of job 097's four rentals; the first two for 099a and the first for 099d, replaced before any measurement; 100d, which stalled downloading the model; 100e, which never reached its job; 103c, which never started its job; and 105c, 105d and 105g below.

Jobs 099–104: Job 099's two untested clauses are one machine's replay check, whose step failed when a Python package did not install. Job 100's header gained a ninth prediction, about the two replacement machines, after the first four machines

had started; it compares the replacements with values one of which was already measured. Job 100's untested clauses are the planning time of the deployed-path windows, which the engine does not time, and the paired prediction for a machine that never reached a timed step. Job 101 relaunched the lowest-ratio panel machine and added a second Ryzen 7 9800X3D behind a slower link. In job 103 the machines could not install the plan's solver beside Debian's numpy, so the engine ran the greedy schedule in the plan configurations, as it logs. Job 104 repeated job 103 with the plan in its own environment on three of the same machines: Pf; a second Core Ultra 9 285K, whose link reads at 0.49 of its CPU rate; and Pg. Job 103's 5950X had turned out to be Pg's second launch. Most of job 104's speed predictions held: **Few-1R** gains on Pf, and gains at least 0.03 over **MIN-1R** at gpt-oss 11% on every machine.

Jobs 105 and 106: Two of job 105's clauses predicted quantities already measured: the profiled G (job 069c) and the plan's copies, which are a property of the trace (job 104). They test that the engine reproduces them. Every speed prediction of job 105 held: **LA** gains at least 3% on the machine with a ratio above 0.8 and loses on both below 0.4, and **Few-1R** beats the deployed cache on every machine, beats **MIN-1R** by at least 0.03 on the three with a ratio under 0.7, and copies at most 0.65 of **MIN-1R**'s experts. Three more launched machines produced no results: 105c never started its job, and the job's bandwidth gate stopped 105d and 105g, whose cards were throttled, before any timed run. Job 106 ran on five machines. Two server machines were new to the job; one computed wrong outputs, and the other varied by more than the registered 2% between rounds (Appendix I). The 3 stable machines are Pf, the Threadripper 9960X of job 105 and panel machine Pe, all rented before, and on them the relation held on every cell. **Margin**'s pooled median speed-up of at least 1.02 held only because of the two unstable machines; on the stable ones it is 1.01.

Jobs 107 and 108: Job 107 had validity gates registered before launch. Of 7 machines rented, 3 stopped at the gate and 2 failed the round check. The test had two valid machines of the three it required, which by its registered rule makes it inconclusive; its per-cell clauses failed on both. The profiled GPU compute, **Margin**'s reads and pooled median, both fewest-admission loads and the crossover between them held. Replacing machines that failed the round check, and the rules that chose the last two machines, were amendments committed before each machine started. Jobs 102, 105 and 106 also had their headers edited after launch, each time to substitute a machine, with the predictions unchanged and marked so in the header. Job 108 tested the relation on machines never rented before, with the class drawn after job 107 gated before any download, gpt-oss only, and an 8% band. Of 6 machines rented, 2 stopped at the gate, 1 failed the round check, and three were valid. The profiled G of the EPYC 7543 could not be tested.

Job 109: The job fixed its population, machines never rented before with a link-to-CPU ratio of at least 0.5, before any started; five machines started, and all five passed every gate and round check. A smoke run on gpt-oss-20b, on a machine outside the sample, first checked that the two new combinations run and compute the same model. One listed offer was no longer offered and was skipped as registered. The job's deadline cut the 25% round of the Ryzen 9 5950X, the slowest machine. The header's statement that no listed offer was in the rental ledger was wrong for one: the Core i9-13900KF's offer had been rented twice before (the first rental for 099d, and 100e), both times without a result. Its GPU had never been measured, and it passed the identity gate. Several of the job's bands were wide against what earlier jobs had shown: the mean share closed by **MIN-1R** (0.25 to 0.55, where the panel had shown 33%), the margin's speed (0.98 to 1.06, where it had run $1.021\times$ the deployed cache at 11% and $1.018\times$ at 25%) and the capture bound (at most 0.35). These held, but could hardly have failed. Job 102's floor of 0.50 for the per-layer read time was loose in the same way: its header's analytic estimate, before launch, was 0.85 to 0.99. The clauses that could have failed are the interaction's sign on every machine and the per-machine bounds on the variants without foresight; job 110's predictions 3 to 5 are range checks of the same kind.

Jobs 110 and 111: Job 110 tested the RTX 5090 trend of Section VII, frozen in its header, on RTX 4090 machines; five machines started. Of these, one stopped at its ratio gate (at 0.2499, below 0.25), and every other one failed its output check after one round. That check used the RTX 5090's reference loss. The RTX 4090 computes a loss 3% higher in every configuration of every process, so the check failed by design rather than by fault, and job 110 is void as registered. Job 111 relaunched three of the four machines that had run a round, with the card's own reference (its first round's loss on one of them) and job 110's predictions unchanged. The fourth, 110e, was not relaunched: its GPU was left out of the relaunch list, an oversight we found after the job. Its first round deviated from the trend by -5.6% at 11%, inside the band, so leaving it out does not change the trend's outcome. In that void round the larger margin captured 78% of the read-ahead oracle's gain, a gain of only 5.6%: where the oracle gains little, capture is unstable. The relaunched machines' first rounds were in the repository when the relaunch was committed. Those first rounds had come within 6% of the trend for **MIN-1R** and the read-ahead oracle; for the two single changes, whose band is ± 0.06 in log, one was outside it (**MIN-2R**, 6.4% off, 0.062 in log). The predictions were committed before any RTX 4090 data, and the corrected reference after one round on these machines. Of job 111's first launches, two failed to download the model, before any model run; they were relaunched once, by an amendment committed before the relaunch.

Job 112: The job's predictions were committed before any of its machines started. It asked whether the RTX 5090 trend carries to RTX 4090 machines at higher ratios (job 110's predictions unchanged, the ratio gate at 0.45); whether landing the two-read arm's copies in time rescues it (**Few-2R-early**, a new plan file for an engine path that had run before, so no smoke

TABLE S6

When the copy lands, against how often the expert is read: MIN’s fewest-admission set read once (**Few-1R**), read twice with the copy published two steps after the admission (**Few-2R**), and read twice with each copy issued a step earlier so that it lands for the next use (**Few-2R-early**).

Machine	Ratio	Speed vs deployed			Misses per token		
		1R	2R	early	1R	2R	early
<i>RTX 5090</i>							
Ryzen 9 5950X	0.74	1.41	1.03	0.98	40.7	55.6	54.5
Ryzen 9 7945HX	0.88	1.25	1.03	1.00	40.7	55.7	54.6
<i>RTX 4090</i>							
Ryzen 7 7800X3D	0.57	1.21	1.03	1.00	40.9	54.7	53.0
Core i5-12400	0.77	1.35	1.03	0.98	40.8	55.7	54.5

Speed relative to the deployed cache on the same machine (every 95% interval over problems within ± 0.01) and host misses per token, at gpt-oss 11%; registered before the machines were rented.

TABLE S7

The two-read path with the in-step fetches off (registered; gpt-oss 11%).

Machine	Ratio	vs Dep.	vs Dep. CPU only		vs Dep.		Few-2R misses
		Dep. CPU only	MIN-2R CPU only	Few-2R CPU only	Few-2R CPU only	Few-1R	
Ryzen 9 5950X	0.74	0.98 $\pm$.01	1.12 $\pm$.01	1.26 $\pm$.01	1.23 $\pm$.01	1.40 $\pm$.01	55.6 / 43.0
Core Ultra 9 285K	0.56	1.07 $\pm$.01	1.09 $\pm$.01	1.23 $\pm$.01	1.32 $\pm$.01	1.34 $\pm$.01	54.7 / 43.0

CPU only: the configuration with no in-step fetches, every miss served by the CPU, each admission copied in the background. *Dep.*: the deployed cache. Speeds are relative to the configuration named above them, on the same machine; subscripts are half-widths of 95% intervals over problems. Misses: **Few-2R**’s per token with the fetch table, and CPU only.

run); whether a second probe after the timed runs agrees with the first; and whether two of job 109’s RTX 5090 machines reproduce on a second rental. One of those two, job 109d’s machine, never finished downloading the model in two hours and was destroyed without results; job 109f’s machine replaced it, started while the download was still running, once two rounds could no longer fit before its deadline. The job’s Core i5-12400 RTX 4090 machine had run job 091’s system grid, a different protocol, and its UUID was allowed by name in the header. The timing predictions came from a replay of the engine’s publication rule; in the engine the early copies barely changed the misses (0.97–0.98 of **Few-2R**’s at 11%, against a registered 0.95 or less) and slowed the arm by up to 0.05 of the deployed cache’s speed (against a registered 0.01): the first clause failed on 5 of 8 machine-budgets and the second on 7 of 8. The early arm admitted 11.7–11.8 experts per token at 11% against **Few-2R**’s 6.2–7.4 and read 4–7% more from the host. The clause that the second read costs more than the late landing (T3) held, but only because the early arm was slower than **Few-2R** on every machine: once the manipulation failed, T3 could not fail, so it carries no evidence. **Few-2R-early** stayed at or below 1.10 $\times$, as registered. The first probe ran while the model downloaded on some machines and after it on others; the second ran on the idle machine. On the two relaunched RTX 5090 machines the deployed cache’s time came within 0.7% of job 109’s, while the probe’s link-to-CPU ratio moved by up to 12% between the two rentals. Table S6 gives the timing control machine by machine.

Job 113: The control that job 112’s counters called for: the two-read arms and a deployed cache with the in-step fetches off (`fetch=0`, every miss served by the CPU, as job 093’s all-CPU arm), so that the two-read arms hold the scheduled set. The predictions came from a model of the probe (CPU reads at B_c in series with the GPU’s non-expert time, background copies overlapped, all host traffic at most the probe’s concurrent rate) and were committed before any machine started. Two of the five listed offers were no longer listed and were skipped as registered; both machines that started, a Ryzen 9 5950X and a Core Ultra 9 285K, were new and valid (0.56–0.74 link-to-CPU ratio), and the job stopped, as registered, once two machines were valid. The offers’ absence at launch was not saved at the time: `results/113_offers_at_launch.txt` on the `gpu` branch transcribes the launch session’s query and was committed after the job. Job 114 saved its offer listing with its registration. The manipulation held: **Few-2R**’s misses fell to 0.77–0.79 of their value with the fetches at 11% and 0.72–0.74 at 25% (registered 0.88 or less). MIN’s set paid read twice (**Few-2R** 1.23–1.26 $\times$ and **MIN-2R** 1.09–1.12 $\times$ the deployed cache without fetches; registered at least 1.10 and 1.03), held on both machines with their intervals. Two clauses failed on the Core Ultra 9 285K: turning the fetches off made the deployed cache faster there, not slower (registered within 0.80–1.00), and reading once beat the fetch-free two-read arm by less than 0.05. At 25%, which ran one round, the fetch-free **Few-2R** ran 1.33–1.36 $\times$ the deployed cache against **Few-1R**’s 1.29–1.43 $\times$. Its clauses are per machine; two machines show what can happen on this engine, not how often it happens across machines.

Job 114: The decomposition with the in-step fetches off, which the reviews of job 113 asked for: the 2×2 on the CPU-only path (the deployed cache, Dep-1R and MIN-2R with `fetch=0`; MIN-1R and the read-ahead oracle, whose forced plans replace the fetch table in every decode step) beside the cells of Table IV, all in one process per round: two rounds at 11% and one at 25%. The predictions came from job 113’s machines and a fit of their times to CPU-served misses, background copies and in-step fetches; they and the list of offers were committed before any machine started. Of the 7 machines that started, three stopped at the first gate because their GPUs had been rented before (the gate admits only new GPUs and job 109’s), and one, a Ryzen 9 9950X3D2, failed the round check (its deployed cache’s two rounds were 6.1% apart) and is not scored. The replacements did not follow the registered order exactly: we chose them from offer listings taken minutes before the ones committed with the results, and in the committed listings a cheaper qualifying offer, a Core Ultra 9 285K, was listed at each replacement and skipped. The header’s statement that none of its listed offers had been rented before was wrong for one (job 089’s machine, which the gate then stopped). The choices were made before any of these machines produced data, so they could not select on outcomes, but they changed which machines were sampled. The three valid machines had link-to-CPU ratios 0.57–0.79; one is job 109f’s. MIN-1R’s misses and fetches per token matched the earlier machines’ within 1% at both budgets, so the fetch table played no part in it, and the CPU-only deployed cache and MIN-2R made no in-step fetch and matched job 113’s counters. One valid Ryzen 9 9950X’s probe gave an all-zero fetch table, so its deployed path already ran without in-step fetches; the two clauses that asked the fetches’ removal to change something (Dep-1R’s fetches at most halved, MIN-2R at least $1.03 \times$ faster) failed there, and held on the other two. Of the job’s 33 clauses, 17 held with their intervals, 14 on the point estimate (the counter checks, which have no interval) and 2 failed; its validity condition, at least two valid machines, was met. The interaction now has an interval over problems: on the CPU-only path at 11% it was 0.4–3.2 ms per token, above zero on all 3 machines; at 25% it was -0.1 to 0.1 ms, above zero on one.

The clause-by-clause tables (565 rows for jobs 073–098, 273 for job 099, 216 for job 100, 25 for job 101, 33 for job 102, 45 and 47 for jobs 103 and 104, 31 for job 105, 181 for job 106, 45 for job 107, 20 for job 108, 90 for job 109, 1 for job 110, 37 for job 111, 57 for job 112, 14 for job 113, 33 for job 114; from job 107 on, plus each job’s validity condition) are in the scorecard, a separate document (`paper/supplement.pdf`), generated by the same scripts.

APPENDIX E FORESIGHT IN THE ENGINE

Four jobs ran oracle policies in the engine, each on a rented RTX 5090 machine with weaker host memory than host B’s. Jobs 093–097 report ratios of mean speeds rather than of mean times per token; recomputed either way, no ratio moves by more than 0.006. The fourth (096) is the confirmatory factorial of Fig. 3: ten configurations per cell, each in a fresh context, in an order shuffled per cell with a recorded seed, on hosts O4 (Ryzen 9 9950X: 43, 45 and 51 GB/s by the CPU, over the link and at best) and O5 (Ryzen 9 9950X3D: 47, 27 and 51 GB/s), with ten predictions per machine committed before either started (`prereg/foresight_outcome_096.md`). They held at most cells. The failures: MIN-1R’s gain above the predicted band at Qwen3 25% on O4 and different between the two machines by up to 0.26 against the 0.15 predicted; Belady-1R faster than predicted at Qwen3 12.5% on O4, with reads below the predicted 1.25 – $1.9 \times$ those of the *ahead-only* oracle (below) at Qwen3 12.5% on both machines and at gpt-oss 11% on O4 (down to $1.09 \times$); the bytes share at Qwen3 12.5% on O4 (60% of the gap against a predicted 25–55%); and the reads Dep-1R saves at gpt-oss 40% on O5, just above the predicted 22%. Dep-1R came to 0.97 – $1.03 \times$ the deployed cache at the host-bound cells, at the bottom of the predicted 0.97 – $1.12 \times$ although its reads fell as predicted, and to 0.95 – $1.02 \times$ at the GPU-bound cells. Table S8 lists every cell of the single-read factorial of Section V on host O3 (a Ryzen 9 9950X whose probe reads 53 GB/s by the CPU at the helper count, 43 over the link, 60 with both and 61.5 at best; the model’s FETCH tables $0, 0, 1, 2, 2$ for gpt-oss and $0, 0, 1, 1, 2, 3, 4, 4, 5$ for Qwen3), and Table S9 the milliseconds per token of the engine states the three jobs realised beside each machine’s bound and the model’s values. Table S10 lists the first job (093), Belady-2R with a window sweep on host O1 (a Ryzen 9 9950X: 44, 46, 49 and 50.5 GB/s; tables $0, 0, 1, 2, 3$ and $0, 0, 1, 2, 3, 4, 5, 6, 7$), and Table S11 the second (094), MIN-2R beside Belady-2R on host O2 (a Ryzen 9 7950X: 44, 46, 50 and 50.6 GB/s; the same tables). On O1 and O2 every cell but gpt-oss 40% and Qwen3 43.75% is host-bound in Eq. (1); on O3, whose memory is stronger, the same four.

The single-read oracles: In MIN-1R the step’s plan is computed before the layer runs from the trace’s record of the current step: the misses MIN with bypass would admit are fetched by the engine’s FETCH operation into their victims’ slots and run from there, and every other miss runs on the CPU; the oracle refuses an admission when no resident’s next use is further than the candidate’s. Reading ahead admits, each step, the non-resident experts whose first use is at least d steps away ($d=2$ unpaced, 3 paced), provided the victim’s next use is further than the candidate’s use after that first one; an expert used sooner than d steps away misses on the CPU (the *ahead-only* oracle), or is fetched as in MIN-1R (the read-ahead oracle; paced with $d=3$, the form of Table II). The same policies simulated on the trace with the engine’s semantics (the lookahead within the sequence, residents carried over with their next use recomputed at each sequence start) predict the engine’s reads and hit rates: MIN-1R’s hit rate within 0.4 points and its reads within 0 to 4%, the read-ahead oracle’s hits within two points (Table S8). Simulated, MIN-1R reads R_{eng}^* experts per token, MIN’s count in the engine: its window ends at the problem’s end, and an

TABLE S8

The single-read oracles (job 095; RTX 5090, Ryzen 9 9950X, host memory 61.5 GB/s at best) beside Belady-2R. Per cell: the deployed cache's speed and its share of this machine's bound; for each oracle with the whole sequence in view, its ratio to the deployed cache (paired, 95% interval), hit rate, host reads per token (CPU misses + fetches + admissions) as a multiple of R^* (each read counted once), share of the bound and the calibrated model on its own counters against its measured time. MIN-1R: MIN with bypass, admitted misses fetched into their slot this step; *ahead only*: MIN's admissions copied when their first use is at least two steps away, other misses served by the CPU; *read-ahead*: MIN-1R and ahead only together; *read-ahead, paced*: the same on the paced copy path with a three-step lead (the read-ahead oracle of Table II); Belady-2R: Belady's rule without bypass, unpaced. The simulation columns give the trace's prediction of reads and hits for MIN-1R (its reads are R_{eng}^* , Appendix E) and the read-ahead oracle.

Cell	deployed		ratio	MIN-1R				ratio	read-ahead				ratio	Belady-2R				ahead only		read-ahead, paced		sim. R_{eng}^*	MIN-1R		sim. read-ahead	
	tok/s	% bound		hits	reads/ R^*	% bound	model		hits	reads/ R^*	% bound	ratio		hits	reads/ R^*	% bound	ratio	% bound	ratio	% bound	hits		reads	hits		
gpt-oss 11%	55.5	46	1.29 [1.28, 1.29]	72	1.04	59	-8%	1.39 [1.39, 1.40]	86	1.23	64	0.91 [0.90, 0.92]	98	1.77	42	1.26	58	1.45	67	39	73	46	87			
gpt-oss 25%	90.3	30	1.32 [1.31, 1.33]	89	1.08	39	-13%	1.49 [1.48, 1.50]	97	1.45	44	1.14 [1.12, 1.15]	99	2.11	34	1.45	43	1.66	50	16	89	22	99			
gpt-oss 40%	128.8	25	1.25 [1.24, 1.27]	95	1.14	32	-14%	1.39 [1.37, 1.41]	99	1.60	35	1.27 [1.25, 1.28]	99	2.20	32	1.38	35	1.49	38	7	95	12	100			
Qwen3 12.5%	32.3	48	1.34 [1.34, 1.35]	74	1.04	64	-9%	1.42 [1.41, 1.43]	84	1.20	68	0.69 [0.68, 0.71]	77	2.47	33	1.22	58	1.41	68	100	74	115	84			
Qwen3 25%	51.8	34	1.38 [1.37, 1.39]	88	1.04	48	-15%	1.48 [1.46, 1.49]	96	1.42	51	1.06 [1.05, 1.08]	99	2.00	37	1.41	48	1.67	57	45	88	61	97			
Qwen3 43.75%	91.4	30	1.26 [1.25, 1.28]	96	1.08	38	-17%	1.41 [1.39, 1.42]	99	1.62	42	1.27 [1.26, 1.29]	100	2.29	38	1.40	42	1.50	45	15	96	23	100			

TABLE S9

The accounting, measured: milliseconds per token (1000 over the mean speed; the main text averages times per token, which gives values up to 1.3% higher on these runs) of the engine's states that the oracle jobs realised, on each oracle machine, beside that machine's bound and the model's values for the states it prices; O1–O3 ran jobs 093–095 in a fixed order, O4 and O5 job 096 in a shuffled order. *Deployed*: the deployed cache; *Dep-1R*: the same with each admission fetched into its slot in the step that misses it (one read, no foresight); *no ovl.*: the deployed cache with its within-layer overlap off; *all CPU*: with no FETCH table; MIN-2R: MIN with bypass with background admissions (two reads per admitted expert); Belady-2R: Belady's rule without bypass, which admits what MIN bypasses, unpaced and on the paced copy path (O1 and O2 ran only the paced form); Belady-1R: Belady's rule within a two-step lead with each admission read once; MIN-1R: MIN with bypass, admitted misses fetched into their slot (one read, serialised: the model's foresight-only state $v(F)$); *read-ahead*: MIN-1R with MIN's admissions also copied two steps ahead (one read, overlapped: the model's $v(O, F)$), unpaced and on the paced copy path with a three-step lead (the read-ahead oracle of Table II); *bound*: Eq. (1) on that machine; $v(F)$, $v(O, F)$, $v(\text{all})$: the model.

Host	Cell	Deployed	Dep-1R	no ovl.	all CPU	MIN-2R	Belady-2R		Belady-1R	MIN-1R	read-ahead		bound	$v(F)$	$v(O, F)$	$v(\text{all})$
							unpaced	paced			unpaced	paced				
O1	gpt-oss 11%	21.12	–	21.01	22.91	–	–	22.88	–	–	–	–	10.06	14.54	11.27	10.06
O1	gpt-oss 25%	12.89	–	12.91	14.21	–	–	9.22	–	–	–	–	4.03	8.04	5.10	4.03
O1	gpt-oss 40%	8.90	–	8.95	–	–	–	5.99	–	–	–	–	1.97	5.75	3.84	1.97
O1	Qwen3 12.5%	36.92	–	36.91	43.29	–	–	40.29	–	–	–	–	18.04	24.90	19.29	18.04
O1	Qwen3 25%	22.73	–	22.76	26.60	–	–	18.22	–	–	–	–	8.11	14.04	9.22	8.11
O1	Qwen3 43.75%	12.71	–	12.67	–	–	–	8.12	–	–	–	–	3.32	8.25	5.51	3.32
O2	gpt-oss 11%	20.85	–	–	–	20.17	–	21.57	–	–	–	–	10.04	14.27	11.96	10.04
O2	gpt-oss 25%	12.72	–	–	–	10.77	–	7.80	–	–	–	–	4.02	7.86	5.86	4.02
O2	gpt-oss 40%	8.79	–	–	–	7.13	–	5.62	–	–	–	–	1.97	5.63	3.78	1.97
O2	Qwen3 12.5%	36.64	–	–	–	36.58	–	38.74	–	–	–	–	18.00	24.70	20.06	18.00
O2	Qwen3 25%	22.24	–	–	–	19.67	–	15.80	–	–	–	–	8.09	13.58	9.90	8.09
O2	Qwen3 43.75%	12.02	–	–	–	10.11	–	7.42	–	–	–	–	3.32	7.56	4.90	3.32
O3	gpt-oss 11%	18.01	–	–	–	–	19.79	–	–	14.00	12.91	12.39	8.26	12.63	9.49	8.26
O3	gpt-oss 25%	11.08	–	–	–	–	9.74	–	–	8.42	7.44	6.66	3.31	7.05	4.28	3.31
O3	gpt-oss 40%	7.77	–	–	–	–	6.14	–	–	6.20	5.58	5.20	1.95	5.24	3.56	1.95
O3	Qwen3 12.5%	30.95	–	–	–	–	44.54	–	–	23.02	21.78	21.89	14.81	20.35	16.06	14.81
O3	Qwen3 25%	19.32	–	–	–	–	18.17	–	–	14.01	13.09	11.58	6.66	11.68	7.76	6.66
O3	Qwen3 43.75%	10.94	–	–	–	–	8.59	–	–	8.66	7.78	7.28	3.30	7.24	4.76	3.30
O4	gpt-oss 11%	20.59	20.18	–	–	20.11	19.78	22.03	19.64	15.13	14.06	13.78	10.02	13.77	11.63	10.02
O4	gpt-oss 25%	12.43	12.07	–	–	10.55	9.76	8.13	9.61	8.72	7.65	6.87	4.01	7.40	5.34	4.01
O4	gpt-oss 40%	8.39	8.19	–	–	6.91	6.15	5.38	6.03	6.28	5.59	5.21	1.97	5.12	3.20	1.97
O4	Qwen3 12.5%	36.36	35.83	–	–	36.54	48.80	39.66	31.64	25.29	24.01	24.58	17.97	23.98	19.88	17.97
O4	Qwen3 25%	22.09	21.75	–	–	20.06	17.74	16.93	17.57	14.60	13.50	12.24	8.08	13.12	9.54	8.08
O4	Qwen3 43.75%	11.94	11.74	–	–	10.09	8.44	7.17	8.33	8.69	7.72	7.29	3.32	7.32	4.56	3.32
O5	gpt-oss 11%	21.26	21.31	–	–	21.47	33.34	25.57	33.03	18.34	17.97	16.39	9.94	15.22	11.20	9.94
O5	gpt-oss 25%	13.01	13.35	–	–	11.18	15.87	11.17	15.60	10.60	10.54	8.68	3.98	9.13	5.06	3.98
O5	gpt-oss 40%	8.73	9.22	–	–	7.16	8.34	6.25	8.13	7.27	6.87	5.91	1.97	6.26	3.62	1.97
O5	Qwen3 12.5%	37.63	37.42	–	–	40.11	69.42	44.89	49.11	31.37	31.27	30.01	17.82	27.95	19.27	17.82
O5	Qwen3 25%	23.24	23.27	–	–	21.12	30.78	21.76	30.47	18.51	19.64	16.03	8.01	16.77	9.15	8.01
O5	Qwen3 43.75%	12.49	12.92	–	–	10.32	12.08	9.00	11.86	10.21	9.61	8.23	3.31	8.94	5.02	3.31

expert served in a step stays for that step. The simulator the job's predictions were registered against gave a resident carried over from the previous sequence no next use, so it was evicted first; that over-counted MIN-1R's reads by 17% at the two GPU-bound cells, which is why two of its reads clauses failed (Appendix D). Each configuration of llama-ec-bench starts with a cold cache in a fresh context. The paced copy path copies one 16 MB piece at a time on a single copier thread and publishes a landed copy at the start of the next step; the unpaced path copies each admission as one transfer. The background admissions of Belady-2R drive host memory at 71–82% of the probe's best rate at the host-bound cells, the single-read oracles that read ahead at 64–85% (the paced read-ahead oracle 71–82%), the deployed cache itself at 66–82%, MIN-1R at 42–67%. The predictions and their outcomes are in `prereg/foresight_outcome_095.md`: of 56 clauses, 44 held and 12 failed, all of size (Appendix D).




Fig. S1. The first oracle job (093, host O1): speed as a percentage of this machine’s bound against the window W of future routing Belady-2R sees, at the four host-bound cells. “none” is the deployed cache; “all” the whole sequence; points are connected for the trend only. The tick at the right is the model’s foresight-only term for the cell, the speed it predicts with that one fix; the oracle’s host traffic sits at 89–94% of the probe’s best rate at every window, so its speed is R^* over its reads times that utilisation, 44–45% of the bound with the whole sequence in view.

The double-reading oracles: Belady-2R and MIN-2R read the routing of the next W decode steps; their copies ride the paced admission path, and an expert still in flight at its step runs on the CPU like any other miss. Belady-2R, the *hit-optimal* oracle, admits the experts those steps will route to, soonest first, evicting the resident whose next use is furthest (Belady within the window), for one future use as readily as for ten. MIN-2R, the *bytes-optimal* oracle, is MIN with bypass with its admissions made through the background path: a miss is served where it is, by the CPU or a demand copy, and admitted only if its next use comes before the furthest next use among the residents. With the whole sequence in view Belady-2R runs $1.40\times$ [1.38, 1.41] the deployed cache at gpt-oss 25% and $1.25\times$ [1.24, 1.26] at Qwen3 25% on host O1, $1.63\times$ and $1.41\times$ on O2, and 1.48 – $1.62\times$ at the two GPU-bound cells; half of its gain at gpt-oss 25% arrives between $W=4$ and 16 (9.2 tokens by log interpolation; Qwen3 25%: 6.1 against the trace’s 3.9). Where the host binds hardest it loses, 3–8% at gpt-oss 11% and Qwen3 12.5% on both machines, and at every window on O1 (0.78 – $0.92\times$): it reads more host bytes than the deployed cache (82 and 195 experts per token against 63 and 160), 2.02 – $2.18\times R^*$, because Belady admits an expert used once where the optimum bypasses it and re-copies what it evicts (its admissions, each one read, are most of its reads; a miss served on the CPU before its copy lands is the second read of that expert). An unpaced copy path, run at the four host-bound cells of O1, reaches 98–99% hits at three of them and still loses at gpt-oss 11% and Qwen3 12.5% (where it hits 75%), so copies that arrive in time do not remove the loss there: the extra bytes cause it. MIN-2R gains 13–23% at the four cells of 25% and above and 0–3% at the two below (6% with a 16-token window); its misses fall (30 to 20 and 85 to 64 per token at 25%) but its total reads, misses plus the admissions that read each admitted expert again, are 1.73 – $2.04\times R^*$ and exceed the deployed cache’s by 4–5% at the two lowest budgets; the model on its own counters predicts its time within 9% at every cell, and the bytes it saves account for 4–21 points of its 13–23% gain at the four cells (most of it at gpt-oss 40% and Qwen3 43.75%, a third at Qwen3 25%). Belady-2R reads 1.01 – $1.13\times$ the host bytes of MIN-2R on O2 and gains 41–63% at the same four cells, and the model, with its GPU term calibrated on the deployed cache’s run, over-predicts those runs’ time per token by 37–53% (on O1, 24–39%): its gain is the overlap of copies issued ahead of their use with the step, which the model charges in sequence. Pacing the deployed cache’s own copies changes it by under 1% (0.997 – $1.005\times$). With no FETCH table (every miss on the CPU) the deployed cache runs 8–15% slower at the four host-bound cells; with the engine’s within-layer overlap of a layer’s CPU misses with its GPU hits turned off it changes by under 1% at every cell (0.5% faster to 0.6% slower), so the overlap the accounting charges is across the step and in rebalancing, not within a layer. MIN-2R with a 16-token window keeps 34–176% of the whole-sequence gain at the five cells where that gain’s interval excludes zero (more than all of it at gpt-oss 11%, where the shorter window admits less).

Why the misses of MIN-2R stay above R^ :* Most of the excess comes from the engine’s in-step fetches, which put misses the schedule did not choose into slots: with them off (job 113, Section IV), Few-2R’s misses fall from 54.7–55.6 to 43.0 per token, against Few-1R’s 40.7. Besides that and the second read, MIN-2R’s misses alone are 1.31 – $1.47\times R^*$ at the four host-bound cells (at the two GPU-bound ones the ideal simulation below is itself 6–10% above R^* , which sees across sequences). The engine’s paced copy path moves one 16 MB piece at a time from pageable host memory on a single copier thread, behind the demand copies, publishes a landed copy at the start of the next step, and admits nothing while 64 copies are queued; at the 16–30 admissions per token (210–285 MB) of the two lowest budgets it lags the admissions by steps, and an expert still in flight is a miss the CPU reads again. The last four columns of Table S11 simulate MIN-2R on the measured routing trace with the engine’s own semantics (the lookahead within the sequence, residents carried over, no queue cap) and a

TABLE S10

Belady-2R, the hit-optimal oracle, in the engine (job 093; RTX 5090, Ryzen 9 9950X, host memory 50.5 GB/s at best). tok/s of the deployed cache and of the oracle with the whole sequence in view; their ratio, paired by sequence, with its 95% interval; hit rate and misses per token of each (experts read from host memory by the CPU or by a demand copy), the optimum's hit rate and reads (R^*) in parentheses, and the oracle's background admissions (Belady's rule without bypass: it admits what MIN bypasses); speed as a percentage of this machine's bound; the accounting's foresight-only term $v(F)$ for the cell in tok/s; the share of that term the oracle recovered and the share of the gap to the bound it closed.

Cell	deployed	oracle	ratio	hits % (opt.)	misses/token (opt.)	oracle adm.	% of bound	$v(F)$	recovered %	closed %
gpt-oss 11%	47.4	43.7	0.92 [0.91, 0.93]	58 → 67 (73)	60 → 47 (38)	35	47.6 → 44.0	68.8	−27	−16
gpt-oss 25%	77.6	108.5	1.40 [1.38, 1.41]	79 → 92 (89)	30 → 11 (15)	20	31.2 → 43.7	124.4	76	41
gpt-oss 40%	112.4	166.9	1.48 [1.46, 1.51]	89 → 98 (95)	15 → 3 (7)	12	22.1 → 32.8	174.0	92	42
Qwen3 12.5%	27.1	24.8	0.92 [0.91, 0.92]	60 → 66 (75)	155 → 132 (97)	63	48.9 → 44.8	40.2	−28	−18
Qwen3 25%	44.0	54.9	1.25 [1.24, 1.26]	78 → 89 (89)	85 → 42 (43)	46	35.7 → 44.5	71.2	52	31
Qwen3 43.75%	78.7	123.1	1.56 [1.53, 1.59]	92 → 99 (96)	32 → 5 (14)	25	26.1 → 40.8	121.3	103	49

TABLE S11

MIN-2R, the bytes-optimal oracle, beside Belady-2R, the hit-optimal one, on the second oracle machine (job 094; RTX 5090, Ryzen 9 7950X, host memory 50.6 GB/s at best). Per cell: the deployed cache's speed; each oracle's ratio to it with the whole sequence in view (paired, 95% interval), its hit rate, its misses per token (experts read from host memory by the CPU or by a demand copy); its background admissions in parentheses: those of MIN-2R are this step's CPU-served misses, each read a second time by the copy, those of Belady-2R are prefetches of what MIN would bypass as readily as of what it would admit) and the calibrated model on its own counters against its measured time; the optimum's reads (R^*) and hit rate from the trace; and MIN-2R simulated on the trace with the engine's lookahead (within the sequence) under a copy latency of 1 (ideal), 2 and 3 steps, with the latency at which the simulation meets the engine's misses.

Cell	deployed	ratio	MIN-2R, W =all				Belady-2R, W =all				optimum R^* hits	simulated misses/token			
			hits	misses (adm.)	model	ratio	hits	misses (adm.)	model	$d=1$		2	3	fit	
gpt-oss 11%	48.0	1.03 [1.03, 1.04]	65	50 (16)	+7%	0.97 [0.96, 0.98]	71	42 (39)	+19%	38.3 73	39	43	47	3.9	
gpt-oss 25%	78.6	1.18 [1.18, 1.19]	86	20 (9)	+6%	1.63 [1.61, 1.66]	95	7 (23)	+53%	15.3 89	16	17	19	4.7	
gpt-oss 40%	113.8	1.23 [1.22, 1.24]	94	9 (4)	+2%	1.56 [1.54, 1.59]	98	3 (12)	+39%	6.8 95	7	8	9	3.5	
Qwen3 12.5%	27.3	1.00 [1.00, 1.01]	64	138 (30)	+5%	0.95 [0.94, 0.95]	67	128 (65)	+11%	96.5 75	97	122	143	2.8	
Qwen3 25%	45.0	1.13 [1.12, 1.14]	83	64 (21)	+9%	1.41 [1.39, 1.43]	91	34 (52)	+37%	43.4 89	44	54	61	3.4	
Qwen3 43.75%	83.2	1.19 [1.18, 1.20]	95	19 (9)	+1%	1.62 [1.58, 1.65]	99	3 (27)	+47%	13.7 96	14	17	19	3.0	

copy that lands d steps after it is issued, during which the expert misses and is not re-admitted: at $d=1$, the ideal, the misses come within 0–4% of R^* at the host-bound cells; the engine's misses meet the simulation at $d=2.8$ –4.7 there and at 3.0–3.5 at the GPU-bound cells, three to five steps. At the fitted latency the simulation admits 1.1 – $1.9\times$ what the engine admitted, so the queue cap, not only the latency, keeps the engine's misses above the ideal. The predictions and their outcomes are in `prereg/foresight_outcome_093.md` and `094.md`: of 36 clauses for the first job, 9 held, 2 held on the point estimate, 21 failed and 4 were untested; of 38 for the second, 12 held, 1 on the point estimate, 24 failed and 1 was untested (Appendix D). The second job's failures are of size, not direction: the misses of MIN-2R were predicted within 15% of R^* and came 29–47% above (the copy latency and the queue cap); its hit rate was predicted within 4 points of the optimum's and fell 5–11 points short at three cells; its gain was predicted at 10–45% at the host-bound cells and came below that at the two lowest budgets; the model was predicted to over-predict it by 10–40% at the host-bound cells and over-predicted by 5–9%; and the two oracles were predicted within 15% of each other at 25% and above, where Belady-2R was 1.24 – $1.38\times$ faster.

Where a token's time goes at a GPU-bound budget: On Qwen3 at 43.75% Nsight Systems splits a token of our cache (FreeToken in parentheses) into 4.93 (5.20) ms of expert copies over PCIe, 2.29 (2.21) of resident-expert products, 1.65 (1.34) of attention projections, 0.49 (0.37) of attention, 0.38 of output head and 0.97 (0.40) of norms, small kernels and cache control, with neither system overlapping its copies with GPU compute. Turning off the within-layer overlap of a layer's CPU misses with its GPU hits changes speed by under 1% at every budget, and predicting the next layer's experts [S3], [S4] never lowers host bytes.

APPENDIX F

THE TRACE STUDY IN DETAIL

Policies: Table S12 scores LRU, LFU, decayed frequency, ARC, S3-FIFO and a profiled static cache against MIN and against Belady within a window of W future tokens, at three budgets per model, on each model's own sampled text (9 traces, the cache carried across conversations). Without admitting every miss, the online policies are close to one another and far from MIN: the best reads 35–110% more, decayed frequency is best at 12 of 27 budget points and S3-FIFO at 10. Where the cache turns over every token ($C/k \leq 1.5$) the six are within 4% of one another; from $C/k=2$ LRU reads 2–11% more than decayed frequency on eight of nine models, and classic LFU fails at large budgets. Adding *admit every miss* (every miss copied into the slot of the lowest-scored resident not requested now) changes the ranking: it is best at 18 of 27 points, ties at 5, and loses by at most 2.7%; it saves up to 5.8% over the previous best, and the best online policy still reads 35–110% more than MIN (`prereg/policy_aa.json`). At the small budgets almost every resident turns over within a few tokens, and the expert just

TABLE S12

Host reads per token relative to the exact optimum (Belady’s MIN with bypass over the whole future), median over nine models at three budgets, and the range over all 27 model–budget cells. Event-atomic replay of each model’s own sampled text; every online policy is scored on the last 90% of the trace. Mixtral (8 experts) uses $C = 2, 3$ and 4.

Policy	$C=E/8$	$E/4$	$3E/8$	range
LRU	1.44	1.77	2.03	1.37–2.26
LFU	1.44	1.93	2.40	1.37–6.15
Decayed frequency	1.45	1.70	1.94	1.35–2.10
with hysteresis κ	1.46	1.75	2.00	1.36–2.55
ARC	1.44	1.76	1.99	1.37–2.23
S3-FIFO	1.44	1.69	1.94	1.38–2.11
Static, profiled	1.92	2.74	3.48	1.43–13.24
Belady, $W=1$	1.18	1.45	1.69	1.07–2.06
Belady, $W=4$	1.01	1.13	1.31	1.00–1.95
Belady, $W=16$	1.00	1.00	1.01	1.00–1.66

TABLE S13

Predicting a held-out model’s horizon W_{50} (fitted on the other eight models, tested on the ninth, over all 26 budget points): error as a factor, $\max(\hat{W}/W, W/\hat{W})$.

Rule	Constants	Median	90th pct.	Max
$W_{50} \propto C/k$	1	1.22	1.57	2.42
$W_{50} = a(C/k)^b$	2	1.19	1.51	1.67
$D(W_{50}) = cC$	1, and $D(W)$	1.16	1.42	1.61

requested is a good bet. In our engine admitting every miss reads all of its misses over the link, so its value in time depends on the link (Section V). A static cache profiled on other text reads $1.4\text{--}13.2\times$ MIN because the hot set moves: the C most requested experts of a layer change by 22–58% between windows of 100 tokens and 33–75% between windows of 1,000.

The horizon, before and after the stronger baseline: Against decayed frequency with each admission fetched once (the study’s original baseline), MIN reads 26–56% fewer experts (median 41%), and the window that closes half of the gap fitted $W_{50} \approx 0.59(C/k)^{1.33}$; admitting every miss lowers the baseline at 23 of 26 points by up to 6.7%, raises W_{50} by a median 9% (at most 42%), and refits it to $0.67(C/k)^{1.31}$, with a leave-one-model-out exponent range of 1.25–1.36 and, without the points below one token, 1.35 (it was 1.51 against the original baseline). Ninety percent of the gap needs a median of $2.7C/k$ tokens. Table S13 compares the rule in distinct experts with the power law in tokens and with $W_{50} \propto C/k$, each fitted on eight models and tested on the ninth: the first two are not distinguished, and both beat the third. The rule in distinct experts inverts the held-out model’s own curve of distinct experts against window length, which needs that model’s routing trace. Across the 9 models, the window that closes half of the gap holds a median $0.65C$ distinct experts per layer (quartiles 0.61–0.72).

A learned admission order: We trained a logistic model of whether an expert is requested within the next C/k tokens on features the engine has when it plans a step (decayed request counts at five half-lives, the time since the last request, whether the expert is requested now, and transition scores from the previous token’s experts at the same layer and, through an 8-bit table, at every layer), fitted on each model’s mixed-domain trace and replayed on the AIME routing, so no test text is seen in training (Table S14). With each admission read once it reads 6–12% fewer experts than decayed frequency with the same single-read admission at the four host-bound budgets, closing 17–30% of that policy’s excess over MIN; admitting every miss alone closes most of that at Qwen3 12.5% and little at gpt-oss (prereg/learned_offline.json and the research checks). In the engine (job 097, predictions committed before launch) it reads 6–13% fewer experts than Dep-1R at the host-bound budgets on both machines, but runs $1.01\text{--}1.03\times$ it on O4’s machine and $0.90\text{--}0.94\times$ on host O6: the reads it saves are CPU reads, bought with more fetches in the step, each a copy its layer waits on, and with $237\text{--}585\mu\text{s}$ of host time per token (over all six budgets).

Speculation is not foresight: Speculative decoding verifies K draft tokens in one forward pass, so a layer’s experts for K positions are requested together and a missed expert is read once for all of them. On the same traces, with drafts accepted at rate α until the first rejection, batched verification saves an online cache a median 15–32% of its reads at $K=8$ when rejected drafts cost nothing, the equivalent of a window of two to four tokens; once the rejected drafts’ experts are counted, the saving is 8–12% at $\alpha=0.9$ and becomes a loss of 18–20% at $\alpha=0.8$, because the union of K positions’ expert sets is 20–79% of Kk . MIN itself gains little from batching (–7 to 15%). Systems that use drafts as a look-ahead sensor [S5] compare against a baseline that already pays for verification; ours charges the rejected drafts to a cache that does not speculate.

TABLE S14

A learned admission order, offline: host reads per token on the AIME-25 routing the engine runs, of the decayed frequency with each admission read once (Dep-1R; within 10% of the engine’s reads on hosts O4 and O5), the learned order with each admission read once (trained on the model’s mixed-domain trace; 8-bit cross-layer table, as in the engine), Belady within one and four tokens of perfect foresight, and MIN with bypass in its no-evict form (Table S24). *closed*: the share of Dep-1R’s excess over MIN that the learned order removes. [†]host-bound.

cell	Dep-1R	learned	W=1	W=4	MIN	closed
gpt-oss 11% [†]	59.4	56.0	54.4	47.3	39.0	17%
gpt-oss 25% [†]	28.3	26.1	27.3	25.1	15.5	17%
gpt-oss 40%	14.3	12.6	14.0	13.2	6.8	23%
Qwen3 12.5% [†]	166.2	146.3	134.1	109.1	100.4	30%
Qwen3 25% [†]	84.2	74.8	74.1	62.8	44.3	24%
Qwen3 43.75%	28.8	26.3	27.7	25.3	13.8	17%

TABLE S15

Partial foresight on the AIME routing the engine runs: host reads per token of admitting every miss and of MIN-1R with the whole problem in view (R_{eng}^* , Appendix E), and the share of the gap between them that each window closes (Table II). Recall 0.5: each future expert kept with probability 0.5 and replaced by a random one (fill) or left out (drop); means of three seeds. Linear: the ridge forecaster of Appendix G with a 4-token window. $\kappa=1$: the same exact 4-token window with the deployed margin instead of admitting every miss beyond it, as a share of its own gap.

Budget	Reads per token		Exact window W (tokens)						Recall 0.5, fill / drop		Linear	$\kappa=1$
	every miss	R_{eng}^*	1	2	4	8	16	$W=8$	all	$W=4$	$W=4$	
gpt-oss 11%	58.6	39.3	0.18	0.31	0.52	0.81	0.98	0.29 / 0.29	0.43 / 0.48	0.13	0.62	
gpt-oss 25%	28.2	16.0	0.07	0.14	0.25	0.42	0.66	0.14 / 0.13	0.46 / 0.47	0.05	0.34	
Qwen3 12.5%	150.0	100.3	0.31	0.51	0.77	0.96	0.99	0.42 / 0.43	0.48 / 0.52	0.19	0.88	
Qwen3 25%	79.6	44.9	0.16	0.28	0.46	0.70	0.93	0.29 / 0.28	0.49 / 0.53	0.11	0.60	

APPENDIX G

PARTIAL FORESIGHT IN DETAIL

The window policies: At step g the engine plans the step’s misses before it runs, from the trace’s record of step g (the routers give it layer by layer in a deployable system) and from a view of records $g+1..g+W$ of the same problem. A miss is processed in request order: it takes a free slot; otherwise the victim is a resident not requested now that is not seen in the window (lowest decayed score), or, if every resident is seen, the one seen furthest; a miss seen in the window is admitted unless the victim is seen sooner, and a miss not seen is admitted if the victim is not seen either (admit every miss beyond the window; with the deployed margin $\kappa=1$, only if its score beats the victim’s by one). On the panel admissions are fetched in the step, so every miss is read once; on the deployed path of the follow-up the machine’s fetch table splits the misses between the CPU and an in-step copy, and an admission served by the CPU is copied afterwards, a second read. In a degraded view each expert of a future record is kept with probability r , decided by a hash of (seed, g , record, layer, position), so step g sees one view throughout and step $g+1$ an independent one; a dropped expert is replaced by a uniformly drawn one (fill: precision and recall r) or left out (drop: recall r , precision 1). `scripts/value_map.py` replays the plan bit for bit: on a CPU build the engine’s reads equal the replay’s in 30 configurations at three budgets (`cputest099`), and on the panel’s GPUs, where the routing in the run can differ from the recorded trace by numerics, within 1.7% at both gpt-oss budgets. The panel ran the first 20 of the 30 problems; the table’s reads are on all 30. Table S16 lists the windows on the deployed path and in the step on the machines of job 100. Copied in the step, a window beats the deployed path where the link reads at least 0.92 of the CPU rate and loses to it at 0.54 or below; on the deployed path the engine misses 7–33% more than an instant replay (the deployed state matches its replay within 2.4%), and background copies that land late are the likely cause.

What the table says: Horizon and accuracy trade: at gpt-oss 11% an 8-token window at recall 0.5 closes 0.29 of the gap between admitting every miss and R_{eng}^* , about what two exact tokens close (0.31), and the whole future at recall 0.5 closes 0.43. Fill and drop cost about the same for windows of up to 16 tokens; with the whole future in view at low recall, dropping costs less (the largest difference, 0.12, is there). A larger cache needs a longer window: at 25% the same exact 4-token window closes 0.25 of the gap, against 0.52 at 11%.

A realisable forecaster: Ridge regression from the all-layer routing of step t to that of step $t+h$ ($h=1..8$), fitted on each model’s mixed-domain own text and applied to the AIME routing, forecasts the top k experts of each layer with a precision of 0.48 one token ahead and 0.26 four ahead on gpt-oss (0.54 and 0.32 on Qwen3); a decayed request count, the deployed cache’s own signal, scores 0.30 and 0.28. Driving the window policy, the ridge forecasts close 3–15% of the gap on gpt-oss and 6–19% on Qwen3 over windows of 1–8 tokens (`prereg/forecaster_linear.json`). Its precision falls with the horizon, so over a window it is closer to 0.3 than to its one-token value; at that average precision a window with random errors closes about as much on gpt-oss, and more on Qwen3, where the forecaster’s correct forecasts are mostly experts the

TABLE S16

Foresight on each read path (job 100): speed relative to the deployed cache on the same machine. *In step*: admit every miss with a window of W exact tokens, each admission copied in the step (as on the panel). *Deployed path*: the deployed policy with the same window, its misses split between the CPU and an in-step copy by the machine’s fetch table and its admissions copied in the background; $W=8$ at recall 0.5 in the column marked r . The last three columns: MIN-2R, MIN-1R and the read-ahead oracle. Machines sorted by the probe’s link-to-CPU ratio.

Machine	Link/ CPU Budget		In step		Deployed path				MIN-2R	MIN-1R	Read- ahead
			$W=4$	16	4	16	r				
Pd again	0.41	gpt-oss 11%	0.66	0.89	1.02	1.03	0.94	1.00	1.05	1.11	
Pd again	0.41	gpt-oss 25%	0.69	0.81	1.03	1.12	1.01	1.19	1.11	1.29	
9950X, slower link	0.54	gpt-oss 11%	0.65	0.89	1.01	1.01	0.96	0.99	1.07	1.23	
9950X, slower link	0.54	gpt-oss 25%	0.69	0.80	0.96	1.06	1.00	1.10	1.11	1.37	
5700X3D	0.92	gpt-oss 11%	1.09	1.32	0.99	1.00	0.95	0.99	1.37	1.44	
5700X3D	0.92	gpt-oss 25%	1.06	1.23	1.03	1.07	1.00	1.16	1.48	1.80	
Ph again	0.96	gpt-oss 11%	0.99	1.20	1.01	1.03	0.98	1.03	1.27	1.45	
Ph again	0.96	gpt-oss 25%	0.97	1.09	1.03	1.07	1.01	1.16	1.29	1.64	

cache already holds. A sequence model reports recall above 90% one token ahead with $k+3$ candidates [S6]; its precision is then at most $k/(k+3)$, 0.57 for gpt-oss and 0.73 for Qwen3. A recurrent forecaster in that spirit (a GRU of 512 units over the same all-layer vectors, one output head per horizon, trained on the same text; `scripts/forecaster_gru.py`) reaches a precision of 0.40 and 0.26 one and four tokens ahead on gpt-oss (0.46 and 0.31 on Qwen3), a recall at $k+3$ of 0.53 and 0.54, and closes 2–11% and 6–17% of the gap: no better than the linear model. Our training text is smaller than SeqMoE’s and differs from the test text, so this bounds what such a model does across text, not what it can do in domain.

APPENDIX H

A CLOSED-FORM ACCOUNT OF THE TIME, AND ITS REGISTERED TESTS

This appendix holds the account of the deployed cache’s time that the main text uses once (Section III): what it is, how it was found, how it fared in the tests registered on machines new to it, and the split of the time it implies. The measured decomposition of Section IV does not rely on it.

A. The relation

For the deployed cache, we describe the time per token T by

$$T = G + \frac{(M + (1 - G/T) A) S}{B_{\text{host}}}, \quad (\text{S1})$$

solved for T : it is the larger root of $T^2 - (G + t_M + t_A)T + t_A G = 0$, with $t_M = MS/B_{\text{host}}$ and $t_A = AS/B_{\text{host}}$. M and A are the misses and the background admissions per token (the engine’s counters), S the expert’s bytes, B_{host} the machine’s best probed rate, and G the GPU’s own kernel time per token in an Nsight profile of the same engine (4.1–4.7 ms for gpt-oss; where a machine was not profiled or its profile captured no decode kernels, the median of a fixed set of 7 profiles; Appendix I). G exceeds T_{GPU} of Eq. (2) because it also counts the expert kernels the GPU runs for resident and fetched experts. No constant is fitted to the times. Every host read costs host time, except the background copies that run while the GPU computes: they spread over the next step, and a share G/T of them falls under the GPU’s compute. The *plain* form drops that share. We found the plain form on 30 launches of gpt-oss (exploratory): it is within 6% on 28 of them at 11% but over-predicts at 25%, where admissions are a larger share of the reads. The overlap term came from where those copies run, after the plain form failed its first registered test; over 34 launches it brings 33 instead of 18 within 6% at 25%, at the cost of 29 instead of 32 at 11%.

B. Registered tests

After the plain form failed, we registered Eq. (S1) before three more experiments, each predicting it within a band at every cell, with G profiled on each machine before its timed runs (Table S17). None passed as registered.

The first test: held on machines rented before, failed on new ones: On new runs of 3 machines we had rented before, it came within 6% at all 12 cells of both models (median error 2.3%). This does not single out the overlap term: discounting admissions by half passes 11 of those cells. The same test failed on the two server machines new to it, by 10–52%; one computed wrong outputs and the other ran its rounds 52% apart (slowest over fastest), so the following tests set validity gates before launch (Appendix I).

TABLE S17

The four registered tests of Eq. (S1). A *cell* is one model at one budget on one machine. *New*: machines no earlier launch had rented. Each test registered a band for every cell and a median error of at most 4%; job 108 also allowed at most one cell beyond 6%. Jobs 107 and 108 count only the machines that passed their registered validity checks; job 107 registered that with fewer than three it is inconclusive. *New machines*': their cells within the band. The last column sets aside, after the fact, the hosts the text discusses: in job 106 the two that computed wrong outputs or ran unsteadily, in job 108 the one whose GPU compute could not be profiled.

Job	Form	Models	Band	Machines (new)	Cells within	New machines'	Median	As registered	Set aside
105	plain	gpt-oss	6%	4 (2)	5/8	2/4	4.2%	failed	–
106	overlap	both	6%	5 (2)	12/20	0/8	4.7%	failed	held (12/12, 2.3%)
107	overlap	both	6%	2 (2)	1/8	1/8	31.7%	inconclusive	–
108	overlap	gpt-oss	8%	3 (3)	4/6	4/6	6.7%	failed	failed (4/4, 5.3%)

Machines never rented before: not passed: The next test admitted only machines whose GPU no earlier launch had recorded, with little host memory in use, correct outputs and rounds within 2% of each other. Only two of 7 passed, short of the three required, so by its own rule the test is inconclusive; the relation missed the band at 7 of their 8 cells, by up to 59% on a dual-socket server (51% against the probe's second-highest reading) and 11% on a Ryzen 9 5900XT. A last test restricted the machines in advance to the class the relation had been found on (one NUMA node, at most 32 usable cores) and widened the band to 8%. It failed on an EPYC 7543 inside that class, by 34.3%, and without that machine still failed two pooled clauses (Appendix I).

The error depends on the budget: On the three consumer machines new to the last two tests the relation under-predicted every gpt-oss cell. Over all 25 consumer machines its mean error is -2.2% at gpt-oss 11% (95% *t*-interval over machines -4.0 to -0.4) and $+1.6\%$ at 25% (0.0 to 3.1): what Eq. (S1) leaves out changes sign with the budget.

C. After the fact: the processor

The processor separates most of the machines the relation fits from those it misses, not all; we drew this split after the tests. On 39 valid launches on 25 machines with consumer processors, the read rate the deployed cache implies, $(M + (1 - G/T)A)S/(T - G)$, is 0.88–1.16 of the probe's best rate at gpt-oss 11%, and the relation is within 8% on 37 of them (G was profiled on the machine for 9 of these launches). On server processors, 5 of 8 launches failed a validity check, and the three valid ones read at AMD engineering sample 0.27, EPYC 7543 0.58, EPYC 7402P 1.04 of it: the last fits the relation, the other two do not. We have not established why. On the dual-socket server the probe's own threads read slower with 128 threads, about the engine's 126 helpers, than with 16 (118 against 151 GB/s), and on an EPYC 9655 rented earlier the per-layer timers showed a latency floor no consumer machine has (the scorecard).

D. The split it implies (exploratory)

On consumer processors, then, the relation splits the deployed cache's time into three parts (Fig. S2). At gpt-oss 11% the GPU's compute takes 12–45% of the time, MIN's reads, which are the bound, 31–54%, and the reads beyond MIN's 22–35%. The first is serialisation, the GPU's compute in series with the reads, and makes up 23–66% of the gap to Eq. (1); the second is the reads MIN would not make. With G removed, the time scales with the machine's read rate with an elasticity of -0.95 (95% interval -1.06 to -0.84 over the machines of the exploratory set), as reads at B_{host} would. Foresight can attack both parts: fewer reads need MIN's choices, and overlap needs reads issued ahead. Reading below the probed rate takes at most 11% of the time on the consumer machines; on the server machines it is a third part (EPYC 7402P 0%; EPYC 7543 34%; AMD ES, 2 sockets 59% of the time at 11%).

APPENDIX I

MIN'S SCHEDULES AND THE MACHINE, IN DETAIL

The fewest-admission schedule, machine by machine: Tables S18 and S19 list jobs 104 and 105: MIN's set loaded by the CPU (MIN-2R, Few-2R) or in the step (MIN-1R, Few-1R), as the greedy rule picks it or as the schedule with the fewest admissions does. On the engine's 20 problems the fewest-admission schedule copies 0.60–0.73 times as many experts in the step as the greedy one, for 2.5–5.3% more misses (more than its replay predicts; its measured gains include that cost). On every launch it gains over the deployed cache with either load and at both budgets, each 95% interval above 1; on the machine with the fastest link, Few-2R is 0.07 behind MIN-2R at 25%.

Job 106, every machine: Table S20 adds the two server machines new to job 106. On the EPYC 7302 machine the engine computed wrong outputs: its teacher-forced loss on the deployed cache was 0.21–0.45 nats per token across rounds and models, against 0.19 (gpt-oss) and 0.085 (Qwen3) on every other machine and in its own untimed routing pass, and its counters changed between rounds. On the Xeon 8347C machine the deployed cache ran its gpt-oss 11% rounds 52% apart (slowest over fastest), against at most 2% registered, and the greedy copy gained in two of its three rounds where we had




Fig. S2. The split implied by Eq. (S1) (exploratory): the deployed cache’s time per token at the two host-bound gpt-oss budgets, machine by machine, split into the GPU’s compute (G , profiled on the machine, or the median profile where none was taken or it captured no decode kernels), MIN’s host reads at the machine’s best rate (the bound of Eq. (1)) and the reads the cache makes beyond MIN’s. The dot is the measured time; hatched, the measured time beyond Eq. (S1), reading below the machine’s best rate; a bar that ends past the dot is an over-prediction. Machines with the same processor are numbered in order of rental. One launch per machine that ran stably: 25 with consumer processors (above the dashed line) and 3 with server processors (below). *: first rented for a registered test of Eq. (S1).

TABLE S18

Which of MIN’s hit-optimal sets the oracle follows (job 104): speed relative to the deployed cache on the same machine, with MIN’s set loaded by the CPU (two reads: MIN-2R, Few-2R) or in the step (one read: MIN-1R, Few-1R), as the greedy rule picks it or as the schedule with the fewest admissions does; and the copies per token each makes in the step. Machines sorted by the probe’s link-to-CPU ratio.

Machine	Link/ CPU	Budget	By the CPU		In the step		Copies/token	
			MIN-2R	Few-2R	MIN-1R	Few-1R	MIN-1R	Few-1R
Pf again	0.29	gpt-oss 11%	0.94	1.11	0.86	1.04	19.8	11.8
Pf again	0.29	gpt-oss 25%	1.13	1.22	0.96	1.06	9.7	7.0
285K, second	0.49	gpt-oss 11%	1.03	1.08	1.05	1.18	19.8	11.8
285K, second	0.49	gpt-oss 25%	1.20	1.21	1.12	1.17	9.7	7.0
Pg again (5950X)	0.63	gpt-oss 11%	1.02	1.06	1.26	1.34	19.8	11.8
Pg again (5950X)	0.63	gpt-oss 25%	1.20	1.15	1.32	1.35	9.7	7.0

predicted it would lose. When their jobs reached the gate, before any download, `free` reported 131–132 GB of host memory in use (12–15 GB on the other machines), and their probes’ CPU read rates fell when more threads read than the container had usable physical cores. On both, Eq. (S1) under-predicted the deployed cache by 10–52%. We report their clauses (in the scorecard, `supplement.pdf`) and draw nothing from them about the relation on those runs; both are servers, and job 107’s one valid server was under-predicted as well (Appendix H). The three machines that ran stably were all rented before: Pf, the Threadripper 9960X of job 105, and panel host Pe (same GPU).

Job 107, every machine: Table S21 lists the four machines that passed the gate before any download. The EPYC 9754 and the EPYC 7663 ran their gpt-oss 11% rounds 3.5% and 6.0% apart, beyond the registered 2%, and stopped there, as registered; the relation under-predicted them by 31% and 21%. The others stopped at the gate with 50–96 GB of host memory in use. The registration replaced only machines that failed the gate, from a list of offers in order; replacing machines that failed the round check was added during the job, after the results of the first four machines were in, in an amendment committed before the

TABLE S19

Job 105, registered before the runs, nothing fitted: the deployed cache’s time per token against the plain form of Eq. (S1) (without the overlap term, which came after this job), the GPU’s compute profiled on the same machine before the timed runs plus the run’s counted host reads at the machine’s best rate; and the speed, relative to the deployed cache, of LA, the deployed cache with the engine’s layer-ahead copy (each layer’s top predicted expert of the next layer copied ahead of use), and of MIN-1R and Few-1R, MIN’s greedy and fewest-admission sets copied in the step. Machines sorted by the probe’s link-to-CPU ratio.

Machine	Link/ CPU	Budget	G_{prof} (ms)	Deployed (ms)			LA	MIN-1R	Few-1R
				plain	Eq. S1	measured			
Pf again	0.28	gpt-oss 11%	4.1	13.6	13.9	0.72	0.86	1.03	
Pf again	0.28	gpt-oss 25%	4.4	9.8	9.5	0.69	0.97	1.06	
Threadripper 9960X	0.32	gpt-oss 11%	4.1	9.3	9.1	0.91	0.91	1.04	
Threadripper 9960X	0.32	gpt-oss 25%	4.4	7.3	6.8	0.87	0.98	1.04	
EPYC 7402P	0.45	gpt-oss 11%	4.3	19.1	18.2	0.81	1.01	1.15	
EPYC 7402P	0.45	gpt-oss 25%	4.5	12.9	11.8	0.78	1.09	1.15	
O4 again (9950X)	1.04	gpt-oss 11%	4.4	20.8	20.5	1.04	1.36	1.34	
O4 again (9950X)	1.04	gpt-oss 25%	4.4	13.2	12.4	1.01	1.42	1.39	

TABLE S20

Job 106, registered before the runs, every machine (*: deployed cache varied by more than the registered 2% between rounds): the deployed cache’s time per token against Eq. (S1) (GPU compute profiled on the same machine before the timed runs; the run’s own counters), and the speed relative to the deployed cache of the two online admission rules (Margin: the decayed count with a larger admission margin; *lrm*: the learned reuse predictor with a margin; both chosen on other text) and of Few-1R, MIN’s fewest-admission set copied in the step, with round-level 95% intervals (rounds, then problems: three rounds at gpt-oss 11%, two at 25%; Qwen3 ran one round, so its intervals are over problems). Machines sorted by the probe’s link-to-CPU ratio.

Machine	Link/ CPU	Budget	G (ms)	Deployed (ms)			lrm	Few-1R
				Eq. S1	meas.	Margin		
EPYC 7302*	0.14	gpt-oss 11%	4.3	12.2	17.7	1.06 [1.01, 1.10]	0.97 [0.92, 1.01]	0.85 [0.82, 0.88]
EPYC 7302*	0.14	gpt-oss 25%	4.7	8.7	12.2	1.08 [1.00, 1.17]	1.10 [1.06, 1.14]	–
EPYC 7302*	0.14	Qwen3 12.5%	5.8	20.4	25.4	1.07 [1.02, 1.13]	0.93 [0.86, 1.00]	–
EPYC 7302*	0.14	Qwen3 25%	6.3	14.3	17.2	1.01 [0.93, 1.08]	0.84 [0.76, 0.94]	–
Xeon 8347C*	0.23	gpt-oss 11%	4.1	12.0	19.5	1.28 [0.99, 1.56]	1.08 [0.80, 1.45]	1.33 [1.00, 1.61]
Xeon 8347C*	0.23	gpt-oss 25%	4.5	8.6	10.1	1.06 [1.05, 1.08]	1.02 [0.99, 1.04]	–
Xeon 8347C*	0.23	Qwen3 12.5%	5.3	21.4	45.0	1.96 [1.67, 2.24]	1.29 [0.94, 1.74]	–
Xeon 8347C*	0.23	Qwen3 25%	5.8	14.9	16.6	0.86 [0.71, 1.02]	0.73 [0.59, 0.88]	–
Pf again (285K)	0.29	gpt-oss 11%	4.3	13.8	13.7	1.02 [1.01, 1.04]	1.03 [1.02, 1.03]	1.02 [1.01, 1.03]
Pf again (285K)	0.29	gpt-oss 25%	4.7	9.6	9.5	1.06 [1.04, 1.07]	1.05 [1.04, 1.07]	–
Pf again (285K)	0.29	Qwen3 12.5%	5.4	24.8	24.4	1.02 [1.00, 1.03]	1.02 [1.01, 1.03]	–
Pf again (285K)	0.29	Qwen3 25%	5.9	16.9	16.2	1.04 [1.03, 1.04]	1.03 [1.02, 1.03]	–
Threadripper 9960X again	0.32	gpt-oss 11%	4.1	9.1	9.4	1.00 [0.99, 1.02]	0.98 [0.97, 1.00]	1.06 [1.05, 1.06]
Threadripper 9960X again	0.32	gpt-oss 25%	4.4	6.9	7.0	1.02 [1.00, 1.04]	0.99 [0.97, 1.00]	–
Threadripper 9960X again	0.32	Qwen3 12.5%	5.4	15.1	14.4	1.02 [1.01, 1.02]	0.99 [0.99, 1.00]	–
Threadripper 9960X again	0.32	Qwen3 25%	5.7	11.2	10.5	1.01 [1.01, 1.01]	0.98 [0.98, 0.99]	–
Pe again (9950X)	1.02	gpt-oss 11%	4.5	20.4	20.5	1.01 [1.01, 1.01]	1.01 [1.01, 1.02]	1.34 [1.33, 1.35]
Pe again (9950X)	1.02	gpt-oss 25%	4.5	12.8	12.4	1.01 [1.01, 1.01]	1.01 [1.00, 1.02]	–
Pe again (9950X)	1.02	Qwen3 12.5%	5.8	35.0	36.7	1.01 [1.01, 1.01]	1.02 [1.02, 1.02]	–
Pe again (9950X)	1.02	Qwen3 25%	5.9	22.2	22.5	1.01 [1.01, 1.01]	1.01 [1.01, 1.02]	–

machine it added started. By then one listed offer was gone and the next would have exceeded the remaining credit, so the last two machines were chosen by rules committed before each started (the cheapest verified offer not rented before with cheap traffic and downloads of at least 1000 and 900Mb/s), and both stopped at the gate. On the two valid machines, the relation predicted the admission margin’s speed from its own counters within a median 0.016 (up to 0.067, on the server), against 0.021 for predicting no change. The profiled GPU compute was 4.1–4.5 ms on gpt-oss and 5.6–5.9 ms on Qwen3, within the range of earlier machines.

Job 108, every machine: Table S22 lists the machines that ran. The EPYC 7K62 ran its two gpt-oss 11% rounds 7.0% apart and stopped, as registered; the relation under-predicted it by 26.8%. The gate stopped a Ryzen 7 5700X3D whose GPU job 100 had rented, and an EPYC 9354 machine with 323GB of memory in use. Machines were rented from the registered list in order until three were valid. On both EPYC machines the Nsight trace held almost no kernels (driver 570; the Ryzen machines ran drivers 580 and later), so their G is the median of the 7 profiles of jobs 069c and 105, as everywhere else (4.28 and 4.50 ms). This is a deviation from the registration that does not change the outcome: without the EPYC 7543 the median error over the 4 cells is 5.3% and all are within 8%, but two pooled clauses still fail. The registered offer list said none of its offers was in the rental ledger, but three were: the 5700X3D’s (job 100; the gate turned it away), the 9950X3D’s (rented for a few minutes in job 099 without producing data; its GPU matched no earlier launch’s) and an i9-13900KF’s that was never reached.

TABLE S21

Job 107, registered before the runs, on the machines never rented before that passed the registered gate before any download ($\dagger$: rounds further apart than the registered 2%; such a machine stopped after its gpt-oss 11% rounds and is outside the predictions): the deployed cache’s time per token against Eq. (S1) (GPU compute profiled on the same machine before the timed runs; the run’s own counters), and the speed relative to the deployed cache of Margin and of MIN’s fewest-admission set copied in the step or loaded by the CPU, with round-level 95% intervals (rounds, then problems: two rounds of gpt-oss, one of Qwen3, whose intervals are over problems). Machines sorted by the probe’s link-to-CPU ratio.

Machine	Link/ CPU	Budget	G (ms)	Deployed (ms)				Few-1R, Few-2R			
				Eq. S1	meas.	Margin		in the step		by the CPU	
EPYC 7663 $\dagger$	0.21	gpt-oss 11%	4.3	11.3	14.3	1.03	[1.00, 1.05]	0.97	[0.94, 1.01]	1.13	[1.07, 1.18]
AMD engineering sample	0.44	gpt-oss 11%	4.1	8.2	20.0	1.05	[1.04, 1.06]	1.16	[1.14, 1.17]	1.13	[1.11, 1.14]
AMD engineering sample	0.44	gpt-oss 25%	4.2	6.4	13.4	0.99	[0.98, 1.00]	–		–	
AMD engineering sample	0.44	Qwen3 12.5%	5.6	13.6	32.7	1.03	[1.02, 1.04]	–		–	
AMD engineering sample	0.44	Qwen3 25%	5.9	10.4	23.7	1.07	[1.06, 1.09]	–		–	
EPYC 9754 $\dagger$	0.51	gpt-oss 11%	4.2	10.1	14.7	0.97	[0.95, 0.99]	1.18	[1.15, 1.20]	1.05	[1.02, 1.08]
Ryzen 9 5900XT	0.75	gpt-oss 11%	4.4	27.4	29.5	1.02	[1.02, 1.02]	1.42	[1.41, 1.43]	1.04	[1.04, 1.04]
Ryzen 9 5900XT	0.75	gpt-oss 25%	4.5	16.6	17.0	1.02	[1.01, 1.02]	–		–	
Ryzen 9 5900XT	0.75	Qwen3 12.5%	5.8	48.3	54.4	1.02	[1.02, 1.02]	–		–	
Ryzen 9 5900XT	0.75	Qwen3 25%	5.9	29.7	31.7	1.02	[1.02, 1.02]	–		–	

TABLE S22

Job 108, registered before the runs and restricted in advance to desktop-class machines never rented before: the deployed cache’s time per token (gpt-oss, mean of two rounds) against Eq. (S1), with G profiled on the same machine before the timed runs (on the two EPYC hosts, whose traces held no decode kernels, the median of the profiles of jobs 069c and 105), the run’s own counters and the probe’s best rate, and the read rate the time implies as a share of that rate. $\dagger$: rounds further apart than the registered 2%; outside the predictions.

Machine	Cores	B_{host}	Budget	G (ms)	Eq. (S1)	Measured	Implied/best
EPYC 7K62 $\dagger$	23	22	11%	4.3	46.2	63.1	0.71
Ryzen 9 5950X	16	36	11%	4.5	27.2	29.1	0.92
Ryzen 9 5950X	16	36	25%	4.6	16.5	16.7	0.99
Ryzen 9 9950X3D	16	65	11%	4.2	16.8	18.1	0.91
Ryzen 9 9950X3D	16	65	25%	4.2	10.8	11.2	0.94
EPYC 7543	30	84	11%	4.3	14.9	22.7	0.58
EPYC 7543	30	84	25%	4.5	10.1	14.0	0.61

The probe’s rate, and machines as the unit: B_{host} is the highest single reading of one probe run. Its second-highest reading is a median 0.8% lower and more than 5% lower on 1 of 42 launches (28 machines): launch 107d (Table S21), by 26% (Table S23). With the second-highest reading, Eq. (S1) is within 6% on 34 instead of 31 of the 39 consumer launches at gpt-oss 11%; with the best CPU-only reading, which ignores that the link reads at the same time, on 22. By Eq. (S1)’s own accounting the engine reads more than 5% faster than the highest reading on 12 of 84 launch-budgets (up to $1.22\times$), so the bounds are relative to the probe. Taking the machine as the unit, one value per machine: **Few-1R** gains a geometric mean $1.26\times$ (95% t -interval over machines, on the log scale, 1.19–1.33) on the 15 machines that ran stably; **Margin**, a geometric mean of $1.021\times$ (1.000–1.043) at gpt-oss 11% and $1.018\times$ (0.987–1.049) at 25% on 5. Equation (S1)’s error is a mean -2.2% (-4.0 to -0.4%) at gpt-oss 11% over the 25 consumer machines, and $+1.6\%$ (0.0 to 3.1%) at 25%; on the 21 machines it was found on, -1.5% (-3.5 to 0.4%) at 11%: of the 30 launches it was found on, it under-predicts 25 at 11% and over-predicts 23 at 25%. G was profiled on the machine for 9 of the 39 consumer launches; the others use the median of the profiles. The median |error| is 3.2% on the launches with their own G and 1.7% on the others, but the profiled launches all come from the registered tests.

Admitting less, in detail: The margin κ of **Margin** was chosen as the value that minimises the deployed path’s reads on each model’s mixed-domain text, not the AIME text (scripts/online_admit.py): 3 at the smallest budget and 2 at 25%. It cuts the deployed cache’s host reads by 4–6% at gpt-oss 11% and 6–10% at 25%. On the three machines of job 106 that ran stably it ran 1.00 – $1.02\times$ and 1.01 – $1.06\times$ the deployed cache (Qwen3: 1.01 – $1.04\times$; Table S20). We had registered at least 1.01 at gpt-oss on every machine: it fell short on two of the three stable ones, and the registered pooled median of at least 1.02 held only through the two unstable machines (1.01 on the stable ones). On the two valid machines of job 107 it ran 0.987 – $1.074\times$ (median 1.021), below the registered 0.995 at one cell, the server at gpt-oss 25%. Predicted from each configuration’s own counters, Eq. (S1) gives these ratios within a median 0.013 on job 106’s stable machines, but predicting no change does almost as well (0.016), and on job 107’s server it missed the registered 0.03 at two cells. A learned reuse predictor with the same kind of margin (job 097’s, fitted on other text) reads less still (7–14% fewer) but costs 267–439 μs of host time per token on gpt-oss and runs 0.98 – $1.05\times$: it pays where host memory is slow and loses where it is fast, against the registered prediction that it would not lose.

TABLE S23

How the probe’s rate changes the picture, over the 84 launch-budgets at gpt-oss 11% and 25% of the machines that passed our checks: Eq. (S1) with the highest probe reading (as in the paper), the second-highest, and the best CPU-only reading. *Faster than the rate*: launch-budgets whose implied read rate exceeds the rate by more than 5%. The error columns are the consumer launches at 11%.

Rate	Faster than the rate (max)	Within 6%	Within 8%	Median error
highest reading	12 (1.22 $\times$)	31/39	37/39	1.7%
second-highest	15 (1.26 $\times$)	34/39	37/39	1.5%
best CPU-only	46 (1.53 $\times$)	22/39	29/39	5.6%

The slowest link: On the EPYC 7302 machine of job 106, whose link moves 13 GB/s, **Few-1R** lost in each round (0.82–0.88 $\times$); that machine computed wrong outputs throughout, so we draw nothing from it, but below some link rate even the fewest-admission set should not be copied in the step.

The machine decides how much: Across machines, the gain of every state that moves reads from the CPU to the link follows the link-to-CPU ratio (Spearman over the machines that ran it through job 104, 95% bootstrap interval): MIN-1R 0.89 [0.63, 0.97] (19 machines), the read-ahead oracle 0.87 [0.54, 0.99] (17), admitting every miss 0.95 [0.75, 0.99] (13). MIN-2R, which reads every admission twice and copies it in the background, tracks the link’s absolute rate more closely (0.79 [0.49, 0.92], 17 machines, against 0.45 for the ratio; the two intervals overlap, and the feature was picked from a search over all probe features). MIN-1R gains 1.02–1.05 $\times$ at 11% on 2 further machines with ratios 0.40 and 0.48, and loses only on the machines whose link reads at under a third of their CPU rate; issued ahead of use, as in the read-ahead oracle, the same copies lose beyond their interval on no machine they ran on (at worst 0.997 $\times$, on Pf).

Machines differ more than problems: Problems rank largely the same on every machine: Kendall’s coefficient of concordance of MIN-1R’s per-problem gain over 19 machines and 20 problems is 0.77 at 11% and 0.82 at 25%, and a problem’s gain follows how far the deployed policy’s reads exceed MIN’s on its trace (Spearman 0.88 at 11%, 0.68 at 25%). Yet machines carry 96% of the variance of its log gain and problems 1–3%.

APPENDIX J

THE EXPERT CACHE IN MORE DETAIL

This section holds the measurements of Section II that the main text summarises: where the cache’s gain comes from, the choice of the FETCH split, the machines with half the link, long outputs, FreeToken’s own headline model, two other llama.cpp-based systems, output parity, and the cost of prompts and batches. Job numbers refer to the scripts on the gpu branch.

Against the pooled bound their designs allow, FreeToken (one cache for all layers) and llama.cpp (whole layers pinned) stand at 18–40% and 8–20% in Table III’s configurations. With every weight in GPU memory (an RTX PRO 6000), stock llama.cpp decodes gpt-oss at 255.3 and Qwen3 at 177.8 tok/s.

Where the gain comes from: Figure S3 builds the system up from stock llama.cpp on one machine. Placing each layer’s hot experts is the first gain: a static cache of the 32 experts per layer most requested on other text (the models’ own MATH-500 outputs) runs 1.31 $\times$ and 1.89 $\times$ stock, which pins whole layers instead. Its hit rate (52% and 67%) depends on how well the profile matches the workload: profiled in hindsight on the measured problems it hits 75% and 72% and runs faster (hatched). Replacing experts online closes that gap without a profile, if admissions are rationed: LRU admits every miss, and on Qwen3 its copies raise host reads from 127 to 188 experts per token, so it runs 0.88 $\times$ the static cache (1.30 $\times$ on gpt-oss); the decayed-frequency policy reads 41 and 118 per token, runs 1.58 $\times$ and 1.12 $\times$ the profiled static cache, and the hindsight cache runs 0.98 [0.95, 1.01] and 1.01 [0.98, 1.04] times as fast as it. On traces LRU misses only 2–11% more than decayed frequency at these budgets (Appendix F); the rest of its cost here is the copies it admits. On host S, with everything else as in the full system, LRU runs 6–21% slower than decayed frequency across the six cells of Table III and trails FreeToken at four of them: the admission policy is most of the lead over FreeToken (77–84% of it at gpt-oss 25 and 40%, more than the whole of it at gpt-oss 11% and Qwen3 12.5 and 25%). The GPU-signalled CPU helpers, GPU slot maps, GPU sampling and FETCH add 2–9% each, and the model’s FETCH table 5–8% over the fixed one. The full system runs 2.68 $\times$ (gpt-oss) and 2.81 $\times$ (Qwen3) stock llama.cpp.

How tight is the bound?: Table S24 changes one assumption at a time on host B. Pooled slots need 4.5–18.6% fewer reads and raise the bound at the host-bound budgets by 5–10%; starting MIN warm changes it by at most 1.4%; the median probe sample instead of the highest lowers it there by 8%. The GPU term moves most: at batch size 1 the all-in-VRAM runs reach only 52% (gpt-oss) and 61% (Qwen3) of the datasheet rate, and with the measured rate the bounds at 25% and above fall to 279 and 192 tok/s, 8–9% above the all-in-VRAM speed. In the running example (gpt-oss at 25%) the bound on host B is 430 tok/s (279 with the measured GPU rate) and the cache runs 109 tok/s. With every expert in its slots the cache matches stock llama.cpp (0.997 $\times$ and 1.047 $\times$). The bound assumes an expert run on the CPU is read from DRAM (MIN’s shortest




Fig. S3. Ablation at 25% of experts on the GPU (15 AIME-25 problems, 256 tokens, host A). Each step adds one component to the one above; the hatched row is not a step: a static cache of each layer’s 32 most requested experts counted on the measured problems themselves.

TABLE S24

The bound under each tightening, tokens per second on host B, and our cache as a percentage of the bound as published (exact optimum, highest probed host rate, datasheet GPU rate) and with every tightening at once. *no-evict*: Belady’s MIN with bypass that keeps a served expert for the step; *exact*: the same, free to evict an expert after serving it; *pool*: one pool of *LC* slots shared across layers; *median*: B_{host} the median of six concurrent probe samples instead of the highest; *GPU*: the GPU term from the measured all-in-VRAM time instead of the datasheet rate; *layers*: the per-layer sum, which no system without cross-layer prefetch can beat; *all*: exact, pool, median and GPU together (a warm start, not included, changes the bound by at most 1.4%).

Model	Experts on GPU	no-evict	exact	pool	median	GPU	layers	all	Ours, % of exact	% of all
gpt-oss-120b	11%	169	172	182	159	172	163	168	41	41
gpt-oss-120b	25%	427	430	474	397	279	365	278	25	39
gpt-oss-120b	40%	519	519	519	517	279	500	278	29	55
Qwen3-30B-A3B	12.5%	92	96	101	89	96	93	93	42	43
Qwen3-30B-A3B	25%	209	214	234	197	192	193	191	30	33
Qwen3-30B-A3B	43.75%	308	308	308	307	192	298	191	35	57

re-read distance is 0.5 GB of host traffic) and that computation runs at memory speed. The routing traces are each model’s own greedy text on the 30 prompts; replayed through a simulator of the deployed policy they reproduce the engine’s hit rates within 5.3 points.

Reaching the read time: The microbenchmark of Section III (`jobs/ec2/readsched.cu`) takes MIN’s misses per layer and step from `scripts/readsched.py` and runs no model. Usable cores minus two CPU threads read their share of host memory from an ordinary buffer, as the engine’s helpers do, and the copy engine copies the rest, 13.25 MB at a time, from pinned memory into GPU slots; both buffers rotate through 4 GiB so that no cache holds them. *Per layer* splits each layer’s experts, whole, between the two paths at the probe’s rates and waits for both before the next layer; *per token* sums the same split over the token and waits once; *link only* sends every read over PCIe, as admitting every miss does (MIN-1R sends only its admissions); *CPU only* reads everything on the CPU. Table S25 gives each machine. Waiting per layer instead of per token costs at most 1 point of the host term on the two AMD machines and 7 on the Intel machine, whose link is the slowest; there, reading everything on the CPU beats the per-layer split. The split ignores which reads are admissions; the link-aware bound sends every admission over the link and raises the bound by at most 1% at 11% (`scripts/linkaware.py`, `prereg/linkaware.json`).

Choosing the split per machine: On host B, the model’s table beats the fixed table we used before at every budget of both models, by 3.3–8.0% (launch 1). Table S26 shows why at one budget: a single missed expert should run on the CPU (about 155 μs , hidden behind the GPU’s other experts) rather than be copied (about 225 μs , on the critical path); that entry alone is worth 3.7%. Fetching everything, FreeToken’s offload strategy expressed in our engine, is the worst table, and running everything on the CPU the second worst. The model’s table was within 1% of the best of five alternatives on launch 1 and

TABLE S25

Can the bound’s read time be reached (job 102)? MIN’s per-layer host reads of gpt-oss (first 4,000 steps of the AIME routing), replayed without the model through CPU threads and the copy engine. Each mode’s share of the bound’s host term (R^*S/B_{host} , ms per token in the column *Term*); *Model*: the per-layer mode predicted from the probe’s rates. Two repetitions of each mode agree within 0.9%.

Machine	Link/ CPU	C	Term (ms)	Split per layer	Split per token	Link only	CPU only	Model
Pd again	0.40	14	7.04	86%	93%	35%	92%	98%
Pd again	0.40	32	2.84	89%	93%	35%	92%	98%
9950X	1.02	14	9.70	96%	97%	95%	86%	94%
9950X	1.02	32	3.91	96%	97%	94%	86%	92%
9800X3D	0.55	14	9.56	95%	95%	53%	91%	95%
9800X3D	0.55	32	3.85	93%	94%	53%	91%	94%

TABLE S26

FETCH tables at 43.75% on Qwen3-30B-A3B (launch 1): entry n is how many of a layer’s n missed experts are copied over PCIe; the rest run on the CPU. Ratio to the fixed table, paired, 95% CI.

Table	$f(0..8)$	tok/s	vs fixed
Lite	0, 0, 1, 1, 1, 2, 2, 3, 3	109.0	1.057 [1.052, 1.062]
Model	0, 0, 1, 1, 2, 3, 3, 4, 5	108.0	1.048 [1.041, 1.055]
Model, fetch a single miss	0, 1, 1, 1, 2, 3, 3, 4, 5	104.1	1.010 [1.007, 1.012]
Fixed	0, 1, 1, 2, 3, 3, 4, 5, 6	103.1	1
None (CPU only)	0, . . . , 0	100.7	0.977 [0.970, 0.984]
All (fetch every miss)	0, 1, . . . , 8	98.7	0.957 [0.953, 0.962]

tied with it on launch 2. Given each run’s measured misses, the calibrated model of the scorecard ranks the tables correctly except the model’s table against the lite one, which are within 1%, and overstates their differences about $1.8\times$.

Machines with half the link: On RTX 5090s whose links we measured at 27–30 GB/s instead of 53 (jobs 085, 085b and 088: gpt-oss next to a Core i9-14900K, Qwen3 next to a Ryzen 9 7900 in a PCIe 4.0 slot), the model’s tables fetch less (gpt-oss 0, 0, 0, 1, 1, Qwen3 0, 0, 1, 1, 2, 2, 3, 3, 4) and beat the fixed table by 21–34% on gpt-oss and 9–18% on Qwen3, about twice the gain on host B for Qwen3 and four to six times for gpt-oss: the per-machine split matters most where the link is slow. On the Ryzen, ours leads FreeToken by $1.04\text{--}1.08\times$; with the fixed table it would trail at 12.5%. On the Core i9 the first comparison (job 085) was confounded: FreeToken’s hybrid backend ran gpt-oss at 12–66 tok/s, below its own offload backend (25–86), because its CPU executor spread 23 threads across the 8 performance and 16 efficiency cores and its calibration then sent two thirds of the misses over the slow link, the same failure the scorecard reports for llama.cpp; the $1.4\text{--}2.0\times$ lead that gave us measured a thread configuration. We reran it (job 088, a second rental of the same CPU model) with FreeToken’s executor on 8 threads pinned to the performance cores: its hybrid backend then runs $1.4\text{--}3.1\times$ faster than before (37.5, 62.1 and 92.5 tok/s; its calibration now copies 51.7% of misses), beating its own offload backend at 11 and 25% and tying it at 40%, and ours leads its better configuration by 1.58, 1.52 and $1.51\times$ (lowest interval bound 1.47). Restricted to the CPU alone, FreeToken’s 8-thread executor decodes 13.7 tok/s at every budget. llama.cpp on the same rental runs 40.0 tok/s with one thread per performance core (−9.4% from the model’s prediction) and 32.8 with all 24 threads (+25.9%); the hybrid-core failure is a thread-count choice, and the model holds once the choice is right. The model predicted llama.cpp on the Ryzen within +2.7%. A figure in the scorecard puts every same-machine comparison on one axis.

Long outputs: Our lead does not come from the first 256 tokens. On gpt-oss, on ten held-out hard problems (AIME 2022, problems 11–15 of both exams) that neither system has finished reasoning about after 2,048 tokens (at most 1 of 10 reach an answer), ours runs 1.36 [1.34, 1.38], 1.32 [1.30, 1.34] and 1.19 [1.17, 1.21] times FreeToken at 11, 25 and 40% over all 2,048 tokens, and 1.37 [1.34, 1.39], 1.33 [1.31, 1.35] and 1.19 [1.17, 1.22] over tokens 257–2,048 (job 083, host A); both engines stream one event per token. On Qwen3, over tokens 257–2,048 of ten held-out MATH-500 problems, ours runs 1.043 [1.020, 1.070], 1.056 [1.036, 1.077] and 1.054 [1.025, 1.086] times FreeToken at 12.5, 25 and 43.75%. Both leads are larger than over the same runs’ first 256 tokens. Turning off FreeToken’s prefill overlap, which lends it cache slots during each prompt, changes its overall speed by 1% or less.

FreeToken’s own headline model: On FreeToken’s own headline model, Qwen3.6-35B-A3B in BF16 (256 experts of 6.3 MB per layer, hybrid linear attention), the two systems are close: ours runs 1.030 [1.023, 1.038] times FreeToken at 12.5% and ties at 25 and 37.5% (1.010 [0.998, 1.021], 1.003 [0.994, 1.013]), at $1.98\text{--}2.48\times$ llama.cpp (job 086, host A). FreeToken as shipped, with its automatic backend and cache size, runs this model at 93.8 tok/s on that machine, above the 77–83 tok/s its paper reports on an RTX 5090, so our setup does not handicap it; FreeToken’s fetch cap and thread count were left at their defaults, while our FETCH table was swept. Our lead is therefore model-dependent: largest with few large experts (gpt-oss,

top-4 of 128), smallest with many small ones.

FreeToken tuned: On an RTX 5090 with a Ryzen 9 9950X (job 098, one launch, configurations in a shuffled order per cell) FreeToken ran in five settings per cell: hybrid with its calibrated fetch split (its default), hybrid with the fetch cap fixed at one and at two, hybrid with 8 CPU threads instead of one per physical core, and offload. Its best setting was the calibrated hybrid with 8 threads at three cells and offload at the other three (gpt-oss 25 and 40%, Qwen3 43.75%), at most $1.04\times$ its Table III setting; the fixed caps were never better than its own split by more than 0.3%. Ours over FreeToken’s best: 1.17, 1.17 and 1.09 on gpt-oss, 1.02, 1.04 and 0.98 on Qwen3, within 0.04 of host S’s ratios at every cell. With every expert on the CPU, stock llama.cpp decodes gpt-oss $1.13\times$ and Qwen3 $0.99\times$ as fast as FreeToken’s CPU backend.

Other systems and outputs: Two newer llama.cpp-based systems, Pipelined Sharding [S7] and leloch’s expert cache [S8], run gpt-oss-120b at $1.28\text{--}1.45\times$ stock llama.cpp; ours ran $1.42\text{--}2.34\times$ the better of them (job 077). KTransformers [S9] (kt-kernel 0.7.0, exact BF16 CPU experts, the same number of experts statically on the GPU) ran Qwen3 in BF16 at $0.98\text{--}1.01\times$ stock llama.cpp on an RTX 5090 with a Ryzen 9 9950X, and ours at $2.2\text{--}4.4\times$ it (job 078); both static placements read the same CPU expert bytes from the same DRAM. Its published speeds rely on AMX Xeons and INT4/INT8 CPU weights, which that machine lacks and which change the model, so this is not its best setting. Teacher-forced on fixed text (1536 steps per model) against stock llama.cpp with every weight in the VRAM of an RTX PRO 6000 (job 090), our cache at 25% has a mean KL divergence over the full vocabulary of 0.0019 nats on gpt-oss and 0.0005 on Qwen3 (99.9th percentile 0.121 and 0.059, maximum 0.21 and 0.13), picks the same next token on 98.5% and 99.3% of steps, and changes the loss by +0.16% and +0.19%. Stock llama.cpp’s own `--n-cpu-moe` placement diverges from the same reference by as much (mean KL 0.0017 and 0.0008, agreement 98.8% and 99.3%): the differences are those of running experts on the CPU (the helpers quantise activations to Q8_0 where the GPU path uses Q8_1), not of the cache. Task accuracy was not measured.

Prompts and batches: The cache serves single-token steps; prompts and batches above one run llama.cpp’s stock graph with every expert in host memory, and pay for the cache’s resident experts and its one-thread CPU side. In `llama-batched-bench` at 25% on host S (job 089, both models) our prefill at 512 and 2,048 tokens runs 79–83% of stock llama.cpp’s and our decode at 2 and 4 parallel sequences 77–79% of it (the tool feeds both engines the same random tokens).

APPENDIX K RELATED WORK IN FULL

Offloading systems: Demand-paging systems cache experts per layer and copy misses over PCIe [S10]–[S18], or use predicted expert activations to decide what to keep [S19]; an upstream llama.cpp pull request adds such a cache, admitting every miss into an LRU [S20]. CPU-execution systems run non-resident experts on the host [S9], [S21]–[S25], as PowerInfer runs a dense model’s cold neurons [S26]; a set-associative GPU cache with CPU-executed misses [S27], a llama.cpp fork [S8], [S28] and two llama.cpp pull requests [S29], [S30] are the closest designs to ours. FreeToken [S31] and DALI [S32] split each layer’s misses between CPU and PCIe. Predictors guess upcoming experts from the previous layer’s state [S3], [S4], several tokens ahead with a learned forecaster [S6], [S33], or from draft tokens [S5], [S34], [S35]; Read-ME decouples the router so routing is known in advance [S36]. Batched-throughput systems [S37]–[S39] and flash or SSD tiers [S40]–[S42] are outside our scope.

Bounds, oracles and caching theory: WiSP [S43] bounds decode by PCIe turnover, Budgeting Bytes [S44] by bytes over bandwidth for a given fetch set, and MoE-CAP [S45] measures bandwidth utilisation against a single peak; the roofline [S46] is the single-memory ancestor of our bound. Zhang [S47] propose a trace-level evaluation contract and report a 44–46% gap between causal policies and MIN with bypass; Liang et al. [S48] measure, across 20 models, the hit rate of an oracle cache that sees m tokens ahead. We add the time domain: oracles inside a real engine, a measured decomposition of the gap, the cost of how foresight is spent, and a horizon rule in distinct experts. SAEM [S49] gives its own stage-level planner perfect knowledge of the coming stage and finds that worth at most about 14% on an A100; SpecMD [S50] benchmarks expert-cache and speculative prefetch policies across hardware; Angelopoulos et al. [S51] give competitive bounds for layered expert paging. With bypass, MIN maximises hits exactly [S52], [S53]; when bypass and admission cost differ, the time-optimal placement is a different problem [S54], which is why we use MIN’s reads as an input to a bound rather than as a time-optimal policy. Choosing among hit-optimal schedules by what admissions cost has precedents: Demand-MIN for prefetching [S55], offline-optimal caching as a min-cost flow over reuse intervals [S56], and admission rationing in flash caches [S57]. Our fewest-admission schedule is that interval formulation with admissions as a second objective, measured in an engine. Learning-augmented paging is most robust with predictions of whether a page will be evicted before its next use [S58]; learned replacement imitates Belady in CPU and content-delivery caches [S53], [S59] and in an SSD-backed expert cache [S60].

APPENDIX L LIMITATIONS IN DETAIL

- **Hardware.** Power limits differed between machines. Rented machines are often shared or unsteady: of 8 launches on server processors only three passed every check. In the two tests of the time relation on new machines, 3 and 2 machines stopped




Fig. S4. Published MoE-offloading measurements (52 rows from 13 sources) as a percentage of their own bound, computed from the paper’s hardware at datasheet rates and, where a routing trace of the model exists, the per-layer Belady optimum (filled); hollow points assume uniform routing, which overstates the optimum’s reads and so the percentage; the dotted lines are the class medians (no GPU expert capacity, 49.0%; every expert resident, 4.0%). Right: our cache, FreeToken and stock llama.cpp on host B against the probed per-layer bound of Table III (the pooled bound FreeToken’s and llama.cpp’s designs allow is 0–10% higher).

at the gate and 2 and 1 ran unsteadily. The first RTX 4090 test deviated from its registration in three ways: its output check used the RTX 5090’s reference loss, so its machines were relaunched with the RTX 4090’s; two relaunches that failed to download the model were repeated once; and one of the four machines eligible for relaunch was not relaunched. In job 112 one machine never finished downloading the model, and its replacement started before it was given up. Job 113 ran on two machines and stopped there, as registered. Job 114 started seven machines, its registered limit: three had GPUs rented before and one ran unsteadily, so three were valid, not the five it aimed for; its replacements skipped a cheaper listed offer each time, a deviation from its registered order (Appendix D).

- **The bounds.** MIN’s count lets an expert served in a step be evicted later in that step; keeping it for the step raises the bound slightly (Table S24). R^* is counted on the 30-problem trace; on the first 20 problems, which the later experiments time, it changes by -0.4% . The rate check of Table S23 uses the closed-form account of Appendix H. T_{GPU} is not a property of the machine alone; with the RTX 4090’s own profile (3.5 ms) the residual on the slow-link RTX 4090s is 27–31% of the gap at 11%. Measured against Eq. (3), the tighter bound on slow links, the read-ahead oracle still leaves 83–99% of the gap to it on the two slow-link panel machines and 73–81% on the three slow-link RTX 4090s, against a median of 45% on fast links: the oracles fetch their misses over the link, the slow path there.
- **The decomposition.** The GPU’s non-expert time is profiled, not removed. The machines pool runs of 20 and 30 problems and three versions of the engine. Equation (3) assumes that the CPU reads an expert’s weights when it runs the expert. The shares depend on the probe through Eq. (1): with every machine’s B_{host} 10% (22%) higher, the fast-link shares of Table IV at 11% move by at most 5 (9) points for *Both* and 6 (11) for *left*. The CPU-only column rests on three machines, and two of its three are Ryzen 9 5950Xs.
- **The oracles.** In the engine the fewest-admission schedule misses 2.5–5.3% more than the greedy one, more than its replay predicts; its measured gains include that cost. Degraded windows draw errors independently, while a real predictor’s errors correlate with what the cache holds.

APPENDIX M THE AUDIT’S ROWS

Table S27 lists the 52 published measurements of Section III with their bounds. The sources are Fiddler [S21], ProMoE [S12], KTransformers [S9], HybriMoE [S22], Fate [S61], SP-MoE [S34], MoE-SpeQ [S62], the CPU–GPU collaborative inference of Huang et al. [S27], FreeToken [S31], SeqMoE [S6], NVIDIA’s pipelined sharding [S7], a llama.cpp fork [S8] and a llama.cpp pull request [S30]. The in-class subset of Section III (20 rows) is the trace-scored rows less the lossy (Fate) and speculative (SP-MoE, MoE-SpeQ) ones; over all 29 trace-scored rows the median is 9.5%. These systems mostly read on demand, so Eq. (2) would place them higher. Three caveats apply to every row: the ceilings are datasheet ones where ours are probed; the trace is the model’s own sampled text, sometimes of a sibling variant (an FP8 trace for a Q4 deployment); and systems whose

cache is pooled across layers, whose placement pins whole layers, or which verify several tokens per pass are scored against the per-layer one-token bound, which is not theirs.

TABLE S27

The audit's rows (Section III): every published batch-1 measurement we scored, with the system's stated hardware, its experts per layer resident on the GPU (C), the model's experts and top- k , the published speed, the bound at datasheet-physical ceilings (GPU datasheet bandwidth, the top of the host-DRAM band, datasheet PCIe, no latencies; `scripts/audit.py`), how the optimum's reads were obtained (*trace*: Belady's MIN with bypass on an on-policy routing trace of the model; *uniform*: independent uniform routing, which overstates the reads and so the percentage; *no slots*: no GPU expert capacity, the CPU-only rate is the bound; *all fit*: every expert resident, nothing offloaded) and the percentage of the bound. Medians: trace 9.5% ($n=29$), uniform 34.8% (11), no slots 49.0% (8), all fit 4.0% (4), all rows 14.5% (52). Rows below 5% are more likely configuration mismatches with the bound's assumptions (precision, dense layers on the CPU, end-to-end timing) than inefficiency, and are not adjudicated.

System	Model	GPU	C	E/k	tok/s	bound	reads	%
Fiddler [S21]	Mixtral-8x7B	Quadro RTX 6000	1	8/2	2.5	17	trace	14.6
	Mixtral-8x7B	RTX 6000 Ada	3	8/2	8.3	57	trace	14.4
ProMoE [S12]	DeepSeek-V2-Lite	RTX 4090	45	64/6	49.9	216	trace	23.1
	Qwen1.5-MoE-A2.7B	RTX 4090	46	60/4	67.2	216	trace	31.1
	Qwen2-57B-A14B	RTX 4090	42	64/8	26.9	136	trace	19.8
	Mixtral-8x7B	RTX 4090	6	8/2	27.0	152	trace	17.8
KTransformers [S9]	DeepSeek-V3-0324	A100 40GB	0	256/8	5.9	12	no slots	48.8
	DeepSeek-V2.5-1210	A100 40GB	0	160/6	11.9	24	no slots	49.2
	Qwen2-57B-A14B	A100 40GB	0	64/8	22.9	35	no slots	64.9
	DeepSeek-V3-0324	RTX 4080	0	256/8	16.1	33	no slots	48.3
	DeepSeek-V2.5-1210	RTX 4080	0	160/6	17.6	32	no slots	55.3
	Qwen2-57B-A14B	RTX 4080	0	64/8	37.0	48	no slots	76.9
HybriMoE [S22]	DeepSeek-V2-Lite-Chat	RTX A6000	16	64/6	14.5	470	trace	3.1
	DeepSeek-V2-Lite-Chat	RTX A6000	48	64/6	15.3	655	trace	2.3
	Mixtral-8x7B-Instruct	RTX A6000	2	8/2	4.4	53	trace	8.2
	Mixtral-8x7B-Instruct	RTX A6000	6	8/2	7.3	133	trace	5.5
	Qwen2-57B-A14B-Instruct	RTX A6000	16	64/8	9.6	86	trace	11.1
	Qwen2-57B-A14B-Instruct	RTX A6000	48	64/8	9.0	116	trace	7.7
Fate [S61]	Qwen1.5-MoE-A2.7B	RTX 3090	46	60/4	14.0	198	trace	7.1
	Qwen1.5-MoE-A2.7B	GTX 1080 Ti	15	60/4	3.7	76	trace	4.9
SP-MoE [S34]	Mixtral-8x7B	RTX 3090	1	8/2	1.2	16	trace	7.7
	Mixtral-8x7B	RTX 4090	1	8/2	1.3	13	trace	9.5
	Mixtral-8x7B	A100 40GB	3	8/2	1.5	32	trace	4.5
	Phi-3.5-MoE	RTX 4090	3	16/2	3.3	57	trace	5.8
	DeepSeek-V2-Lite	A100 40GB	7	64/6	5.6	142	trace	3.9
MoE-SpeQ [S62]	DeepSeek-V2-Lite	A100 40GB	64	64/6	10.1	335	all fit	3.0
	DeepSeek-V2-Lite	A100 40GB	64	64/6	13.8	335	all fit	4.1
	Qwen1.5-MoE-A2.7B	A100 40GB	60	60/4	13.2	335	all fit	3.9
	Qwen1.5-MoE-A2.7B	A100 40GB	60	60/4	28.3	335	all fit	8.5
	Phi-3.5-MoE	A100 40GB	7	16/2	6.0	129	trace	4.7
	Phi-3.5-MoE	A100 40GB	7	16/2	6.1	129	trace	4.7
CPU-GPU collab. [S27]	Mixtral-8x7B	RTX 4090	1	8/2	4.8	13	trace	36.5
	Phi-3.5-MoE	RTX 4090	3	16/2	10.4	56	trace	18.6
FreeToken [S31]	Qwen3.6-35B-A3B	RTX 5090	94	256/8	77.1	227	uniform	34.0
	DeepSeek-V4-Flash	RTX 5090	28	256/6	24.9	81	uniform	30.8
	Qwen3.6-35B-A3B	RTX 5090	104	256/8	73.8	205	uniform	36.1
	Qwen3.6-35B-A3B	RTX 4090	72	256/8	42.9	123	uniform	34.8
	Qwen3.6-35B-A3B	RTX 4060 Laptop	60	256/8	39.3	97	uniform	40.4
	GLM-5.2	RTX PRO 6000 Blackwell	36	256/8	14.9	30	uniform	49.9
SeqMoE [S6]	Qwen3-30B-A3B-FP8	RTX 4090	57	128/8	104.1	374	trace	27.9
	Qwen3-30B-A3B-FP8	RTX 4090	19	128/8	47.9	374	trace	12.8
	Qwen3.6-35B-A3B-BF16	RTX 5090	102	256/8	118.3	305	uniform	38.8
	Qwen3.6-35B-A3B-BF16	RTX 5090	38	256/8	70.2	212	uniform	33.2
	GPT-OSS-120B-MXFP4	RTX PRO 6000 Blackwell	57	128/4	139.2	378	trace	36.8
	DeepSeek-V4-Flash	RTX PRO 6000 Blackwell	115	256/6	57.4	183	uniform	31.3
Pipelined sharding [S7]	Qwen3-30B-A3B-Instruct-2507	RTX 5090	0	128/8	25.7	199	no slots	12.9

continued

System	Model	GPU	C	E/k	tok/s	bound	reads	%
	Qwen3-30B-A3B-Instruct-2507	RTX 5090	44	128/8	32.1	1032	trace	3.1
	Qwen3-235B-A22B-Instruct-2507	RTX 5090	0	128/8	7.7	40	no slots	19.3
	Qwen3-235B-A22B-Instruct-2507	RTX 5090	42	128/8	11.5	108	uniform	10.7
	Qwen3-30B-A3B-Instruct-2507	RTX 5070 Ti	107	128/8	54.9	511	trace	10.8
llama.cpp fork (leloch) [S8]	Laguna S 2.1	RTX 5090	66	256/10	31.3	88	uniform	35.5
llama.cpp PR 27861 [S30]	Qwen3-30B-A3B Q4_K_M derivative	Radeon RX 7600 8GB	32	128/8	16.5	169	trace	9.8

APPENDIX N ARTIFACT DESCRIPTION

- **System.** A patch to llama.cpp commit 4da6337 (the expert cache, FETCH, the mailbox and the llama-ec-bench instrumentation: teacher-forced loss, logit dumps, routing lookahead records); built with CUDA 13.0 for compute capability 12.0.
- **Jobs.** One shell script per job; each installs FreeToken at commit 0d652e7, downloads the models from Hugging Face, runs the bandwidth probes, writes any on-machine prediction, runs the client, and stores raw rows (one JSON line per problem), server logs and engine counters. A small runner rents a Vast instance, streams a heartbeat, fetches the results and only then destroys the instance.
- **Analysis.** scripts/*_stats.py compute every interval from the raw rows; scripts/perlayer_model.py, speed_limit_v2.py, shapley_gap.py, policy_study.py, foresight_exact.py, batchk_trace.py, fig_audit_sol.py and scorecard.py produce the model comparison, the bound, the accounting, the trace studies, the audit figure and the scorecard; scripts/wsg_tables.py and wsg_numbers2.py generate every table, figure and number in this paper from their outputs. The audit’s 52 published rows and the physical speed-of-light model they are scored against are data/published_measurements.json, data/audit/ and scripts/audit.py.
- **Traces.** R^* is computed by scripts/speed_limit_v2.py from the routing traces of the AIME-25 text the engine runs, on the gpu branch: results/084c_gptoss_trace@vast/route_aime25_gptoss.npz and results/084b_vram_rerun@vast/route_aime25_qwen3.npz; prereg/speed_limit_v2.json holds its output; data/traces/ holds traces of further models.
- **Cost.** The 154 rentals (90 distinct offers) of jobs 058–114 cost 112.7 US dollars at their hourly GPU prices, by the ledger in the repository; storage and download charges, which the ledger omits, add to this.

REFERENCES

- [S1] T. Hoeffler and R. Belli, “Scientific benchmarking of parallel computing systems: Twelve ways to tell the masses when reporting performance results,” in *Proc. Int. Conf. for High Performance Comput., Networking, Storage and Anal. (SC)*, 2015.
- [S2] T. Kalibera and R. Jones, “Rigorous benchmarking in reasonable time,” in *Int. Symp. Memory Manage. (ISMM)*, 2013.
- [S3] R. Hwang, J. Wei, S. Cao, C. Hwang, X. Tang, T. Cao, and M. Yang, “Pre-gated MoE: An algorithm-system co-design for fast and scalable mixture-of-expert inference,” in *Proc. 51st Annu. Int. Symp. Comput. Archit. (ISCA)*, 2024, arXiv:2308.12066.
- [S4] S. Zhu, S. Bohl, R. Oester, and G. Alonso, “Pre-attention expert prediction and prefetching for mixture-of-experts large language models,” *arXiv preprint arXiv:2511.10676*, 2025.
- [S5] S. Li, J. Lin, D. Ge, and Y. Ye, “MoE-SpAc: Efficient MoE inference based on speculative activation utility in heterogeneous edge scenarios,” *arXiv preprint arXiv:2603.09983*, 2026.
- [S6] Z. Wang, Y. Wang, L. Gong, C. Tang, W. Lou, T. Wang, C. Wang, and X. Zhou, “SeqMoE: Toward full-load performance via predictive and graph-compatible MoE offloading,” *arXiv preprint arXiv:2609.12978*, 2026.
- [S7] A. Ukarande, D. Shekhar, M. Blackstein, and R. Rangan, “Efficient, VRAM-constrained xLM inference on clients,” in *Proc. Mach. Learn. and Syst. (MLSys)*, vol. 8, 2026, industry track. arXiv:2604.26334. Artifact: <https://github.com/deepshnv/pipeshard-mlsys26-ae>, tag v2.0.3-mlsys26.
- [S8] leloch, “MoE expert cache for llama.cpp: VRAM caching of hot CPU-resident experts with hybrid hit/miss execution,” <https://github.com/leloch/llama.cpp/tree/moe-cache-v2-pr>, 2026, branch moe-cache-v2-pr, commit e3096b0 (6 August 2026). Proposed upstream as llama.cpp pull request #24524 and RFC discussion #24528; not merged. Accessed September 2026.
- [S9] H. Chen, W. Xie, B. Zhang, J. Tang, J. Wang, J. Dong, S. Chen, Z. Yuan, C. Lin, C. Qiu, Y. Zhu, Q. Ou, J. Liao, X. Chen, Z. Ai, Y. Wu, and M. Zhang, “KTransformers: Unleashing the full potential of CPU/GPU hybrid inference for MoE models,” in *Proc. 31st ACM Symp. Oper. Syst. Princ. (SOSP)*, 2025.
- [S10] A. Eliseev and D. Mazur, “Fast inference of mixture-of-experts language models with offloading,” *arXiv preprint arXiv:2312.17238*, 2023.
- [S11] L. Xue, Y. Fu, Z. Lu, L. Mai, and M. Marina, “MoE-Infinity: Efficient MoE inference on personal machines with sparsity-aware expert cache,” *arXiv preprint arXiv:2401.14361*, 2024.
- [S12] X. Song, Z. Zhong, R. Chen, and H. Chen, “ProMoE: Fast MoE-based LLM serving using proactive caching,” *arXiv preprint arXiv:2410.22134*, 2024.
- [S13] P. Tang, J. Liu, X. Hou, Y. Pu, J. Wang, P.-A. Heng, C. Li, and M. Guo, “HOBbit: A mixed precision expert offloading system for fast MoE inference,” *arXiv preprint arXiv:2411.01433*, 2024.
- [S14] R. Kong, Y. Li, Q. Feng, W. Wang, X. Ye, Y. Ouyang, L. Kong, and Y. Liu, “SwapMoE: Serving off-the-shelf MoE-based large language models with tunable memory budget,” in *Proc. 62nd Annu. Meeting of the Assoc. for Comput. Linguistics (ACL)*, 2024.
- [S15] S. Zhong, L. Liang, Y. Wang, R. Wang, R. Huang, and M. Li, “AdapMoE: Adaptive sensitivity-based expert gating and management for efficient MoE inference,” in *Proc. IEEE/ACM Int. Conf. Comput.-Aided Des. (ICCAD)*, 2024.
- [S16] R. Yi, L. Guo, S. Wei, A. Zhou, S. Wang, and M. Xu, “EdgeMoE: Empowering sparse large language models on mobile devices,” *arXiv preprint arXiv:2308.14352*, 2023.
- [S17] X. He, S. Zhang, K. Tang, S. Shi, Y. Wang, Z. Zeng, Z. Tang, X. Chu, H. Yin, I. W. Tsang, and Y. S. Ong, “ExpertFlow: Efficient mixture-of-experts inference via predictive expert caching and token scheduling,” *arXiv preprint arXiv:2410.17954*, 2024, accepted at DAC 2026 (per arXiv v2).
- [S18] M. Shams, “Paging the experts: A reproducible characterization of flash-backed MoE inference on iPhone,” *arXiv preprint arXiv:2609.29032*, 2026.
- [S19] Z. Du, S. Li, Y. Wu, X. Jiang, J. Sun, Q. Zheng, Y. Wu, A. Li, H. Li, and Y. Chen, “SiDA: Sparsity-inspired data-aware serving for efficient and scalable large mixture-of-experts models,” in *Proc. Mach. Learn. and Syst. (MLSys)*, 2024.
- [S20] A. Gupta, “llama : add a GPU cache for MoE experts kept in host memory,” llama.cpp pull request 29887, <https://github.com/ggml-org/llama.cpp/pull/29887>, 2026, opened 3 October 2026.
- [S21] K. Kamahori, T. Tang, Y. Gu, K. Zhu, and B. Kasicki, “Fiddler: CPU-GPU orchestration for fast inference of mixture-of-experts models,” *arXiv preprint arXiv:2402.07033*, 2024.
- [S22] S. Zhong, Y. Sun, L. Liang, R. Wang, R. Huang, and M. Li, “HybriMoE: Hybrid CPU-GPU scheduling and cache management for efficient MoE inference,” in *Des. Autom. Conf. (DAC)*, 2025, arXiv:2504.05897.
- [S23] Y. Zhang, S. Aggarwal, and T. Mitra, “DAOP: Data-aware offloading and predictive pre-calculation for efficient MoE inference,” in *Des., Autom. & Test in Eur. Conf. (DATE)*, 2025, pp. 1–7, arXiv:2501.10375.
- [S24] W. Wang, Y. Hou, Y. Ji, P. Qu, and Y. Zhang, “Achieving cloud-grade SLOs for local mixture-of-experts inference through CPU-GPU hybrid design,” *arXiv preprint arXiv:2606.10493*, 2026.
- [S25] “ATSIInfer: Automated tensor scheduling for hybrid CPU-GPU LLM inference on consumer devices,” 2026, arXiv:2607.10183.
- [S26] Y. Song, Z. Mi, H. Xie, and H. Chen, “PowerInfer: Fast large language model serving with a consumer-grade GPU,” in *Proc. ACM Symp. Oper. Syst. Princ. (SOSP)*, 2024.
- [S27] E.-M. Huang, L.-S. Lin, and C.-Y. Lee, “Efficient CPU-GPU collaborative inference for MoE-based LLMs on memory-limited systems,” *arXiv preprint arXiv:2512.16473*, 2025.
- [S28] leloch, “RFC: MoE expert cache, VRAM caching of hot CPU-resident experts with hybrid hit/miss execution,” llama.cpp discussion #24528, <https://github.com/ggml-org/llama.cpp/discussions/24528>, 2026, posted 12 June 2026; accessed September 2026.
- [S29] milto22, “llama.cpp pull request #26563: Expert caching that greatly increases performance (self-contained, off by default),” <https://github.com/ggml-org/llama.cpp/pull/26563>, 2026, accessed September 2026.
- [S30] csantiago78, “llama.cpp pull request #27861: GPU-resident LRU cache for host-offloaded MoE expert weights,” <https://github.com/ggml-org/llama.cpp/pull/27861>, 2026, draft, accessed September 2026.
- [S31] S. Yang, X. Fan, M. Pan, H. Xi, Z. Wang, S. Sun, K. Keutzer, S. Han, M. Zaharia, C. Xu, and I. Stoica, “FreeToken: Efficient edge-native MoE serving with bandwidth-adaptive execution,” *arXiv preprint arXiv:2608.16157*, 2026.
- [S32] Z. Zhu, G. Li, P. Wang, Z. Mo, M. Pei, Z. Song, X. Liang, and J. Cheng, “DALI: A workload-aware offloading framework for efficient MoE inference on local PCs,” *arXiv preprint arXiv:2602.03495*, 2026.
- [S33] S. Yadav and B. Asgari, “Mira: Memory-efficient MoE inference using adaptive caching and predictive expert staging,” *arXiv preprint arXiv:2609.38090*, 2026.
- [S34] L. Chen, Z. Wen, T. Wu, X. Zhang, and C. Wu, “SP-MoE: Speculative decoding and prefetching for accelerating MoE-based model inference,” *arXiv preprint arXiv:2510.10302*, 2025.
- [S35] A. Saxena, P.-A. Tsai, H. Taneja, A. Jaleel, and M. Qureshi, “Cascade: Utility-driven speculative decoding for mixture-of-experts,” in *Proc. Mach. Learn. and Syst. (MLSys)*, 2026, arXiv:2506.20675 (2025).
- [S36] R. Cai, Y. Ro, G.-W. Kim, P. Wang, B. Ehteshami Bejnordi, A. Akella, and Z. Wang, “Read-ME: Refactorizing LLMs as router-decoupled mixture of experts with system co-design,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, vol. 37, 2024, pp. 116 126–116 148, arXiv:2410.19123.
- [S37] Y. Sheng, L. Zheng, B. Yuan, Z. Li, M. Ryabinin, B. Chen, P. Liang, C. Ré, I. Stoica, and C. Zhang, “FlexGen: High-throughput generative inference of large language models with a single GPU,” in *Proc. 40th Int. Conf. Mach. Learn. (ICML)*, ser. Proceedings of Machine Learning Research, vol. 202, 2023, pp. 31 094–31 116, arXiv:2303.06865.

- [S38] S. Cao, S. Liu, T. Griggs, P. Schafhalter, X. Liu, Y. Sheng, J. E. Gonzalez, M. Zaharia, and I. Stoica, “MoE-Lightning: High-throughput MoE inference on memory-constrained GPUs,” *arXiv preprint arXiv:2411.11217*, 2024.
- [S39] Y. Yuan, L. Ma, and N. Talati, “MoE-Lens: Towards the hardware limit of high-throughput MoE LLM serving under resource constraints,” *arXiv preprint arXiv:2504.09345*, 2025.
- [S40] K. Alizadeh, I. Mirzadeh, D. Belenko, K. Khatamifard, M. Cho, C. C. Del Mundo, M. Rastegari, and M. Farajtabar, “LLM in a flash: Efficient large language model inference with limited memory,” in *Proc. 62nd Annu. Meeting of the Assoc. for Comput. Linguistics (ACL)*, 2024.
- [S41] F. Liang, Y. Shen, J. Hu, J. Xu, H. Gao, M. Xu, and Z. Meng, “SSD-LLaMA: SSD-native inference for trillion-parameter MoE at 1+ token/s on a consumer PC,” *arXiv preprint arXiv:2609.18110*, 2026.
- [S42] Y. Lin, Y. Wang, R. Cai, H. Liu, and X. Zeng, “The other half of the memory wall: Serving 35B MoEs from SSD with trained routing prediction,” *arXiv preprint arXiv:2609.18063*, 2026.
- [S43] J. Zhang, L. Wu, M. Darbari, and L. Hong, “WiSP: A working-set view of mixture-of-experts serving on extremely low-resource hardware,” *arXiv preprint arXiv:2606.21868*, 2026.
- [S44] H. Zhang, “Budgeting bytes: A windowed storage roofline and dual-budget architecture ablations for storage-bound LLM decoding,” *arXiv preprint arXiv:2609.04238*, 2026.
- [S45] Y. Jiang *et al.*, “MoE-CAP: Cost-accuracy-performance benchmarking for mixture-of-experts systems,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2025, arXiv:2412.07067.
- [S46] S. Williams, A. Waterman, and D. Patterson, “Roofline: An insightful visual performance model for multicore architectures,” *Commun. ACM*, vol. 52, no. 4, pp. 65–76, 2009.
- [S47] Y. Zhang, “Reproducible evaluation of MoE expert caching: Replay semantics, workload contamination, and operating regimes,” *arXiv preprint arXiv:2608.07911*, 2026.
- [S48] J. Liang, S. Wang, M. Tian, Y. Li, D. Tang, and Z. Wei, “Not all models suit expert offloading: On local routing consistency of mixture-of-expert models,” in *Int. Conf. Learn. Represent. (ICLR)*, 2026, arXiv:2505.16056.
- [S49] Zhang, Gao, and Mitra, “SAEM: Stage-aware expert management for memory-efficient MoE inference in chain-of-thought reasoning,” 2026, arXiv:2608.21614.
- [S50] D. Hoang, A. Jaiswal, M. Samragh, and M. Cho, “SpecMD: A comprehensive study on speculative expert prefetching,” in *Int. Conf. Mach. Learn. (ICML)*, 2026, arXiv:2602.03921.
- [S51] S. Angelopoulos, L. Marchal, A. Obrecht, and B. Simon, “Cache management for mixture-of-experts LLMs,” in *Euro-Par*, 2025, arXiv:2509.02408.
- [S52] M. C. Carlisle and E. L. Lloyd, “On the k -coloring of intervals,” *Discrete Appl. Math.*, vol. 59, 1995.
- [S53] A. Jain and C. Lin, “Back to the future: Leveraging Belady’s algorithm for improved cache replacement,” in *Proc. 43rd Int. Symp. Comput. Archit. (ISCA)*, 2016.
- [S54] L. Zhang, R. Karimi, I. Ahmad, and Y. Vigfusson, “Optimal data placement for heterogeneous cache, memory, and storage systems,” *Proc. ACM Meas. and Anal. of Comput. Syst. (POMACS)*, vol. 4, no. 1, pp. 6:1–6:27, 2020.
- [S55] A. Jain and C. Lin, “Rethinking Belady’s algorithm to accommodate prefetching,” in *Int. Symp. Comput. Archit. (ISCA)*, 2018.
- [S56] D. S. Berger, N. Beckmann, and M. Harchol-Balter, “Practical bounds on optimal caching with variable object sizes,” *Proc. ACM Meas. and Anal. of Comput. Syst. (SIGMETRICS)*, vol. 2, no. 2, 2018.
- [S57] A. Eisenman, A. Cidon, E. Pergament, O. Haimovich, R. Stutsman, M. Alizadeh, and S. Katti, “Flashield: a hybrid key-value cache that controls flash write amplification,” in *16th USENIX Symp. Netw. Syst. Des. and Implement. (NSDI)*, 2019.
- [S58] A. Antoniadis, J. Boyar, M. Eliáš, L. M. Favrholt, R. Hoeksma, K. S. Larsen, A. Polak, and B. Simon, “Paging with succinct predictions,” in *Int. Conf. Mach. Learn. (ICML)*, 2023.
- [S59] Z. Song, D. S. Berger, K. Li, and W. Lloyd, “Learning relaxed Belady for content distribution network caching,” in *17th USENIX Symp. Netw. Syst. Des. and Implement. (NSDI)*, 2020.
- [S60] B. Kim, J. Lee, D. Han, H.-J. Yoo, and S. Kim, “FlashMoE: Reducing SSD I/O bottlenecks via ML-based cache replacement for mixture-of-experts inference on edge devices,” *arXiv preprint arXiv:2601.17063*, 2026.
- [S61] Z. Fang, Z. Hong, Y. Huang, Y. Lyu, W. Chen, Y. Yu, F. Yu, and Z. Zheng, “Fate: Fast edge inference of mixture-of-experts models via cross-layer gate,” *arXiv preprint arXiv:2502.12224*, 2025.
- [S62] W. Wang, J. Liu, X. Hou, X. Xia, P. Tang, M. Zhang, C. Li, and M. Guo, “MoE-SpeQ: Speculative quantized decoding with proactive expert prefetching and offloading for mixture-of-experts,” *arXiv preprint arXiv:2511.14102*, 2025.