

WHERE THE SECONDS GO: BOUNDING MIXTURE-OF-EXPERTS DECODE FROM HOST MEMORY, AND MEASURING WHAT FORESIGHT IS WORTH

Anonymous Authors¹

ABSTRACT

A Mixture-of-Experts model too large for one consumer GPU keeps most of its experts in host memory, and decoding one token at a time then waits on host reads. We ask how fast such decoding could be on a given machine, where the rest of the time goes, and what knowing the future routing is worth. Two lower bounds follow from the fewest host reads any cache of C experts per layer can make (Belady’s MIN with bypass) and the machine’s measured read rates. On 23 rented RTX 5090 machines, our llama.cpp expert cache takes the GPU’s profiled compute plus its host reads at the machine’s best rate, with nothing fitted; registered before new runs, the relation held within 6% at both models and two budgets on three machines rented before, and missed on two new server machines that did not run stably. Its gap to the bound is that serialisation plus the reads MIN would not make. Oracles inside the engine show that foresight pays when spent as MIN spends it, and that among MIN’s equally good schedules the one with the fewest admissions also wins where the PCIe link is slow. Carried online, that lesson, admit less, gains 0–6% on gpt-oss. Half of MIN’s saving over admitting every miss needs the routing of about 0.65 C distinct experts per layer ahead, which no forecaster we tested approaches.

1 INTRODUCTION

Open-weight Mixture-of-Experts (MoE) models now carry 30–120 B parameters and activate a few billion per token; a desktop GPU has 24–32 GB. Running such a model for one user therefore means keeping most experts in host DRAM. Each decoded token reads its routed experts from GPU memory if they are resident, and otherwise from host memory: either the CPU runs the expert, or the expert is copied over PCIe into a GPU slot (Ivanov et al., 2021; Hoefler et al., 2021). Many systems build on this: GPU expert caches (Eliseev & Mazur, 2023; Xue et al., 2024; Song et al., 2024a), CPU execution of misses (Kamahori et al., 2024; Chen et al., 2025a; Zhong et al., 2025), a per-miss choice between the two (Yang et al., 2026; Zhu et al., 2026) and expert prediction (Hwang et al., 2024; Zhu et al., 2025; Wang et al., 2026b). Each reports a speed-up over a baseline of its choosing. None states how fast the same machine *could* decode the model within the same GPU memory, or what the rest of the time is made of. We answer three questions for batch-1 decode of gpt-oss-120b (OpenAI, 2025) and Qwen3-30B-A3B (Yang et al., 2025) on rented RTX 5090 hosts.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

How fast could it be? (Section 3) No cache holding C experts per layer reads fewer experts from host memory than Belady’s MIN with bypass. Those reads at the machine’s measured read rate bound the time per token of any system; a system that reads an expert only once its router has chosen it also pays the GPU’s own compute in series, which gives a tighter bound for every system we know of. The read time of both bounds is nearly reachable: a microbenchmark that replays MIN’s reads reaches 86–96% of it.

Where does the time go? (Section 4, Fig. 1) Our cache’s time per token is the GPU’s profiled compute plus its counted host reads at the machine’s best rate, with nothing fitted, less the background copies that fall under the GPU’s compute. Found on 34 launches of gpt-oss and registered before new runs on both models at two budgets, the relation came within 6% on all 12 model-budget cells of the 3 machines that ran stably, all of them rented before (median error 2.3%); on the two server machines new to that job, one of which computed wrong outputs, it missed by up to 52%. Its gap to the bound therefore has two parts: the GPU’s compute, which runs in series with the reads, and the reads it makes beyond MIN’s.

What is knowing the future worth? (Sections 5 and 6) Oracles of the future routing built into the engine show that foresight pays when it is spent as MIN spends it, caching MIN’s set and reading each expert once: 16–51% faster than our deployed cache on three hosts at six budgets. Which of MIN’s many equally good schedules is followed matters:

the greedy one admits so often that it loses where the link is slow, while the one with the fewest admissions gains on every machine that ran stably; where the link is slow, its copies belong in the background. Carried online, that lesson, admit less, gains 0–6% on gpt-oss. Foresight also has a price: half of MIN’s saving needs the routing of about 0.65 C distinct experts per layer ahead, which no online policy, learned admission order, speculative batch or forecaster we tested approaches.

Contributions.

- **Two bounds on time per token** for exact-routing systems with C slots per layer, from MIN’s host reads and the machine’s probed rates, with evidence that their read time is nearly reachable (Section 3).
- **A relation with nothing fitted** for where the time goes, found on 34 launches and tested, with its predictions committed before the runs, on new machines and a second model (Section 4).
- **An in-engine study of what foresight is worth** and how it must be spent, on 23 machines, including MIN’s fewest-admission schedule and an online rule that carries its lesson (Section 5).
- **A price on foresight:** its value by horizon and accuracy, and a one-constant rule for the horizon in distinct experts across nine models (Section 6).
- An artifact in which every prediction was committed before the machine it concerns started, and every number is generated by a script (Appendices B and M).

2 SETTING

Models and workload. gpt-oss-120b (36 MoE layers of 128 experts, top-4, MXFP4, 13.25 MB per expert) and Qwen3-30B-A3B in BF16 (48 layers of 128 experts, top-8, 9.44 MB per expert), on the 30 AIME-25 problems; the cache carries over between problems, as in one user’s session. The system comparison decodes 256 tokens per problem greedily through each system’s server. Everything else replays each model’s own greedy output teacher-forced, so that every configuration sees the same routing. A *budget* is the share C/E of each layer’s experts that fit in GPU slots: 11%, 25% and 40% for gpt-oss, 12.5%, 25% and 43.75% for Qwen3.

The engine. Our cache is a 3,500-line patch to llama.cpp (Gerganov & contributors, 2023–2026). Each MoE layer has C GPU slots. Routing is decided on the GPU, and resident experts run there. A *miss* is read from host memory in one of two ways: CPU helper threads run it, or it is *fetched*, copied into a slot over PCIe and run on the GPU

in the same step. A per-machine table, computed from the machine’s bandwidth probe, sets how many of a layer’s misses to fetch. The *deployed* policy admits a miss into a slot when its decayed request count (half-life 16 tokens) beats the lowest resident’s by a margin $\kappa=1$; the CPU serves the admitted miss and a background copy then fills its slot, so the expert is read from host memory twice.

Terms. *MIN* is Belady’s optimum with bypass (Belady, 1966; Zhang et al., 2020): it *admits* a miss only if some resident is needed later than the miss, and otherwise serves it without admitting it; this maximises hits under C slots (Carlisle & Lloyd, 1995; Jain & Lin, 2016). MIN has many hit-optimal schedules: the *greedy* one admits whenever it may; the *fewest-admission* one admits as rarely as any schedule with the most hits. An admitted expert is loaded *in the step*, copied over PCIe before the GPU runs it (one host read, on the critical path), or *by the CPU*, served there and copied in the background (two reads). The oracles read the routing of the rest of the problem from a trace recorded on the same machine; they decide, they do not predict (Table 1). A budget is *host-bound* when host reads, not the GPU, bound the time; the two smallest budgets of each model are. Each host’s probe measures how fast it reads its memory from CPU threads (B_c), over PCIe (B_p), and the highest rate by any method (B_{host}); B_p/B_c is its *link-to-CPU ratio*.

Statistics. Time per token is averaged over problems; ratios are ratios of means paired by problem, with 95% bootstrap intervals over problems. Machines differ far more than problems: two hosts with the same CPU model differ by 16–29% in the deployed cache’s time, while machines rented again came within 1.9% (job 101) and 3.2% (job 106) of their earlier launch. The machine is therefore the unit: claims are about orderings that hold on every machine of a stated class, and pooled intervals resample machines and then problems (Kalibera & Jones, 2013). A *launch* is one rental; in job 106 each configuration also ran in two or three *rounds*, separate processes within one launch, so that its small effects carry intervals over rounds. Each job’s predictions were committed before its machine started; Appendix B scores those of jobs 073 onward.

3 TWO BOUNDS ON TIME PER TOKEN

Consider any system that executes the model’s exact routing with at most C experts of each layer resident on the GPU. Per token each of the Lk routed experts is served by a resident copy or read from host memory, by the CPU that runs it or by the copy that brings it into a slot. Let R^* be MIN’s host reads per token, summed over layers, on the routing trace; no policy, even one that knows the future, reads fewer. If c experts per token run on the CPU, host memory supplies at least $\max(R^*, c) S$ bytes of experts of

Table 1. Cache configurations. *Reads*: how many times an admitted expert is read from host memory. All run in the same engine; the oracles read the future routing from a trace of the same text recorded on the same machine.

Name	What it admits, and how	Reads
<i>No foresight</i> deployed	decayed request count beats the lowest resident’s by $\kappa=1$; the CPU serves the miss, then a background copy fills the slot	2
single read	the same admissions, fetched into the slot in the step	1
admit every miss	every miss fetched into the lowest-scored slot (demand fetch)	1
<i>Oracles</i>		
MIN, 2 reads	MIN with bypass (greedy rule); the CPU serves the miss, then a copy	2
MIN, 1 read	MIN with bypass (greedy rule); admissions fetched in the step	1
MIN, fewest adm.	the hit-optimal schedule with the fewest admissions (an LP per host), with 2 or 1 reads	2 or 1
MIN, prefetched	MIN, 1 read, plus copies issued three steps ahead	1
Belady	every expert of the coming steps, into the slot needed furthest ahead (also with 1 read, and prefetched)	2 or 1
window W	admit every miss unless the next W tokens show the victim needed sooner; victim the lowest-scored resident not seen in the window, else the one seen furthest	1
window W , recall r	the same, each future expert kept with probability r , else replaced by a random one	1
deployed + window W	the deployed policy with the same window and victim choice; the CPU serves the miss, then a background copy	2

S bytes each, and the GPU reads the dense weights D plus $(Lk - c)S$. With everything overlapped,

$$\bar{T} \geq \min_{0 \leq c \leq Lk} \max \left(\frac{D + (Lk - c)S}{B_{\text{gpu}}}, \frac{\max(R^*, c)S}{B_{\text{host}}} \right). \quad (1)$$

At a host-bound budget the minimum is at $c=0$ and the bound is MIN’s read time, R^*S/B_{host} . The bound covers exact routing, whole experts of unmodified bytes, at most C residents per layer and one token per forward pass; it drops all latencies and lets the two read paths share B_{host} freely. (Its MIN runs over the problems in sequence and may evict an expert in the step that served it; the oracles restart at each problem and keep each step’s experts, which reads a few percent more.)

A tighter bound for systems that read on demand. Equation (1) overlaps the GPU’s work with the host reads.

At batch size 1 a system that reads an expert only after its layer’s router has chosen it cannot: layer l ’s misses are known only after layer l ’s attention and router, and layer $l+1$ starts only after layer l ’s experts. Its time is at least the GPU’s non-expert work plus the reads, in series:

$$\bar{T}_{\text{demand}} \geq T_{\text{GPU}} + \frac{R^*S}{B_{\text{host}}}, \quad (2)$$

where T_{GPU} , the GPU’s attention, dense layers, norms and routers per token, is 2.9–3.4 ms for gpt-oss in Nsight profiles of our engine on 7 hosts. Every system we compare reads on demand; only reading before routing, that is foresight, can approach Eq. (1).

Where systems stand. Table 2 places three systems against the bounds on two hosts. Our cache leads FreeToken (Yang et al., 2026) at 11 of 12 configurations and runs 2.0–4.0 $\times$ stock llama.cpp at equal GPU expert memory; tuned over five settings on a third host, FreeToken still trails at 5 of six cells (Appendix K). On the 21 machines of jobs 093–104, at gpt-oss 11%, our cache stands at 38–54% of Eq. (1) and at 55–70% of Eq. (2). Scored by the same procedure at their papers’ datasheet rates, published batch-1 systems in the bound’s class reach a median 13.6% of Eq. (1) (Appendix L).

The read time is nearly reachable. A microbenchmark (job 102) replays MIN’s per-layer reads of gpt-oss at 11% and 25% on three hosts without the model, CPU threads reading part of each layer’s experts while the copy engine copies the rest into slots; waiting for every layer, it reaches 86–96% of the bound’s read time (Table 21). An expert that is admitted and read once must cross the link; among MIN’s hit-optimal schedules the fewest admissions any of them needs is 12.4 per token at gpt-oss 11%, against 21.8 for the greedy one, and sending those over the link at the probe’s link rate raises the bound by at most 1% at 11% and 7% at 25% (scripts/linkaware.py). The GPU term of Eq. (1), by contrast, is loose (batch-1 decode reaches 52% of the datasheet rate), but it binds only at the largest budgets; relaxing every assumption at once is in Table 20.

4 WHERE THE SECONDS GO

Time is GPU compute plus host reads. For the deployed cache, one relation with nothing fitted describes the time per token T :

$$T = G + \frac{(M + (1 - G/T)A)S}{B_{\text{host}}}, \quad (3)$$

solved for T , with M the misses and A the admissions copied in the background per token (the engine’s counters), S the expert’s bytes, B_{host} the machine’s best probed read

Table 2. Decode speed at equal GPU expert memory on two RTX 5090 hosts (30 AIME-25 problems, the first 256 greedy tokens, one session). Ratios are of mean speeds, paired by problem, with 95% bootstrap intervals. Our cache uses each machine’s fetch table; FreeToken uses its faster backend per budget, picked on host B on a separate launch. *Bound*: Eq. (1) on that machine (MIN’s reads, the host’s highest probed rate) with the GPU at its datasheet rate, and at the rate all-in-VRAM decode measures at batch size 1 (52% of datasheet for gpt-oss, 61% for Qwen3; Table 20 varies the other assumptions). Against the pooled bound their designs allow, FreeToken (one cache for all layers) and llama.cpp (whole layers pinned) stand at 18–40% and 8–20%. With every weight in GPU memory (an RTX PRO 6000) stock llama.cpp decodes gpt-oss at 255 and Qwen3 at 178 tok/s. [†]Selection launch only, which favours FreeToken.

Model	Experts on GPU	llama.cpp (tok/s)	FreeToken (tok/s)	Ours (tok/s)	Ours ÷ FreeToken (95% CI)	Ours ÷ llama.cpp	Bound, datasheet GPU tok/s	Bound, measured GPU tok/s	Ours
Host B: Ryzen 9 9950X3D, card at 17001 MHz memory clock, CPU 72 GB/s, link 53, together 78, highest (the bound’s rate) 88									
gpt-oss-120b	11%	34.9	54.0 (hybrid)	69.9	1.294 [1.278, 1.312]	2.00×	172	41%	172
gpt-oss-120b	25%	40.0	85.6 (hybrid)	109.2	1.275 [1.254, 1.295]	2.73×	430	25%	279
gpt-oss-120b	40%	47.7	132.2 (offload)	152.5	1.154 [1.134, 1.172]	3.20×	519	29%	279
Qwen3-30B-A3B	12.5%	19.5	38.7 (hybrid)	40.0	1.032 [1.022, 1.042]	2.05×	96	42%	96
Qwen3-30B-A3B	25%	22.1	54.7 (hybrid)	63.1	1.153 [1.134, 1.173] [†]	2.86×	214	30%	192
Qwen3-30B-A3B	43.75%	28.4	103.1 (offload)	108.2	1.049 [1.036, 1.064]	3.81×	308	35%	192
Host S: Ryzen 9 9950X at the stock clock, card at 14001 MHz, CPU 56 GB/s, link 53, together 66, highest 71; FreeToken’s backend carried over from host B									
gpt-oss-120b	11%	28.9	48.0 (hybrid)	57.9	1.207 [1.195, 1.219]	2.00×	140	41%	140
gpt-oss-120b	25%	33.6	78.7 (hybrid)	94.1	1.196 [1.163, 1.224]	2.80×	351	27%	275
gpt-oss-120b	40%	40.1	122.2 (offload)	133.6	1.093 [1.077, 1.111]	3.34×	514	26%	275
Qwen3-30B-A3B	12.5%	16.1	32.7 (hybrid)	33.5	1.027 [1.018, 1.036]	2.08×	78	43%	78
Qwen3-30B-A3B	25%	18.5	51.4 (hybrid)	53.9	1.048 [1.037, 1.059]	2.91×	174	31%	174
Qwen3-30B-A3B	43.75%	23.8	98.3 (offload)	95.8	0.974 [0.962, 0.987]	4.02×	305	31%	189

rate, and G the GPU’s own kernel time per token from Nsight profiles of the same engine (4.1–4.7 ms for gpt-oss, and 5.4–5.9 ms for Qwen3 on the hosts of job 106 that ran stably). Every host read takes host time except the background copies that run while the GPU computes: they spread over the next step, and a share G/T of them falls in the GPU’s compute. The *plain* form drops that share, $T = G + (M+A)S/B_{\text{host}}$.

Found, then registered. We found the plain form on the 30 launches of jobs 093–104 at gpt-oss: with G the profiles’ median, it is within 6% of the measured time on 28 of them at 11%, but over-predicts at 25% (median 6%), where background admissions are a larger share of the reads, and it failed its registered test there (job 105). The overlap term follows from where those copies run, not from a fit. On the 34 launches of jobs 093–105 it brings 33 instead of 18 within 6% at 25%, at the cost of 29 instead of 32 at 11%. We then registered the overlap form before job 106, on both models at two budgets each, with G profiled on each host before its timed runs (Table 3). On new runs of three machines rented before (Pf, the Threadripper 9960X of job 105 and panel host Pe), it came within 6% on all 12 model-budget cells (median error 2.3%, largest 6.0%), where the plain form missed by up to 10.1%. The test does not single out the overlap term: discounting every admission by half passes 11 of the 12 cells, and counting misses alone 10. The two server machines new to the job did not run as the others did: one computed wrong outputs (its teacher-forced loss was 0.21–0.45 nats per token, against 0.19 on gpt-oss and 0.085 on Qwen3 everywhere else), and the other’s deployed cache varied by 39% between rounds. There the relation under-predicted by 10–52%, which takes its registered pooled

median error to 4.7% against the 4% registered; we cannot tell how much of that is the relation and how much the machines.

It reads at the machine’s best rate. With G removed, the deployed cache’s time scales with the host’s read rate with an elasticity of -0.95 (95% interval over machines -1.06 to -0.84), as reads at B_{host} would: it executes the reads it decides to make at the machine’s best rate.

So the gap to the bound has two parts. Figure 1 splits each machine’s time by Eq. (3): the GPU’s compute (12–47% of the time at gpt-oss 11%), MIN’s reads, which are the bound (31–54%), and the reads beyond MIN’s (22–35%). The first is serialisation: the GPU’s compute in series with the reads, 23–68% of the gap to Eq. (1). The second is the reads MIN would not make. Foresight can attack both: fewer reads need MIN’s choices, and overlap needs reads issued ahead (with MIN’s copies made a few steps ahead, the implied G falls to 2.7 ms). Copies in the step do the opposite: they add latency on the critical path, and the implied G rises to 7.4 ms for MIN’s set and 16.4 ms when admitting every miss, so Eq. (3) describes the deployed path, not states that copy in the step.

5 WHAT KNOWING THE FUTURE IS WORTH

The deployed cache does not know which experts the coming tokens will need. Oracles that read the future routing from a trace measure what that knowledge is worth and how it must be spent: on three hosts (O3–O5) at all six budgets, on a 10-host panel at the two host-bound gpt-oss budgets,




Figure 1. Where the seconds go, machine by machine: the deployed cache’s time per token at the two host-bound gpt-oss budgets split by Eq. (3) into the GPU’s own compute (the median of the profiles at that budget), MIN’s host reads at the machine’s best rate (the bound of Eq. (1)) and the reads the cache makes beyond MIN’s; the dot is the measured time. One launch per machine (23 machines). Nothing is fitted: G comes from profiles, the reads from the engine’s counters and the rate from each machine’s probe.

and on follow-up launches. A running example: on O4 (a Ryzen 9 9950X) at gpt-oss 11%, Eq. (1) gives 10.0 ms per token and Eq. (2) 13.0; the deployed cache takes 20.6, MIN copying each admission in the step 15.2, and MIN with the copies made a few steps ahead 13.8.

Spent as MIN spends it, foresight pays. MIN with each admission copied once, in the step, is faster than the deployed cache at every budget on O3–O5, by 16–51%, and by 25–81% with the copies made ahead (Fig. 2). The usual ways of spending foresight waste reads: Belady prefetch admits experts MIN would bypass, and serving a miss on the CPU before copying it reads every admission twice. They read 1.6–2.5 $\times$ MIN’s bytes and, at the smallest budget of each model on O4 and O5, gain at most 15% or lose; the trade-off is classical for disks (Cao et al., 1995) and CPU caches (Jain & Lin, 2018).

Which of MIN’s schedules is followed matters. The greedy schedule admits a miss whenever some resident is needed later, and may evict it again before its next use; the fewest-admission schedule copies 0.60–0.73 times as many experts in the step as the greedy one. Given to the engine on 10 launches on 7 machines that ran stably (jobs 104–106), it

gains over the deployed cache on every one (1.02–1.34 $\times$ at gpt-oss 11%), and over the greedy schedule by 0.08–0.17 in speed ratio on the 5 machines whose link reads at 0.28–0.63 of their CPU rate. On the 2 machines whose link reads at 0.28–0.32 of it, the greedy schedule loses (0.85–0.93 $\times$) and the fewest-admission one gains (1.02–1.06 $\times$); rented again in job 106 and run in three rounds, separate processes, it gained in every round, with round-level 95% intervals (resampling rounds, then problems) that start at 1.01 and 1.05. On those machines, though, reading once in the step is not the best way to spend that set: loaded by the CPU, with the copy in the background, it gains more (1.11–1.22 $\times$ at the two budgets, jobs 104–105). On a server host with a 13 GB/s link the fewest-admission set copied in the step lost in each round (0.82–0.88 $\times$), but that host computed wrong outputs throughout (Appendix J), so we draw nothing from it. Where the link is as fast as the CPU, copies in the step are cheap and the two schedules tie. Host by host, both loads and both budgets are in Appendix J.

The machine decides how much. Across machines, the gain of every state that moves reads from the CPU to the link follows the link-to-CPU ratio: for MIN copied in the step, Spearman 0.89 [0.63, 0.97] over 19 machines (Fig. 3).




Figure 2. What foresight is worth in the engine, by how it is spent: speed of each configuration of Table 1 relative to the deployed cache on the same host, at the six budgets ($\dagger$ host-bound). Hosts O3, O4 and O5 at every budget; the 10 panel hosts (grey ticks) at the two gpt-oss host-bound budgets, and 3 of them at Qwen3 12.5%. Within-host 95% intervals are mostly narrower than the markers. Admitting what MIN admits and reading it once (*MIN, 1 read*) gains at every budget on these hosts; the usual ways of spending foresight (*Belady, MIN, 2 reads*) gain at some budgets and lose at others.

Machines also differ more than problems: they carry 96% of the variance of that state’s log gain, problems 1–3% (Appendix J).

The lesson, carried online: admit less. The fewest-admission schedule wins because an admission costs a read. The deployed cache’s admissions are second reads, so the lesson applies to it without foresight. Raising its admission margin from $\kappa=1$ to the value that minimises its reads on other text (2–3; `scripts/online_admit.py`) cuts its host reads by 4–6% at gpt-oss 11% and 6–10% at 25%, and on the three hosts of job 106 that ran stably it runs 1.00–1.02 $\times$ and 1.01–1.06 $\times$ the deployed cache (Qwen3: 1.01–1.04 $\times$; Table 3). The gain is small, as Eq. (3) says it must be: admissions are a small share of the deployed cache’s reads, and part of them is hidden under the GPU’s compute. We had registered at least 1.01 at gpt-oss on every host: it fell short on two of the three stable ones, and the registered pooled median of at least 1.02 held only through the two unstable machines (1.01 on the stable ones). Predicted from each configuration’s own counters, Eq. (3) gives these ratios within a median 0.013 (at most 0.045, beyond the registered 0.03 on Pf at 25%), but predicting no change does almost as well (0.016), so this tests it little. A learned reuse predictor with the same kind of margin (job 097’s, fitted on other text) reads less still (7–14% fewer) but costs 267–439 μ s of host time per token on gpt-oss on those hosts, and runs 0.98–1.05 $\times$: it pays where host memory is slow and loses where it is fast, against the registered prediction that it would not lose, and its host time exceeded the registered

Table 3. Job 106, registered before the runs, the hosts that ran stably: the deployed cache’s time per token against Eq. (3) (GPU compute profiled on the same host before the timed runs; the run’s own counters), and the speed relative to the deployed cache of the two online admission rules (*dk*: the decayed count with a larger admission margin; *lrn*: the learned reuse predictor with a margin; both chosen on other text) and of MIN’s fewest-admission set copied in the step, with round-level 95% intervals (rounds, then problems: three rounds at gpt-oss 11%, two at 25%; Qwen3 ran one round, so its intervals are over problems). Hosts sorted by the probe’s link-to-CPU ratio.

Host	Link/ CPU Budget	G Deployed (ms)			lrn			MIN, fewest in the step
		low	meas.	dk	low	meas.	dk	
Pf again (285K)	0.29 gpt-oss 11%	4.3	13.8	13.7	1.02 [1.01, 1.04]	1.03 [1.02, 1.03]	1.02 [1.01, 1.03]	1.02 [1.01, 1.03]
Pf again (285K)	0.29 gpt-oss 25%	4.7	9.6	9.5	1.06 [1.04, 1.07]	1.05 [1.04, 1.07]	–	–
Pf again (285K)	0.29 Qwen3 12.5%	5.4	24.8	24.4	1.02 [1.00, 1.03]	1.02 [1.01, 1.03]	–	–
Pf again (285K)	0.29 Qwen3 25%	5.9	16.9	16.2	1.04 [1.03, 1.04]	1.03 [1.02, 1.03]	–	–
Threadripper 9960X again	0.32 gpt-oss 11%	4.1	9.1	9.4	1.00 [0.99, 1.02]	0.98 [0.97, 1.00]	1.06 [1.05, 1.06]	1.06 [1.05, 1.06]
Threadripper 9960X again	0.32 gpt-oss 25%	4.4	6.9	7.0	1.02 [1.00, 1.04]	0.99 [0.97, 1.00]	–	–
Threadripper 9960X again	0.32 Qwen3 12.5%	5.4	15.1	14.4	1.02 [1.01, 1.02]	0.99 [0.99, 1.00]	–	–
Threadripper 9960X again	0.32 Qwen3 25%	5.7	11.2	10.5	1.01 [1.01, 1.01]	0.98 [0.98, 0.99]	–	–
Pe again (9950X)	1.02 gpt-oss 11%	4.5	20.4	20.5	1.01 [1.01, 1.01]	1.01 [1.01, 1.02]	1.34 [1.33, 1.35]	1.34 [1.33, 1.35]
Pe again (9950X)	1.02 gpt-oss 25%	4.5	12.8	12.4	1.01 [1.01, 1.01]	1.01 [1.00, 1.02]	–	–
Pe again (9950X)	1.02 Qwen3 12.5%	5.8	35.0	36.7	1.01 [1.01, 1.01]	1.02 [1.02, 1.02]	–	–
Pe again (9950X)	1.02 Qwen3 25%	5.9	22.2	22.5	1.01 [1.01, 1.01]	1.01 [1.01, 1.02]	–	–

caps on Pe.

6 HOW MUCH FORESIGHT IS NEEDED

The oracles know the whole future; a forecaster would know part of it, imperfectly. We give the cache a *window*, the routing of the next W tokens, exact or degraded (each future expert kept with probability r and replaced otherwise, so that precision and recall are both r), and ask what it is worth in the engine and, on nine models’ traces, how long it must be. The window policies and their engine runs are detailed




Figure 3. Which way of spending foresight pays depends on the machine: speed relative to the deployed cache on the same host against the host’s link-to-CPU read-rate ratio from its own probe, at the two gpt-oss host-bound budgets. Stars: MIN’s fewest-admission set, copied in the step (jobs 104–106). Crosses: the deployed cache with the engine’s layer-ahead copy (job 105). Lines: least-squares fits in the log of the ratio, for the trend only.

in Appendix I.

A short window pays little against the deployed cache.

Given to the deployed policy, so that it changes only what is cached, a 16-token window runs $1.00\text{--}1.03\times$ the deployed cache at gpt-oss 11% on 4 hosts, about what MIN with two reads gains; shorter or half-right windows lose up to 6%. Measured against admitting every miss, the online policy that reads least at these budgets, windows recover a share of MIN’s gain quickly (Fig. 4): exact windows of 4 and 16 tokens recover 0.32–0.47 and 0.78–0.93 of it in time at gpt-oss 11%. Accuracy matters as much as horizon: an 8-token window right half the time recovers 0.20–0.26, about what 2 exact tokens buy.

The horizon is measured in distinct experts. On nine models’ own sampled text (33,000–152,000 tokens), the window W_{50} that closes half of the read gap between admitting every miss and MIN ranges from 0.6 to 34 tokens. Paging theory measures lookahead in distinct pages (Albers, 1997; Breslauer, 1995), and so does the rule that fits: the experts a layer requests within W_{50} number about $0.65C$ (quartiles 0.61–0.72). Given a model’s routing trace, this one-constant rule predicts its horizon as well as a two-constant power law in tokens (Appendix H). A forecaster’s target is thus the next $\approx 0.65C$ distinct experts of each layer.

Nothing realisable reaches it. Forecasters trained on routing histories, linear or recurrent, close at most 15% of the read gap on gpt-oss: their precision falls with the horizon (0.48 one token ahead, 0.26 four ahead), and over a window

they do no better than random errors of the same precision (Fig. 4, dotted). Speculative batches save an online cache 8–12% of its reads at 90% draft acceptance and cost 18–20% at 80% (Chen et al., 2025b; Li et al., 2026; Saxena et al., 2026). Layer-ahead predictors (Hwang et al., 2024; Zhu et al., 2025) reach 84% recall in our engine but see less than one token ahead, so they can only hide reads under the GPU’s compute: the engine’s layer-ahead copy, which copies one predicted expert per layer ahead of use, gains $1.04\times$ where the link is as fast as the CPU and loses elsewhere, down to $0.72\times$ (job 105, Table 18).

7 RELATED WORK

Offloading systems. Demand-paging systems cache experts per layer and copy misses over PCIe (Eliseev & Mazur, 2023; Xue et al., 2024; Song et al., 2024a; Tang et al., 2024; Kong et al., 2024; Zhong et al., 2024; Yi et al., 2023; He et al., 2024; Shams, 2026), or use predicted expert activations to decide what to keep (Du et al., 2024); an upstream llama.cpp pull request adds such a cache, admitting every miss into an LRU (Gupta, 2026). CPU-execution systems run non-resident experts on the host (Kamahori et al., 2024; Chen et al., 2025a; Zhong et al., 2025; Zhang et al., 2025; Wang et al., 2026a; ATSTInfer), as PowerInfer runs a dense model’s cold neurons (Song et al., 2024b); a set-associative GPU cache with CPU-executed misses (Huang et al., 2025), a llama.cpp fork (leloch, 2026a;b) and two llama.cpp pull requests (miltos22, 2026; csantiago78, 2026) are the closest designs to ours. FreeToken (Yang et al., 2026) and DALI (Zhu et al., 2026) split each layer’s misses between CPU




Figure 4. The value of partial foresight at the four host-bound budgets: the share of MIN’s gain over *admit every miss* recovered by exact routing of the next W tokens (solid), by a window whose experts are each right with probability 0.5 (dashed), and by a linear forecaster fitted on other text (dotted), in host reads on the AIME routing the engine runs; markers: the same share in time per token on the panel hosts (mean, and range over hosts; 10 hosts for gpt-oss, 3 for Qwen3 at 12.5%; no Qwen3 run at 25%). The grey line marks half of the gain.

and PCIe. Predictors guess upcoming experts from the previous layer’s state (Hwang et al., 2024; Zhu et al., 2025), several tokens ahead with a learned forecaster (Wang et al., 2026b; Yadav & Asgari, 2026), or from draft tokens (Chen et al., 2025b; Li et al., 2026; Saxena et al., 2026); Read-ME decouples the router so routing is known in advance (Cai et al., 2024). Batched-throughput systems (Sheng et al., 2023; Cao et al., 2024; Yuan et al., 2025) and flash or SSD tiers (Alizadeh et al., 2024; Liang et al., 2026a; Lin et al., 2026) are outside our scope.

Bounds, oracles and caching theory. WiSP (Zhang et al., 2026b) bounds decode by PCIe turnover, Budgeting Bytes (Zhang, 2026a) by bytes over bandwidth for a given fetch set, and MoE-CAP (Jiang et al., 2025) measures bandwidth utilisation against a single peak; the roofline (Williams et al., 2009) is the single-memory ancestor of our bound. Zhang (2026b) proposes a trace-level evaluation contract and reports a 44–46% gap between causal policies and MIN with bypass; Liang et al. (2026b) measure, across 20 models, the hit rate of an oracle cache that sees m tokens ahead against the cache ratio C/k . We add the time domain: an oracle inside a real engine, the cost of how foresight is spent, and a horizon rule in distinct experts. SAEM (Zhang et al., 2026a) gives its own stage-level planner perfect knowledge of the coming stage in its engine and finds that worth at most about 14% on an A100; SpecMD (Hoang et al., 2026) benchmarks expert-cache and speculative prefetch policies across hardware; Angelopoulos et al. (2025) give competitive bounds for layered expert paging. With bypass, MIN maximises hits exactly (Carlisle & Lloyd, 1995; Jain & Lin, 2016); when bypass and admission cost differ, the time-optimal placement is a different problem (Zhang et al., 2020), which is

why we use MIN’s reads as an input to a bound rather than as a time-optimal policy. Choosing among hit-optimal schedules by what admissions cost has precedents: Demand-MIN for prefetching (Jain & Lin, 2018), offline-optimal caching as a min-cost flow over reuse intervals (Berger et al., 2018), and admission rationing in flash caches (Eisenman et al., 2019). Our fewest-admission schedule is that interval formulation with admissions as a second objective, measured in an engine. Lookahead in paging is valuable when measured in distinct pages (Albers, 1997; Breslauer, 1995), and learning-augmented paging is most robust with predictions of whether a page will be evicted before its next use (Antoniadis et al., 2023); learned replacement imitates Belady in CPU and content-delivery caches (Jain & Lin, 2016; Song et al., 2020) and in an SSD-backed expert cache (Kim et al., 2026).

8 LIMITATIONS

- **Hardware and sample.** Rented consumer cards in desktop hosts; clocks could not be locked and power limits differed. The panel has 10 hosts, one launch each. Few hosts have a link slower than half their CPU rate: the greedy copy loses on 2 machines, so where the crossovers lie rests on few points. Of the two server machines new to job 106, one computed wrong outputs and the other varied by 39% between rounds; we report their clauses and draw nothing from them, and the summaries of job 106 cover the three machines that ran stably, all rented before. Within-host intervals resample problems; round-level intervals exist only for job 106.

- **Workload.** Batch-1 decode of two models on AIME

prompts, also used while developing the system, teacher-forced on each model’s own greedy text. Prompts and batches above one bypass the cache.

- **The bounds.** Exact routing and per-layer slots (pooling slots across layers lowers MIN’s reads by 4.5–18.6%). They drop latencies and let the read paths share the best probed rate, so a faster reader than our probe could exceed them: by Eq. (3)’s own accounting the engine read more than 5% faster than the probe’s highest rate on 13 of 68 launch-budgets (up to $1.22\times$), and the GPU term uses a datasheet rate that a card can exceed. The GPU term of Eq. (2) is our engine’s kernel time, not a property of the machine alone. The read time is nearly reachable in a microbenchmark that runs no compute.
- **The relation.** Equation (3) describes the deployed path; states that copy in the step add latency it does not model. It assumes the machine delivers its probed rate during the run. Its registered success is on new runs of machines rented before; on the two machines new to that job it missed by up to 52%, and we cannot tell how much of that is the relation. The overlap term assumes background copies spread evenly over the next step; we have not measured that directly, and it makes the fit at gpt-oss 11% slightly worse. It was tested on two models only.
- **The oracles.** Oracles read a recorded trace. In the engine the fewest-admission schedule misses 2.5–5.3% more than the greedy one, more than its replay predicts; its measured gains include that cost. Degraded windows draw errors independently; a real forecaster’s errors are correlated with what the cache holds.
- **Outputs.** Parity is measured by teacher-forced KL divergence against an all-VRAM reference (mean 0.0019 and 0.0005 nats), not by task accuracy.

9 CONCLUSION

With experts in host memory, batch-1 MoE decode is the GPU’s compute plus host bytes. A cache that reads on demand pays them in series, apart from the background copies that fall under the GPU’s compute, and our cache executes its reads at the machine’s best rate: what separates it from the bound is that serialisation and the reads it makes beyond MIN’s. Knowing the future reduces both when it is spent as MIN spends it, and which of MIN’s schedules is followed matters: the fewest-admission one gains on every machine that ran stably, while the greedy one loses where the link is slow, and there the copies belong in the background. Admitting less carries that lesson online, for 0–6%. Foresight has a measurable price: about $0.65C$ distinct experts of each layer’s coming routing to close half of the read gap between admitting every miss and MIN, at accuracies no forecaster we know of or trained reaches.

REFERENCES

- Albers, S. On the influence of lookahead in competitive paging algorithms. *Algorithmica*, 18(3):283–305, 1997. doi: 10.1007/PL00009158.
- Alizadeh, K., Mirzadeh, I., Belenko, D., Khatamifard, K., Cho, M., Del Mundo, C. C., Rastegari, M., and Farajtabar, M. LLM in a flash: Efficient large language model inference with limited memory. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- Angelopoulos, S., Marchal, L., Obrecht, A., and Simon, B. Cache management for mixture-of-experts LLMs. In *Euro-Par*, 2025. arXiv:2509.02408.
- Antoniadis, A., Boyar, J., Eliáš, M., Favrholt, L. M., Hoeksma, R., Larsen, K. S., Polak, A., and Simon, B. Paging with succinct predictions. In *International Conference on Machine Learning (ICML)*, 2023.
- ATSInfer. ATSInfer: Automated tensor scheduling for hybrid CPU-GPU LLM inference on consumer devices, 2026. arXiv:2607.10183.
- Belady, L. A. A study of replacement algorithms for a virtual-storage computer. *IBM Systems Journal*, 5(2): 78–101, 1966.
- Berger, D. S., Beckmann, N., and Harchol-Balter, M. Practical bounds on optimal caching with variable object sizes. *Proceedings of the ACM on Measurement and Analysis of Computing Systems (SIGMETRICS)*, 2(2), 2018.
- Breslauer, D. On competitive on-line paging with lookahead. Technical report, BRICS, 1995.
- Cai, R., Ro, Y., Kim, G.-W., Wang, P., Ehteshami Bejnordi, B., Akella, A., and Wang, Z. Read-ME: Refactorizing LLMs as router-decoupled mixture of experts with system co-design. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 37, pp. 116126–116148, 2024. arXiv:2410.19123.
- Cao, P., Felten, E. W., Karlin, A. R., and Li, K. A study of integrated prefetching and caching strategies. In *Proceedings of ACM SIGMETRICS*, 1995.
- Cao, S., Liu, S., et al. MoE-Lightning: High-throughput MoE inference on memory-constrained GPUs. *arXiv preprint arXiv:2411.11217*, 2024.
- Carlisle, M. C. and Lloyd, E. L. On the k -coloring of intervals. *Discrete Applied Mathematics*, 59, 1995.
- Chen, H., Xie, W., Zhang, B., Tang, J., Wang, J., Dong, J., Chen, S., Yuan, Z., Lin, C., Qiu, C., Zhu, Y., Ou, Q.,

- Liao, J., Chen, X., Ai, Z., Wu, Y., and Zhang, M. KTrans-formers: Unleashing the full potential of CPU/GPU hybrid inference for MoE models. In *Proceedings of the 31st ACM Symposium on Operating Systems Principles (SOSP)*, 2025a. doi: 10.1145/3731569.3764843.
- Chen, L., Wen, Z., Wu, T., Zhang, X., and Wu, C. SP-MoE: Speculative decoding and prefetching for accelerating MoE-based model inference. *arXiv preprint arXiv:2510.10302*, 2025b.
- csantiago078. llama.cpp pull request #27861: GPU-resident LRU cache for host-offloaded MoE expert weights. <https://github.com/ggml-org/llama.cpp/pull/27861>, 2026. Draft, accessed September 2026.
- Du, Z., Li, S., Wu, Y., Jiang, X., Sun, J., Zheng, Q., Wu, Y., Li, A., Li, H., and Chen, Y. SiDA: Sparsity-inspired data-aware serving for efficient and scalable large mixture-of-experts models. In *Proceedings of Machine Learning and Systems (MLSys)*, 2024.
- Eisenman, A., Cidon, A., Pergament, E., Haimovich, O., Stutsman, R., Alizadeh, M., and Katti, S. Flashield: a hybrid key-value cache that controls flash write amplification. In *16th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, 2019.
- Eliseev, A. and Mazur, D. Fast inference of mixture-of-experts language models with offloading. *arXiv preprint arXiv:2312.17238*, 2023.
- Fang, Z., Hong, Z., Huang, Y., Lyu, Y., Chen, W., Yu, Y., Yu, F., and Zheng, Z. Fate: Fast edge inference of mixture-of-experts models via cross-layer gate. *arXiv preprint arXiv:2502.12224*, 2025.
- Fawcett, C. and Hoos, H. H. Analysing differences between algorithm configurations through ablation. *Journal of Heuristics*, 22:431–458, 2016.
- Gerganov, G. and contributors. llama.cpp: LLM inference in C/C++. <https://github.com/ggml-org/llama.cpp>, 2023–2026.
- Gupta, A. llama : add a GPU cache for MoE experts kept in host memory. llama.cpp pull request 29887, <https://github.com/ggml-org/llama.cpp/pull/29887>, 2026. Opened 3 October 2026.
- He, X., Zhang, S., Tang, K., Shi, S., Wang, Y., Zeng, Z., Tang, Z., Chu, X., Yin, H., Tsang, I. W., and Ong, Y. S. ExpertFlow: Efficient mixture-of-experts inference via predictive expert caching and token scheduling. *arXiv preprint arXiv:2410.17954*, 2024. Accepted at DAC 2026 (per arXiv v2).
- Hoang, D., Jaiswal, A., Samragh, M., and Cho, M. SpecMD: A comprehensive study on speculative expert prefetching. In *International Conference on Machine Learning (ICML)*, 2026. arXiv:2602.03921.
- Hoefer, T. and Belli, R. Scientific benchmarking of parallel computing systems: Twelve ways to tell the masses when reporting performance results. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC)*, 2015.
- Hoefer, T., Alistarh, D., Ben-Nun, T., Dryden, N., and Peste, A. Sparsity in deep learning: Pruning and growth for efficient inference and training in neural networks. *Journal of Machine Learning Research*, 22(241), 2021.
- Huang, E.-M., Lin, L.-S., et al. Efficient CPU-GPU collaborative inference for MoE-based LLMs on memory-limited systems. *arXiv preprint arXiv:2512.16473*, 2025.
- Hwang, R., Wei, J., et al. Pre-gated MoE: An algorithm-system co-design for fast and scalable mixture-of-expert inference. In *Proceedings of the 51st Annual International Symposium on Computer Architecture (ISCA)*, 2024. arXiv:2308.12066.
- Ivanov, A., Dryden, N., Ben-Nun, T., Li, S., and Hoefer, T. Data movement is all you need: A case study on optimizing transformers. In *Proceedings of Machine Learning and Systems (MLSys)*, 2021.
- Jain, A. and Lin, C. Back to the future: Leveraging Belady’s algorithm for improved cache replacement. In *Proceedings of the 43rd International Symposium on Computer Architecture (ISCA)*, 2016.
- Jain, A. and Lin, C. Rethinking Belady’s algorithm to accommodate prefetching. In *International Symposium on Computer Architecture (ISCA)*, 2018.
- Jain, R. *The Art of Computer Systems Performance Analysis*. Wiley, 1991.
- Jiang, Y., Fu, Y., Huang, Y., et al. MoE-CAP: Cost-accuracy-performance benchmarking for mixture-of-experts systems. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. arXiv:2412.07067.
- Kalibera, T. and Jones, R. Rigorous benchmarking in reasonable time. In *International Symposium on Memory Management (ISMM)*, 2013.
- Kamahori, K., Tang, T., et al. Fiddler: CPU-GPU orchestration for fast inference of mixture-of-experts models. *arXiv preprint arXiv:2402.07033*, 2024.
- Kim, B., Lee, J., Han, D., Yoo, H.-J., and Kim, S. Flash-MoE: Reducing SSD I/O bottlenecks via ML-based cache

- replacement for mixture-of-experts inference on edge devices. *arXiv preprint arXiv:2601.17063*, 2026.
- Kong, R., Li, Y., Feng, Q., Wang, W., Ye, X., Ouyang, Y., Kong, L., and Liu, Y. SwapMoE: Serving off-the-shelf MoE-based large language models with tunable memory budget. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- leloch. MoE expert cache for llama.cpp: VRAM caching of hot CPU-resident experts with hybrid hit/miss execution. <https://github.com/leloch/llama.cpp/tree/moe-cache-v2-pr>, 2026a. Branch moe-cache-v2-pr, commit e3096b0 (6 August 2026). Proposed upstream as llama.cpp pull request #24524 and RFC discussion #24528; not merged. Accessed September 2026.
- leloch. RFC: MoE expert cache, VRAM caching of hot CPU-resident experts with hybrid hit/miss execution. llama.cpp discussion #24528, <https://github.com/ggml-org/llama.cpp/discussions/24528>, 2026b. Posted 12 June 2026; accessed September 2026.
- Li, S., Lin, J., Ge, D., and Ye, Y. MoE-SpAc: Efficient MoE inference based on speculative activation utility in heterogeneous edge scenarios. *arXiv preprint arXiv:2603.09983*, 2026.
- Li, Y., Tan, J., Wu, B., He, X., and Li, F. ShapleyIQ: Influence quantification by Shapley values for performance debugging of microservices. In *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), Volume 4*, pp. 287–323, 2024. doi: 10.1145/3623278.3624771.
- Liang, F., Shen, Y., Hu, J., Xu, J., Gao, H., Xu, M., and Meng, Z. SSD-LLaMA: SSD-native inference for trillion-parameter MoE at 1+ token/s on a consumer PC. *arXiv preprint arXiv:2609.18110*, 2026a.
- Liang, J. et al. Not all models suit expert offloading: On local routing consistency of mixture-of-expert models. In *International Conference on Learning Representations (ICLR)*, 2026b. arXiv:2505.16056.
- Lin, Y., Wang, Y., Cai, R., Liu, H., and Zeng, X. The other half of the memory wall: Serving 35B MoEs from SSD with trained routing prediction. *arXiv preprint arXiv:2609.18063*, 2026.
- miltos22. llama.cpp pull request #26563: Expert caching that greatly increases performance (self-contained, off by default). <https://github.com/ggml-org/llama.cpp/pull/26563>, 2026. Accessed September 2026.
- OpenAI. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025.
- Saxena, A., Tsai, P.-A., Taneja, H., Jaleel, A., and Qureshi, M. Cascade: Utility-driven speculative decoding for mixture-of-experts. In *Proceedings of Machine Learning and Systems (MLSys)*, 2026. arXiv:2506.20675 (2025).
- Shams, M. Paging the experts: A reproducible characterization of flash-backed MoE inference on iPhone. *arXiv preprint arXiv:2609.29032*, 2026.
- Sheng, Y., Zheng, L., Yuan, B., Li, Z., Ryabinin, M., Chen, B., Liang, P., Ré, C., Stoica, I., and Zhang, C. FlexGen: High-throughput generative inference of large language models with a single GPU. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, volume 202 of *Proceedings of Machine Learning Research*, pp. 31094–31116, 2023. arXiv:2303.06865.
- Song, X., Zhong, Z., et al. ProMoE: Fast MoE-based LLM serving using proactive caching. *arXiv preprint arXiv:2410.22134*, 2024a.
- Song, Y., Mi, Z., Xie, H., and Chen, H. PowerInfer: Fast large language model serving with a consumer-grade GPU. In *Proceedings of the ACM Symposium on Operating Systems Principles (SOSP)*, 2024b.
- Song, Z., Berger, D. S., Li, K., and Lloyd, W. Learning relaxed Belady for content distribution network caching. In *17th USENIX Symposium on Networked Systems Design and Implementation (NSDI)*, 2020.
- Tang, P., Liu, J., et al. HOBbit: A mixed precision expert offloading system for fast MoE inference. *arXiv preprint arXiv:2411.01433*, 2024.
- Ukarande, A., Shekhar, D., Blackstein, M., and Rangan, R. Efficient, VRAM-constrained xLM inference on clients. In *Proceedings of Machine Learning and Systems (MLSys)*, volume 8, 2026. Industry track. arXiv:2604.26334. Artifact: <https://github.com/deepshnv/pipeshard-mlsys26-ae>, tag v2.0.3-mlsys26.
- Wang, W., Liu, J., Hou, X., Xia, X., Tang, P., Zhang, M., Li, C., and Guo, M. MoE-SpeQ: Speculative quantized decoding with proactive expert prefetching and offloading for mixture-of-experts. *arXiv preprint arXiv:2511.14102*, 2025.
- Wang, W., Hou, Y., et al. Achieving cloud-grade SLOs for local mixture-of-experts inference through CPU-GPU hybrid design. *arXiv preprint arXiv:2606.10493*, 2026a.
- Wang, Z., Wang, Y., et al. SeqMoE: Toward full-load performance via predictive and graph-compatible MoE offloading. *arXiv preprint arXiv:2609.12978*, 2026b.

- Williams, S., Waterman, A., and Patterson, D. Roofline: An insightful visual performance model for multicore architectures. *Communications of the ACM*, 52(4):65–76, 2009.
- Xue, L., Fu, Y., et al. MoE-Infinity: Efficient MoE inference on personal machines with sparsity-aware expert cache. *arXiv preprint arXiv:2401.14361*, 2024.
- Yadav, S. and Asgari, B. Mira: Memory-efficient MoE inference using adaptive caching and predictive expert staging. *arXiv preprint arXiv:2609.38090*, 2026.
- Yang, A. et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- Yang, S., Fan, X., Pan, M., Xi, H., Wang, Z., Sun, S., Keutzer, K., Han, S., Zaharia, M., Xu, C., and Stoica, I. FreeToken: Efficient edge-native MoE serving with bandwidth-adaptive execution. *arXiv preprint arXiv:2608.16157*, 2026.
- Yi, R., Guo, L., Wei, S., Zhou, A., Wang, S., and Xu, M. EdgeMoE: Empowering sparse large language models on mobile devices. *arXiv preprint arXiv:2308.14352*, 2023.
- Yuan, Y., Ma, L., et al. MoE-Lens: Towards the hardware limit of high-throughput MoE LLM serving under resource constraints. *arXiv preprint arXiv:2504.09345*, 2025.
- Zhang, Gao, and Mitra. SAEM: Stage-aware expert management for memory-efficient MoE inference in chain-of-thought reasoning, 2026a. *arXiv:2608.21614*.
- Zhang, H. Budgeting bytes: A windowed storage roofline and dual-budget architecture ablations for storage-bound LLM decoding. *arXiv preprint arXiv:2609.04238*, 2026a.
- Zhang, J., Wu, L., Darbari, M., and Hong, L. WiSP: A working-set view of mixture-of-experts serving on extremely low-resource hardware. *arXiv preprint arXiv:2606.21868*, 2026b.
- Zhang, L., Karimi, R., Ahmad, I., and Vigfusson, Y. Optimal data placement for heterogeneous cache, memory, and storage systems. *Proceedings of the ACM on Measurement and Analysis of Computing Systems (POMACS)*, 4(1):6:1–6:27, 2020. doi: 10.1145/3379472.
- Zhang, Y. Reproducible evaluation of MoE expert caching: Replay semantics, workload contamination, and operating regimes. *arXiv preprint arXiv:2608.07911*, 2026b.
- Zhang, Y., Aggarwal, S., and Mitra, T. DAOP: Data-aware offloading and predictive pre-calculation for efficient MoE inference. In *Design, Automation & Test in Europe Conference (DATE)*, pp. 1–7, 2025. doi: 10.23919/DATE64628.2025.10992741. *arXiv:2501.10375*.
- Zhong, S., Liang, L., Wang, Y., Wang, R., Huang, R., and Li, M. AdapMoE: Adaptive sensitivity-based expert gating and management for efficient MoE inference. In *Proceedings of the IEEE/ACM International Conference on Computer-Aided Design (ICCAD)*, 2024.
- Zhong, S. et al. HybriMoE: Hybrid CPU-GPU scheduling and cache management for efficient MoE inference. In *Design Automation Conference (DAC)*, 2025. *arXiv:2504.05897*.
- Zhu, S., Bohl, S., et al. Pre-attention expert prediction and prefetching for mixture-of-experts large language models. *arXiv preprint arXiv:2511.10676*, 2025.
- Zhu, Z., Li, G., Wang, P., Mo, Z., Pei, M., Song, Z., Liang, X., and Cheng, J. DALI: A workload-aware offloading framework for efficient MoE inference on local PCs. *arXiv preprint arXiv:2602.03495*, 2026.

A SCIENTIFIC BENCHMARKING CHECKLIST

How this paper follows the twelve rules of [Hoeffler & Belli \(2015\)](#).

Rule	Where and how
1. Report the base case of a speed-up and its absolute performance.	Every ratio in Table 2 is next to both absolute speeds; llama.cpp is configured with per-layer expert offload at equal expert memory, the setting its users run.
2. Explain subsets of benchmarks or resources.	Two models chosen for size and precision (MXFP4, BF16), a third that is FreeToken’s own headline model; budgets span what fits a 32 GB card; the 43.75% budget is omitted where it does not fit. All 30 AIME-25 problems are used.
3. Arithmetic mean for costs, harmonic mean for rates.	The accounting, the panel and the value of foresight average time per token over problems (a cost) and quote speeds as the harmonic mean this implies; in Table 2 every problem decodes the same number of tokens, and ratios of total time agree with the reported ratios of mean rates within 0.003.
4. Avoid summarising ratios.	Ratios of means are never averaged across budgets or models; ranges list the per-budget values, and the medians reported across models (Table 12) or across published rows (Appendix L) are medians of per-cell quantities, labelled as such.
5. Report whether measurements are deterministic; give confidence intervals.	Decode is greedy and deterministic per engine; timing is not. Every comparison has a 95% interval over problems; the panel (job 099) adds intervals over hosts, the unit that varies most.
6. Do not assume normality.	Intervals are bootstrap percentile intervals (10,000 resamples, paired by problem); no distributional assumption.
7. Compare nondeterministic data soundly.	Paired bootstrap of the ratio of means within a host; across the panel’s hosts, a two-stage bootstrap (hosts, then problems) (Kalibera & Jones, 2013); the variant a system is scored with is picked on a separate launch (except the Qwen3 25% row of Table 2).
8. Check that central tendency is the right summary.	Mean speed over 256 tokens is FreeToken’s own metric; per-token p50 and p99 latencies are recorded for every problem in the released rows.
9. Document all factors and the setup.	Machines, probes, power limits, memory clocks (Section 2), software commits (llama.cpp 4da6337, FreeToken 0d652e7), flags and job scripts are released; Appendix M.
10. Report measurement, synchronisation and summarisation.	One client times every system of Table 2 from streamed token arrivals; the oracle comparisons of Section 4 are timed in process by llama-ec-bench over teacher-forced text, every policy by the same timer (Section 2).
11. Show upper performance bounds.	The bound (Section 3) is a column of Table 2; Table 20 tightens it four ways; Section 3 applies it to the published literature and to our cache with the same procedure.
12. Plot enough to interpret; connect points only for trends.	Figure 2 plots every host’s ratio with its interval; Fig. 4 connects measured windows for the trend only and marks engine measurements separately; Figs. 7 and 9 are bar charts and Figs. 5, 6 and 10 scatter plots.

B PREDICTION LOG

Each job’s predictions are in the header of its script on the public `gpu` branch, committed and pushed before the machine started; git history orders every prediction before its measurement. Outcome notes with the statistics are in `prereg/`. The scorecard, generated from `prereg/scorecard.clauses.json`, splits every prediction of jobs 073–098 into its separable clauses and scores each under one rule (jobs before 073, which built and debugged the engine, also registered predictions in their headers; their outcome notes are in `prereg/`, and they are not scored here; job 072’s, for one, failed or did not apply on two of its three): *held* only when the point estimate is on the predicted side and the 95% paired-bootstrap interval excludes the threshold (for a band, lies inside it); *held (point)* when only the point estimate is, or no interval exists

(except in the hand scoring of jobs 073–098, which holds deterministic counts such as reads outright, since they have no sampling error); *failed* when the point estimate is on the wrong side; *untested* when the clause could not be evaluated; *void* when the job was voided. Of 565 clauses, 218 held, 193 held on the point estimate, 142 failed, 8 were untested and 4 void: 7 of 19 scored clauses held in jobs 073–075, when the system was being built, 19 of 23 in jobs 076–081, 11 of them sign-only predictions made after a same-host pilot, 24 of 38 in jobs 082–087, and 168 of 473 in jobs 088 onward, where 42 of 82 sign clauses held (75 on the point estimate) and 90 of 301 bands failed. Most bands that failed are the sizes of effects whose direction held: our slow-link lead landed at $1.51\text{--}1.58\times$ against a predicted $1.10\text{--}1.50$; the rerun of Table 2 on a second host (job 089) gave ratios $0.06\text{--}0.11$ below host B’s at five of six cells against a predicted ± 0.06 , host-bound speeds $14\text{--}17\%$ lower against a predicted $\pm 12\%$, and LRU $6\text{--}8\%$ slower than decayed frequency at the two GPU-heavy cells against a predicted $10\text{--}30\%$; on the 4090 and 3090 hosts (jobs 091, 092; Appendix D) our lead over FreeToken and its fraction of the limit overshot their bands, FreeToken’s carried-over backend fell far below its band, and two launch-order checks exceeded 3% (7% and, at gpt-oss 11% on the 3090 host, 11%); in the first oracle job (093; Section 4) the oracle’s gain fell outside $20\text{--}80\%$ at the two lowest budgets (it lost), its speed fell just below the predicted $45\text{--}65\%$ of the limit at all four host-bound cells ($43.7\text{--}44.8\%$), its hit rate and admissions missed the optimum’s by more than the predicted 3 points and 30% at three and two of the four, no-overlap cost 0% against a predicted $5\text{--}25\%$, and all-CPU cost 15% on Qwen3 against a predicted $0\text{--}10\%$; in the bytes-oracle job (094) the bands missed as Appendix G describes; in the single-read job (095) 44 of 56 clauses held, and the 12 that failed are of size: the model on the fetch oracle’s counters was predicted within 10% and missed by $13\text{--}17\%$ at four cells, the fetch oracle’s reads were predicted within 10% of the simulation’s and came 17% below at the two GPU-bound cells (the registered simulator left a resident carried over from the previous sequence with no next use, and so evicted it first; corrected, the simulation is within 4% of the engine’s reads at every cell, Appendix G), the prefetching oracle was predicted at $55\text{--}85\%$ of the limit at the host-bound cells and reached 44 and 51% at the two 25% cells, the model’s over-prediction of it was predicted at $15\text{--}50\%$ and came to 7 and 13% at the two lowest budgets, and the paced variant, predicted within 10% of the unpaced one, was $12\text{--}13\%$ faster at the two 25% cells; in the confirmatory factorial (096) the clauses held at most cells of both hosts and the failures are listed in Appendix G, the largest being the difference between the hosts; and in the FreeToken tuning job (098) its best setting stayed within the predicted 10% of its default and our ratios within the predicted bands, while the fixed fetch caps, the thread count and the CPU-only comparison missed as `prereg/freetoken_outcome_098.md` describes. Among the 23 threshold failures are ours at $1.56\text{--}1.78\times$ llama.cpp on the slow hosts against a predicted 1.8 , the Mixtral tie against a predicted $1.3\times$, the hit-optimal oracle’s gain at the two GPU-bound cells (48 and 56%) against a predicted ceiling of 15% , the bytes-optimal oracle short of beating the hit-optimal one by 15% at the two lowest budgets, and its 16-token window short of keeping 70% of its gain at three cells, fetch’s gain differing between hosts O4 and O5 by more than 0.15 at five cells, Belady’s single-read prefetch above $1.05\times$ the online policy at Qwen3 12.5% on O4, and FreeToken’s CPU backend running close to stock llama.cpp’s where we predicted at least $2\times$ and $1.2\times$; the 3 equality failures are two of the 4090 host’s FETCH tables, which copy one more expert than predicted, and FreeToken’s best backend at gpt-oss 25% in job 098 (offload, not hybrid); the 7 sign failures are LRU behind FreeToken at gpt-oss 11% , where we predicted it would still lead, and the hit-optimal oracle slower than the online policy at gpt-oss 11% and Qwen3 12.5% . Against the running log kept during the study, two clauses are downgraded (a llama.cpp clause of job 074b that the log counted as held was never run; job 085b’s own restatement of its fourth prediction, “at least 2 of 3”, failed at 1 of 3, though the pooled 4 of 6 held), job 073’s “identical outputs” is scored on fully identical outputs ($0\text{--}3\%$) rather than on agreement of the first 160 characters ($27\text{--}47\%$), and ten clauses the log counted as held, among them job 087’s first, hold on the point estimate only under the interval rule; `prereg/scorecard_outcome.md` lists every change. Job 085’s Qwen3 half ran as 085b on a second machine with a slow link after a hung conversion. Job 084 ran on a throttled card and was rerun, with its predictions unchanged, as 084b/084c. The batched-verification variants of Section 6 (rejected drafts free, routed like the true token, or routed at random) are in `prereg/batchchk_outcome.md`.

The clauses were scored by hand from the outcome notes. A script (`scripts/scorecard_auto.py`) re-scores every clause whose measured value is a number and whose threshold parses, 469 of 565: it agrees with the hand score on 432; of the 37 differences, 36 are the deterministic counts (reads, divergences, table entries) that the hand scoring holds outright and the script scores held (point), and 1 is a boundary value the hand scoring read unrounded (1.249 against a band starting at 1.25). Of the clauses the script scores held on the point estimate, 213 have no interval and 7 have one that crosses the threshold. Job 099’s clauses are scored by machine from the start (`scripts/panel_099.py`): of 273, 190 held, 42 held on the point estimate, 39 failed and 2 were untested (`prereg/scorecard_099.json`); the two untested clauses are one host’s replay check, whose step failed when a Python package did not install. The failures concentrate on the two hosts whose link reads at less than half their CPU’s rate, where the bands, set from four earlier hosts none of which was link-starved, did not encode host dependence; and on the planning time of the noisy whole-future window (`prereg/panel_outcome_099.md`). Job

100 (`scripts/job100.py`): of 216, 41 held, 122 held on the point estimate, 27 failed and 26 were untested. Most failures are one prediction, that the deployed-path windows miss within 10% as often as a replay that admits at once (their copies land late); the rest are the probe-only model on the half-width-link host and four band or ordering edges. The untested clauses are the planning time of the deployed-path windows, which the engine does not time, and the paired prediction for a host that never reached a timed step (`prereg/window_outcome_100.md`). Job 101 (`scripts/job101.py`; the lowest-ratio panel machine relaunched, and a second Ryzen 7 9800X3D behind a slower link): of 25, 6 held, 19 held on the point estimate, 0 failed and 0 were untested (`prereg/crossover_outcome_101.md`). Job 102 (`scripts/job102.py`; the read-schedule microbenchmark on three hosts): of 33, 0 held, 28 held on the point estimate (none has an interval), 5 failed and 0 were untested. Four failures are one prediction, that waiting per layer would cost at least 0.03 of the host term against the analytic schedule; on two of the three hosts it costs nothing measurable. The fifth is the ordering of the two budgets on the Intel host (`prereg/readsched_outcome_102.md`). Rentals that produced no results were destroyed and are in the rental ledger (`gpu/vast_ledger.json`): the first two for 099a, replaced before any measurement; 100d, which stalled downloading the model; 100e, which never reached its job; and 103c, which never started its job. Job 103 (`scripts/job103.py`; MIN’s fewest-admission schedule against the greedy one): the hosts could not install the plan’s solver beside Debian’s numpy, so the engine ran the greedy schedule in the plan configurations, as it logs; of 45 clauses 4 held on the point estimate (the relaunch of Pf) and 41 are untested. Job 104 repeated it with the plan in its own environment on three of the same machines (Pf, a second Core Ultra 9 285K whose link reads at 0.49 of its CPU rate, and Pg, whose second launch job 103’s 5950X turned out to be): of 47, 1 held, 22 held on the point estimate and 24 failed. Most speed predictions held (the fewest-admission set gains on Pf, and gains at least 0.03 over the greedy set at gpt-oss 11% on every machine; with two reads it is at most 0.01 behind the greedy set at 5 of 6 cells, and 0.047 behind on Pg at gpt-oss 25%); the other failures are the counters at gpt-oss 25% (copies 0.73 of greedy’s against at most 0.70, misses 5% more against at most 3%), the engine’s misses against the host’s replay, the two-read reads, and the interaction, which we predicted would shrink with the plan’s set and which grew (`prereg/minadm_outcome_104.md`). Job 105 (`scripts/job105.py`; the plain form of Eq. (3) with G profiled on each host before its timed runs, the engine’s layer-ahead copy, and the fewest-admission set, on an EPYC 7402P, a Threadripper 9960X, Pf and O4): of 31, 7 held, 20 held on the point estimate and 4 failed. Every failure is the law at gpt-oss 25%, which over-predicts the deployed cache by 6.9–9.3% (against at most 6%) on 3 of the four hosts (3.8% on Pf) and so takes the pooled median error (4.2%) past the 4% registered; at 11% it is within -1.5 to 4.7% . Every speed prediction held: the layer-ahead copy gains at least 3% on the host with a ratio above 0.8 and loses on both below 0.4, and the fewest-admission set beats the deployed cache on every host, beats the greedy set by at least 0.03 on the three with a ratio under 0.7, and copies at most 0.65 of the greedy set’s experts (`prereg/sumlaw_outcome_105.md`). Two clauses predicted quantities already measured, the profiled G (job 069c) and the plan’s copies, which are a property of the trace (job 104); they test that the engine reproduces them. Three more launched hosts produced no results: 105c never started its job, and the job’s bandwidth gate stopped 105d and 105g, whose cards were throttled, before any timed run. Job 106 (`scripts/job106.py`; Eq. (3) with its overlap term on both models at two budgets, the two online admission rules, and MIN’s sets launch by launch, on five hosts): of 181, 21 held, 112 held on the point estimate and 48 failed. 36 of the failures are on the two server machines new to the job, one of which computed wrong outputs and the other of which varied by more than the registered 2% between rounds (Appendix J); the three stable hosts are Pf, the Threadripper 9960X of job 105 and panel host Pe, all rented before. On the 3 stable hosts the law held on every cell; the 11 failures there are the margin rule’s speed-up of at least 1.01 at gpt-oss on two hosts, the learned order’s speed-up of at least 1.00 on the fast-memory host, its host-time caps on the fastest-link host, and the law’s prediction of both rules’ ratios on Pf at 25% (off by 0.030 and 0.045 against at most 0.03). The pooled clauses: the law’s median error over all hosts failed (1), and the margin rule’s pooled median of at least 1.02 held only because of the two unstable machines (on the stable ones it is 1.01) (`prereg/onlineadmit_outcome_106.md`).

The clause-by-clause tables (565 rows for jobs 073–098, 273 for job 099, 216 for job 100, 25 for job 101, 33 for job 102, 45 and 47 for jobs 103 and 104, 31 for job 105, 181 for job 106) are in the supplement (`paper/supplement.pdf`), generated by the same scripts.

C THE CALIBRATED TIME MODEL

Our cache’s one per-machine decision, how many of a layer’s misses to fetch rather than run on the CPU, comes from a calibrated model of time per token. With X_c the bytes the CPU reads per token and X_p the bytes the link reads,

$$T = G + \max\left(\frac{X_c}{B_c}, \frac{X_p}{B_p}, \frac{X_c + X_p}{B_{cp}}\right), \quad (4)$$

with one constant G per engine and model, fitted once on one host and frozen. It is a calibrated model, not a law: G is fitted, and it exceeds the whole all-in-VRAM token (5.17 against 3.92 ms) because it absorbs per-token host work and the CPU helpers' shortfall against the probe (below). As a predictor of time it ties a naive one-path model with the same G (time = G + all host bytes over the best probed rate):

Held-out set	Eq. (4)		one path	
	median	90th pct.	median	90th pct.
blind test, 33 rows on 7 hosts	3.4%	12%	3.6%	31%
frozen, 30 rows on 4 later hosts	5.6%	16%	6.1%	11%

It over-predicts Qwen3's speed on the later hosts by +5 to +21%, a budget-dependent bias whose G was borrowed from llama.cpp. Its value is the per-machine fetch table it picks (Eq. (5) below), which gains 3.3–8.0% over a fixed table on host B and 9–34% on hosts with half the link (Appendix K).

The blind test in full. On every machine of the blind test the job wrote the model's prediction from its own probe before any model ran (Table 5; all of its cache rows are gpt-oss-120b on RTX 5090 hosts). Over 29 configurations on 7 machines (33 measurements; one machine ran twice) the median error is 3.7% per configuration; without the 12-channel EPYC 9655 it is 3.3% over 24 configurations (90th percentile 7.5%). llama.cpp --n-cpu-moe is predicted within 4% on the Ryzen hosts, including two Qwen3 configurations on host B (+3.6%, +4.0%; for Qwen3, $G=4.3$ ms, the mean residual of three llama.cpp runs on another host), and misses by +10–+119% on the hybrid-core and server hosts. With its constants frozen, the model scores the cache on the 4 RTX 5090 hosts rented later (30 measurements at the six cells of Table 2, not written before the runs) at a median 5.6%, 90th percentile 16%, worst 21% (Qwen3-30B-A3B 43.75% on oracle host O1). Refit under leave-one-host-out validation, the token-level max has a median error of 3.1% (90th percentile 9.4%), the additive form $G + X_c/B_c + X_p/B_p$ 4.5% (11.3%) and a per-layer max summed over layers 3.2% (10.2%); the max beats the sum on the three rows that copy misses on demand behind a 53 GB/s link, where both paths carry most of the bytes, and on the EPYC 9655 rows, where both fail, the sum is closer on the two rows that copy behind a slow link, and elsewhere the two are within a few points (Appendix C). Each failure has a measured cause (Appendix C): on-demand copies behind a slow link share DRAM less evenly than the probe's combined rate says; llama.cpp splits every matrix evenly across threads and runs at 46% of the model on a hybrid-core host; and on a 12-channel EPYC host memory stops being the limit and a per-layer latency floor takes over. G (5.17 ms refit) exceeds the whole all-in-VRAM token (3.92 ms) because it absorbs a shortfall the probe does not see: the helpers read at $1.09\text{--}1.31\times$ the probe's time per expert, and a fitted rate factor is collinear with G at this sample size.

What the model says, and the split it picks. Equation (4) is a calibrated model of execution that reads host memory in sequence with its GPU work, as our cache and llama.cpp do. For such an engine every mechanism acts only through the bytes read from host memory per token and the number of read paths in use: a second path adds bandwidth on every host of the blind test but the 12-channel EPYC (B_{cp} against $\max(B_c, B_p)$): 8–17% on its Ryzen hosts, 1% on the Core Ultra, 33% on the Zen 2 EPYC; the RTX 3090 host's Core i9-11900KF, where the two paths together read less than the CPU alone, is a second exception), and that is the whole gain available to fetching or splitting misses between the CPU and PCIe. Prefetching never reduces the bytes; what it can buy is overlap of the reads with the step, which the model does not price: charging every read in sequence, it over-predicts the time per token of MIN prefetched by +1 to +38% (Section 4) and of a Belady prefetch by up to 53% (Appendix G). The model's job in the system is the split. Its per-layer form,

$$T(n, f) = \max\left(\frac{fS}{B_p} + g, [n > f](L_c + \frac{(n-f)S}{B_c}), [0 < f < n] \frac{nS}{B_{cp}}\right), \quad (5)$$

with g the GPU's expert work per layer (48 μ s for Qwen3, 37 μ s for gpt-oss, from profiles) and $L_c=20$ μ s the hand-off, picks for each count n of missed experts the number f to copy with the smallest T . On host B the model's table beats the fixed table we used before at every budget of both models, by 3.3–8.0%; on hosts whose link we measured at 27–30 GB/s instead of 53 it fetches less and gains 21–34% on gpt-oss and 9–18% on Qwen3, the per-machine split mattering most where the link is slow (Appendix K). In the running example on host B the table is 0, 0, 1, 1, 2 ($f(0..4)$): one missed expert runs on the CPU, hidden behind the GPU's other experts; from two misses the model copies one. The lead over FreeToken at that cell, $1.27\times$, is mostly the admission policy (77% of it by the LRU ablation of Appendix K), the rest the CPU path that FreeToken's bandwidth model declines to use there.

Table 5. The calibrated model predicted blind: each machine wrote its prediction from its own bandwidth probe, with constants frozen in the public repository, before any model run. Error of predicted over measured speed, gpt-oss-120b (cache rows) and llama.cpp --n-cpu-moe (last column, %).

Host CPU	Job	n	Median err %	Max err %	llama.cpp err %
Core Ultra 7 270K (8P+16E)	069c	5	1.5	11.9	+119
EPYC 7352 (Zen 2)	069c	5	3.2	11.7	+10
EPYC 9655 (12-ch. DDR5)	069c	5	18.7	31.9	+16
Ryzen 9 7900	072	5	4.9	8.2	+4
Ryzen 9 9950X3D (A)	073	8	1.6	4.5	—
Ryzen 9 7950X	074	1	4.0	4.0	—
Ryzen 9 9950X	076	4	2.5	5.0	—
Ryzen 9 9950X3D (B)	082	0	—	—	+4, +4

Is the max the right form? In 19 of the 29 configurations nothing is copied, so they test one roof. We therefore refit the model and its alternatives on the same 33 measurements under leave-one-host-out validation (each of the 7 hosts predicted with constants fitted on the other six): the token-level max of Eq. (4) with G refit has a median error of 3.1% (90th percentile 9.4%); the additive form $G + X_c/B_c + X_p/B_p$, 4.5% (11.3%); a per-layer max summed over layers from the engine’s per-layer miss histograms, 3.2% (10.2%). The additive form ties in sample but misses the three rows that copy misses on demand behind a 53 GB/s link, where both paths carry most of the bytes, by -11%, while the max is within $\pm 0.2\%$ there; on the slow-link copy rows the additive form is the closer of the two. The per-layer form neither beats the max nor explains its failures: the Jensen gap between a sum of per-layer maxima and the maximum of sums is a fifth of the Core Ultra’s residual, none of the EPYC 7352’s, and has the wrong sign on the EPYC 9655 and on the prefetch rows. We keep Eq. (4), with its evidence stated: the max beats the sum on the three fast-link copy rows and on the EPYC 9655 rows, where both fail, the sum is closer on the two slow-link copy rows, and elsewhere the two are within a few points, so the form is settled only where both paths carry most of the bytes. FreeToken’s own runs are not predicted, because its counters are not exposed; the model covers our engine and llama.cpp.

Where it fails, and why. Each failure now has a measured cause. (i) Copying misses on demand (FETCH) at 25% is over-predicted by +11.7% on the Core Ultra 7 and +11.4% on the EPYC 7352 (links of 36 and 26 GB/s); +8.8% and +8.5% with G refit. The engine’s per-layer timers show why: a CPU miss served beside a copy takes 20–62% longer than one served alone, so the two paths share DRAM less evenly than the probe’s combined rate says; a variant of the model that uses the probe’s per-path rates under contention and a per-layer hand-off constant predicts both rows within +3.1% and +4.1%, at a worse median (4.2%). (ii) llama.cpp on a Core Ultra 7 with 8 performance and 16 efficiency cores runs at 46% of the model’s prediction at its default thread count, because it splits every matrix evenly across its threads; our helpers claim work dynamically and stay on the model on the same machine. We did not test llama.cpp restricted to the performance cores. (iii) On a 12-channel EPYC 9655 (559 GB/s) host memory stops being the limit (+16 to +29% with G refit under leave-one-host-out validation; Table 5 gives the blind errors): the per-layer timers show a floor of 55–70 μ s per layer with CPU misses that no desktop host has (their intercepts are at or below zero), the latency of the hand-off at a bandwidth where a 13 MB expert takes 24 μ s. (iv) Prefetch configurations are under-predicted by -8.7 to -3.4%: their copies overlap attention, which the model charges in full.




Figure 5. Ours (filled) and FreeToken (hollow) as a percentage of each machine’s own bound (exact optimum, the machine’s highest probed host rate, the card’s datasheet rate) at every cell of Table 6: two RTX 5090 hosts, an RTX 4090 and an RTX 3090 (FreeToken’s 3090 cells, its carried-over backend running below llama.cpp, are omitted). The limit ranges from 41 to 519 tok/s across these cells; ours stays at 25–46% of it.




Figure 6. Ours over FreeToken’s better backend (carried over from host B on host S) in every same-machine comparison of the current engine on the two models of Table 2 (eight hosts), against the machine’s CPU-to-PCIe read ratio (FreeToken’s own probe). Filled: FETCH table from the model on that machine; hollow: the fixed table. The rightmost gpt-oss points are a Core i9-14900K with FreeToken’s executor on its performance cores; the rightmost Qwen3 points are a Ryzen 9 7900.

D THE GRID: OTHER CARDS, AND THE SCOPE OF THE GAIN

Table 6 lists every cell of Table 2’s protocol on the four machines we measured, each against its own limit, and Fig. 5 plots the fractions. The protocol was repeated on an RTX 4090 next to a Core i5-12400 (the host reads 37 GB/s, the link 25) and an RTX 3090 next to a Core i9-11900KF (45 GB/s; CPU and link together read *less* than the CPU alone, so the model’s table copies nothing), at the cells that fit 24 GB. On the 4090 ours runs $1.03\text{--}1.37\times$ FreeToken and $1.8\text{--}2.6\times$ llama.cpp at half host B’s absolute speeds, the host being the reason; on the 3090 it runs $1.6\text{--}2.2\times$ llama.cpp (FreeToken’s carried-over backend ran below llama.cpp there and is not counted). The cards’ own bandwidth never enters: every one of the seven cells on the 24 GB cards is host-bound at the limit, and ours stands at 30–46% and 27–42% of it. Mixtral-8x7B (8 experts per layer, top-2; 26 GB of experts on the 3090) marks the scope: with 2 of 8 experts resident the cache hits 31% of reads where whole-layer pinning hits 25%, so there is little locality to earn, and ours ties llama.cpp (1.015 [1.008, 1.024]); on the 128-expert models the same cache hits 53–79% at 11–25% resident. Across eight hosts the gpt-oss lead over FreeToken grows with the host’s CPU-to-link ratio, from $1.09\text{--}1.22\times$ where the CPU reads $1.05\text{--}1.14\times$ its link’s rate to $1.51\text{--}1.58\times$ behind a link of half the bandwidth, in the direction the model’s split implies (the model does not predict FreeToken’s runs), and the Qwen3 lead stays at $0.95\text{--}1.15\times$ (Fig. 6, Appendix K).

The 24 GB cards run the cells that fit (the VRAM arithmetic, written on the machine from the GGUF header, skipped Qwen3 at 43.75%: 28.0 GB of slots, dense weights and KV against a 23 GB limit). The predictions for jobs 091 and 092 and their outcomes are in `prereg/grid_outcome_091.md` and `092.md`: on the 4090 our lead over FreeToken at gpt-oss 11%

Table 6. The grid: Table 2’s protocol on every card and host we measured, each cell against its own machine’s limit (the exact optimum’s reads on the same trace, the machine’s highest probed host rate, the card’s datasheet rate). *Bound*: what binds the limit; *host* when the optimum’s reads alone take longer than the GPU’s whole read, so the GPU has slack, *both* when the limit runs experts on the CPU until the two paths take equally long, which is where the all-in-VRAM speed can be exceeded. Ours and FreeToken are launch 2 (FreeToken’s backend carried over from host B’s selection); the 40 and 43.75% budgets do not fit a 24 GB card. [†]On the RTX 3090 FreeToken’s carried-over hybrid backend ran below llama.cpp; its own calibration recommends its offload backend there, which we did not run, so its speed is reported and no ratio against it is printed. Mixtral-8x7B (26 GB of experts, top-2 of 8) is ours against llama.cpp only, on the RTX 3090, launch 1.

Model	Experts on GPU	llama.cpp (tok/s)	FreeToken (tok/s)	Ours (tok/s)	Ours ÷ FreeToken (95% CI)	Ours ÷ llama.cpp	Limit (tok/s)	Bound	Ours, % of limit
RTX 5090 , 9950X3D (host B): B_{host} 88 GB/s, B_{gpu} 1792 GB/s; tables 0, 0, 1, 1, 2 / 0, 0, 1, 1, 2, 3, 3, 4, 5									
gpt-oss-120b	11%	34.9	54.0	69.9	1.29 [1.28, 1.31]	2.00	172	host	41
gpt-oss-120b	25%	40.0	85.6	109.2	1.27 [1.25, 1.29]	2.73	430	host	25
gpt-oss-120b	40%	47.7	132.2	152.5	1.15 [1.13, 1.17]	3.20	519	both	29
Qwen3-30B-A3B	12.5%	19.5	38.7	40.0	1.03 [1.02, 1.04]	2.05	96	host	42
Qwen3-30B-A3B	25%	22.1	54.7	63.1	1.15 [1.13, 1.17]	2.86	214	host	30
Qwen3-30B-A3B	43.75%	28.4	103.1	108.2	1.05 [1.04, 1.06]	3.81	308	both	35
RTX 5090 , 9950X (host S): B_{host} 71 GB/s, B_{gpu} 1792 GB/s; tables 0, 0, 1, 2, 3 / 0, 0, 1, 2, 2, 3, 4, 5, 6									
gpt-oss-120b	11%	28.9	48.0	57.9	1.21 [1.19, 1.22]	2.00	140	host	41
gpt-oss-120b	25%	33.6	78.7	94.1	1.20 [1.16, 1.22]	2.80	351	host	27
gpt-oss-120b	40%	40.1	122.2	133.6	1.09 [1.08, 1.11]	3.34	514	both	26
Qwen3-30B-A3B	12.5%	16.1	32.7	33.5	1.03 [1.02, 1.04]	2.08	78	host	43
Qwen3-30B-A3B	25%	18.5	51.4	53.9	1.05 [1.04, 1.06]	2.91	174	host	31
Qwen3-30B-A3B	43.75%	23.8	98.3	95.8	0.97 [0.96, 0.99]	4.02	305	both	31
RTX 4090 , i5-12400 (job 091): B_{host} 37 GB/s, B_{gpu} 1008 GB/s; tables 0, 0, 1, 2, 2 / 0, 0, 1, 1, 2, 3, 4, 4, 5									
gpt-oss-120b	11%	18.1	23.8	32.7	1.37 [1.35, 1.39]	1.81	74	host	44
gpt-oss-120b	25%	21.0	42.5	55.0	1.30 [1.26, 1.33]	2.62	184	host	30
Qwen3-30B-A3B	12.5%	10.5	18.2	18.8	1.03 [1.02, 1.05]	1.78	41	host	46
Qwen3-30B-A3B	25%	12.3	30.0	31.3	1.05 [1.03, 1.06]	2.54	91	host	34
RTX 3090 , i9-11900KF (job 092): B_{host} 45 GB/s, B_{gpu} 936 GB/s; tables 0, 0, 0, 0, 0 / 0, 0, 0, 0, 0, 0, 0, 0, 0									
gpt-oss-120b	11%	22.8	8.4 [†]	37.3	—	1.64	88	host	42
Qwen3-30B-A3B	12.5%	12.3	9.0 [†]	19.2	—	1.56	49	host	39
Qwen3-30B-A3B	25%	13.8	11.2 [†]	30.1	—	2.18	110	host	27
Mixtral-8x7B Q4_K_M	$C=2$ (25%)	8.22	—	8.34	—	1.015 [1.008, 1.024]	—	—	—

(37%) exceeded the predicted 10–35% band, ours ran $1.78 \times$ llama.cpp at Qwen3 12.5% against a predicted 1.8, the model’s tables copied one more expert than host B’s at one entry each (the prediction reasoned from the link alone; the CPU is slower in the same proportion), and ours’ fraction of the limit at Qwen3 12.5% (45.7%) exceeded the predicted 25–45%; on the 3090 FreeToken’s carried-over backend ran below llama.cpp, ours ran 1.56 – $1.64 \times$ llama.cpp at two cells against a predicted 1.8, Mixtral tied at $C=2$ against a predicted $1.3 \times$ and its $C=4$ cell was lost to the job’s deadline, and one launch-order check exceeded 3% (7%).

E THE ORDER-FREE ACCOUNTING OF THE GAP

Two choices separate the deployed cache from MIN with one read: *what to cache* (the deployed policy’s admissions or MIN’s) and *how to load* (by the CPU, two reads, or in the step, one). The engine ran all four combinations on O4, O5, the panel and the three new machines of jobs 100–101. In the running example, changing only what to cache closes 4.5% of the 10.6 ms gap, changing only how to load 3.9%, and both 51.5% (rounded): the choices interact, because MIN caches 4.6 times as many experts per token as the deployed policy and, loaded by the CPU, each is read twice. We credit each choice with its effect averaged over the two orders, its Shapley value (Fawcett & Hoos, 2016; Jain, 1991) (26.1% and 25.5% here); Table 7 also lists each alone and the interaction. Copying ahead is nested under MIN’s set (12.7%); what no oracle touches is the rest (35.7%). On the 13 hosts whose link reads at least half as fast as the CPU, either change alone closes at most 26% of the gap and both together 14–60%; on the 2 below that line, loading in the step costs time at 3 of 4 host-budgets, and at 3 of 4 on 2 more (job 103). These use the greedy set. With the fewest-admission set (job 104), both together close 6–49% at 11% on its three machines, against –26–39% with the greedy set, and the interaction remains (1–33%).

Table 7. An order-free accounting of the gap between the deployed cache and the bound at the host-bound budgets, as percentages of the gap (ms per token). The engine ran both choices of what to cache (the deployed policy’s admissions or MIN’s) with both ways to load a cached expert (two reads: the CPU serves it, then a copy; one read: a copy in the step). *Load alone* and *cache alone* are each choice’s effect with the other at the deployed level; the *interaction* is how much more MIN’s set is worth loaded once than loaded twice. The Shapley values average each choice’s effect over the other’s two levels and add up, with the prefetch (nested under MIN’s set, since it needs foresight) and the rest, to 100%. Panel: mean over its hosts (number in parentheses) and, in small type, the range over its hosts (not an interval); a fast link reads at least half as fast as the host’s CPU, a slow one less. Last rows: the new machines of the follow-up launches, with their link-to-CPU ratio.

Host	Budget	Gap (ms)	Load alone	Cache alone	Interaction	Shapley value cache	load	Prefetch	Rest
O4	gpt-oss 11%	10.6	4	4	43	26	25	13	36
O4	gpt-oss 25%	8.5	4	22	17	31	13	22	34
O4	Qwen3 12.5%	18.4	3	−1	58	28	32	4	36
O4	Qwen3 25%	14.1	2	14	37	33	21	17	30
O5	gpt-oss 11%	11.4	0	−2	28	12	14	17	57
O5	gpt-oss 25%	9.1	−4	20	10	25	1	21	52
O5	Qwen3 12.5%	19.8	1	−13	43	9	23	7	62
O5	Qwen3 25%	15.3	0	14	17	22	9	16	53
Panel, fast link (8)	gpt-oss 11%	11.6	1 (−2 to 4)	3 (−2 to 6)	38 (30 to 47)	22 (14 to 26)	20 (14 to 27)	13 (−4 to 20)	45 (38 to 55)
Panel, fast link (8)	gpt-oss 25%	9.0	0 (−4 to 3)	22 (18 to 26)	15 (10 to 19)	29 (25 to 35)	8 (1 to 12)	20 (9 to 23)	43 (38 to 49)
Panel, slow link (2)	gpt-oss 11%	9.7	−6 (−9 to −3)	−4 (−9 to 0)	3 (−6 to 11)	−3 (−11 to 6)	−5 (−12 to 2)	17 (9 to 24)	91 (83 to 99)
Panel, slow link (2)	gpt-oss 25%	7.8	−17 (−22 to −11)	19 (16 to 23)	3 (3 to 3)	21 (17 to 25)	−15 (−21 to −10)	22 (17 to 26)	73 (68 to 78)
5700X3D (0.92)	gpt-oss 11%	14.9	4	−2	55	25	31	8	36
5700X3D (0.92)	gpt-oss 25%	11.5	4	22	24	34	17	18	31
9950X, x8 (0.54)	gpt-oss 11%	8.3	0	−3	17	5	8	28	59
9950X, x8 (0.54)	gpt-oss 25%	7.4	−15	13	17	22	−6	26	59
9800X3D, slower link (0.61)	gpt-oss 11%	10.8	−2	−4	29	11	12	20	57
9800X3D, slower link (0.61)	gpt-oss 25%	8.9	−5	20	12	26	1	20	54

F THE ACCOUNTING AS A MODEL

Between the measured time of a cell and its limit lie five modelled fixes: overlap of host reads with GPU work (in Eq. (1)’s form, which also rebalances experts onto the CPU), the optimum’s reads instead of the online policy’s (*foresight*), the best online policy’s reads at the probe’s highest rate instead of the engine’s through Eq. (4) (policy and read paths), the engine’s per-token host work, and GPU kernels at the datasheet rate. Adding them in a fixed order makes the first one added look largest, so we report each fix’s Shapley value, its average marginal cost over every order (Fig. 7), a method with a precedent in performance debugging of microservices (Li et al., 2024). Two of the five terms are definitions rather than measurements: the GPU term is the residual, its factor calibrated per cell so that the model with no fix applied reproduces the measured time, and “overlap” is Eq. (1)’s form. On host B the model has two regimes. In the four host-bound cells missing foresight is the largest cost, 37–53% of the gap (26–31% of the token; in the running example 37%, where it edges overlap by 0.14 ms and is the largest marginal cost in 70 of the 120 orders); in the two GPU-bound cells host reads serialised with GPU work is largest, 39–46% of the gap, and foresight second at 23%. Policy and read paths cost 12–16%, per-token host work 3–10%, and GPU kernels below datasheet, net of the overlap the engine achieves, 4–16%. The regimes are a property of the host: recomputed for the three oracle hosts, whose host memory is weaker, foresight edges overlap at three of their six GPU-bound cells. The model prices foresight and overlap as nearly independent players, and the interior states it values are not realised by any online policy. The rest of this section measures them.

G FORESIGHT IN THE ENGINE

Four jobs ran oracle policies in the engine, each on a rented RTX 5090 host with weaker host memory than host B’s. Jobs 093–097 report ratios of mean speeds rather than of mean times per token; recomputed either way, no ratio moves by more than 0.006. The fourth (096) is the confirmatory factorial of Fig. 2: ten configurations per cell, each in a fresh context, in an order shuffled per cell with a recorded seed, on hosts O4 (Ryzen 9 9950X: 43, 45 and 51 GB/s by the CPU, over the link and at best) and O5 (Ryzen 9 9950X3D: 47, 27 and 51 GB/s), with ten predictions per host committed before either started (prereg/foresight_outcome_096.md). They held at most cells. The failures: fetch’s gain above the predicted band at Qwen3 25% on O4 and different between the hosts by up to 0.26 against the 0.15 predicted; Belady’s single-read prefetch faster than predicted at Qwen3 12.5% on O4, with reads below the predicted 1.25–1.9× the MIN prefetch’s at Qwen3 12.5%




Figure 7. The accounting as a model: Shapley attribution of the gap between our cache’s measured time per token and the bound, for every cell of Table 2 on host B, averaged over all 120 orders in which the five fixes can be applied. Fixes: overlapping host reads with GPU work, which in Eq. (1)’s form also moves experts onto the CPU until the two paths balance; the optimum’s reads instead of the best online policy’s (foresight); the best online policy’s reads at the probe’s highest rate instead of the engine’s reads through Eq. (4) (policy and read paths); the engine’s per-token host work; and GPU kernels at the datasheet rate, the term that closes the model on the measured time. Whiskers: the range of foresight’s marginal cost over the orders.

on both hosts and at gpt-oss 11% on O4 (down to $1.09\times$); the bytes share at Qwen3 12.5% on O4 (60% of the gap against a predicted 25–55%); and the reads fetch-on-admit saves at gpt-oss 40% on O5, just above the predicted 22%. Reading the online policy’s admissions once came to $0.97\text{--}1.03\times$ at the host-bound cells, at the bottom of the predicted $0.97\text{--}1.12\times$ although its reads fell as predicted, and to $0.95\text{--}1.02\times$ at the GPU-bound cells. Table 8 lists every cell of the single-read factorial of Section 4 on host O3 (a Ryzen 9 9950X whose probe reads 53 GB/s by the CPU at the helper count, 43 over the link, 60 with both and 61.5 at best; the model’s FETCH tables 0, 0, 1, 2, 2 for gpt-oss and 0, 0, 1, 1, 2, 3, 4, 4, 5 for Qwen3), and Table 9 the milliseconds per token of the engine states the three jobs realised beside each host’s limit and the model’s values. Table 10 lists the first job (093), the hit-optimal oracle with a window sweep on host O1 (a Ryzen 9 9950X: 44, 46, 49 and 50.5 GB/s; tables 0, 0, 1, 2, 3 and 0, 0, 1, 2, 3, 4, 5, 6, 7), and Table 11 the second (094), MIN with bypass with background admissions beside the hit-optimal oracle on host O2 (a Ryzen 9 7950X: 44, 46, 50 and 50.6 GB/s; the same tables). On O1 and O2 every cell but gpt-oss 40% and Qwen3 43.75% is host-bound at the limit; on O3, whose memory is stronger, the same four.

The single-read oracles. In the fetch oracle the step’s plan is computed before the layer runs from the trace’s record of the current step: the misses MIN with bypass would admit are fetched by the engine’s FETCH operation into their victims’

Table 8. The single-read oracles (job 095; RTX 5090, Ryzen 9 9950X, host memory 61.5 GB/s at best) beside the hit-optimal one. Per cell: the online policy’s speed and its share of this host’s limit; for each oracle with the whole sequence in view, its ratio to the online policy (paired, 95% interval), hit rate, host reads per token (CPU misses + fetches + admissions) as a multiple of the optimum’s (each read counted once), share of the limit and the calibrated model on its own counters against its measured time. *fetch*: MIN with bypass, admitted misses fetched into their slot this step; *lead*: the scheduled single-read prefetch (first use at least two steps away); *both*: fetch and lead; *paced*: both on the paced copy path with a three-step lead; *hit-opt.*: the hit-optimal oracle (Belady prefetch without bypass, unpaced). The simulation columns give the trace’s prediction of reads and hits for fetch and both.

Cell	online		ratio	fetch			model	both			hit-opt.			lead			paced			simulated			
	tok/s	% lim.		hits	reads/R*	% lim.		ratio	hits	reads/R*	% lim.	ratio	hits	reads/R*	% lim.	ratio	% lim.	ratio	% lim.	fetch reads	hits	both reads	hits
gpt-oss 11%	55.5	46	1.29 [1.28, 1.29]	72	1.04	59	-8%	1.39 [1.39, 1.40]	86	1.23	64	0.91 [0.90, 0.92]	98	1.77	42	1.26	58	1.45	67	39	73	46	87
gpt-oss 25%	90.3	30	1.32 [1.31, 1.33]	89	1.08	39	-13%	1.49 [1.48, 1.50]	97	1.45	44	1.14 [1.12, 1.15]	99	2.11	34	1.45	43	1.66	50	16	89	22	99
gpt-oss 40%	128.8	25	1.25 [1.24, 1.27]	95	1.14	32	-14%	1.39 [1.37, 1.41]	99	1.60	35	1.27 [1.25, 1.28]	99	2.20	32	1.38	35	1.49	38	7	95	12	100
Qwen3 12.5%	32.3	48	1.34 [1.34, 1.35]	74	1.04	64	-9%	1.42 [1.41, 1.43]	84	1.20	68	0.69 [0.68, 0.71]	77	2.47	33	1.22	58	1.41	68	100	74	115	84
Qwen3 25%	51.8	34	1.38 [1.37, 1.39]	88	1.04	48	-15%	1.48 [1.46, 1.49]	96	1.42	51	1.06 [1.05, 1.08]	99	2.00	37	1.41	48	1.67	57	45	88	61	97
Qwen3 43.75%	91.4	30	1.26 [1.25, 1.28]	96	1.08	38	-17%	1.41 [1.39, 1.42]	99	1.62	42	1.27 [1.26, 1.29]	100	2.29	38	1.40	42	1.50	45	15	96	23	100




Figure 8. The first oracle job (093, host O1): speed as a percentage of this host’s limit against the window W of future routing the hit-optimal oracle sees, at the four host-bound cells. “none” is the online policy; “all” the whole sequence; points are connected for the trend only. The tick at the right is the model’s foresight-only term for the cell, the speed it predicts with that one fix; the oracle’s host traffic sits at 89–94% of the probe’s best rate at every window, so its speed is R^* over its reads times that utilisation, 44–45% of the limit with the whole sequence in view.

slots and run from there, and every other miss runs on the CPU; the oracle refuses an admission when no resident’s next use is further than the candidate’s. The single-read prefetch admits, each step, the non-resident experts whose first use is at least d steps away ($d=2$ unpaced, 3 paced), provided the victim’s next use is further than the candidate’s use after that first one; an expert used sooner than d steps away misses on the CPU, or is fetched when the two are combined. The same policies simulated on the trace with the engine’s semantics (the lookahead within the sequence, residents carried over with their next use recomputed at each sequence start) predict the engine’s reads and hit rates: the fetch oracle’s hit rate within 0.4 points and its reads within 0 to 4%, the combined oracle’s hits within two points (Table 8). The simulator the job’s predictions were registered against gave a resident carried over from the previous sequence no next use, so it was evicted first; that over-counted the fetch oracle’s reads by 17% at the two GPU-bound cells, which is why two of its reads clauses failed (Appendix B). Each configuration of `llama-ec-bench` starts with a cold cache in a fresh context. The paced copy path copies one 16 MB piece at a time on a single copier thread and publishes a landed copy at the start of the next step; the unpaced path copies each admission as one transfer. The hit-optimal oracle’s background admissions drive host memory at 71–82% of the probe’s best rate at the host-bound cells, the prefetching single-read oracles at 64–85% (the paced one 71–82%), the online policy itself at 66–82%, the fetch oracle at 42–67%. The predictions and their outcomes are in `prereg/foresight_outcome_095.md`: of 56 clauses, 44 held and 12 failed, all of size (Appendix B).

The double-reading oracles. Both read the routing of the next W decode steps; their copies ride the paced admission path, and an expert still in flight at its step runs on the CPU like any other miss. The *hit-optimal* oracle admits the experts those steps will route to, soonest first, evicting the resident whose next use is furthest (Belady within the window), for one future use as readily as for ten. The *bytes-optimal* oracle is MIN with bypass with its admissions made through the background path: a miss is served where it is, by the CPU or a demand copy, and admitted only if its next use comes before the furthest

Table 9. The accounting, measured: milliseconds per token of the engine’s states that the oracle jobs realised, on each oracle host, beside that host’s limit and the model’s values for the states it prices; O1–O3 ran jobs 093–095 in a fixed order, O4 and O5 job 096 in a shuffled order. *online*: the deployed policy; *foa*: the same with each admission fetched into its slot in the step that misses it (one read, no foresight); *no ovl.*: its within-layer overlap off; *all CPU*: no FETCH table; *bypass*: MIN with bypass with background admissions (two reads per admitted expert); *hit-opt.*: Belady prefetch without bypass, which admits what the optimum bypasses, unpaced, and the same on the paced copy path (the first *paced* column; O1 and O2 ran only that form); *nb2*: Belady within a two-step lead with each admission read once; *fetch*: MIN with bypass, admitted misses fetched into their slot (one read, serialised: the model’s foresight-only state $v(F)$); *both*: fetch and the scheduled single-read prefetch (one read, overlapped: the model’s $v(O, F)$); the second *paced*: both on the paced copy path with a three-step lead; *limit*: the limit on that host; $v(F)$, $v(O, F)$, $v(\text{all})$: the model.

Host	Cell	online	foa	no ovl.	all CPU	bypass	hit-opt.	paced	nb2	fetch	both	paced	limit	$v(F)$	$v(O, F)$	$v(\text{all})$
O1	gpt-oss 11%	21.12	–	21.01	22.91	–	–	22.88	–	–	–	–	10.06	14.54	11.27	10.06
O1	gpt-oss 25%	12.89	–	12.91	14.21	–	–	9.22	–	–	–	–	4.03	8.04	5.10	4.03
O1	gpt-oss 40%	8.90	–	8.95	–	–	–	5.99	–	–	–	–	1.97	5.75	3.84	1.97
O1	Qwen3 12.5%	36.92	–	36.91	43.29	–	–	40.29	–	–	–	–	18.04	24.90	19.29	18.04
O1	Qwen3 25%	22.73	–	22.76	26.60	–	–	18.22	–	–	–	–	8.11	14.04	9.22	8.11
O1	Qwen3 43.75%	12.71	–	12.67	–	–	–	8.12	–	–	–	–	3.32	8.25	5.51	3.32
O2	gpt-oss 11%	20.85	–	–	–	20.17	–	21.57	–	–	–	–	10.04	14.27	11.96	10.04
O2	gpt-oss 25%	12.72	–	–	–	10.77	–	7.80	–	–	–	–	4.02	7.86	5.86	4.02
O2	gpt-oss 40%	8.79	–	–	–	7.13	–	5.62	–	–	–	–	1.97	5.63	3.78	1.97
O2	Qwen3 12.5%	36.64	–	–	–	36.58	–	38.74	–	–	–	–	18.00	24.70	20.06	18.00
O2	Qwen3 25%	22.24	–	–	–	19.67	–	15.80	–	–	–	–	8.09	13.58	9.90	8.09
O2	Qwen3 43.75%	12.02	–	–	–	10.11	–	7.42	–	–	–	–	3.32	7.56	4.90	3.32
O3	gpt-oss 11%	18.01	–	–	–	–	19.79	–	–	14.00	12.91	12.39	8.26	12.63	9.49	8.26
O3	gpt-oss 25%	11.08	–	–	–	–	9.74	–	–	8.42	7.44	6.66	3.31	7.05	4.28	3.31
O3	gpt-oss 40%	7.77	–	–	–	–	6.14	–	–	6.20	5.58	5.20	1.95	5.24	3.56	1.95
O3	Qwen3 12.5%	30.95	–	–	–	–	44.54	–	–	23.02	21.78	21.89	14.81	20.35	16.06	14.81
O3	Qwen3 25%	19.32	–	–	–	–	18.17	–	–	14.01	13.09	11.58	6.66	11.68	7.76	6.66
O3	Qwen3 43.75%	10.94	–	–	–	–	8.59	–	–	8.66	7.78	7.28	3.30	7.24	4.76	3.30
O4	gpt-oss 11%	20.59	20.18	–	–	20.11	19.78	22.03	19.64	15.13	14.06	13.78	10.02	13.77	11.63	10.02
O4	gpt-oss 25%	12.43	12.07	–	–	10.55	9.76	8.13	9.61	8.72	7.65	6.87	4.01	7.40	5.34	4.01
O4	gpt-oss 40%	8.39	8.19	–	–	6.91	6.15	5.38	6.03	6.28	5.59	5.21	1.97	5.12	3.20	1.97
O4	Qwen3 12.5%	36.36	35.83	–	–	36.54	48.80	39.66	31.64	25.29	24.01	24.58	17.97	23.98	19.88	17.97
O4	Qwen3 25%	22.09	21.75	–	–	20.06	17.74	16.93	17.57	14.60	13.50	12.24	8.08	13.12	9.54	8.08
O4	Qwen3 43.75%	11.94	11.74	–	–	10.09	8.44	7.17	8.33	8.69	7.72	7.29	3.32	7.32	4.56	3.32
O5	gpt-oss 11%	21.26	21.31	–	–	21.47	33.34	25.57	33.03	18.34	17.97	16.39	9.94	15.22	11.20	9.94
O5	gpt-oss 25%	13.01	13.35	–	–	11.18	15.87	11.17	15.60	10.60	10.54	8.68	3.98	9.13	5.06	3.98
O5	gpt-oss 40%	8.73	9.22	–	–	7.16	8.34	6.25	8.13	7.27	6.87	5.91	1.97	6.26	3.62	1.97
O5	Qwen3 12.5%	37.63	37.42	–	–	40.11	69.42	44.89	49.11	31.37	31.27	30.01	17.82	27.95	19.27	17.82
O5	Qwen3 25%	23.24	23.27	–	–	21.12	30.78	21.76	30.47	18.51	19.64	16.03	8.01	16.77	9.15	8.01
O5	Qwen3 43.75%	12.49	12.92	–	–	10.32	12.08	9.00	11.86	10.21	9.61	8.23	3.31	8.94	5.02	3.31

next use among the residents. With the whole sequence in view the hit-optimal oracle runs $1.40 \times [1.38, 1.41]$ the online policy at gpt-oss 25% and $1.25 \times [1.24, 1.26]$ at Qwen3 25% on host O1, $1.63 \times$ and $1.41 \times$ on O2, and $1.48\text{--}1.62 \times$ at the two GPU-bound cells; half of its gain at gpt-oss 25% arrives between $W=4$ and 16 (9.2 tokens by log interpolation; Qwen3 25%: 6.1 against the trace’s 3.9). Where the host binds hardest it loses, 3–8% at gpt-oss 11% and Qwen3 12.5% on both hosts, and at every window on O1 ($0.78\text{--}0.92 \times$): it reads more host bytes than the online policy (82 and 195 experts per token against 63 and 160), $2.02\text{--}2.18 \times$ the optimum’s R^* , because Belady admits an expert used once where the optimum bypasses it and re-copies what it evicts (its admissions, each one read, are most of its reads; a miss served on the CPU before its copy lands is the second read of that expert). An unpaced copy path, run at the four host-bound cells of O1, reaches 98–99% hits at three of them and still loses at gpt-oss 11% and Qwen3 12.5% (where it hits 75%), so copies that arrive in time do not remove the loss there: the extra bytes cause it. The bytes-optimal oracle gains 13–23% at the four cells of 25% and above and 0–3% at the two below (6% with a 16-token window); its misses fall (30 to 20 and 85 to 64 per token at 25%) but its total reads, misses plus the admissions that read each admitted expert again, are $1.73\text{--}2.04 \times R^*$ and exceed the online policy’s by 4–5% at the two lowest budgets; the model on its own counters predicts its time within 9% at every cell, and the bytes it saves account for 4–21 points of its 13–23% gain at the four cells (most of it at gpt-oss 40% and Qwen3 43.75%, a third at Qwen3 25%). The hit-optimal oracle reads $1.01\text{--}1.13 \times$ the bytes-optimal one’s host bytes on O2 and gains 41–63% at the same four cells, and the model, with its GPU term calibrated on the online run, over-predicts those runs’ time per token by 37–53% (on O1, 24–39%): its gain is the overlap of copies issued ahead of their use with the step, which the model charges in sequence. Pacing the online policy’s own copies changes it by under 1% ($0.997\text{--}1.005 \times$). With no FETCH table (every miss on the CPU) the online policy runs 8–15% slower at the four host-bound cells; with the engine’s within-layer overlap of a layer’s CPU misses with its GPU hits turned off it changes by under 1% at every cell (0.5% faster to 0.6% slower), so the overlap the accounting charges is across the step and in rebalancing, not within a layer. The bytes-optimal oracle with a 16-token window keeps 34–176% of the whole-sequence gain at the five cells where that

Table 10. The hit-optimal oracle in the engine (job 093; RTX 5090, Ryzen 9 9950X, host memory 50.5 GB/s at best). tok/s of the online policy and of the oracle with the whole sequence in view; their ratio, paired by sequence, with its 95% interval; hit rate and misses per token of each (experts read from host memory by the CPU or by a demand copy), the optimum’s reads in parentheses, and the oracle’s background admissions (Belady prefetch without bypass: it admits what the optimum bypasses); speed as a percentage of this host’s limit; the accounting’s foresight-only term $v(F)$ for the cell in tok/s; the share of that term the oracle recovered and the share of the gap to the limit it closed.

Cell	online	oracle	ratio	hits % (opt.)	misses/token (opt.)	oracle adm.	% of limit	$v(F)$	recovered %	closed %
gpt-oss 11%	47.4	43.7	0.92 [0.91, 0.93]	58 → 67 (73)	60 → 47 (38)	35	47.6 → 44.0	68.8	−27	−16
gpt-oss 25%	77.6	108.5	1.40 [1.38, 1.41]	79 → 92 (89)	30 → 11 (15)	20	31.2 → 43.7	124.4	76	41
gpt-oss 40%	112.4	166.9	1.48 [1.46, 1.51]	89 → 98 (95)	15 → 3 (7)	12	22.1 → 32.8	174.0	92	42
Qwen3 12.5%	27.1	24.8	0.92 [0.91, 0.92]	60 → 66 (75)	155 → 132 (97)	63	48.9 → 44.8	40.2	−28	−18
Qwen3 25%	44.0	54.9	1.25 [1.24, 1.26]	78 → 89 (89)	85 → 42 (43)	46	35.7 → 44.5	71.2	52	31
Qwen3 43.75%	78.7	123.1	1.56 [1.53, 1.59]	92 → 99 (96)	32 → 5 (14)	25	26.1 → 40.8	121.3	103	49

gain’s interval excludes zero (more than all of it at gpt-oss 11%, where the shorter window admits less).

Why the bytes-optimal oracle’s misses stay above the optimum’s reads. Besides the second read, its misses alone are $1.31\text{--}1.47\times$ the optimum’s reads at the four host-bound cells (at the two GPU-bound ones the ideal simulation below is itself 6–10% above the optimum, which sees across sequences). The engine’s paced copy path moves one 16 MB piece at a time from pageable host memory on a single copier thread, behind the demand copies, publishes a landed copy at the start of the next step, and admits nothing while 64 copies are queued; at the 16–30 admissions per token (210–285 MB) of the two lowest budgets it lags the admissions by steps, and an expert still in flight is a miss the CPU reads again. The last four columns of Table 11 simulate MIN with bypass on the measured routing trace with the engine’s own semantics (the lookahead within the sequence, residents carried over, no queue cap) and a copy that lands d steps after it is issued, during which the expert misses and is not re-admitted: at $d=1$, the ideal, the misses come within 0–4% of the limit’s optimum at the host-bound cells; the engine’s misses meet the simulation at $d=2.8\text{--}4.7$ there and at 3.0–3.5 at the GPU-bound cells, three to five steps. At the fitted latency the simulation admits $1.1\text{--}1.9\times$ what the engine admitted, so the queue cap, not only the latency, keeps the engine’s misses above the ideal. The predictions and their outcomes are in `prereg/foresight_outcome_093.md` and `094.md`: of 36 clauses for the first job, 9 held, 2 held on the point estimate, 21 failed and 4 were untested; of 38 for the second, 12 held, 1 on the point estimate, 24 failed and 1 was untested (Appendix B). The second job’s failures are of size, not direction: the bypass oracle’s misses were predicted within 15% of the optimum’s reads and came 29–47% above (the copy latency and the queue cap); its hit rate was predicted within 4 points of the optimum’s and fell 5–11 points short at three cells; its gain was predicted at 10–45% at the host-bound cells and came below that at the two lowest budgets; the model was predicted to over-predict it by 10–40% at the host-bound cells and over-predicted by 5–9%; and the two oracles were predicted within 15% of each other at 25% and above, where the hit-optimal one was $1.24\text{--}1.38\times$ faster.

Where a token’s time goes at a GPU-bound budget. On Qwen3 at 43.75% Nsight Systems splits a token of our cache (FreeToken in parentheses) into 4.93 (5.20) ms of expert copies over PCIe, 2.29 (2.21) of resident-expert products, 1.65 (1.34) of attention projections, 0.49 (0.37) of attention, 0.38 of output head and 0.97 (0.40) of norms, small kernels and cache control, with neither system overlapping its copies with GPU compute. Turning off the within-layer overlap of a layer’s CPU misses with its GPU hits changes speed by under 1% at every budget, and predicting the next layer’s experts (Hwang et al., 2024; Zhu et al., 2025) never lowers host bytes.

H THE TRACE STUDY IN DETAIL

Policies. Table 12 scores LRU, LFU, decayed frequency, ARC, S3-FIFO and a profiled static cache against MIN and against Belady within a window of W future tokens, at three budgets per model, on each model’s own sampled text (9 traces, the cache carried across conversations). Without admitting every miss, the online policies are close to one another and far from MIN: the best reads 35–110% more, decayed frequency is best at 12 of 27 budget points and S3-FIFO at 10. Where the cache turns over every token ($C/k \leq 1.5$) the six are within 4% of one another; from $C/k=2$ LRU reads 2–11% more than decayed frequency on eight of nine models, and classic LFU fails at large budgets. Adding *admit every miss* (every miss copied into the slot of the lowest-scored resident not requested now) changes the ranking: it is best at 18 of 27 points, ties at 5, and loses by at most 2.7%; it saves up to 5.8% over the previous best, and the best online policy still reads 35–110% more than MIN (`prereg/policy_aa.json`). At the small budgets almost every resident turns over within a few tokens, and the expert just requested is a good bet. In our engine admitting every miss reads all of its misses over the

Table 11. The bytes-optimal oracle (MIN with bypass) beside the hit-optimal one on the second oracle host (RTX 5090, Ryzen 9 7950X, host memory 50.6 GB/s at best). Per cell: the online policy’s speed; each oracle’s ratio to it with the whole sequence in view (paired, 95% interval), its hit rate, its misses per token (experts read from host memory by the CPU or by a demand copy; its background admissions in parentheses: the bypass oracle’s are this step’s CPU-served misses, each read a second time by the copy, the hit-optimal one’s are prefetches of what the optimum would bypass as readily as of what it would admit) and the calibrated model on its own counters against its measured time; the optimum’s reads and hit rate from the trace; and MIN with bypass simulated on the trace with the engine’s lookahead (within the sequence) under a copy latency of 1 (ideal), 2 and 3 steps, with the latency at which the simulation meets the engine’s misses.

Cell	online	bypass oracle, $W=\text{all}$				hit-optimal oracle, $W=\text{all}$				optimum		simulated misses/token			
		ratio	hits	misses (adm.)	model	ratio	hits	misses (adm.)	model	reads	hits	$d=1$	2	3	fit
gpt-oss 11%	48.0	1.03 [1.03, 1.04]	65	50 (16)	+7%	0.97 [0.96, 0.98]	71	42 (39)	+19%	38.3	73	39	43	47	3.9
gpt-oss 25%	78.6	1.18 [1.18, 1.19]	86	20 (9)	+6%	1.63 [1.61, 1.66]	95	7 (23)	+53%	15.3	89	16	17	19	4.7
gpt-oss 40%	113.8	1.23 [1.22, 1.24]	94	9 (4)	+2%	1.56 [1.54, 1.59]	98	3 (12)	+39%	6.8	95	7	8	9	3.5
Qwen3 12.5%	27.3	1.00 [1.00, 1.01]	64	138 (30)	+5%	0.95 [0.94, 0.95]	67	128 (65)	+11%	96.5	75	97	122	143	2.8
Qwen3 25%	45.0	1.13 [1.12, 1.14]	83	64 (21)	+9%	1.41 [1.39, 1.43]	91	34 (52)	+37%	43.4	89	44	54	61	3.4
Qwen3 43.75%	83.2	1.19 [1.18, 1.20]	95	19 (9)	+1%	1.62 [1.58, 1.65]	99	3 (27)	+47%	13.7	96	14	17	19	3.0

link, so its value in time depends on the link (Section 6). A static cache profiled on other text reads $1.4\text{--}13.2\times$ MIN because the hot set moves: the C most requested experts of a layer change by 22–58% between windows of 100 tokens and 33–75% between windows of 1,000.

The horizon, before and after the stronger baseline. Against decayed frequency with each admission fetched once (the study’s original baseline), MIN reads 26–56% fewer experts (median 41%), and the window that closes half of the gap fitted $W_{50} \approx 0.59 (C/k)^{1.33}$; admitting every miss lowers the baseline at 23 of 26 points by up to 6.7%, raises W_{50} by a median 9% (at most 42%), and refits it to $0.67 (C/k)^{1.31}$, with a leave-one-model-out exponent range of 1.25–1.36 and, without the points below one token, 1.35 (it was 1.51 against the original baseline). Ninety percent of the gap needs a median of $2.7 C/k$ tokens. Table 13 compares the rule in distinct experts with the power law in tokens and with $W_{50} \propto C/k$, each fitted on eight models and tested on the ninth: the first two are not distinguished, and both beat the third. The rule in distinct experts inverts the held-out model’s own curve of distinct experts against window length, which needs that model’s routing trace.

A learned admission order. We trained a logistic model of whether an expert is requested within the next C/k tokens on features the engine has when it plans a step (decayed request counts at five half-lives, the time since the last request, whether the expert is requested now, and transition scores from the previous token’s experts at the same layer and, through an 8-bit table, at every layer), fitted on each model’s mixed-domain trace and replayed on the AIME routing, so no test text is seen in training (Table 14). With each admission read once it reads 6–12% fewer experts than decayed frequency with the same single-read admission at the four host-bound budgets, closing 17–30% of that policy’s excess over MIN; admitting every miss alone closes most of that at Qwen3 12.5% and little at gpt-oss (prereg/learned.offline.json and the research checks). In the engine (job 097, predictions committed before launch) it reads 6–13% fewer experts than the single-read online policy at the host-bound budgets on both hosts, but runs $1.01\text{--}1.03\times$ it on O4’s machine and $0.90\text{--}0.94\times$ on host O6: the reads it saves are CPU reads, bought with more fetches in the step, each a copy its layer waits on, and with $237\text{--}585 \mu\text{s}$ of host time per token (over all six budgets).

Speculation is not foresight. Speculative decoding verifies K draft tokens in one forward pass, so a layer’s experts for K positions are requested together and a missed expert is read once for all of them. On the same traces, with drafts accepted at rate α until the first rejection, batched verification saves an online cache a median 15–32% of its reads at $K=8$ when rejected drafts cost nothing, the equivalent of a window of two to four tokens; once the rejected drafts’ experts are counted, the saving is 8–12% at $\alpha=0.9$ and becomes a loss of 18–20% at $\alpha=0.8$, because the union of K positions’ expert sets is 20–79% of Kk . MIN itself gains little from batching (–7 to 15%). Systems that use drafts as a look-ahead sensor (Li et al., 2026) compare against a baseline that already pays for verification; ours charges the rejected drafts to a cache that does not speculate.

I PARTIAL FORESIGHT IN DETAIL

The window policies. At step g the engine plans the step’s misses before it runs, from the trace’s record of step g (the routers give it layer by layer in a deployable system) and from a view of records $g+1..g+W$ of the same problem. A miss is processed in request order: it takes a free slot; otherwise the victim is a resident not requested now that is not seen in the

Table 12. Host reads per token relative to the exact optimum (Belady’s MIN with bypass over the whole future), median over nine models at three budgets, and the range over all 27 model–budget cells. Event-atomic replay of each model’s own sampled text; every online policy is scored on the last 90% of the trace. Mixtral (8 experts) uses $C = 2, 3$ and 4.

Policy	$C=E/8$	$E/4$	$3E/8$	range
LRU	1.44	1.77	2.03	1.37–2.26
LFU	1.44	1.93	2.40	1.37–6.15
Decayed frequency	1.45	1.70	1.94	1.35–2.10
with hysteresis κ	1.46	1.75	2.00	1.36–2.55
ARC	1.44	1.76	1.99	1.37–2.23
S3-FIFO	1.44	1.69	1.94	1.38–2.11
Static, profiled	1.92	2.74	3.48	1.43–13.24
Belady, $W=1$	1.18	1.45	1.69	1.07–2.06
Belady, $W=4$	1.01	1.13	1.31	1.00–1.95
Belady, $W=16$	1.00	1.00	1.01	1.00–1.66

Table 13. Predicting a held-out model’s horizon W_{50} (fitted on the other eight models, tested on the ninth, over all 26 budget points): error as a factor, $\max(\hat{W}/W, W/\hat{W})$.

Rule	Constants	Median	90th pct.	Max
$W_{50} \propto C/k$	1	1.22	1.57	2.42
$W_{50} = a (C/k)^b$	2	1.19	1.51	1.67
$D(W_{50}) = cC$	1, and $D(W)$	1.16	1.42	1.61

window (lowest decayed score), or, if every resident is seen, the one seen furthest; a miss seen in the window is admitted unless the victim is seen sooner, and a miss not seen is admitted if the victim is not seen either (admit every miss beyond the window; with the deployed margin $\kappa=1$, only if its score beats the victim’s by one). On the panel admissions are fetched in the step, so every miss is read once; on the deployed path of the follow-up the host’s fetch table splits the misses between the CPU and an in-step copy, and an admission served by the CPU is copied afterwards, a second read. In a degraded view each expert of a future record is kept with probability r , decided by a hash of (seed, g , record, layer, position), so step g sees one view throughout and step $g+1$ an independent one; a dropped expert is replaced by a uniformly drawn one (fill: precision and recall r) or left out (drop: recall r , precision 1). `scripts/value_map.py` replays the plan bit for bit: on a CPU build the engine’s reads equal the replay’s in 30 configurations at three budgets (`cputest099`), and on the panel’s GPUs, where the routing in the run can differ from the recorded trace by numerics, within 1.7% at both gpt-oss budgets. The panel ran the first 20 of the 30 problems; the table’s reads are on all 30. Table 16 lists the windows on the deployed path and in the step on the hosts of job 100. Copied in the step, a window beats the deployed path where the link reads at least 0.92 of the CPU rate and loses to it at 0.54 or below; on the deployed path the engine misses 7–33% more than an instant replay (the deployed state matches its replay within 2.4%), and background copies that land late are the likely cause.

What the table says. Horizon and accuracy trade: at gpt-oss 11% an 8-token window at recall 0.5 closes 0.29 of the gap, about what two exact tokens close (0.31), and the whole future at recall 0.5 closes 0.43. Fill and drop cost about the same

Table 14. A learned admission order, offline: host reads per token on the AIME-25 routing the engine runs, of the decayed frequency with each admission read once (single read; within 10% of the engine’s fetch-on-admit reads on hosts O4 and O5), the learned order with single-read admission (trained on the model’s mixed-domain trace; 8-bit cross-layer table, as in the engine), Belady within one and four tokens of perfect foresight, and MIN with bypass in its no-evict form (Table 20). *closed*: the share of the single-read policy’s excess over MIN that the learned order removes. [†]host-bound.

cell	single read	learned	$W=1$	$W=4$	MIN	closed
gpt-oss 11% [†]	59.4	56.0	54.4	47.3	39.0	17%
gpt-oss 25% [†]	28.3	26.1	27.3	25.1	15.5	17%
gpt-oss 40%	14.3	12.6	14.0	13.2	6.8	23%
Qwen3 12.5% [†]	166.2	146.3	134.1	109.1	100.4	30%
Qwen3 25% [†]	84.2	74.8	74.1	62.8	44.3	24%
Qwen3 43.75%	28.8	26.3	27.7	25.3	13.8	17%

Table 15. Partial foresight on the AIME routing the engine runs: host reads per token of admitting every miss and of MIN with one read, and the share of the gap between them that each window closes (Table 1). Recall 0.5: each future expert kept with probability 0.5 and replaced by a random one (fill) or left out (drop); means of three seeds. Linear: the ridge forecaster of Appendix I with a 4-token window. $\kappa=1$: the same exact 4-token window with the deployed margin instead of admitting every miss beyond it, as a share of its own gap.

Budget	Reads per token		Exact window W (tokens)					Recall 0.5, fill / drop		Linear	$\kappa=1$
	every miss	MIN	1	2	4	8	16	$W=8$	all	$W=4$	$W=4$
gpt-oss 11%	58.6	39.3	0.18	0.31	0.52	0.81	0.98	0.29 / 0.29	0.43 / 0.48	0.13	0.62
gpt-oss 25%	28.2	16.0	0.07	0.14	0.25	0.42	0.66	0.14 / 0.13	0.46 / 0.47	0.05	0.34
Qwen3 12.5%	150.0	100.3	0.31	0.51	0.77	0.96	0.99	0.42 / 0.43	0.48 / 0.52	0.19	0.88
Qwen3 25%	79.6	44.9	0.16	0.28	0.46	0.70	0.93	0.29 / 0.28	0.49 / 0.53	0.11	0.60

for windows of up to 16 tokens; with the whole future in view at low recall, dropping costs less (the largest difference, 0.12, is there). A larger cache needs a longer window: at 25% the same exact 4-token window closes 0.25 of the gap, against 0.52 at 11%.

A realisable forecaster. Ridge regression from the all-layer routing of step t to that of step $t+h$ ($h=1..8$), fitted on each model’s mixed-domain own text and applied to the AIME routing, forecasts the top k experts of each layer with a precision of 0.48 one token ahead and 0.26 four ahead on gpt-oss (0.54 and 0.32 on Qwen3); a decayed request count, the online policy’s own signal, scores 0.30 and 0.28. Driving the window policy, the ridge forecasts close 3–15% of the gap on gpt-oss and 6–19% on Qwen3 over windows of 1–8 tokens (prereg/forecaster_linear.json). Its precision falls with the horizon, so over a window it is closer to 0.3 than to its one-token value; at that average precision a window with random errors closes about as much on gpt-oss, and more on Qwen3, where the forecaster’s correct forecasts are mostly experts the cache already holds. A sequence model reports recall above 90% one token ahead with $k+3$ candidates (Wang et al., 2026b); its precision is then at most $k/(k+3)$, 0.57 for gpt-oss and 0.73 for Qwen3. A recurrent forecaster in that spirit (a GRU of 512 units over the same all-layer vectors, one output head per horizon, trained on the same text; scripts/forecaster_gru.py) reaches a precision of 0.40 and 0.26 one and four tokens ahead on gpt-oss (0.46 and 0.31 on Qwen3), a recall at $k+3$ of 0.53 and 0.54, and closes 2–11% and 6–17% of the gap: no better than the linear model. Our training text is smaller than SeqMoE’s and differs from the test text, so this bounds what such a model does across text, not what it can do in domain.

J MIN’S SCHEDULES AND THE MACHINE, IN DETAIL

The fewest-admission schedule, host by host. Tables 17 and 18 list jobs 104 and 105: MIN’s set loaded by the CPU or in the step, as the greedy rule picks it or as the schedule with the fewest admissions does. On the engine’s 20 problems the fewest-admission schedule copies 0.60–0.73 times as many experts in the step as the greedy one, for 2.5–5.3% more misses (more than its replay predicts; its measured gains include that cost). On every launch it gains over the deployed cache with either load and at both budgets, each 95% interval above 1; loaded by the CPU on the machine with the fastest link, it is 0.07 behind the greedy set at 25%.

Table 16. Foresight on each read path (job 100): speed relative to the deployed cache on the same host. *In step*: admit every miss with a window of W exact tokens, each admission copied in the step (as on the panel). *Deployed path*: the deployed policy with the same window, its misses split between the CPU and an in-step copy by the host’s fetch table and its admissions copied in the background; $W=8$ at recall 0.5 in the column marked r . Hosts sorted by the probe’s link-to-CPU ratio.

Host	Link/ CPU	Budget	In step $W=4$ 16		Deployed path 4 16 r			MIN, 2 reads	MIN, 1 read	MIN, pref.
Pd again	0.41	gpt-oss 11%	0.66	0.89	1.02	1.03	0.94	1.00	1.05	1.11
Pd again	0.41	gpt-oss 25%	0.69	0.81	1.03	1.12	1.01	1.19	1.11	1.29
9950X, slower link	0.54	gpt-oss 11%	0.65	0.89	1.01	1.01	0.96	0.99	1.07	1.23
9950X, slower link	0.54	gpt-oss 25%	0.69	0.80	0.96	1.06	1.00	1.10	1.11	1.37
5700X3D	0.92	gpt-oss 11%	1.09	1.32	0.99	1.00	0.95	0.99	1.37	1.44
5700X3D	0.92	gpt-oss 25%	1.06	1.23	1.03	1.07	1.00	1.16	1.48	1.80
Ph again	0.96	gpt-oss 11%	0.99	1.20	1.01	1.03	0.98	1.03	1.27	1.45
Ph again	0.96	gpt-oss 25%	0.97	1.09	1.03	1.07	1.01	1.16	1.29	1.64

Table 17. Which of MIN’s hit-optimal sets the oracle follows (job 104): speed relative to the deployed cache on the same host, with MIN’s set loaded by the CPU (two reads) or in the step (one read), as the greedy rule picks it or as the schedule with the fewest admissions does; and the copies per token each makes in the step. Hosts sorted by the probe’s link-to-CPU ratio.

Host	Link/ CPU	Budget	By the CPU greedy fewest		In the step greedy fewest		Copies/token greedy fewest	
Pf again	0.29	gpt-oss 11%	0.94	1.11	0.86	1.04	19.8	11.8
Pf again	0.29	gpt-oss 25%	1.13	1.22	0.96	1.06	9.7	7.0
285K, second	0.49	gpt-oss 11%	1.03	1.08	1.05	1.18	19.8	11.8
285K, second	0.49	gpt-oss 25%	1.20	1.21	1.12	1.17	9.7	7.0
Pg again (5950X)	0.63	gpt-oss 11%	1.02	1.06	1.26	1.34	19.8	11.8
Pg again (5950X)	0.63	gpt-oss 25%	1.20	1.15	1.32	1.35	9.7	7.0

Job 106, every host. Table 19 adds the two server machines new to job 106. On the EPYC 7302 host the engine computed wrong outputs: its teacher-forced loss on the deployed cache was 0.21–0.45 nats per token across rounds and models, against 0.19 (gpt-oss) and 0.085 (Qwen3) on every other host and in its own untimed routing pass, and its counters changed between rounds. On the Xeon 8347C host the deployed cache varied by 39% between rounds at gpt-oss 11%, against at most 2% registered, and the greedy copy gained in two of its three rounds where we had predicted it would lose. When their jobs reached the gate, before any download, `free` reported 131–132 GB of host memory in use (12–15 GB on the other hosts), and their probes’ CPU read rates fell when more threads read than the container had usable physical cores. On both, Eq. (3) under-predicted the deployed cache by 10–52%. We report their clauses (in the supplement) and draw nothing from them. The three machines that ran stably were all rented before: Pf, the Threadripper 9960X of job 105, and panel host Pe (same GPU).

The machine decides how much. Across machines, the gain of every state that moves reads from the CPU to the link follows the link-to-CPU ratio (Spearman over the machines that ran it through job 104, 95% bootstrap interval): MIN copied in the step 0.89 [0.63, 0.97] (19 machines), copied ahead 0.87 [0.54, 0.99] (17), admitting every miss 0.95 [0.75, 0.99] (13). MIN loaded by the CPU, which reads every admission twice and copies it in the background, tracks the link’s absolute rate more closely (0.79 [0.49, 0.92], 17 machines, against 0.45 for the ratio; the two intervals overlap, and the feature was picked from a search over all probe features). The greedy schedule copied in the step gains 1.02–1.05 $\times$ at 11% on 2 further machines with ratios 0.40 and 0.48, and loses only on the machines whose link reads at under a third of their CPU rate;

Table 18. Job 105, registered before the runs, nothing fitted: the deployed cache’s time per token against Eq. (3), the GPU’s compute profiled on the same host before the timed runs plus the run’s counted host reads at the machine’s best rate; and the speed, relative to the deployed cache, of the deployed cache with the engine’s layer-ahead copy (each layer’s top predicted expert of the next layer copied ahead of use) and of MIN’s greedy and fewest-admission sets copied in the step. Hosts sorted by the probe’s link-to-CPU ratio.

Host	Link/ CPU	Budget	G_{prof} (ms)	Deployed (ms) law	Layer- ahead	MIN, in step greedy	fewest
Pf again	0.28	gpt-oss 11%	4.1	13.6	13.9	0.72	1.03
Pf again	0.28	gpt-oss 25%	4.4	9.8	9.5	0.69	1.06
Threadripper 9960X	0.32	gpt-oss 11%	4.1	9.3	9.1	0.91	1.04
Threadripper 9960X	0.32	gpt-oss 25%	4.4	7.3	6.8	0.87	1.04
EPYC 7402P	0.45	gpt-oss 11%	4.3	19.1	18.2	0.81	1.15
EPYC 7402P	0.45	gpt-oss 25%	4.5	12.9	11.8	0.78	1.15
O4 again (9950X)	1.04	gpt-oss 11%	4.4	20.8	20.5	1.04	1.34
O4 again (9950X)	1.04	gpt-oss 25%	4.4	13.2	12.4	1.01	1.39

made ahead of use, the same copies lose beyond their interval on no machine they ran on (at worst $0.997\times$, on Pf).

Machines differ more than problems. Problems rank largely the same on every machine: Kendall’s coefficient of concordance of MIN-copied-in-the-step’s per-problem gain over 19 machines and 20 problems is 0.77 at 11% and 0.82 at 25%, and a problem’s gain follows how far the deployed policy’s reads exceed MIN’s on its trace (Spearman 0.88 at 11%, 0.68 at 25%). Yet machines carry 96% of the variance of its log gain and problems 1–3%.

K THE EXPERT CACHE IN MORE DETAIL

This section holds the measurements of Section 2 that the main text summarises: where the cache’s gain comes from, the choice of the FETCH split, the hosts with half the link, long outputs, FreeToken’s own headline model, two other llama.cpp-based systems, output parity, and the cost of prompts and batches. Job numbers refer to the scripts on the gpu branch.

Where the gain comes from. Figure 9 builds the system up from stock llama.cpp on one machine. Placing each layer’s hot experts is the first gain: a static cache of the 32 experts per layer most requested on other text (the models’ own MATH-500 outputs) runs $1.31\times$ and $1.89\times$ stock, which pins whole layers instead. Its hit rate (52% and 67%) depends on how well the profile matches the workload: profiled in hindsight on the measured problems it hits 75% and 72% and runs faster (hatched). Replacing experts online closes that gap without a profile, if admissions are rationed: LRU admits every miss, and on Qwen3 its copies raise host reads from 127 to 188 experts per token, so it runs $0.88\times$ the static cache ($1.30\times$ on gpt-oss); the decayed-frequency policy reads 41 and 118 per token, runs $1.58\times$ and $1.12\times$ the profiled static cache, and the hindsight cache runs 0.98 [0.95, 1.01] and 1.01 [0.98, 1.04] times as fast as it. On traces LRU misses only 2–11% more than decayed frequency at these budgets (Section 6); the rest of its cost here is the copies it admits. On host S, with everything else as in the full system, LRU runs 6–21% slower than decayed frequency across the six cells of Table 2 and trails FreeToken at four of them: the admission policy is most of the lead over FreeToken (77–84% of it at gpt-oss 25 and 40%, more than the whole of it at gpt-oss 11% and Qwen3 12.5 and 25%). The GPU-signalled CPU helpers, GPU slot maps, GPU sampling and FETCH add 2–9% each, and the model’s FETCH table 5–8% over the fixed one. The full system runs $2.68\times$ (gpt-oss) and $2.81\times$ (Qwen3) stock llama.cpp.

How tight is the bound? Table 20 changes one assumption at a time on host B. Pooled slots need 4.5–18.6% fewer reads and raise the host-bound limits by 5–10%; starting MIN warm changes it by at most 1.4%; the median probe sample instead of the highest lowers the host-bound limits by 8%. The GPU term moves most: at batch size 1 the all-in-VRAM runs reach only 52% (gpt-oss) and 61% (Qwen3) of the datasheet rate, and with the measured rate the limits at 25% and above fall to 279 and 192 tok/s, 8–9% above the all-in-VRAM speed. In the running example (gpt-oss at 25%) the limit on host B is

Table 19. Job 106, registered before the runs, every host (*: deployed cache varied by more than the registered 2% between launches): the deployed cache’s time per token against Eq. (3) (GPU compute profiled on the same host before the timed runs; the run’s own counters), and the speed relative to the deployed cache of the two online admission rules (*dk*: the decayed count with a larger admission margin; *lrn*: the learned reuse predictor with a margin; both chosen on other text) and of MIN’s fewest-admission set copied in the step, with round-level 95% intervals (rounds, then problems: three rounds at gpt-oss 11%, two at 25%; Qwen3 ran one round, so its intervals are over problems). Hosts sorted by the probe’s link-to-CPU ratio.

Host	Link/ CPU	Budget	G Deployed (ms)				lrn	MIN, fewest in the step
			(ms)	law	meas.	dk		
EPYC 7302*	0.14	gpt-oss 11%	4.3	12.2	17.7	1.06 [1.01, 1.10]	0.97 [0.92, 1.01]	0.85 [0.82, 0.88]
EPYC 7302*	0.14	gpt-oss 25%	4.7	8.7	12.2	1.08 [1.00, 1.17]	1.10 [1.06, 1.14]	–
EPYC 7302*	0.14	Qwen3 12.5%	5.8	20.4	25.4	1.07 [1.02, 1.13]	0.93 [0.86, 1.00]	–
EPYC 7302*	0.14	Qwen3 25%	6.3	14.3	17.2	1.01 [0.93, 1.08]	0.84 [0.76, 0.94]	–
Xeon 8347C*	0.23	gpt-oss 11%	4.1	12.0	19.5	1.28 [0.99, 1.56]	1.08 [0.80, 1.45]	1.33 [1.00, 1.61]
Xeon 8347C*	0.23	gpt-oss 25%	4.5	8.6	10.1	1.06 [1.05, 1.08]	1.02 [0.99, 1.04]	–
Xeon 8347C*	0.23	Qwen3 12.5%	5.3	21.4	45.0	1.96 [1.67, 2.24]	1.29 [0.94, 1.74]	–
Xeon 8347C*	0.23	Qwen3 25%	5.8	14.9	16.6	0.86 [0.71, 1.02]	0.73 [0.59, 0.88]	–
Pf again (285K)	0.29	gpt-oss 11%	4.3	13.8	13.7	1.02 [1.01, 1.04]	1.03 [1.02, 1.03]	1.02 [1.01, 1.03]
Pf again (285K)	0.29	gpt-oss 25%	4.7	9.6	9.5	1.06 [1.04, 1.07]	1.05 [1.04, 1.07]	–
Pf again (285K)	0.29	Qwen3 12.5%	5.4	24.8	24.4	1.02 [1.00, 1.03]	1.02 [1.01, 1.03]	–
Pf again (285K)	0.29	Qwen3 25%	5.9	16.9	16.2	1.04 [1.03, 1.04]	1.03 [1.02, 1.03]	–
Threadripper 9960X again	0.32	gpt-oss 11%	4.1	9.1	9.4	1.00 [0.99, 1.02]	0.98 [0.97, 1.00]	1.06 [1.05, 1.06]
Threadripper 9960X again	0.32	gpt-oss 25%	4.4	6.9	7.0	1.02 [1.00, 1.04]	0.99 [0.97, 1.00]	–
Threadripper 9960X again	0.32	Qwen3 12.5%	5.4	15.1	14.4	1.02 [1.01, 1.02]	0.99 [0.99, 1.00]	–
Threadripper 9960X again	0.32	Qwen3 25%	5.7	11.2	10.5	1.01 [1.01, 1.01]	0.98 [0.98, 0.99]	–
Pe again (9950X)	1.02	gpt-oss 11%	4.5	20.4	20.5	1.01 [1.01, 1.01]	1.01 [1.01, 1.02]	1.34 [1.33, 1.35]
Pe again (9950X)	1.02	gpt-oss 25%	4.5	12.8	12.4	1.01 [1.01, 1.01]	1.01 [1.00, 1.02]	–
Pe again (9950X)	1.02	Qwen3 12.5%	5.8	35.0	36.7	1.01 [1.01, 1.01]	1.02 [1.02, 1.02]	–
Pe again (9950X)	1.02	Qwen3 25%	5.9	22.2	22.5	1.01 [1.01, 1.01]	1.01 [1.01, 1.02]	–

430 tok/s (279 with the measured GPU rate) and the cache runs 109 tok/s. With every expert in its slots the cache matches stock llama.cpp (0.997 $\times$ and 1.047 $\times$). The bound assumes an expert run on the CPU is read from DRAM (MIN’s shortest re-read distance is 0.5 GB of host traffic) and that computation runs at memory speed. The routing traces are each model’s own greedy text on the 30 prompts; replayed through a simulator of the deployed policy they reproduce the engine’s hit rates within 5.3 points.

Reaching the read time. The microbenchmark of Section 3 (`jobs/ec2/readsched.cu`) takes MIN’s misses per layer and step from `scripts/readsched.py` and runs no model. Usable cores minus two CPU threads read their share of host memory from an ordinary buffer, as the engine’s helpers do, and the copy engine copies the rest, 13.25 MB at a time, from pinned memory into GPU slots; both buffers rotate through 4 GiB so that no cache holds them. *Per layer* splits each layer’s experts, whole, between the two paths at the probe’s rates and waits for both before the next layer; *per token* sums the same split over the token and waits once; *link only* sends every read over PCIe, as admitting every miss does (MIN with one read sends only its admissions); *CPU only* reads everything on the CPU. Table 21 gives each host. Waiting per layer instead of per token costs at most 1 point of the host term on the two AMD hosts and 7 on the Intel host, whose link is the slowest; there, reading everything on the CPU beats the per-layer split. The split ignores which reads are admissions; Section 3 gives the link-aware bound, which sends every admission over the link (`scripts/linkaware.py`, `prereg/linkaware.json`).

Choosing the split per machine. On host B, the model’s table beats the fixed table we used before at every budget of both models, by 3.3–8.0% (launch 1). Table 22 shows why at one budget: a single missed expert should run on the CPU (about 155 μ s, hidden behind the GPU’s other experts) rather than be copied (about 225 μ s, on the critical path); that entry alone is worth 3.7%. Fetching everything, FreeToken’s offload strategy expressed in our engine, is the worst table, and running everything on the CPU the second worst. The model’s table was within 1% of the best of five alternatives on launch 1 and tied with it on launch 2. Given each run’s measured misses, Eq. (5) ranks the tables correctly except law against lite, which are within 1%, and overstates their differences about 1.8 $\times$.




Figure 9. Ablation at 25% of experts on the GPU (15 AIME-25 problems, 256 tokens, host A). Each step adds one component to the one above; the hatched row is not a step: a static cache of each layer’s 32 most requested experts counted on the measured problems themselves.

Machines with half the link. On RTX 5090s whose links we measured at 27–30 GB/s instead of 53 (jobs 085, 085b and 088: gpt-oss next to a Core i9-14900K, Qwen3 next to a Ryzen 9 7900 in a PCIe 4.0 slot), the model’s tables fetch less (gpt-oss 0, 0, 0, 1, 1, Qwen3 0, 0, 1, 1, 2, 2, 3, 3, 4) and beat the fixed table by 21–34% on gpt-oss and 9–18% on Qwen3, about twice the gain on host B for Qwen3 and four to six times for gpt-oss: the per-machine split matters most where the link is slow. On the Ryzen, ours leads FreeToken by 1.04–1.08 $\times$; with the fixed table it would trail at 12.5%. On the Core i9 the first comparison (job 085) was confounded: FreeToken’s hybrid backend ran gpt-oss at 12–66 tok/s, below its own offload backend (25–86), because its CPU executor spread 23 threads across the 8 performance and 16 efficiency cores and its calibration then sent two thirds of the misses over the slow link, the same failure Appendix C reports for llama.cpp; the 1.4–2.0 $\times$ lead that gave us measured a thread configuration. We reran it (job 088, a second rental of the same CPU model) with FreeToken’s executor on 8 threads pinned to the performance cores: its hybrid backend then runs 1.4–3.1 $\times$ faster than before (37.5, 62.1 and 92.5 tok/s; its calibration now copies 51.7% of misses), beating its own offload backend at 11 and 25% and tying it at 40%, and ours leads its better configuration by 1.58, 1.52 and 1.51 $\times$ (lowest interval bound 1.47). Restricted to the CPU alone, FreeToken’s 8-thread executor decodes 13.7 tok/s at every budget. llama.cpp on the same rental runs 40.0 tok/s with one thread per performance core (–9.4% from the model’s prediction) and 32.8 with all 24 threads (+25.9%): the hybrid-core failure is a thread-count choice, and the model holds once the choice is right. The law predicted llama.cpp on the Ryzen within +2.7%. Figure 6 puts every same-machine comparison on one axis.

Long outputs. Our lead does not come from the first 256 tokens. On gpt-oss, on ten held-out hard problems (AIME 2022, problems 11–15 of both exams) that neither system has finished reasoning about after 2,048 tokens (at most 1 of 10 reach an answer), ours runs 1.36 [1.34, 1.38], 1.32 [1.30, 1.34] and 1.19 [1.17, 1.21] times FreeToken at 11, 25 and 40% over all 2,048 tokens, and 1.37 [1.34, 1.39], 1.33 [1.31, 1.35] and 1.19 [1.17, 1.22] over tokens 257–2,048 (job 083, host A); both engines stream one event per token. On Qwen3, over tokens 257–2,048 of ten held-out MATH-500 problems, ours runs 1.043 [1.020, 1.070], 1.056 [1.036, 1.077] and 1.054 [1.025, 1.086] times FreeToken at 12.5, 25 and 43.75%. Both leads are larger than over the same runs’ first 256 tokens. Turning off FreeToken’s prefill overlap, which lends it cache slots during each prompt, changes its overall speed by 1% or less.

FreeToken’s own headline model. On FreeToken’s own headline model, Qwen3.6-35B-A3B in BF16 (256 experts of 6.3 MB per layer, hybrid linear attention), the two systems are close: ours runs 1.030 [1.023, 1.038] times FreeToken at 12.5% and ties at 25 and 37.5% (1.010 [0.998, 1.021], 1.003 [0.994, 1.013]), at 1.98–2.48 $\times$ llama.cpp (job 086, host A). FreeToken as shipped, with its automatic backend and cache size, runs this model at 93.8 tok/s on that machine, above the 77–83 tok/s its paper reports on an RTX 5090, so our setup does not handicap it; FreeToken’s fetch cap and thread count were left at their defaults, while our FETCH table was swept. Our lead is therefore model-dependent: largest with few large

Table 20. The bound under each tightening, tokens per second on host B, and our cache as a percentage of the bound as published (exact optimum, highest probed host rate, datasheet GPU rate) and with every tightening at once. *no-evict*: Belady’s MIN with bypass that keeps a served expert for the step; *exact*: the same, free to evict an expert after serving it; *pool*: one pool of LC slots shared across layers; *median*: B_{host} the median of six concurrent probe samples instead of the highest; *GPU*: the GPU term from the measured all-in-VRAM time instead of the datasheet rate; *layers*: the per-layer sum, which no system without cross-layer prefetch can beat; *all*: exact, pool, median and GPU together (a warm start, not included, changes the bound by at most 1.4%).

Model	Experts on GPU	no-evict	exact	pool	median	GPU	layers	all	Ours, % of exact	% of all
gpt-oss-120b	11%	169	172	182	159	172	163	168	41	41
gpt-oss-120b	25%	427	430	474	397	279	365	278	25	39
gpt-oss-120b	40%	519	519	519	517	279	500	278	29	55
Qwen3-30B-A3B	12.5%	92	96	101	89	96	93	93	42	43
Qwen3-30B-A3B	25%	209	214	234	197	192	193	191	30	33
Qwen3-30B-A3B	43.75%	308	308	308	307	192	298	191	35	57

Table 21. Can the bound’s read time be reached (job 102)? MIN’s per-layer host reads of gpt-oss (first 4,000 steps of the AIME routing), replayed without the model through CPU threads and the copy engine. Each mode’s share of the bound’s host term (R^*S/B_{host} , ms per token in the column *Term*); *Model*: the per-layer mode predicted from the probe’s rates. Two repetitions of each mode agree within 0.9%.

Host	Link/ CPU	C	Term (ms)	Split per layer	Split per token	Link only	CPU only	Model
Pd again	0.40	14	7.04	86%	93%	35%	92%	98%
Pd again	0.40	32	2.84	89%	93%	35%	92%	98%
9950X	1.02	14	9.70	96%	97%	95%	86%	94%
9950X	1.02	32	3.91	96%	97%	94%	86%	92%
9800X3D	0.55	14	9.56	95%	95%	53%	91%	95%
9800X3D	0.55	32	3.85	93%	94%	53%	91%	94%

experts (gpt-oss, top-4 of 128), smallest with many small ones.

FreeToken tuned. On an RTX 5090 with a Ryzen 9 9950X (job 098, one launch, configurations in a shuffled order per cell) FreeToken ran in five settings per cell: hybrid with its calibrated fetch split (its default), hybrid with the fetch cap fixed at one and at two, hybrid with 8 CPU threads instead of one per physical core, and offload. Its best setting was the calibrated hybrid with 8 threads at three cells and offload at the other three (gpt-oss 25 and 40%, Qwen3 43.75%), at most $1.04\times$ its Table 2 setting; the fixed caps were never better than its own split by more than 0.3%. Ours over FreeToken’s best: 1.17, 1.17 and 1.09 on gpt-oss, 1.02, 1.04 and 0.98 on Qwen3, within 0.04 of host S’s ratios at every cell. With every expert on the CPU, stock llama.cpp decodes gpt-oss $1.13\times$ and Qwen3 $0.99\times$ as fast as FreeToken’s CPU backend.

Other systems and outputs. Two newer llama.cpp-based systems, Pipelined Sharding (Ukarande et al., 2026) and leloch’s expert cache (leloch, 2026a), run gpt-oss-120b at $1.28\text{--}1.45\times$ stock llama.cpp; ours ran $1.42\text{--}2.34\times$ the better of them (job 077). KTransformers (Chen et al., 2025a) (kt-kernel 0.7.0, exact BF16 CPU experts, the same number of experts statically on the GPU) ran Qwen3 in BF16 at $0.98\text{--}1.01\times$ stock llama.cpp on an RTX 5090 with a Ryzen 9 9950X, and ours at $2.2\text{--}4.4\times$ it (job 078); both static placements read the same CPU expert bytes from the same DRAM. Its published speeds

Table 22. FETCH tables at 43.75% on Qwen3-30B-A3B (launch 1): entry n is how many of a layer’s n missed experts are copied over PCIe; the rest run on the CPU. Ratio to the fixed table, paired, 95% CI.

Table	$f(0..8)$	tok/s	vs fixed
Lite	0, 0, 1, 1, 1, 2, 2, 3, 3	109.0	1.057 [1.052, 1.062]
Model	0, 0, 1, 1, 2, 3, 3, 4, 5	108.0	1.048 [1.041, 1.055]
Model, fetch a single miss	0, 1, 1, 1, 2, 3, 3, 4, 5	104.1	1.010 [1.007, 1.012]
Fixed	0, 1, 1, 2, 3, 3, 4, 5, 6	103.1	1
None (CPU only)	0, . . . , 0	100.7	0.977 [0.970, 0.984]
All (fetch every miss)	0, 1, . . . , 8	98.7	0.957 [0.953, 0.962]




Figure 10. Published MoE-offloading measurements (52 rows from 13 sources) as a percentage of their own bound, computed from the paper’s hardware at datasheet rates and, where a routing trace of the model exists, the per-layer Belady optimum (filled); hollow points assume uniform routing, which overstates the optimum’s reads and so the percentage; the dotted lines are the class medians (no GPU expert capacity, 49.0%; every expert resident, 4.0%). Right: our cache, FreeToken and stock llama.cpp on host B against the probed per-layer limit of Table 2 (the pooled bound FreeToken’s and llama.cpp’s designs allow is 0–10% higher).

rely on AMX Xeons and INT4/INT8 CPU weights, which that host lacks and which change the model, so this is not its best setting. Teacher-forced on fixed text (1536 steps per model) against stock llama.cpp with every weight in the VRAM of an RTX PRO 6000 (job 090), our cache at 25% has a mean KL divergence over the full vocabulary of 0.0019 nats on gpt-oss and 0.0005 on Qwen3 (99.9th percentile 0.121 and 0.059, maximum 0.21 and 0.13), picks the same next token on 98.5% and 99.3% of steps, and changes the loss by +0.16% and +0.19%. Stock llama.cpp’s own `--n-cpu-moe` placement diverges from the same reference by as much (mean KL 0.0017 and 0.0008, agreement 98.8% and 99.3%); the differences are those of running experts on the CPU (the helpers quantise activations to Q8_0 where the GPU path uses Q8_1), not of the cache. Task accuracy was not measured.

Prompts and batches. The cache serves single-token steps; prompts and batches above one run llama.cpp’s stock graph with every expert in host memory, and pay for the cache’s resident experts and its one-thread CPU side. In llama-batched-bench at 25% on host S (job 089, both models) our prefill at 512 and 2,048 tokens runs 79–83% of stock llama.cpp’s and our decode at 2 and 4 parallel sequences 77–79% of it (the tool feeds both engines the same random tokens).

L THE AUDIT’S ROWS

Table 23 lists the 52 published measurements of Section 3 with their bounds. The sources are Fiddler (Kamahori et al., 2024), ProMoE (Song et al., 2024a), KTransformers (Chen et al., 2025a), HybriMoE (Zhong et al., 2025), Fate (Fang et al., 2025), SP-MoE (Chen et al., 2025b), MoE-SpeQ (Wang et al., 2025), the CPU–GPU collaborative inference of Huang et al. (2025), FreeToken (Yang et al., 2026), SeqMoE (Wang et al., 2026b), NVIDIA’s pipelined sharding (Ukarande et al., 2026), a llama.cpp fork (leloch, 2026a) and a llama.cpp pull request (csantiago78, 2026). The in-class subset of Section 3 (20 rows) is the trace-scored rows less the lossy (Fate) and speculative (SP-MoE, MoE-SpeQ) ones; over all 29 trace-scored rows the median is 9.5%. These systems mostly read on demand, so Eq. (2) would place them higher. Three caveats apply to every row: the ceilings are datasheet ones where ours are probed; the trace is the model’s own sampled text, sometimes of a sibling variant (an FP8 trace for a Q4 deployment); and systems whose cache is pooled across layers, whose placement pins whole layers, or which verify several tokens per pass are scored against the per-layer one-token bound, which is not theirs.

Table 23: The audit’s rows (Section 3): every published batch-1 measurement we scored, with the system’s stated hardware, its experts per layer resident on the GPU (C), the model’s experts and top- k , the published speed, the bound at datasheet-physical ceilings (GPU datasheet bandwidth, the top of the host-DRAM band, datasheet PCIe, no latencies; `scripts/audit.py`), how the optimum’s reads were obtained (*trace*: Belady’s MIN with bypass on an on-policy routing trace of the model; *uniform*: independent uniform routing, which overstates the reads and so the percentage; *no slots*: no GPU expert capacity, the CPU-only rate is the bound; *all fit*: every expert resident, nothing offloaded) and the percentage of the bound. Medians: trace 9.5% ($n=29$), uniform 34.8% (11), no slots 49.0% (8), all fit 4.0% (4), all rows 14.5% (52). Rows below 5% are more likely configuration mismatches with the bound’s assumptions (precision, dense layers on the CPU, end-to-end timing) than inefficiency, and are not adjudicated.

System	Model	GPU	C	E/k	tok/s	bound	reads	%
Fiddler (Kamahori et al., 2024)	Mixtral-8x7B	Quadro RTX 6000	1	8/2	2.5	17	trace	14.6
	Mixtral-8x7B	RTX 6000 Ada	3	8/2	8.3	57	trace	14.4
ProMoE (Song et al., 2024a)	DeepSeek-V2-Lite	RTX 4090	45	64/6	49.9	216	trace	23.1
	Qwen1.5-MoE-A2.7B	RTX 4090	46	60/4	67.2	216	trace	31.1
	Qwen2-57B-A14B	RTX 4090	42	64/8	26.9	136	trace	19.8
KTransformers (Chen et al., 2025a)	Mixtral-8x7B	RTX 4090	6	8/2	27.0	152	trace	17.8
	DeepSeek-V3-0324	A100 40GB	0	256/8	5.9	12	no slots	48.8
	DeepSeek-V2.5-1210	A100 40GB	0	160/6	11.9	24	no slots	49.2
	Qwen2-57B-A14B	A100 40GB	0	64/8	22.9	35	no slots	64.9
	DeepSeek-V3-0324	RTX 4080	0	256/8	16.1	33	no slots	48.3
	DeepSeek-V2.5-1210	RTX 4080	0	160/6	17.6	32	no slots	55.3
	Qwen2-57B-A14B	RTX 4080	0	64/8	37.0	48	no slots	76.9
HybriMoE (Zhong et al., 2025)	DeepSeek-V2-Lite-Chat	RTX A6000	16	64/6	14.5	470	trace	3.1
	DeepSeek-V2-Lite-Chat	RTX A6000	48	64/6	15.3	655	trace	2.3
	Mixtral-8x7B-Instruct	RTX A6000	2	8/2	4.4	53	trace	8.2
	Mixtral-8x7B-Instruct	RTX A6000	6	8/2	7.3	133	trace	5.5
	Qwen2-57B-A14B-Instruct	RTX A6000	16	64/8	9.6	86	trace	11.1
	Qwen2-57B-A14B-Instruct	RTX A6000	48	64/8	9.0	116	trace	7.7
Fate (Fang et al., 2025)	Qwen1.5-MoE-A2.7B	RTX 3090	46	60/4	14.0	198	trace	7.1
	Qwen1.5-MoE-A2.7B	GTX 1080 Ti	15	60/4	3.7	76	trace	4.9
SP-MoE (Chen et al., 2025b)	Mixtral-8x7B	RTX 3090	1	8/2	1.2	16	trace	7.7
	Mixtral-8x7B	RTX 4090	1	8/2	1.3	13	trace	9.5
	Mixtral-8x7B	A100 40GB	3	8/2	1.5	32	trace	4.5
	Phi-3.5-MoE	RTX 4090	3	16/2	3.3	57	trace	5.8
	DeepSeek-V2-Lite	A100 40GB	7	64/6	5.6	142	trace	3.9
MoE-SpeQ (Wang et al., 2025)	DeepSeek-V2-Lite	A100 40GB	64	64/6	10.1	335	all fit	3.0
	DeepSeek-V2-Lite	A100 40GB	64	64/6	13.8	335	all fit	4.1
	Qwen1.5-MoE-A2.7B	A100 40GB	60	60/4	13.2	335	all fit	3.9
	Qwen1.5-MoE-A2.7B	A100 40GB	60	60/4	28.3	335	all fit	8.5
	Phi-3.5-MoE	A100 40GB	7	16/2	6.0	129	trace	4.7
	Phi-3.5-MoE	A100 40GB	7	16/2	6.1	129	trace	4.7
CPU–GPU collab. (Huang et al., 2025)	Mixtral-8x7B	RTX 4090	1	8/2	4.8	13	trace	36.5
	Phi-3.5-MoE	RTX 4090	3	16/2	10.4	56	trace	18.6

continued

Where the Seconds Go

System	Model	GPU	C	E/k	tok/s	bound	reads	%
FreeToken (Yang et al., 2026)	Qwen3.6-35B-A3B	RTX 5090	94	256/8	77.1	227	uniform	34.0
	DeepSeek-V4-Flash	RTX 5090	28	256/6	24.9	81	uniform	30.8
	Qwen3.6-35B-A3B	RTX 5090	104	256/8	73.8	205	uniform	36.1
	Qwen3.6-35B-A3B	RTX 4090	72	256/8	42.9	123	uniform	34.8
	Qwen3.6-35B-A3B	RTX 4060 Laptop	60	256/8	39.3	97	uniform	40.4
SeqMoE (Wang et al., 2026b)	GLM-5.2	RTX PRO 6000 Blackwell	36	256/8	14.9	30	uniform	49.9
	Qwen3-30B-A3B-FP8	RTX 4090	57	128/8	104.1	374	trace	27.9
	Qwen3-30B-A3B-FP8	RTX 4090	19	128/8	47.9	374	trace	12.8
	Qwen3.6-35B-A3B-BF16	RTX 5090	102	256/8	118.3	305	uniform	38.8
	Qwen3.6-35B-A3B-BF16	RTX 5090	38	256/8	70.2	212	uniform	33.2
	GPT-OSS-120B-MXFP4	RTX PRO 6000 Blackwell	57	128/4	139.2	378	trace	36.8
	DeepSeek-V4-Flash	RTX PRO 6000 Blackwell	115	256/6	57.4	183	uniform	31.3
	Qwen3-30B-A3B-Instruct-2507	RTX 5090	0	128/8	25.7	199	no slots	12.9
Pipelined sharding (Ukarande et al., 2026)	Qwen3-30B-A3B-Instruct-2507	RTX 5090	44	128/8	32.1	1032	trace	3.1
	Qwen3-235B-A22B-Instruct-2507	RTX 5090	0	128/8	7.7	40	no slots	19.3
	Qwen3-235B-A22B-Instruct-2507	RTX 5090	42	128/8	11.5	108	uniform	10.7
	Qwen3-30B-A3B-Instruct-2507	RTX 5070 Ti	107	128/8	54.9	511	trace	10.8
llama.cpp fork (leloch) (leloch, 2026a)	Laguna S 2.1	RTX 5090	66	256/10	31.3	88	uniform	35.5
llama.cpp PR 27861 (csantiago78, 2026)	Qwen3-30B-A3B Q4_K_M derivative	Radeon RX 7600 8GB	32	128/8	16.5	169	trace	9.8

M ARTIFACT DESCRIPTION

- **System.** A patch to llama.cpp commit 4da6337 (the expert cache, FETCH, the mailbox and the llama-ec-bench instrumentation: teacher-forced loss, logit dumps, routing lookahead records); built with CUDA 13.0 for compute capability 12.0.
- **Jobs.** One shell script per job; each installs FreeToken at commit 0d652e7, downloads the models from Hugging Face, runs the bandwidth probes, writes any on-machine prediction, runs the client, and stores raw rows (one JSON line per problem), server logs and engine counters. A small runner rents a Vast instance, streams a heartbeat, fetches the results and only then destroys the instance.
- **Analysis.** scripts/*_stats.py compute every interval from the raw rows; scripts/perlayer_model.py, speed_limit_v2.py, shapley_gap.py, policy_study.py, foresight_exact.py, batch_k_trace.py, fig_audit_sol.py and scorecard.py produce the model comparison, the limit, the accounting, the trace studies, the audit figure and the scorecard; scripts/wsg_tables.py and wsg_numbers2.py generate every table, figure and number in this paper from their outputs. The audit’s 52 published rows and the physical speed-of-light model they are scored against are data/published_measurements.json, data/audit/ and scripts/audit.py.
- **Cost.** The 88 rentals (53 distinct machines) of jobs 058–101 cost 70.8 US dollars at their hourly GPU prices, by the ledger in the repository; storage and download charges, which the ledger omits, add to this.