

# Where the Seconds Go:

## A Speed Limit and an Accounting for Mixture-of-Experts Decode on One Consumer GPU

Harshith Kantamneni  
Independent Researcher

### Abstract

A Mixture-of-Experts model too large for one consumer GPU keeps most experts in host memory; each decoded token waits for the ones it needs. How fast could it decode, and what separates real systems from that? A *host-memory law*, a GPU term plus host bytes over the probed rate of the paths that read them, predicts decode time blind on 7 machines with 3.7% median error. Belady’s optimum with bypass on the model’s routing gives a *speed limit* in seconds for any exact-routing system with a per-layer expert budget: llama.cpp, FreeToken and our llama.cpp expert cache reach 9–20, 20–40 and 25–43% of it on two RTX 5090 hosts; 52 published measurements, a median 9.5% of theirs. A Shapley *accounting* says the rest is mostly missing routing foresight where the host binds and serialised host reads where the GPU does. In the engine, the limit’s own policy given future routing runs 13–23% faster at expert budgets of 25% and above on the bytes it saves, and a hit-optimal oracle 25–63% faster by overlapping copies with the step; at the lowest budgets the first gains 3% at most and the second loses. Nine models’ traces give the horizon foresight needs,  $W_{50} \approx 0.59 (C/k)^{1.33}$  tokens, which speculative batching does not supply. The cache leads FreeToken at 11 of 12 cells (–3 to +29%) and runs 2.0–4.0× llama.cpp. Every prediction was committed before its measurement; 61% of 243 scored clauses held.

## 1 Introduction

Open-weight Mixture-of-Experts models now carry 30–120B parameters and activate a few billion per token; a desktop GPU has 24–32GB. Running such a model at batch 1 therefore means keeping most experts in host DRAM, and each decoded token then reads its routed experts from one of three places: GPU memory, host memory by the CPU that runs the expert, or host memory across PCIe into a GPU slot [16, 12]. A line of systems has built on this: GPU expert caches with demand fetch [8, 41, 34], CPU execution of misses [18, 5, 51, 14, 50, 53], permiss choices between the two [44], and expert prediction [15, 52, 39]. Each reports a speed-up over a baseline of its choosing. None states how fast its machine *could* decode the model within the same GPU memory, or how

much of the remaining time any design could remove, and the trace-level evaluation contracts proposed for this field [49] count cache misses, not seconds.

We answer those questions on consumer hardware (RTX 5090, 4090 and 3090 hosts with desktop CPUs, rented by the hour) with two models, gpt-oss-120b [28] and Qwen3-30B-A3B [43], routing traces of nine, and 52 published measurements from 13 sources. The method is Hoefer and Belli’s [13]: a performance model with bounds, statistically sound comparisons, and every prediction committed to a public repository before the machine it concerns was started (Appendices A and B).

### Contributions.

- **A host-memory law, validated as a predictor** (Section 3): time per token =  $G + \max(X_c/B_c, X_p/B_p, (X_c+X_p)/B_{cp})$ , host bytes  $X$  from the engine’s counters and rates  $B$  from a probe, predicted blind on 7 machines and tested under leave-one-host-out validation against the additive and per-layer forms, with each failure traced to a measured cause.
- **A speed limit in seconds** (Section 4): a bound on any exact-routing system with  $C$  expert slots per layer, from the exact Belady optimum with bypass, two host read paths and measured ceilings, tightened four ways, reported beside every measured speed on four machines, and applied to 52 published measurements (Section 6).
- **An accounting, and its check** (Section 4): Shapley attribution of the gap to the limit, which no ordering of fixes can bias, with the regime it predicts (foresight where the host binds, overlap where the GPU does), and its largest term tested in the engine with two oracles of future routing (worth 13–23% on the bytes it saves and 25–63% with the overlap it allows, at budgets of 25% and above; at most 3% below).
- **What the traces say** (Section 5): across nine models, which policy earns its reads, how fast the hot set drifts, how many tokens of foresight are worth half the optimum, and what batched speculative verification cannot buy.

- **The instrument** (Section 7): a llama.cpp expert cache whose per-machine miss split comes from the law, measured against FreeToken and llama.cpp on two RTX 5090 hosts, an RTX 4090 and an RTX 3090, with an ablation of its parts and a measurement of its scope.

**What we do not claim.** Running missed experts on the CPU [18, 5], caching hot experts with CPU-executed misses [51, 14, 21], splitting a layer’s misses between CPU and PCIe from a bandwidth model [44, 53], a PCIe turnover ceiling [47], bytes over bandwidth as a bound [46], Belady with bypass as a cache-level oracle [49] and a table of systems against the all-in-VRAM speed [39] are all published. Our system uses the first three; what is new is the limit in seconds, the accounting and its measurement, the trace results, the audit, and the measurements themselves.

## 2 Setting and Method

**Workload.** Single-request decode. gpt-oss-120b (MXFP4 experts; 36 MoE layers of 128 experts, top-4, 13.25 MB per expert) and Qwen3-30B-A3B in BF16 (48 layers of 128 experts, top-8, 9.44 MB per expert). Prompts are the 30 AIME-25 problems in FreeToken’s own benchmark format; each system decodes 256 tokens per problem, greedily, after one warm-up request on a fixed held-out text, in one session in which the cache carries over between problems. One client measures every system with FreeToken’s own metric, (tokens−1) over the time between the first and last token’s arrival.

**Machines and probes.** RTX 5090 (32 GB, 1.79 TB/s datasheet) hosts with desktop CPUs, rented by the hour. Each job first runs a probe (`concur`) that reads host memory from  $t$  CPU threads, over PCIe with zero-copy kernels and with the copy engine, and from both at once; FreeToken’s own probe (`ft bench bw`) runs too. We call the CPU rate at the helper thread count  $B_c$ , the zero-copy PCIe rate  $B_p$  and the combined rate  $B_{cp}$ . GPU clocks cannot be locked on these rentals; we report each card’s power limit. Table 1 is measured on two RTX 5090 hosts: *listing B*, a Ryzen 9 9950X3D rental whose card reports a 17001 MHz memory clock,  $1.21\times$  the bandwidth the limit assumes, against 14001 MHz on 25 of the 26 other RTX 5090 rentals that recorded one; and the *stock-clock host*, a Ryzen 9 9950X rental at the common clock with a 22% slower CPU. The ablation and the long-output runs use *listing A*, a second rental of the 9950X3D model at the common clock.

**Equal memory.** llama.cpp keeps whole layers’ experts on the GPU (`--n-cpu-moe n`), our cache keeps  $C$  experts of every layer, and FreeToken takes a cache rate; we match the number of experts on the GPU (within 2.5%)

and report measured VRAM. FreeToken preallocates 28.4–30.1 GiB whatever its rate.

**Statistics and protocol.** Ratios are of mean speeds, paired by problem, with 95% percentile intervals from 10,000 bootstrap resamples; they cover problem sampling within one launch on one machine, not rentals (two rentals of one CPU model differed by 19–22% in host bandwidth). Where a system has variants (our FETCH tables, FreeToken’s offload and hybrid backends), a first launch picks each system’s faster variant per budget and a second launch, which the comparison uses, runs only those (one row of Table 1 has no second launch; on the stock-clock host FreeToken’s variants were carried over from listing B’s selection rather than re-picked). Every job’s predictions are in the header of its script, committed and pushed before the machine starts; Appendix B scores every clause under one rule.

## 3 A Host-Memory Law

With experts partly in host memory, a decode step does two kinds of work: GPU work that every token needs (attention, dense projections, resident experts, the output head, launches), and reads of the experts that are not resident. The second is limited by host memory. The CPU reads  $X_c$  bytes per token (the experts it runs), PCIe reads  $X_p$  bytes (the experts copied to the GPU, whether on demand or ahead of time), and when both run at once they share the DRAM:

$$T = G + \max\left(\frac{X_c}{B_c}, \frac{X_p}{B_p}, \frac{X_c + X_p}{B_{cp}}\right). \quad (1)$$

$X_c$  and  $X_p$  come from the engine’s counters of misses, fetches and admissions, which are counts, not timings.  $B_c$ ,  $B_p$  and  $B_{cp}$  come from the machine’s probe.  $G$  is one constant per engine and model, fitted once on one host (4.82 ms for our cache on gpt-oss-120b, 5.06 ms for llama.cpp) and then frozen in the repository. The form is FreeToken’s residual-bandwidth split of one layer [44], lifted to the token; the difference is that the three rates are probed rather than assumed to add, and that the law is tested as a predictor.

**Blind test.** On every later machine the job wrote the law’s prediction from its own probe before any model ran (Table 2). Over 29 configurations of our cache on 7 machines (33 measurements; one machine ran twice) the median error is 3.7% per configuration (3.4% per measurement); without the 12-channel EPYC 9655 it is 3.3% over 24 configurations (90th percentile 7.5%). llama.cpp `--n-cpu-moe` is predicted within 4% on the Ryzen hosts, including two Qwen3 configurations predicted on the headline machine (+3.6%, +4.0%; for Qwen3,  $G=4.3$  ms, the mean residual of three llama.cpp runs on another host).

Table 1: Decode speed at equal GPU expert memory on two RTX 5090 hosts. 30 AIME-25 problems, first 256 decode tokens, greedy, session; our cache uses the FETCH split computed from each machine’s bandwidth probe; FreeToken uses its faster backend per budget, picked on listing B on a separate launch and carried over to the second host. Ratios are of mean speeds, paired by problem, with 95% bootstrap intervals. *Speed limit*: the fastest any exact-routing system with the same slots per layer can decode on that machine (Section 4: exact optimum, that host’s highest probed rate, datasheet GPU rate; Table 3 tightens it). With every weight in the VRAM of an RTX PRO 6000 (the RTX 5090’s datasheet bandwidth, 96 GB), stock llama.cpp decodes gpt-oss at 255 and Qwen3 at 178 tok/s. <sup>†</sup>Launch 1 for both systems: FreeToken’s backend was picked on the launch it is scored on, which favours it, and the row has no confirmation launch.

| Model                                                                                                                                         | Experts on GPU | llama.cpp (tok/s) | FreeToken (tok/s)     | Ours (tok/s) | Ours ÷ FreeToken (95% CI)         | Ours ÷ llama.cpp | Speed limit | Ours, % of limit |
|-----------------------------------------------------------------------------------------------------------------------------------------------|----------------|-------------------|-----------------------|--------------|-----------------------------------|------------------|-------------|------------------|
| <b>Listing B:</b> Ryzen 9 9950X3D, card at 17001 MHz memory clock, CPU 72 GB/s, link 53, together 78                                          |                |                   |                       |              |                                   |                  |             |                  |
| gpt-oss-120b                                                                                                                                  | 11%            | 34.9              | 54.0 (hybrid)         | <b>69.9</b>  | 1.294 [1.278, 1.312]              | 2.00×            | 172         | 41%              |
| gpt-oss-120b                                                                                                                                  | 25%            | 40.0              | 85.6 (hybrid)         | <b>109.2</b> | 1.275 [1.254, 1.295]              | 2.73×            | 430         | 25%              |
| gpt-oss-120b                                                                                                                                  | 40%            | 47.7              | 132.2 (offload)       | <b>152.5</b> | 1.154 [1.134, 1.172]              | 3.20×            | 519         | 29%              |
| Qwen3-30B-A3B                                                                                                                                 | 12.5%          | 19.5              | 38.7 (hybrid)         | <b>40.0</b>  | 1.032 [1.022, 1.042]              | 2.05×            | 96          | 42%              |
| Qwen3-30B-A3B                                                                                                                                 | 25%            | 22.1              | 54.7 (hybrid)         | <b>63.1</b>  | 1.153 [1.134, 1.173] <sup>†</sup> | 2.86×            | 214         | 30%              |
| Qwen3-30B-A3B                                                                                                                                 | 43.75%         | 28.4              | 103.1 (offload)       | <b>108.2</b> | 1.049 [1.036, 1.064]              | 3.81×            | 308         | 35%              |
| <b>Stock-clock host:</b> Ryzen 9 9950X, card at 14001 MHz, CPU 56 GB/s, link 53, together 66; FreeToken’s backend carried over from listing B |                |                   |                       |              |                                   |                  |             |                  |
| gpt-oss-120b                                                                                                                                  | 11%            | 28.9              | 48.0 (hybrid)         | <b>57.9</b>  | 1.207 [1.195, 1.219]              | 2.00×            | 140         | 41%              |
| gpt-oss-120b                                                                                                                                  | 25%            | 33.6              | 78.7 (hybrid)         | <b>94.1</b>  | 1.196 [1.163, 1.224]              | 2.80×            | 351         | 27%              |
| gpt-oss-120b                                                                                                                                  | 40%            | 40.1              | 122.2 (offload)       | <b>133.6</b> | 1.093 [1.077, 1.111]              | 3.34×            | 514         | 26%              |
| Qwen3-30B-A3B                                                                                                                                 | 12.5%          | 16.1              | 32.7 (hybrid)         | <b>33.5</b>  | 1.027 [1.018, 1.036]              | 2.08×            | 78          | 43%              |
| Qwen3-30B-A3B                                                                                                                                 | 25%            | 18.5              | 51.4 (hybrid)         | <b>53.9</b>  | 1.048 [1.037, 1.059]              | 2.91×            | 174         | 31%              |
| Qwen3-30B-A3B                                                                                                                                 | 43.75%         | 23.8              | <b>98.3</b> (offload) | 95.8         | 0.974 [0.962, 0.987]              | 4.02×            | 305         | 31%              |

Table 2: The law predicted blind: each machine wrote its prediction from its own bandwidth probe, with constants frozen in the public repository, before any model run. Error of predicted over measured speed, gpt-oss-120b (cache rows) and llama.cpp --n-cpu-moe (last column, %).

| Host CPU                   | Job  | $n$ | Median<br> err  % | Max<br> err  % | llama.cpp<br>err % |
|----------------------------|------|-----|-------------------|----------------|--------------------|
| Core Ultra 7 270K (8P+16E) | 069c | 5   | 1.5               | 11.9           | +119               |
| EPYC 7352 (Zen 2)          | 069c | 5   | 3.2               | 11.7           | +10                |
| EPYC 9655 (12-ch. DDR5)    | 069c | 5   | 18.7              | 31.9           | +16                |
| Ryzen 9 7900               | 072  | 5   | 4.9               | 8.2            | +4                 |
| Ryzen 9 9950X3D (A)        | 073  | 8   | 1.6               | 4.5            | –                  |
| Ryzen 9 7950X              | 074  | 1   | 4.0               | 4.0            | –                  |
| Ryzen 9 9950X              | 076  | 4   | 2.5               | 5.0            | –                  |
| Ryzen 9 9950X3D (B)        | 082  | 0   | –                 | –              | +4, +4             |

**Is the max the right form?** In 19 of the 29 configurations nothing is copied, so they test one roof. We therefore refit the law and its alternatives on the same 33 measurements under leave-one-host-out validation (each of the 7 hosts predicted with constants fitted on the other six): the token-level max of Eq. (1) with  $G$  refit has a median error of 3.1% (90th percentile 9.4%); the additive form  $G + X_c/B_c + X_p/B_p$ , 4.5% (11.3%); a per-layer max summed over layers from the engine’s per-layer miss histograms, 3.2% (10.2%). The additive form ties in sample but misses the three rows that copy misses on demand behind a 53 GB/s link, where both paths carry most of the bytes, by -11%, while the max is within  $\pm 0.2\%$  there; on the slow-link copy rows the additive form is the closer of the two. The per-layer form neither beats the max

nor explains its failures: the Jensen gap between a sum of per-layer maxima and the maximum of sums is a fifth of the Core Ultra’s residual, none of the EPYC 7352’s, and has the wrong sign on the EPYC 9655 and on the prefetch rows. We keep Eq. (1), with its evidence stated: the max beats the sum on three fast-link rows and ties elsewhere, so the form is settled only where both paths carry most of the bytes. FreeToken’s own runs are not predicted, because its counters are not exposed; the law covers our engine and llama.cpp.

**Where it fails, and why.** Each failure now has a measured cause. (i) Copying misses on demand (FETCH) at 25% is over-predicted by +11.7% on the Core Ultra 7 and +11.4% on the EPYC 7352 (links of 36 and 26 GB/s); +8.8% and +8.5% with  $G$  refit. The engine’s per-layer timers show why: a CPU miss served beside a copy takes 20–62% longer than one served alone, so the two paths share DRAM less evenly than the probe’s combined rate says; a variant of the law that uses the probe’s per-path rates under contention and a per-layer hand-off constant predicts both rows within +3.1% and +4.1%, at a worse median (4.2%). (ii) llama.cpp on a Core Ultra 7 with 8 performance and 16 efficiency cores runs at 46% of the law at its default thread count, because it splits every matrix evenly across its threads; our helpers claim work dynamically and stay on the law on the same machine. We did not test llama.cpp restricted to the performance cores. (iii) On a 12-channel EPYC 9655 (559 GB/s) host memory stops being the limit (+16 to +29%): the per-layer timers show a floor of 55–70  $\mu$ s per layer with CPU misses that no desktop host has (their intercepts are at or below zero),

the latency of the hand-off at a bandwidth where a 13 MB expert takes  $24\mu\text{s}$ . (iv) Prefetch configurations are under-predicted by -8.7 to -3.4%: their copies overlap attention, which the law charges in full.

**What the law says.** Every mechanism in this design space acts only through the bytes read from host memory per token and the number of read paths in use. A second path adds bandwidth on every host we measured but the 12-channel EPYC ( $B_{cp}$  against  $\max(B_c, B_p)$ ): 8–17% on the Ryzen hosts, 1% on the Core Ultra, 33% on the Zen 2 EPYC), and that is the whole gain available to fetching, prefetching or splitting misses between the CPU and PCIe. Prefetching never reduces the bytes.  $G$  (5.17 ms refit) exceeds the whole all-in-VRAM token (3.92 ms) because it absorbs a shortfall the probe does not see: the helpers read at  $1.09\text{--}1.31\times$  the probe’s time per expert on desktop hosts, and a fitted rate factor is collinear with  $G$  at this sample size. And the “host lottery” is the DRAM: two rentals of the same CPU differed by 19–22% in speed, which is 62 against 47 GB/s of host bandwidth.

## 4 The Speed Limit, and Where the Seconds Go

**Bound.** Consider any system that executes the model’s exact routing with at most  $C$  experts of each layer resident on the GPU. Per token, each of the  $Lk$  routed experts is either served by a resident copy (the GPU reads it from VRAM) or not, in which case its  $S$  bytes are read from host memory once more: by the CPU that runs it, or by the copy that brings it into a slot. Let  $R^*$  be the reads per token of Belady’s MIN with bypass [2, 9, 48] applied per layer to the routing trace, the fewest reads of any policy, even one that knows the future (an expert may be evicted after serving the current token; a policy that keeps every requested expert until the token ends reads 0.5–4.0% more). If  $c$  experts per token run on the CPU, host memory supplies at least  $\max(R^*, c)S$  bytes and the GPU reads the dense bytes  $D$  plus  $(Lk - c)S$ . With everything overlapped,

$$\bar{T} \geq \min_{0 \leq c \leq Lk} \max\left(\frac{D + (Lk - c)S}{B_{\text{gpu}}}, \frac{\max(R^*, c)S}{B_{\text{host}}}\right), \quad (2)$$

with  $B_{\text{gpu}}$  the GPU’s datasheet rate and  $B_{\text{host}}$  the highest host read rate the probe measured by any method (87.5 GB/s on listing B). The bound is for systems with a per-layer budget; a system free to pool its  $LC$  slots across layers has a looser one, 5–10% higher at the host-bound budgets (Table 3). The bound assumes that an expert run on the CPU is read from DRAM (the optimum’s shortest re-read distance is 0.5 GB of host traffic, far beyond any CPU cache), that computation runs at memory speed, and that the two read paths share  $B_{\text{host}}$  freely; it drops all latencies and starts from an empty cache.  $B_{\text{host}}$

Table 3: The speed limit under each tightening, tokens per second on the headline machine, and our cache as a percentage of the limit as published (exact optimum, highest probed host rate, datasheet GPU rate) and with every tightening at once. *084*: the optimum used in the earlier draft; *exact*: Belady’s MIN with bypass that may evict an expert after serving it; *pool*: one pool of  $LC$  slots shared across layers; *median*:  $B_{\text{host}}$  the median of six concurrent probe samples instead of the highest; *GPU*: the GPU term from the measured all-in-VRAM time instead of the datasheet rate; *layers*: the per-layer sum, which no system without cross-layer prefetch can beat; *all*: exact, pool, median and GPU together (a warm start, not included, changes the limit by at most 1.4%).

| Model         | Experts on GPU | 084 | exact | pool | median | GPU | layers | all | Ours, % of exact | % of all |
|---------------|----------------|-----|-------|------|--------|-----|--------|-----|------------------|----------|
| gpt-oss-120b  | 11%            | 169 | 172   | 182  | 159    | 172 | 163    | 168 | 41               | 41       |
| gpt-oss-120b  | 25%            | 427 | 430   | 474  | 397    | 279 | 365    | 278 | 25               | 39       |
| gpt-oss-120b  | 40%            | 519 | 519   | 519  | 517    | 279 | 500    | 278 | 29               | 55       |
| Qwen3-30B-A3B | 12.5%          | 92  | 96    | 101  | 89     | 96  | 93     | 93  | 42               | 43       |
| Qwen3-30B-A3B | 25%            | 209 | 214   | 234  | 197    | 192 | 193    | 191 | 30               | 33       |
| Qwen3-30B-A3B | 43.75%         | 308 | 308   | 308  | 307    | 192 | 298    | 191 | 35               | 57       |

is the highest rate our probe reached, not a hardware peak, so a system that reads host memory faster than the probe could exceed the limit (a second rental of the same CPU probed 90.3 GB/s). It can exceed the all-in-VRAM speed, because the CPU’s bandwidth adds to the GPU’s. The traces are the models’ own greedy text on the 30 AIME-25 prompts, templated and tokenised exactly as the server does, with routing recorded teacher-forced by our llama.cpp instrumentation; each system’s own greedy text drifts from it within the first few hundred characters on gpt-oss (Section 7).

**Tightening it.** Table 3 changes one assumption at a time. A single pool of  $LC$  slots shared across layers, the true equal-memory constraint, needs 4.5–18.6% fewer reads than per-layer slots and raises the four host-bound limits by 5–10%; it leaves the two largest budgets unchanged, where the limit is the machine’s aggregate bandwidth and read counts no longer matter. Starting the optimum warm, after 1024 tokens, changes it by at most 1.4%.  $B_{\text{host}}$  is one probe sample; the six concurrent samples on this machine read 80.7 GB/s at the median (mean 81.1, 95% interval 78.4–84.3), and the median lowers the host-bound limits by 8%. The GPU term is the one that moves most: at batch 1 the all-in-VRAM runs reach only 52% (gpt-oss) and 61% (Qwen3) of the datasheet rate, and with the measured rate the limit at 25% and above collapses to one aggregate line, 279 and 192 tok/s. A per-layer sum,  $\sum_i \max(\text{GPU}_i, \text{host}_i)$  with the optimum’s reads per layer, which no system without cross-layer prefetch can beat, sits 1–5% of a measured token below the token-level maximum (3–11% if dense work is also serialised). With every tightening at once, our cache stands at 33–57% of the limit, FreeToken at 29–54% and stock llama.cpp






Figure 1: Where the seconds go: Shapley attribution of the gap between our cache’s measured time per token and the speed limit, for every cell of Table 1, averaged over all 120 orders in which the five fixes can be applied. Fixes: overlapping host reads with GPU work, which in Eq. (2)’s form also moves experts onto the CPU until the two paths balance; the optimum’s reads instead of the best online policy’s (foresight); the best online policy’s reads at the probe’s highest rate instead of the engine’s reads through Eq. (1) (policy and read paths); the engine’s per-token host work; and GPU kernels at the datasheet rate, the term that closes the model on the measured time. Whiskers: the range of foresight’s marginal cost over the orders.

at 12–21%; against the limit as published, 25–42%, 20–40% and 9–20%. Even with every weight in VRAM, our cache with every expert in its slots matches or beats stock llama.cpp (0.997 $\times$  and 1.047 $\times$ ): the cache costs nothing when the model fits. The traces represent the measured runs: replayed through a simulator of our deployed policy, they give hit rates within 1.4 points of the engine’s own counters on gpt-oss and within 0.6–5.3 points on Qwen3.

**The accounting.** An earlier draft added one cost at a time to the limit in a fixed order and found foresight largest in every cell; costed after GPU efficiency instead, it is smaller than that in two cells, and costed after overlap it vanishes at the largest budgets. Shapley values remove the choice: each fix’s share is its average marginal cost over every order (Fig. 1), a method with a precedent in performance debugging of microservices [23]. The result has two regimes, and Eq. (2) says which one a cell is in. In the four cells where the limit binds on host reads ( $c^*=0$ : gpt-oss at 11 and 25%, Qwen3 at 12.5 and 25%), missing foresight is the largest cost, 37–53% of the gap (26–31% of the token); in the two where it binds on the GPU (gpt-oss 40%, Qwen3 43.75%, balanced by running 13 and 30 experts per token on the CPU), host reads serialised with GPU work is largest, 39–46% of the gap

(28–30% of the token), and foresight second at 23%. Policy and read paths cost 12–16% of the gap, of which most is the difference between the probe’s highest rate and the 77.7 GB/s the engine’s own law uses (at that rate the term falls to 0–10%); per-token host work 3–10%; and GPU kernels below datasheet, net of the overlap the engine achieves, 4–16%. Under this accounting, in every order of every cell the largest cost is foresight or overlap; GPU efficiency, policy and host work never are. Two of the five terms are definitions rather than measurements, and the reader should weigh them as such. The GPU term is the residual: its factor is calibrated per cell so that the model with no fix applied reproduces the measured time, so the five shares close the gap by construction; charging the all-in-VRAM shortfall in full instead raises it to 26–28% of the gap in the GPU-bound cells, where it is then the largest cost in 40 and 15 of the 120 orders, and changes no host-bound cell. And “overlap” is Eq. (2)’s form, which both overlaps host reads with GPU work and rebalances 13–30 experts per token onto the CPU; the profile below finds that overlap alone, with the engine’s present split, would save 5–10%. The interior states of the model are not measured, with one exception: Section 5 and the oracle measurement below test the foresight term, the one no engineering of read paths can recover, and that measurement probes two other parts of the model on the way. Foresight is largest in every one of the 120 orders at Qwen3 12.5%, in 115 of them at gpt-oss 11% and Qwen3 25%, and ties overlap within 0.14 ms at gpt-oss 25%.

**One token, profiled.** On Qwen3 at 43.75% with our earlier fixed FETCH table (listing A), Nsight Systems splits a token of our cache (FreeToken) into 4.93 (5.20) ms of expert copies over PCIe, 2.29 (2.21) of resident-expert products, 1.65 (1.34) of attention projections, 0.49 (0.37) of attention, 0.38 of output head and 0.97 (0.40) of norms, small kernels and cache control; our host-side time is 0.27 ms. In that configuration neither system overlaps its PCIe copies with GPU compute; doing so would save 5–10%, against the 28–30% of the token the accounting charges to its overlap-and-rebalance fix at these cells: the rest of that term is the rebalancing. Our extra 0.5 ms of GPU work comes from unfused batch-1 kernels: we launch 1.6 $\times$  as many.

**Foresight, measured.** We then gave the engine foresight of future routing, in two forms, and measured what each is worth. An oracle policy reads the routing of the next  $W$  decode steps from a trace recorded on the same machine over the same teacher-forced text; its copies ride the paced admission path (16 MB pieces; pacing the online policy itself changes it by under 1%), and an expert still in flight at its step runs on the CPU like any other miss. The *hit-optimal* oracle admits the experts those steps will route to, soonest first, evicting the resident whose next use is furthest (Belady within the window); it admits an expert for one future use as readily as for ten. The *bytes-optimal*




Figure 2: Foresight in the engine, measured on another rented RTX 5090 (Ryzen 9 9950X): speed as a percentage of this host’s limit against the window  $W$  of future routing the hit-optimal oracle sees, at the four host-bound cells. “none” is the online policy; “all” the whole sequence; points are connected for the trend only. The tick at the right is the accounting’s foresight-only term for the cell, the speed the model predicts with that one fix.

oracle is the policy whose reads define the limit, MIN with bypass: a miss runs on the CPU and is admitted only if its next use comes before the furthest next use among the residents, which it then evicts. Each ran on a rented RTX 5090 host with weak host memory (a Ryzen 9 9950X and a Ryzen 9 7950X, both reading 44–44 GB/s by the CPU and 46–46 over the link; Fig. 2, Appendix D). With the whole sequence in view the hit-optimal oracle runs  $1.40\times$  [1.38, 1.41] the online policy at gpt-oss 25% and  $1.25\times$  [1.24, 1.26] at Qwen3 25% on the first host,  $1.63\times$  and  $1.41\times$  on the second, and  $1.48$ – $1.62\times$  at the two GPU-bound cells; half of the gain at gpt-oss 25% arrives between  $W=4$  and 16 (9.2 tokens by log interpolation), near the 10 its trace gives in Section 5. Where the host binds hardest it *loses*, 3–8% at gpt-oss 11% and Qwen3 12.5% on both hosts, and at every window on the first: it reads more host bytes than the online policy (82 and 195 experts per token, misses plus admissions, against 63 and 160), because Belady admits an expert used once where the optimum, at 38 and 97 reads per token, bypasses it. The bytes-optimal oracle gains 13–23% at the four cells of 25% and above and 0–3% at the two below, and the law on its own counters predicts its time within 9%: fewer host bytes, the mechanism the accounting’s foresight term describes, and nothing else. Its reads fall (30 to 20 and 85 to 64 misses per token at 25%) but stay  $1.29$ – $1.47\times$  the optimum’s, because a copy lands two to three steps after it is issued and the expert misses again meanwhile: the same policy simulated on the trace reproduces the engine’s misses at a copy latency of 2.6–3.6 steps at the host-bound cells and comes within 2–14% of the optimum’s reads at one (Table 9). The hit-optimal oracle, which reads as many host bytes as the online policy, gains 41–63% at

Table 4: Host reads per token relative to the exact optimum (Belady’s MIN with bypass over the whole future), median over nine models at three budgets, and the range over all 27 model–budget cells. Event-atomic replay of each model’s own sampled text; every online policy is scored on the last 90% of the trace. Mixtral (8 experts) uses  $C = 2, 3$  and 4.

| Policy                   | $C=E/8$ | $E/4$ | $3E/8$ | range      |
|--------------------------|---------|-------|--------|------------|
| LRU                      | 1.44    | 1.77  | 2.03   | 1.37–2.26  |
| LFU                      | 1.44    | 1.93  | 2.40   | 1.37–6.15  |
| Decayed frequency        | 1.45    | 1.70  | 1.94   | 1.35–2.10  |
| with hysteresis $\kappa$ | 1.46    | 1.75  | 2.00   | 1.36–2.55  |
| ARC                      | 1.44    | 1.76  | 1.99   | 1.37–2.23  |
| S3-FIFO                  | 1.44    | 1.69  | 1.94   | 1.38–2.11  |
| Static, profiled         | 1.92    | 2.74  | 3.48   | 1.43–13.24 |
| Belady, $W=1$            | 1.18    | 1.45  | 1.69   | 1.07–2.06  |
| Belady, $W=4$            | 1.01    | 1.13  | 1.31   | 1.00–1.95  |
| Belady, $W=16$           | 1.00    | 1.00  | 1.01   | 1.00–1.66  |

the same four cells on the same host, and the law, with its GPU term calibrated on the online run, over-predicts those runs by 37–53%: its gain is the overlap of copies issued ahead of their use with the step, which the law charges in sequence. In this engine, then, the accounting’s two largest terms are one mechanism. The bytes foresight saves are worth 13–23% where the link has slack; the rest of foresight’s measured value is the overlap it makes possible, and where the host binds hardest neither form gains more than 3%. The engine’s within-layer overlap of a layer’s CPU misses with its GPU hits, turned off, changes speed by under 1% at every cell, so the overlap the accounting charges is across the step, not within a layer.

## 5 What the Traces Say

The accounting names foresight as the largest cost where the host binds. What would recover it, and what would not? We replay the models’ own sampled text for 9 MoE models [27, 43, 28] (9 traces of 33,000–152,000 tokens, the cache carried across conversations) through an event-atomic simulator: each token requests its  $k$  experts per layer as one set, and no expert of the current token is evicted to serve it, the accounting that cache-level studies of this problem call for [49].

**Policies.** Table 4 scores LRU, LFU, decayed frequency, ARC, S3-FIFO and a profiled static cache against the exact optimum and against Belady within a window of  $W$  future tokens, at three budgets per model. Three things hold across models. First, the online policies are close to one another and far from the optimum: the best of them reads 35–110% more than it, while decayed frequency is

the best online policy in 12 of 27 cells and S3-FIFO in 10 (within 1% of the best in 20 and 18). Where the cache turns over every token ( $C/k \leq 1.5$ , 7 cells) the six online policies are within 4% of one another; from  $C/k=2$  LRU reads 2–11% more than decayed frequency on eight of nine models, Qwen2-57B being the exception, and classic LFU fails at large budgets ( $6.2\times$  the optimum). The literature’s disagreement on recency against frequency [33, 19, 42, 11] is consistent with this regime dependence. Second, a static cache profiled on other text reads  $1.4\text{--}13.2\times$  the optimum, because the hot set moves: the  $C$  most requested experts of a layer change by 22–58% between consecutive windows of 100 tokens and 33–75% between windows of 1,000 (median 51%), more over the longer window on 8 of nine models, so popularity follows the conversation; the skew is mild, the top quarter of experts taking 31–56% of requests. Third, four tokens of foresight beat every online policy in every cell (median  $1.13\times$  the optimum at  $E/4$ ), so foresight is worth several times the choice of policy. Our deployed hysteresis costs 1–8% reads over plain decayed frequency (up to 22% on Qwen3); it was tuned for time, where it rations copies, not for reads.

**How far ahead.** The optimum reads 26–56% fewer experts from host memory than the best online policy (median 41%). The lookahead needed scales with the cache’s turnover time  $C/k$ , slots per layer over experts per token: half of the gap closes at  $W_{50} \approx 0.59 (C/k)^{1.33}$  tokens ( $r=0.971$  in log-log over 26 points; the exponent’s 95% interval over a bootstrap of models is 1.09–1.47, its leave-one-model-out range 1.24–1.37, and without the 6 points whose  $W_{50}$  lies below one token it is 1.51), and 90% needs a median of  $2.7 C/k$ . Small caches ( $C/k \leq 2$ ) recover half of the gap with one or two tokens of foresight; gpt-oss-120b at 25% ( $C/k=8$ ) needs about 10 tokens, and at 50% about 34 (both measured on its trace; the fit gives 9 and 24); in the engine, half of the oracle’s gain at 25% arrived by  $W \approx 9.2$  (Section 4). The law adds the condition: foresight is worth these reads only where the host binds; where the bus is already saturated, as on the slow devices of [47, 46], it is worth nothing. Predicting the next layer’s experts, the usual form of prediction [15, 52], gives no cross-token foresight: it can raise the hit rate by prefetching but never lowers host bytes, which is what we measured when we built it.

**Speculation is not foresight.** Speculative decoding verifies  $K$  draft tokens in one forward pass, so a layer’s experts for  $K$  positions are requested together and a missed expert is read once for all of them, which several systems propose as the source of foresight for a cache [6, 22, 30]. On the same traces, with  $K$  positions per step and drafts accepted at rate  $\alpha$  until the first rejection, batched verification saves an online cache a median 15–32% of its reads across models at  $K=8$  when rejected drafts cost nothing (48% on the model that gains most), the equivalent of a window of two to four tokens; once the rejected drafts’




Figure 3: Published MoE-offloading measurements (52 rows from 13 sources) as a percentage of their own speed of light, computed from the paper’s hardware at datasheet rates and, where a routing trace of the model exists, the per-layer Belady optimum (filled); hollow points assume uniform routing, which overstates the optimum’s reads and so the percentage. Right: our cache, FreeToken and stock llama.cpp on listing B against the probed limit of Table 1.

experts are counted, the saving is 8–12% at  $\alpha=0.9$  and becomes a loss of 18–20% at  $\alpha=0.8$ , because the union of  $K$  positions’ expert sets is 20–79% of  $Kk$ , several times  $C$  (Appendix B lists the variants). The optimum gains little from batching (–7 to 15%). Foresight must come from knowing future routing, as a router decoupled from the backbone [3] or a forecaster [39] would provide, and  $W_{50}$  says how far ahead.

## 6 Published Systems Against Their Own Limit

Rule 11 of scientific benchmarking asks for an upper bound beside every measurement [13]. We applied a bound of the same construction as Eq. (2) (Belady’s optimum with bypass per layer, CPU and PCIe paths, a maximum per layer step; `scripts/audit.py`), with each paper’s stated hardware at datasheet rates and a routing trace of its model where we have one, to 52 published measurements of MoE offloading from 13 sources, collected for an earlier audit of this literature (Appendix F). Figure 3 shows the result. The 29 trace-based measurements reach a median 9.5% of their own speed of light (quartiles 4.9–17.8%, best 36.8%); the 11 that needed an i.i.d. routing assumption, 34.8%, an overestimate; measurements with no GPU expert capacity (the CPU-only rate is their bound) and with every expert resident are plotted apart. Our own reference points are not like for like: their ceilings are probed, and

a probed ceiling is lower than a datasheet one by the ratio of achieved to nominal host bandwidth, which we did not record for the published machines and so do not quantify; it would raise the published percentages by that ratio, which on our own hosts would be a factor of about 1.1–1.4. Even so, the trace-scored measurements sit at a median tenth, and three quarters of them below a fifth, of what their machines could deliver, and Fig. 1 says what the missing part is made of.

## 7 The Expert Cache as the Instrument

**Design.** Our cache is a 3,700-line patch to llama.cpp. Each MoE layer has  $C$  GPU slots. Routing is decided on the GPU; resident experts run there, and each miss either runs on CPU helper threads, which the GPU signals through a mailbox in mapped memory, or is copied into a slot and run on the GPU (FETCH). Admission follows a decayed-frequency score (half-life 16 tokens, an admission margin  $\kappa=1$ ; a llama.cpp pull request admits by a decayed heat with hysteresis in the same spirit [26]) and copies in the background. Sampling runs on the GPU. The cache serves single-token steps; prompts and batches above one run llama.cpp’s stock graph with every expert in host memory and pay for the resident experts and a one-thread CPU side (79–83% of stock llama.cpp’s prefill speed, 77–79% of its decode at 2–4 sequences; Appendix E). The one per-machine decision is the FETCH table  $f(n)$ : of a layer’s  $n$  missed experts, how many to copy. A per-layer form of Eq. (1) gives it,

$$T(n, f) = \max\left(\frac{fS}{B_p} + g, [n > f](L_c + \frac{(n-f)S}{B_c}), [0 < f < n]\frac{nS}{B_{cp}}\right), \quad (3)$$

where  $g$  is the GPU’s expert work per layer ( $48\mu s$  for Qwen3,  $37\mu s$  for gpt-oss, from profiles) and  $L_c=20\mu s$  the hand-off; the table picks the  $f$  with the smallest  $T$  and is computed from the probe when the job starts. On listing B the law’s table beats the fixed table we used before at every budget of both models, by 3.3–8.0%; on hosts whose link we measured at 27–30 GB/s instead of 53 it fetches less and gains 21–34% on gpt-oss and 9–18% on Qwen3, the per-machine split mattering most where the link is slow (Appendix E, Table 10).

**Where the gain comes from.** Figure 4 builds the system up from stock llama.cpp on one machine. Placing each layer’s hot experts is the first gain: a static cache of the 32 experts per layer most requested on other text (the models’ own MATH-500 outputs) runs  $1.31\times$  and  $1.89\times$  stock, which pins whole layers instead. Its hit rate (52% and 67%) depends on how well the profile matches the workload: profiled in hindsight on the measured problems it hits 75% and 72% and runs faster (hatched). Replacing experts online closes that gap without a profile, if admissions are rationed: LRU admits every miss, and on




Figure 4: Ablation at 25% of experts on the GPU (15 AIME-25 problems, 256 tokens, listing A). Each step adds one component to the one above; the hatched row is not a step: a static cache of each layer’s 32 most requested experts counted on the measured problems themselves.

Qwen3 its copies raise host reads from 127 to 188 experts per token, so it runs  $0.88\times$  the static cache ( $1.30\times$  on gpt-oss); the decayed-frequency policy reads 41 and 118 per token, runs  $1.58\times$  and  $1.12\times$  the profiled static cache, and the hindsight cache runs  $0.98$  [0.95, 1.01] and  $1.01$  [0.98, 1.04] times as fast as it. On traces LRU misses only 2–11% more than decayed frequency at these budgets (Section 5); the rest of its cost here is the copies it admits. On the stock-clock host, with everything else as in the full system, LRU runs 6–21% slower than decayed frequency across the six cells of Table 1 and trails FreeToken at four of them: the admission policy is most of the lead over FreeToken (77–84% of it at gpt-oss 25 and 40%, more than the whole of it at gpt-oss 11% and Qwen3 12.5 and 25%). The GPU-signalled CPU helpers, GPU slot maps, GPU sampling and FETCH add 2–9% each, and the law’s FETCH table 5–8% over the fixed one. The full system runs  $2.68\times$  (gpt-oss) and  $2.81\times$  (Qwen3) stock llama.cpp.

**Two hosts.** Table 1 compares the three systems at six budgets on two RTX 5090 hosts. On listing B ours leads FreeToken at every budget, by 15–29% on gpt-oss and 3–15% on Qwen3; the lead is largest when many experts miss and the CPU path carries them (gpt-oss at 11–25%), smallest at Qwen3 12.5%, where our cache already reads host memory at 80% of the machine’s combined rate, and at Qwen3 43.75%, where few experts miss and FreeToken’s fused GPU kernels count. On the stock-clock host, whose CPU reads 22% slower (56 against 72 GB/s) behind the same link, ours leads at five cells (gpt-oss 9–21%; Qwen3 12.5 and 25%  $1.03$  and  $1.05\times$ ), and FreeToken leads at Qwen3 43.75% ( $0.97$ , interval  $0.962$ – $0.987$ ). Every



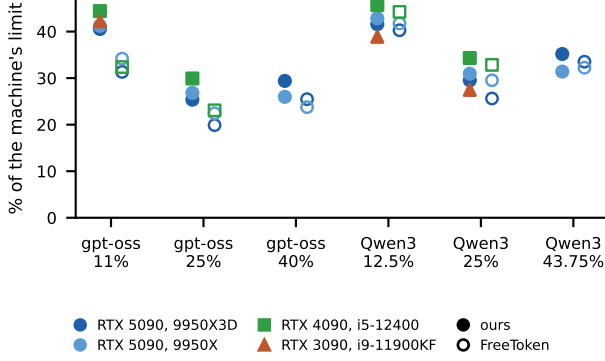


Figure 5: Ours (filled) and FreeToken (hollow) as a percentage of each machine’s own speed limit (exact optimum, the machine’s highest probed host rate, the card’s datasheet rate) at every cell of Table 7: two RTX 5090 hosts, an RTX 4090 and an RTX 3090 (FreeToken’s 3090 cells, its carried-over backend running below llama.cpp, are omitted). The limit ranges from 41 to 519 tok/s across these cells; ours stays at 25–46% of it.

ratio is 0.01–0.11 below listing B’s, beyond the  $\pm 0.06$  we predicted at five of six cells. The card explains little of it: FreeToken’s offload backend, which puts no expert on the CPU, lost only 5–8% at the two cells it runs, while ours lost 12–17% and llama.cpp 16–17%. With the CPU slower the law’s tables copy more of each layer’s misses (0,0,1,2,3 and 0,0,1,2,2,3,4,5,6 against 0,0,1,1,2 and 0,0,1,1,2,3,3,4,5), and the CPU path, which is our lead over FreeToken, is worth less. The limit moves with the host and the fraction much less: this host’s four host-bound limits are 19% lower, and ours stands at 26–43% of them against 25–42% on listing B. On FreeToken’s own headline model, Qwen3.6-35B-A3B, the two systems tie within 3% at three budgets (Appendix E): the lead is model-dependent, largest with few large experts and smallest with many small ones.

**Other cards, and the scope of the gain.** Table 7 (Appendix C) repeats the protocol on an RTX 4090 next to a Core i5-12400 (the host reads 37 GB/s, the link 25) and an RTX 3090 next to a Core i9-11900KF (45 GB/s; CPU and link together read *less* than the CPU alone, so the law’s table copies nothing), at the cells that fit 24 GB. On the 4090 ours runs 1.03–1.37 $\times$  FreeToken and 1.8–2.6 $\times$  llama.cpp at half listing B’s absolute speeds, the host being the reason; on the 3090 it runs 1.6–2.2 $\times$  llama.cpp (FreeToken’s carried-over backend ran below llama.cpp there and is not counted). The cards’ own bandwidth never enters: every one of the seven cells on the 24 GB cards is host-bound at the limit, and ours stands at 30–46% and 27–42% of it (Fig. 5). Mixtral-8x7B (8 experts per layer, top-2; 26 GB of experts on the 3090) marks the scope: with 2 of 8 experts resident the cache hits 31% of reads where whole-layer pinning hits 25%,

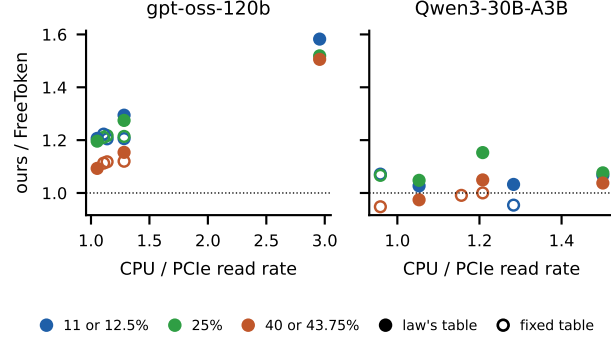


Figure 6: Ours over FreeToken’s better backend (carried over from listing B on the stock-clock host) in every same-machine comparison of the current engine on the two models of Table 1 (eight hosts), against the machine’s CPU-to-PCIe read ratio (FreeToken’s own probe). Filled: FETCH table from the law on that machine; hollow: the fixed table. The rightmost gpt-oss points are a Core i9-14900K with FreeToken’s executor on its performance cores; the rightmost Qwen3 points are a Ryzen 9 7900.

so there is little locality to earn, and ours ties llama.cpp (1.015 [1.008, 1.024]); on the 128-expert models the same cache hits 53–79% at 11–25% resident. Across eight hosts the gpt-oss lead over FreeToken grows with the host’s CPU-to-link ratio, from 1.09–1.22 $\times$  where the CPU reads 1.05–1.14 $\times$  its link’s rate to 1.51–1.58 $\times$  behind a link of half the bandwidth, as the law predicts, and the Qwen3 lead stays at 0.95–1.15 $\times$  (Fig. 6, Appendix E). Teacher-forced against an all-VRAM reference, the cache’s output differs by a mean KL of 0.0019 and 0.0005 nats over the full vocabulary (top-1 agreement 98.5% and 99.3%), the same divergence as stock llama.cpp’s own CPU placement (Appendix E).

## 8 Related Work

**Offloading systems.** Demand-paging systems cache experts per layer and copy misses over PCIe [8, 41, 34, 36, 31]. CPU-execution systems run non-resident experts on the host [18, 5, 51, 50, 38]; a set-associative GPU cache with CPU-executed misses [14], a llama.cpp fork [20, 21] and two llama.cpp pull requests [26, 7] are the closest designs to ours, and FreeToken [44] and DALI [53] split each layer’s misses between CPU and PCIe, FreeToken from a closed-form bandwidth model. Pipelined Sharding [37] plans sub-layer placement under a VRAM budget. Predictors guess upcoming experts from the previous layer’s state [15, 52], several steps ahead with a learned forecaster [39, 42], or a batch’s routes across layers [10]; README [3] makes routing known in advance by decoupling the router and then caches in Belady’s manner. Speculative decoding has been combined with expert offloading as a source of batching [6, 22, 30]. Batched-throughput

systems [32, 4, 45] and SSD tiers [24, 25] are outside our single-request, host-DRAM scope.

**Models, bounds and audits.** FreeToken’s split and our Eq. (3) are bandwidth models of one layer; Eq. (1) is a whole-token model tested blind across machines and against alternatives. Max-of-paths cost models go back to FlexGen [32] and MoE-Lightning [4]; WiSP [47] bounds decode by PCIe turnover, Budgeting Bytes [46] by bytes over bandwidth, and MoE-CAP [17] measures bandwidth utilisation against a single peak. Qiu et al. [29] predict llama.cpp throughput from GGUF metadata for models held entirely in GPU or unified memory; Suram [35] matches one offloaded configuration to bytes per token over bandwidth; Si et al. [33] derive when a copy pays off and find recency competitive with frequency at large budgets. The roofline [40] is the single-memory ancestor of our bound. Zhang [49] proposes a trace-level evaluation contract for expert caching, decomposes the online-to-optimal gap in cache blocks in one fixed order, and audits the reporting of ten papers; SeqMoE [39] reports four systems against the all-in-VRAM speed. Trace studies of policies cover two or three models each [49, 33, 19, 11]. Shapley attribution of a performance gap has a precedent in ShapleyIQ [23].

**Lookahead caching.** Limited lookahead in paging goes back to Albers [1]; Hasslinger et al. [9] combine Belady within a window with LRU or LFU beyond it for web caches, the policy family we use to measure the horizon.

## 9 Limitations

- **Hardware.** Consumer cards (RTX 5090 for the headline comparisons, a 4090 and a 3090 at the cells that fit 24 GB; an RTX PRO 6000 for all-in-VRAM runs) and desktop Ryzen and Intel hosts, rented; GPU clocks could not be locked, power limits differed (350–575 W), and listing B’s memory runs  $1.21\times$  faster than the other cards’; Table 1 therefore reports a stock-clock host beside it, where our lead is 0.01–0.11 smaller and FreeToken leads at one cell. A 12-channel server CPU breaks the law’s bandwidth assumption (Section 3). Intervals cover problem sampling on one machine, not machine-to-machine variance, which Fig. 6 shows across eight hosts.
- **Workload.** Single-request decode; AIME prompts, which were also used while developing the system; greedy decoding; 130–186-token prompts. Long outputs are measured on held-out problems. Prompts and batches above one bypass the cache and run 79–83% and 77–79% of stock llama.cpp’s speed (Section 7); speculative verification bypasses it and is not measured.
- **Competitor settings.** FreeToken ran its faster backend per budget with its own bandwidth calibration,

default thread count and fetch cap, while our FETCH table was swept; its CPU backend failed on a hybrid-core host at its default thread count and was rerun on the performance cores. Each comparison uses one confirmation launch, except Qwen3 at 25%, which is from the selection launch for both systems.

- **The bound, the accounting and the audit.** Exact routing only: expert skipping, substitution or deferral change the model and fall outside it; a per-layer budget, which a pooled cache can beat by 5–10%. Two of the accounting’s five terms are definitions (Section 4). The oracles read a recorded trace, not a predictor, and ran on two rented hosts with weak host memory; their copies land two to three steps after they are issued, so what they measure is this engine’s foresight, not the limit’s. The published rows use datasheet ceilings and, for eleven of them, i.i.d. routing.
- **Constants and outputs.**  $G$ ,  $g$  and  $L_c$  are fitted once and frozen; the per-layer model overstates differences between FETCH tables about  $1.8\times$ . Output parity is measured by teacher-forced KL divergence, agreement and loss against an all-VRAM reference, not by task accuracy.

## 10 Conclusion

On a consumer machine, MoE decode time is a fixed GPU term plus host-memory bytes over the rate of the paths that read them. That law, predicted blind and tested against its alternatives, is accurate to a few percent on desktop hosts; it tells a system which path to use for a miss on a given machine; and with the optimum’s bytes it gives a speed limit in seconds that published systems reach a median tenth of. An accounting that no order of fixes can bias says what the rest is: where the host binds, the policy’s ignorance of future routing, and where the GPU binds, host reads that wait for it. Nine models’ traces say that neither an online policy nor batching alone buys that foresight; it must come from knowing the routing a few to tens of tokens ahead. Given it by an oracle, the engine runs 13–23% faster on the bytes the limit’s policy saves and 41–63% faster when the same foresight copies experts ahead of their use: in the engine the accounting’s two largest terms are one mechanism, because foresight is what makes the overlap possible. Where the host binds hardest neither form gains more than 3%: the copies land late and compete with the helpers for the same memory.

## References

- [1] Susanne Albers. On the influence of lookahead in competitive paging algorithms. *Algorithmica*, 18(3):283–305, 1997.

- [2] Laszlo A. Belady. A study of replacement algorithms for a virtual-storage computer. *IBM Systems Journal*, 5(2):78–101, 1966.
- [3] Ruisi Cai, Yeonju Ro, Geon-Woo Kim, Peihao Wang, Babak Ehteshami Bejnordi, Aditya Akella, and Zhangyang Wang. Read-ME: Refactorizing LLMs as router-decoupled mixture of experts with system co-design. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 37, pages 116126–116148, 2024. arXiv:2410.19123.
- [4] Shiyi Cao, Shu Liu, et al. MoE-Lightning: High-throughput MoE inference on memory-constrained GPUs. *arXiv preprint arXiv:2411.11217*, 2024.
- [5] Hongtao Chen, Weiyu Xie, Boxin Zhang, Jingqi Tang, Jiahao Wang, Jianwei Dong, Shaoyuan Chen, Ziwei Yuan, Chen Lin, Chengyu Qiu, Yuening Zhu, Qingliang Ou, Jiaqi Liao, Xianglin Chen, Zhiyuan Ai, Yongwei Wu, and Mingxing Zhang. KTransformers: Unleashing the full potential of CPU/GPU hybrid inference for MoE models. In *Proceedings of the 31st ACM Symposium on Operating Systems Principles (SOSP)*, 2025.
- [6] Liangkun Chen, Zijian Wen, Tian Wu, Xiaoxi Zhang, and Chuan Wu. SP-MoE: Speculative decoding and prefetching for accelerating MoE-based model inference. *arXiv preprint arXiv:2510.10302*, 2025.
- [7] csantiago78. llama.cpp pull request #27861: GPU-resident LRU cache for host-offloaded MoE expert weights. <https://github.com/ggml-org/llama.cpp/pull/27861>, 2026. Draft, accessed September 2026.
- [8] Artyom Eliseev and Denis Mazur. Fast inference of mixture-of-experts language models with offloading. *arXiv preprint arXiv:2312.17238*, 2023.
- [9] Gerhard Hasslinger, Juho Heikkinen, Konstantinos Ntougias, Frank Hasslinger, and Oliver Hohlfeld. Optimum caching versus LRU and LFU: Comparison and combined limited look-ahead strategies. In *Proceedings of the 16th International Symposium on Modeling and Optimization in Mobile, Ad Hoc, and Wireless Networks (WiOpt)*, pages 1–6, 2018.
- [10] Xin He, Shunkang Zhang, Kaijie Tang, Shaohuai Shi, Yuxin Wang, Zihao Zeng, Zhenheng Tang, Xiaowen Chu, Haiyan Yin, Ivor W. Tsang, and Yew Soon Ong. ExpertFlow: Efficient mixture-of-experts inference via predictive expert caching and token scheduling. *arXiv preprint arXiv:2410.17954*, 2024. Accepted at DAC 2026 (per arXiv v2).
- [11] Duc Hoang, Ajay Jaiswal, Mohammad Samragh, and Minsik Cho. SpecMD: A comprehensive study on speculative expert prefetching. In *International Conference on Machine Learning (ICML)*, 2026. arXiv:2602.03921.
- [12] Torsten Hoefer, Dan Alistarh, Tal Ben-Nun, Nikoli Dryden, and Alexandra Peste. Sparsity in deep learning: Pruning and growth for efficient inference and training in neural networks. *Journal of Machine Learning Research*, 22(241), 2021.
- [13] Torsten Hoefer and Roberto Belli. Scientific benchmarking of parallel computing systems: Twelve ways to tell the masses when reporting performance results. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC)*, 2015.
- [14] En-Ming Huang, Li-Shang Lin, et al. Efficient CPU-GPU collaborative inference for MoE-based LLMs on memory-limited systems. *arXiv preprint arXiv:2512.16473*, 2025.
- [15] Ranggi Hwang, Jianyu Wei, et al. Pre-gated MoE: An algorithm-system co-design for fast and scalable mixture-of-expert inference. In *Proceedings of the 51st Annual International Symposium on Computer Architecture (ISCA)*, 2024. arXiv:2308.12066.
- [16] Andrei Ivanov, Nikoli Dryden, Tal Ben-Nun, Shigang Li, and Torsten Hoefer. Data movement is all you need: A case study on optimizing transformers. In *Proceedings of Machine Learning and Systems (MLSys)*, 2021.
- [17] Yinsicheng Jiang, Yao Fu, Yeqi Huang, et al. MoE-CAP: Cost-accuracy-performance benchmarking for mixture-of-experts systems. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. arXiv:2412.07067.
- [18] Keisuke Kamahori, Tian Tang, et al. Fiddler: CPU-GPU orchestration for fast inference of mixture-of-experts models. *arXiv preprint arXiv:2402.07033*, 2024.
- [19] Byeongju Kim, Jungwan Lee, Donghyeon Han, Hoi-Jun Yoo, and Sangyeob Kim. FlashMoE: Reducing SSD I/O bottlenecks via ML-based cache replacement for mixture-of-experts inference on edge devices. *arXiv preprint arXiv:2601.17063*, 2026.
- [20] leloch. MoE expert cache for llama.cpp: VRAM caching of hot CPU-resident experts with hybrid hit/miss execution. <https://github.com/leloch/llama.cpp/tree/moe-cache-v2-pr>, 2026. Branch moe-cache-v2-pr, commit e3096b0 (6 August 2026). Proposed upstream as llama.cpp pull request #24524 and RFC discussion #24528; not merged. Accessed September 2026.
- [21] leloch. RFC: MoE expert cache, VRAM caching of hot CPU-resident experts with hybrid hit/miss execution. llama.cpp discussion #24528, <https://github.com/ggml-org/llama.cpp/discussions/24528>, 2026. Posted 12 June 2026; accessed September 2026.
- [22] Shuhuai Li, Jianghao Lin, Dongdong Ge, and Yinyu Ye. MoE-SpAc: Efficient MoE inference based on speculative activation utility in heterogeneous edge scenarios. *arXiv preprint arXiv:2603.09983*, 2026.
- [23] Ye Li, Jian Tan, Bin Wu, Xiao He, and Feifei Li. ShapleyIQ: Influence quantification by Shapley values for performance debugging of microservices. In *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, Volume 4, pages 287–323, 2024.
- [24] Fangzhou Liang, Yibin Shen, Jianmin Hu, Jiayang Xu, Hanchi Gao, Minxian Xu, and Zili Meng. SSD-LLaMA: SSD-native inference for trillion-parameter MoE at 1+ token/s on a consumer PC. *arXiv preprint arXiv:2609.18110*, 2026.
- [25] Yu Lin, Yiming Wang, Runyuan Cai, Hanze Liu, and Xiaodong Zeng. The other half of the memory wall: Serving 35B MoEs from SSD with trained routing prediction. *arXiv preprint arXiv:2609.18063*, 2026.
- [26] milto22. llama.cpp pull request #26563: Expert caching that greatly increases performance (self-contained, off by default). <https://github.com/ggml-org/llama.cpp/pull/26563>, 2026. Accessed September 2026.



- [27] Niklas Muennighoff, Luca Soldaini, et al. OLMoE: Open mixture-of-experts language models. *arXiv preprint arXiv:2409.02060*, 2024.
- [28] OpenAI. gpt-oss-120b & gpt-oss-20b model card. *arXiv preprint arXiv:2508.10925*, 2025.
- [29] Xinyu Qiu, Chuhong Xu, et al. GGUF-metadata prediction of single-sequence llama.cpp throughput across three systems. *arXiv preprint arXiv:2609.14864*, 2026.
- [30] Anish Saxena, Po-An Tsai, Hritvik Taneja, Aamer Jaleel, and Moinuddin Qureshi. Cascade: Utility-driven speculative decoding for mixture-of-experts. In *Proceedings of Machine Learning and Systems (MLSys)*, 2026. arXiv:2506.20675 (2025).
- [31] Musa Shams. Paging the experts: A reproducible characterization of flash-backed MoE inference on iPhone. *arXiv preprint arXiv:2609.29032*, 2026.
- [32] Ying Sheng, Lianmin Zheng, Binhang Yuan, Zhuohan Li, Max Ryabinin, Beidi Chen, Percy Liang, Christopher Ré, Ion Stoica, and Ce Zhang. FlexGen: High-throughput generative inference of large language models with a single GPU. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, volume 202 of *Proceedings of Machine Learning Research*, pages 31094–31116, 2023. arXiv:2303.06865.
- [33] Yuan Si, Yufeng Lin, et al. Who should own the expert cache? kernel-managed tiering for trillion-parameter MoE inference. *arXiv preprint arXiv:2608.12103*, 2026.
- [34] Xiaoni Song, Zihang Zhong, et al. ProMoE: Fast MoE-based LLM serving using proactive caching. *arXiv preprint arXiv:2410.22134*, 2024.
- [35] Shriniwas Ramesh Suram. Cacheable by design? training mixture-of-experts routers for locality against the edge memory-bandwidth wall: A pre-registered negative result, with a systems measurement study. *arXiv preprint arXiv:2608.18261*, 2026.
- [36] Peng Tang, Jiacheng Liu, et al. HOBbit: A mixed precision expert offloading system for fast MoE inference. *arXiv preprint arXiv:2411.01433*, 2024.
- [37] Aditya Ukarande, Deep Shekhar, Marc Blackstein, and Ram Rangan. Efficient, VRAM-constrained xLM inference on clients. In *Proceedings of Machine Learning and Systems (MLSys)*, volume 8, 2026. Industry track. arXiv:2604.26334. Artifact: <https://github.com/deepshnv/pipeshard-mlsys26-ae>, tag v2.0.3-mlsys26.
- [38] Wenxin Wang, Yule Hou, et al. Achieving cloud-grade SLOs for local mixture-of-experts inference through CPU-GPU hybrid design. *arXiv preprint arXiv:2606.10493*, 2026.
- [39] Zihan Wang, Yuqi Wang, et al. SeqMoE: Toward full-load performance via predictive and graph-compatible MoE offloading. *arXiv preprint arXiv:2609.12978*, 2026.
- [40] Samuel Williams, Andrew Waterman, and David Patterson. Roofline: An insightful visual performance model for multicore architectures. *Communications of the ACM*, 52(4):65–76, 2009.
- [41] Leyang Xue, Yao Fu, et al. MoE-Infinity: Efficient MoE inference on personal machines with sparsity-aware expert cache. *arXiv preprint arXiv:2401.14361*, 2024.
- [42] Sanjali Yadav and Bahar Asgari. Mira: Memory-efficient MoE inference using adaptive caching and predictive expert staging. *arXiv preprint arXiv:2609.38090*, 2026.
- [43] An Yang et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [44] Shuo Yang, Xiaoze Fan, Melissa Pan, Haocheng Xi, Zhe Wang, Shanlin Sun, Kurt Keutzer, Song Han, Matei Zaharia, Chenfeng Xu, and Ion Stoica. FreeToken: Efficient edge-native MoE serving with bandwidth-adaptive execution. *arXiv preprint arXiv:2608.16157*, 2026.
- [45] Yichao Yuan, Lin Ma, et al. MoE-Lens: Towards the hardware limit of high-throughput MoE LLM serving under resource constraints. *arXiv preprint arXiv:2504.09345*, 2025.
- [46] Hanhaodi Zhang. Budgeting bytes: A windowed storage roofline and dual-budget architecture ablations for storage-bound LLM decoding. *arXiv preprint arXiv:2609.04238*, 2026.
- [47] Jiamu Zhang, Liang Wu, Mayank Darbari, and Liangjie Hong. WiSP: A working-set view of mixture-of-experts serving on extremely low-resource hardware. *arXiv preprint arXiv:2606.21868*, 2026.
- [48] Lei Zhang, Reza Karimi, Irfan Ahmad, and Ymir Vigfusson. Optimal data placement for heterogeneous cache, memory, and storage systems. *Proceedings of the ACM on Measurement and Analysis of Computing Systems (PO-MACS)*, 4(1):6:1–6:27, 2020.
- [49] Yu Zhang. Reproducible evaluation of MoE expert caching: Replay semantics, workload contamination, and operating regimes. *arXiv preprint arXiv:2608.07911*, 2026.
- [50] Yujie Zhang, Shivam Aggarwal, and Tulika Mitra. DAOP: Data-aware offloading and predictive pre-calculation for efficient MoE inference. In *Design, Automation & Test in Europe Conference (DATE)*, pages 1–7, 2025. arXiv:2501.10375.
- [51] Shuzhang Zhong et al. HybriMoE: Hybrid CPU-GPU scheduling and cache management for efficient MoE inference. In *Design Automation Conference (DAC)*, 2025. arXiv:2504.05897.
- [52] Shien Zhu, Samuel Bohl, et al. Pre-attention expert prediction and prefetching for mixture-of-experts large language models. *arXiv preprint arXiv:2511.10676*, 2025.
- [53] Zeyu Zhu, Gang Li, Peisong Wang, Zitao Mo, Minnan Pei, Zhuoran Song, Xiaoyao Liang, and Jian Cheng. DALI: A workload-aware offloading framework for efficient MoE inference on local PCs. *arXiv preprint arXiv:2602.03495*, 2026.



## A Scientific benchmarking checklist

How this paper follows the twelve rules of Hoefer and Belli [13].

| Rule                                                                         | Where and how                                                                                                                                                                                                                      |
|------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 1. Report the base case of a speed-up and its absolute performance.          | Every ratio in Table 1 is next to both absolute speeds; llama.cpp is configured with per-layer expert offload at equal expert memory, the setting its users run.                                                                   |
| 2. Explain subsets of benchmarks or resources.                               | Two models chosen for size and precision (MXFP4, BF16), a third that is FreeToken’s own headline model; budgets span what fits a 32 GB card; the 43.75% budget is omitted where it does not fit. All 30 AIME-25 problems are used. |
| 3. Arithmetic mean for costs, harmonic mean for rates.                       | Every problem decodes the same number of tokens, so the harmonic mean of rates is total tokens over total time; ratios of total time agree with the reported ratios of mean rates within 0.003.                                    |
| 4. Avoid summarising ratios.                                                 | Ratios are computed from mean speeds, never averaged across budgets or models; ranges list the per-budget values.                                                                                                                  |
| 5. Report whether measurements are deterministic; give confidence intervals. | Decode is greedy and deterministic per engine; timing is not. Every comparison has a 95% interval.                                                                                                                                 |
| 6. Do not assume normality.                                                  | Intervals are bootstrap percentile intervals (10,000 resamples, paired by problem); no distributional assumption.                                                                                                                  |
| 7. Compare nondeterministic data soundly.                                    | Paired bootstrap of the ratio of means; the variant a system is scored with is picked on a separate launch (except the Qwen3 25% row of Table 1).                                                                                  |
| 8. Check that central tendency is the right summary.                         | Mean speed over 256 tokens is FreeToken’s own metric; per-token p50 and p99 latencies are recorded for every problem in the released rows.                                                                                         |
| 9. Document all factors and the setup.                                       | Machines, probes, power limits, memory clocks (Section 2), software commits (llama.cpp 4da6337, FreeToken 0d652e7), flags and job scripts are released; Appendix F.                                                                |
| 10. Report measurement, synchronisation and summarisation.                   | One client times every system from streamed token arrivals; no server-side timers enter a comparison (Section 2).                                                                                                                  |
| 11. Show upper performance bounds.                                           | The speed limit (Section 4) and all-in-VRAM speed are columns of Table 1; Table 3 tightens the limit four ways; Section 6 applies it to the published literature.                                                                  |
| 12. Plot enough to interpret; connect points only for trends.                | Figures 1 and 4 are bar charts; Figs. 3, 5 and 6 are scatter plots; Fig. 2 connects measured windows for the trend only, and its $W_{50}$ is a stated log interpolation between two measured windows.                              |

## B Prediction log

Each job’s predictions are in the header of its script on the public `gpu` branch, committed and pushed before the machine started; git history orders every prediction before its measurement. Outcome notes with the statistics are in `prereg/`. Table 6, generated from `prereg/scorecard.clauses.json`, splits every prediction of jobs 073–094 into its separable clauses and scores each under one rule: *held* only when the point estimate is on the predicted side and the 95% paired-bootstrap interval excludes the threshold (for a band, lies inside it); *held (point)* when only the point estimate is, or no interval exists; *failed* when the point estimate is on the wrong side; *untested* when the clause could not be evaluated; *void* when the job was voided. Of 243 clauses, 112 held, 28 held on the point estimate, 91 failed, 8 were untested and 4 void: 7 of 19 scored clauses held in jobs 073–075, when the system was being built, 19 of 23 in jobs 076–081, 11 of them sign-only predictions made after a same-host pilot, 24 of 38 in jobs 082–087, and 62 of 151 in the jobs run after the reviews (088 onward), where 23 of 26 sign clauses held and 56 of 91 bands failed. The bands that failed are the sizes of effects whose direction held: our slow-link lead landed at  $1.51\text{--}1.58\times$  against a predicted  $1.10\text{--}1.50$ ; the rerun of Table 1 on a second host (job 089) gave ratios  $0.06\text{--}0.11$  below the headline machine’s at five of six cells against a predicted  $\pm 0.06$ , host-bound speeds  $14\text{--}17\%$  lower against a predicted  $\pm 12\%$ , and LRU  $6\text{--}8\%$  slower than decayed frequency at the two GPU-heavy cells against a predicted  $10\text{--}30\%$ ; on the 4090 and 3090 hosts (jobs 091, 092; Appendix C) our lead over FreeToken and its fraction of the limit overshoot their bands, FreeToken’s carried-over

backend fell far below its band, and two launch-order checks exceeded 3%; in the oracle job (093; Section 4) the oracle’s gain fell outside 20–80% at the two lowest budgets (it lost), its speed fell just below the predicted 45–65% of the limit at all four host-bound cells (43.7–44.8%), its hit rate and admissions missed the optimum’s by more than the predicted 3 points and 30%, no-overlap cost 0% against a predicted 5–25%, and all-CPU cost 15% on Qwen3 against a predicted 0–10%; in the bytes-oracle job (094) the bands missed as Appendix D describes. The 11 threshold failures are ours at  $1.56\text{--}1.78\times$  llama.cpp on the slow hosts against a predicted 1.8, the Mixtral tie against a predicted  $1.3\times$ , the hit-optimal oracle’s gain at the two GPU-bound cells (48 and 56%) against a predicted ceiling of 15%, the bytes-optimal oracle short of beating the hit-optimal one by 15% at the two lowest budgets, and its 16-token window short of keeping 70% of its gain at three cells; the 2 equality failures are the 4090 host’s FETCH tables, which copy one more expert than predicted; the 3 sign failures are LRU behind FreeToken at gpt-oss 11%, where we predicted it would still lead, and the hit-optimal oracle slower than the online policy at gpt-oss 11% and Qwen3 12.5%. Against the earlier draft’s log, two clauses are downgraded (a llama.cpp clause of job 074b that the log counted as held was never run; job 085b’s own restatement of its fourth prediction, “at least 2 of 3”, failed at 1 of 3, though the pooled 4 of 6 held), job 073’s “identical outputs” is scored on fully identical outputs (0–3%) rather than on agreement of the first 160 characters (27–47%), and ten clauses the log counted as held, among them job 087’s first, hold on the point estimate only under the interval rule; `prereg/scorecard.outcome.md` lists every change. Job 085’s Qwen3 half ran as 085b on a second machine with a slow link after a hung conversion. Job 084 ran on a throttled card and was rerun, with its predictions unchanged, as 084b/084c. The batched-verification variants of Section 5 (rejected drafts free, routed like the true token, or routed at random) are in `prereg/batchchk.outcome.md`.

Table 6: Prediction scorecard, jobs 073–094: every committed prediction split into its separable clauses and scored under one rule. *Held*: the point estimate is on the predicted side and the 95% paired-bootstrap interval excludes the threshold (for a band, lies inside it). *Held (point)*: the point estimate is on the predicted side but the interval includes the threshold, or no interval exists. *Untested*: the clause could not be evaluated. *Void*: the job was voided (084, a throttled card; its predictions are scored on 084b). Types: sign (a direction), threshold (a numeric bar), band (an interval), equality (a deterministic outcome). Measured values are ratios of mean speeds with 95% intervals unless the clause says otherwise; `prereg/scorecard.clauses.json` names the source of every number. Totals: 243 clauses, 112 held, 28 held (point), 91 failed, 8 untested, 4 void.

| Job                                                                                    | Clause                                                               | Type      | Measured                                                                                                           | Status          |
|----------------------------------------------------------------------------------------|----------------------------------------------------------------------|-----------|--------------------------------------------------------------------------------------------------------------------|-----------------|
| <b>Jobs 073–075:</b> 20 clauses, 7 held, 3 held (point), 9 failed, 1 untested, 0 void  |                                                                      |           |                                                                                                                    |                 |
| 073 P1a                                                                                | ours+FETCH $\geq 5\%$ ahead of FreeToken, 11%                        | threshold | 1.109 [1.098, 1.120]                                                                                               | held            |
| 073 P1b                                                                                | ours+FETCH $\geq 5\%$ ahead of FreeToken, 25%                        | threshold | 1.084 [1.076, 1.092]                                                                                               | held            |
| 073 P1c                                                                                | ours+FETCH within $\pm 5\%$ of FreeToken, 40%                        | band      | 0.941 [0.933, 0.951]                                                                                               | <b>failed</b>   |
| 073 P2                                                                                 | ours+FETCH $\geq 1.8\times$ llama.cpp, every budget                  | threshold | 1.832 [1.808, 1.856]; 2.358 [2.322, 2.396]; 2.647 [2.605, 2.686]                                                   | held            |
| 073 P3a                                                                                | same-problem warm-up reads 5-15% faster, the two caches              | band      | 0.991 [0.962, 1.028]; 0.997 [0.986, 1.008]                                                                         | <b>failed</b>   |
| 073 P3b                                                                                | same-problem warm-up within $\pm 2\%$ , llama.cpp                    | band      | 0.962 [0.934, 0.979]                                                                                               | <b>failed</b>   |
| 073 P4                                                                                 | $\geq 80\%$ of greedy outputs identical to llama.cpp’s, every system | equality  | ours (6 configurations): 0.000-0.033; FreeToken (6 configurations): 0.000; llama.cpp -ncmoe 32 / 22: 0.700 / 0.400 | <b>failed</b>   |
| 074 P1                                                                                 | ec-bench launch time $\leq 0.25$ ms, C51                             | threshold | 0.201                                                                                                              | held (point)    |
| 074 P2a                                                                                | ours v2 $\geq 127$ tok/s, C51                                        | threshold | 95.1                                                                                                               | <b>failed</b>   |
| 074 P2b                                                                                | ours v2 ahead of FreeToken offload, C51                              | sign      | 0.831 [0.817, 0.845]                                                                                               | <b>failed</b>   |
| 074 P3a                                                                                | ours v2 $\geq 59$ tok/s, C14                                         | threshold | 41.5                                                                                                               | <b>failed</b>   |
| 074 P3b                                                                                | ours v2 $\geq 91$ tok/s, C32                                         | threshold | 65.3                                                                                                               | <b>failed</b>   |
| 074 P4                                                                                 | half-life 64 within $\pm 2\%$ of half-life 16, C51                   | band      | 1.012 [0.997, 1.027]                                                                                               | held (point)    |
| 074b P1                                                                                | GPU sampling cuts between-step host time by $\geq 1$ ms              | threshold | 1.6                                                                                                                | held (point)    |
| 074b P2a                                                                               | ours v2 $\geq 10\%$ faster with GPU sampling, C51                    | threshold | 1.117 [1.109, 1.125]                                                                                               | held            |
| 074b P2b                                                                               | stock llama.cpp $\geq 3\%$ faster with GPU sampling                  | threshold | –                                                                                                                  | <i>untested</i> |
| 074b P3                                                                                | ours v2+FETCH with GPU sampling $\geq$ FreeToken offload, C51        | sign      | 1.074 [1.063, 1.085]                                                                                               | held            |
| 075 P1                                                                                 | ours ahead of FreeToken at 40% (RTX PRO 6000)                        | sign      | 1.122 [1.107, 1.137]                                                                                               | held            |
| 075 P2                                                                                 | ours within -3 to +10% of FreeToken at 60%                           | band      | 1.118 [1.109, 1.127]                                                                                               | <b>failed</b>   |
| 075 P3                                                                                 | ours at 60% $\geq 55\%$ of the all-in-VRAM speed                     | threshold | 0.726 [0.717, 0.735]                                                                                               | held            |
| <b>Jobs 076–081:</b> 23 clauses, 19 held, 4 held (point), 0 failed, 0 untested, 0 void |                                                                      |           |                                                                                                                    |                 |
| 076 P1                                                                                 | ours ahead of FreeToken at 11/25/40%                                 | sign      | 1.205 [1.190, 1.220]; 1.217 [1.201, 1.234]; 1.118 [1.108, 1.129]                                                   | held            |

*continued*

| Job                                                                                     | Clause                                                          | Type      | Measured                                                                                                                          | Status       |
|-----------------------------------------------------------------------------------------|-----------------------------------------------------------------|-----------|-----------------------------------------------------------------------------------------------------------------------------------|--------------|
| 076 P2                                                                                  | ours $\geq 1.8\times$ llama.cpp, every budget                   | threshold | 1.924 [1.894, 1.956]; 2.683 [2.631, 2.735];<br>3.191 [3.130, 3.250]                                                               | held         |
| 076 P3                                                                                  | the law within 8% on own-text runs (four configurations)        | band      | 1.009 [0.962, 1.066]; 1.006 [0.947, 1.086];<br>1.052 [0.960, 1.164]; 0.960 [0.909, 1.024]                                         | held (point) |
| 077 P1a                                                                                 | both entrants run end to end on gpt-oss-120b                    | equality  | 30 of 30 problems at every budget for both                                                                                        | held         |
| 077 P1b                                                                                 | both entrants faster than llama.cpp at 25/40%                   | sign      | pipeshard 25 / 40%: 1.348 [1.336, 1.360];<br>1.417 [1.399, 1.435]; leloch 25 / 40%:<br>1.446 [1.431, 1.461]; 1.280 [1.269, 1.291] | held         |
| 077 P2                                                                                  | ours ahead of both entrants, every budget                       | sign      | 1.421 [1.403, 1.439]; 1.896 [1.871, 1.921];<br>2.343 [2.293, 2.391]                                                               | held         |
| 078 P1                                                                                  | ours ahead of FreeToken at 12.5/25% (Qwen3)                     | sign      | 1.071 [1.063, 1.080]; 1.067 [1.043, 1.094]                                                                                        | held         |
| 078 P2                                                                                  | ours ahead of KTransformers, every budget                       | sign      | 2.168 [2.140, 2.197]; 3.059 [2.975, 3.162];<br>4.420 [4.310, 4.550]                                                               | held         |
| 078 P3                                                                                  | ours $\geq 1.5\times$ llama.cpp, every budget                   | threshold | 2.122 [2.092, 2.151]; 3.051 [2.967, 3.152];<br>4.470 [4.350, 4.609]                                                               | held         |
| 079 P1a                                                                                 | host-side time per step < 1 ms                                  | threshold | 0.266                                                                                                                             | held (point) |
| 079 P1b                                                                                 | host-side time < 10% of the token time                          | threshold | 0.027                                                                                                                             | held (point) |
| 079 P2                                                                                  | $\geq 1.5\times$ FreeToken’s kernels per token                  | threshold | 1.606                                                                                                                             | held (point) |
| 079 P3                                                                                  | ours within $\pm 5\%$ of FreeToken at 43.75%                    | band      | 0.991 [0.975, 1.009]                                                                                                              | held         |
| 080 P1                                                                                  | law table $\geq 3\%$ over the current table, 43.75%             | threshold | 1.048 [1.041, 1.055]                                                                                                              | held         |
| 080 P2                                                                                  | law table within 2% of the sweep’s best, 43.75%                 | band      | 0.991 [0.986, 0.996]                                                                                                              | held         |
| 080 P3                                                                                  | best table leads FreeToken, 43.75%                              | sign      | 1.049 [1.036, 1.064]                                                                                                              | held         |
| 080 P4                                                                                  | law table beats the current table, 25%                          | sign      | 1.067 [1.062, 1.071]                                                                                                              | held         |
| 081 P1                                                                                  | law table beats current, gpt-oss, every budget                  | sign      | 1.079 [1.072, 1.085]; 1.054 [1.041, 1.066];<br>1.033 [1.021, 1.044]                                                               | held         |
| 081 P2                                                                                  | ours leads FreeToken, gpt-oss 11/25/40%                         | sign      | 1.294 [1.278, 1.312]; 1.275 [1.254, 1.295];<br>1.154 [1.134, 1.172]                                                               | held         |
| 081 P3a                                                                                 | law table beats current, Qwen3 12.5%                            | sign      | 1.080 [1.073, 1.084]                                                                                                              | held         |
| 081 P3b                                                                                 | ours leads FreeToken, Qwen3 12.5%                               | sign      | 1.032 [1.022, 1.042]                                                                                                              | held         |
| 081 P4a                                                                                 | ours $\geq 1.8\times$ llama.cpp, gpt-oss, every budget          | threshold | 2.001 [1.964, 2.041]; 2.731 [2.668, 2.795];<br>3.196 [3.133, 3.259]                                                               | held         |
| 081 P4b                                                                                 | ours $\geq 2\times$ llama.cpp, Qwen3 12.5%                      | threshold | 2.051 [2.026, 2.077]                                                                                                              | held         |
| <b>Jobs 082–087:</b> 43 clauses, 24 held, 4 held (point), 10 failed, 1 untested, 4 void |                                                                 |           |                                                                                                                                   |              |
| 082 P1                                                                                  | steady-state lead over FreeToken, gpt-oss, held-out prompts     | sign      | –                                                                                                                                 | untested     |
| 082 P2                                                                                  | FreeToken gains more after 256 tokens than ours, every budget   | sign      | FreeToken: 1.017 / 1.065 / 1.071; ours:<br>1.038 / 1.094 / 1.117                                                                  | failed       |
| 082 P3                                                                                  | FreeToken ahead in steady state, Qwen3 43.75%                   | sign      | 1.054 [1.025, 1.086]                                                                                                              | failed       |
| 082 P4a                                                                                 | every ablation step $\geq -1\%$                                 | threshold | 0.723 [0.712, 0.733]; 0.744 [0.740, 0.748]                                                                                        | failed       |
| 082 P4b                                                                                 | ours $\geq 2\times$ stock at 25%, both models                   | threshold | 2.658 [2.566, 2.749]; 2.800 [2.743, 2.858]                                                                                        | held         |
| 082 P5a                                                                                 | parity: $\geq 98\%$ top-1 agreement, both models                | threshold | 0.986 / 0.993                                                                                                                     | held         |
| 082 P5b                                                                                 | parity: $ \Delta\text{NLL}  \leq 1\%$ , both models             | threshold | -0.0022 / 0.0031                                                                                                                  | held         |
| 082 P6                                                                                  | llama.cpp Qwen3 25/43.75% within 5% of the law                  | band      | 0.966 [0.962, 0.969]; 0.962 [0.958, 0.965]                                                                                        | held         |
| 083 P1                                                                                  | still reasoning at 2,048 tokens on $\geq 8$ of 10, both systems | threshold | ours: 9 / 9 / 10; FreeToken: 10 / 9 / 10                                                                                          | held         |
| 083 P2                                                                                  | ours within 15% of its Table 1 rate                             | band      | 1.046 [1.019, 1.072]; 1.037 [1.013, 1.060];<br>1.033 [1.003, 1.061]                                                               | held         |
| 083 P3a                                                                                 | ours leads FreeToken, all tokens, 11/25/40%                     | sign      | 1.358 [1.338, 1.378]; 1.322 [1.303, 1.342];<br>1.191 [1.172, 1.212]                                                               | held         |
| 083 P3b                                                                                 | ours leads FreeToken, tokens 257-2,048                          | sign      | 1.366 [1.343, 1.389]; 1.328 [1.310, 1.350];<br>1.194 [1.175, 1.215]                                                               | held         |
| 083 P4                                                                                  | lead $\geq 15\%$ at 11 and 25%                                  | threshold | 1.358 [1.338, 1.378]; 1.322 [1.303, 1.342]                                                                                        | held         |
| 084 P1                                                                                  | Qwen3 all in VRAM at 120-190 tok/s                              | band      | –                                                                                                                                 | void         |
| 084 P2                                                                                  | gpt-oss all in VRAM within 5% of 261.4 tok/s                    | band      | –                                                                                                                                 | void         |
| 084 P3                                                                                  | ours with 128 slots within 5% of all in VRAM, both              | band      | –                                                                                                                                 | void         |
| 084 P4                                                                                  | sim. hit rate within 5 points; ours at 20-50% of the limit      | band      | –                                                                                                                                 | void         |
| 084b P1                                                                                 | Qwen3 all in VRAM at 120-190 tok/s                              | band      | 177.8 [177.1, 178.4]                                                                                                              | held         |
| 084b P2                                                                                 | gpt-oss all in VRAM within 5% of 261.4 tok/s                    | band      | 0.977 [0.971, 0.982]                                                                                                              | held         |
| 084b P3                                                                                 | ours with 128 slots within 5% of all in VRAM, both              | band      | 0.997 [0.991, 1.004]; 1.047 [1.043, 1.052]                                                                                        | held (point) |
| 084b P4a                                                                                | sim. hit rate within 5 points of the engine’s, every budget     | band      | -0.014 / -0.008 / -0.0003 / -0.053 / -0.026<br>/ -0.006                                                                           | failed       |
| 084b P4b                                                                                | ours at 20-50% of the speed limit, every budget                 | band      | 0.413 / 0.256 / 0.294 / 0.433 / 0.302 /<br>0.352                                                                                  | held (point) |

continued

| Job                                                                                       | Clause                                                                                                            | Type      | Measured                                                                                                                                            | Status        |
|-------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------|-----------|-----------------------------------------------------------------------------------------------------------------------------------------------------|---------------|
| 084c                                                                                      | <i>No predictions of its own: it supplies the gpt-oss routing trace that 084b-P4a and 084b-P4b are scored on.</i> | –         | –                                                                                                                                                   | –             |
| 085 P1                                                                                    | law’s tables fetch less than on the headline machine                                                              | equality  | gpt-oss: 0,0,0,1,1 (two entries lower); Qwen3: 0,0,0,1,1,2,2,3,3 (six entries lower)                                                                | held          |
| 085 P2a                                                                                   | law beats fixed, gpt-oss, every budget                                                                            | sign      | 1.340 [1.327, 1.355]; 1.274 [1.256, 1.292]; 1.211 [1.187, 1.235]                                                                                    | held          |
| 085 P2b                                                                                   | law’s gain larger than headline at $\geq 2$ of 3 (pro rata of 4 of 6)                                             | threshold | 3 of 3 (every interval excludes the headline value); pooled with 085b: 6 of 6                                                                       | held          |
| 085 P3                                                                                    | ours leads FreeToken, gpt-oss, every budget                                                                       | sign      | 2.030 [1.993, 2.068]; 1.699 [1.662, 1.737]; 1.428 [1.391, 1.463]                                                                                    | held          |
| 085 P4                                                                                    | lead larger than headline at $\geq 2$ of 3 (pro rata of 4 of 6)                                                   | threshold | 3 of 3 (every interval excludes the headline value); pooled with 085b (1 of 3): 4 of 6                                                              | held          |
| 085 P5                                                                                    | llama.cpp at 25% within 5% of the law, gpt-oss                                                                    | band      | 0.818 [0.817, 0.819]                                                                                                                                | <b>failed</b> |
| 085b P1                                                                                   | law’s table fetches less than on the headline machine                                                             | equality  | 0,0,1,1,2,2,3,3,4 (three entries lower)                                                                                                             | held          |
| 085b P2a                                                                                  | law beats fixed, Qwen3, every budget                                                                              | sign      | 1.184 [1.177, 1.190]; 1.149 [1.142, 1.157]; 1.094 [1.086, 1.101]                                                                                    | held          |
| 085b P2b                                                                                  | law’s gain larger than headline at $\geq 2$ of 3                                                                  | threshold | 3 of 3 (every interval excludes the headline value)                                                                                                 | held          |
| 085b P3                                                                                   | ours leads FreeToken, Qwen3, every budget                                                                         | sign      | 1.067 [1.058, 1.076]; 1.076 [1.063, 1.089]; 1.038 [1.017, 1.060]                                                                                    | held          |
| 085b P4                                                                                   | lead larger than headline at $\geq 2$ of 3                                                                        | threshold | 1 of 3: 12.5% 1.067 [1.058, 1.076] > 1.032; 25% 1.076 [1.063, 1.089] < 1.153; 43.75% 1.038 [1.017, 1.060] below 1.049 with the interval covering it | <b>failed</b> |
| 085b P5                                                                                   | llama.cpp at 25% within 5% of the law, Qwen3                                                                      | band      | 0.974 [0.972, 0.976]                                                                                                                                | held          |
| 086 P1                                                                                    | FreeToken as shipped within 15% of its paper’s rate                                                               | band      | 93.8 [91.3, 96.5]                                                                                                                                   | held (point)  |
| 086 P2                                                                                    | ours leads FreeToken at 12.5/25/37.5%                                                                             | sign      | 1.030 [1.023, 1.038]; 1.010 [0.998, 1.021]; 1.003 [0.994, 1.013]                                                                                    | <b>failed</b> |
| 086 P3                                                                                    | ours $\geq 2\times$ llama.cpp, every budget                                                                       | threshold | 1.976 [1.950, 2.005]; 2.311 [2.264, 2.361]; 2.481 [2.441, 2.521]                                                                                    | <b>failed</b> |
| 086 P4                                                                                    | law beats fixed, every budget                                                                                     | sign      | 1.071 [1.068, 1.074]; 1.060 [1.057, 1.063]; 1.057 [1.052, 1.061]                                                                                    | held          |
| 087 P1                                                                                    | profiled static cache $\geq 1.3\times$ stock, both models                                                         | threshold | 1.307 [1.236, 1.379]; 1.889 [1.803, 1.969]                                                                                                          | held (point)  |
| 087 P2                                                                                    | replacement adds less than placement, both models                                                                 | sign      | 0.268 [0.122, 0.434]; -0.765 [-0.898, -0.619]                                                                                                       | <b>failed</b> |
| 087 P3                                                                                    | decayed frequency beats the profiled static cache, both                                                           | sign      | 1.575 [1.497, 1.674]; 1.124 [1.070, 1.187]                                                                                                          | held          |
| 087 P4a                                                                                   | every step after the static cache $\geq -1\%$                                                                     | threshold | 0.877 [0.835, 0.924]                                                                                                                                | <b>failed</b> |
| 087 P4b                                                                                   | ours $\geq 2\times$ stock, both models                                                                            | threshold | 2.680 [2.595, 2.763]; 2.815 [2.754, 2.877]                                                                                                          | held          |
| <b>Jobs 088–094:</b> 157 clauses, 62 held, 17 held (point), 72 failed, 6 untested, 0 void |                                                                                                                   |           |                                                                                                                                                     |               |
| 088 P1                                                                                    | FreeToken hybrid on 8 threads $\geq 1.3\times$ job 085’s hybrid, every budget                                     | threshold | 3.12 / 2.04 / 1.4                                                                                                                                   | held (point)  |
| 088 P2                                                                                    | FreeToken’s best improves on job 085’s best at $\geq 2$ of 3 budgets                                              | threshold | 1.52 / 1.29 / 1.08                                                                                                                                  | held (point)  |
| 088 P3a                                                                                   | ours leads FreeToken’s best, 11% (CI > 1)                                                                         | sign      | 1.582 [1.564, 1.603]                                                                                                                                | held          |
| 088 P3b                                                                                   | ours leads FreeToken’s best, 25% (CI > 1)                                                                         | sign      | 1.519 [1.472, 1.559]                                                                                                                                | held          |
| 088 P3c                                                                                   | ours leads FreeToken’s best, 40% (CI > 1)                                                                         | sign      | 1.506 [1.476, 1.537]                                                                                                                                | held          |
| 088 P3d                                                                                   | ours / FreeToken between 1.10 $\times$ and 1.50 $\times$ , 11%                                                    | band      | 1.582 [1.564, 1.603]                                                                                                                                | <b>failed</b> |
| 088 P3e                                                                                   | ours / FreeToken between 1.10 $\times$ and 1.50 $\times$ , 25%                                                    | band      | 1.519 [1.472, 1.559]                                                                                                                                | <b>failed</b> |
| 088 P3f                                                                                   | ours / FreeToken between 1.10 $\times$ and 1.50 $\times$ , 40%                                                    | band      | 1.506 [1.476, 1.537]                                                                                                                                | <b>failed</b> |
| 088 P4a                                                                                   | llama.cpp -t 8 within 10% of the law                                                                              | band      | -9.4                                                                                                                                                | held (point)  |
| 088 P4b                                                                                   | llama.cpp -t 8 faster than -t 24                                                                                  | sign      | 1.219 [1.213, 1.225]                                                                                                                                | held          |
| 088 P5                                                                                    | the law’s table is 0,0,0,1,1 or fetches less                                                                      | equality  | 0,0,0,1,1                                                                                                                                           | held          |
| 089 P1a                                                                                   | the card reports a 14,001 MHz memory clock                                                                        | equality  | 14001 MHz                                                                                                                                           | held          |
| 089 P1b                                                                                   | device read $\geq 1,500$ GB/s (not throttled)                                                                     | threshold | 1694                                                                                                                                                | held          |
| 089 P2a                                                                                   | ours / FreeToken within $\pm 0.06$ of Table 1, gpt-oss 11%                                                        | band      | -0.088 [-0.099, -0.075]                                                                                                                             | <b>failed</b> |
| 089 P2b                                                                                   | ours / FreeToken within $\pm 0.06$ of Table 1, gpt-oss 25%                                                        | band      | -0.078 [-0.112, -0.051]                                                                                                                             | <b>failed</b> |
| 089 P2c                                                                                   | ours / FreeToken within $\pm 0.06$ of Table 1, gpt-oss 40%                                                        | band      | -0.061 [-0.077, -0.043]                                                                                                                             | <b>failed</b> |
| 089 P2d                                                                                   | ours / FreeToken within $\pm 0.06$ of Table 1, Qwen3 12.5%                                                        | band      | -0.006 [-0.014, 0.004]                                                                                                                              | held          |

continued



| Job     | Clause                                                              | Type      | Measured                            | Status        |
|---------|---------------------------------------------------------------------|-----------|-------------------------------------|---------------|
| 089 P2e | ours / FreeToken within $\pm 0.06$ of Table 1, Qwen3 25%            | band      | -0.105 [-0.116, -0.094]             | <b>failed</b> |
| 089 P2f | ours / FreeToken within $\pm 0.06$ of Table 1, Qwen3 43.75%         | band      | -0.075 [-0.087, -0.062]             | <b>failed</b> |
| 089 P3a | ours' speed within $\pm 12\%$ of listing B's, gpt-oss 11%           | band      | -17.1                               | <b>failed</b> |
| 089 P3b | ours' speed within $\pm 12\%$ of listing B's, gpt-oss 25%           | band      | -13.8                               | <b>failed</b> |
| 089 P3c | ours' speed 5-20% below listing B's, gpt-oss 40%                    | band      | -12.4                               | held (point)  |
| 089 P3d | ours' speed within $\pm 12\%$ of listing B's, Qwen3 12.5%           | band      | -16.0                               | <b>failed</b> |
| 089 P3e | ours' speed within $\pm 12\%$ of listing B's, Qwen3 25%             | band      | -14.6                               | <b>failed</b> |
| 089 P3f | ours' speed 5-20% below listing B's, Qwen3 43.75%                   | band      | -11.5                               | held (point)  |
| 089 P4  | ours-first and FreeToken-first orders agree within 3%, every cell   | band      | -1.0 / 0.2 / 0.7 / 0.3 / -0.2 / 0.1 | held (point)  |
| 089 P5a | LRU 10-30% slower than decayed frequency, gpt-oss 11%               | band      | 18.5 [17.7, 19.3]                   | held          |
| 089 P5b | LRU 10-30% slower than decayed frequency, gpt-oss 25%               | band      | 12.8 [11.5, 14.1]                   | held          |
| 089 P5c | LRU 10-30% slower than decayed frequency, gpt-oss 40%               | band      | 7.8 [6.5, 9.1]                      | <b>failed</b> |
| 089 P5d | LRU 10-30% slower than decayed frequency, Qwen3 12.5%               | band      | 20.9 [20.6, 21.3]                   | held          |
| 089 P5e | LRU 10-30% slower than decayed frequency, Qwen3 25%                 | band      | 15.4 [14.9, 15.9]                   | held          |
| 089 P5f | LRU 10-30% slower than decayed frequency, Qwen3 43.75%              | band      | 6.4 [5.9, 6.9]                      | <b>failed</b> |
| 089 P5g | LRU trails FreeToken at Qwen3 12.5%                                 | sign      | 0.814 [0.808, 0.821]                | held          |
| 089 P5h | LRU still leads FreeToken at gpt-oss 11%                            | sign      | 0.973 [0.962, 0.986]                | <b>failed</b> |
| 089 P5i | LRU still leads FreeToken at gpt-oss 25%                            | sign      | 1.045 [1.016, 1.068]                | held          |
| 089 P6a | prefill at 2,048 tokens within $\pm 20\%$ of stock, gpt-oss         | band      | 0.800                               | held (point)  |
| 089 P6b | prefill at 2,048 tokens within $\pm 20\%$ of stock, Qwen3           | band      | 0.825                               | held (point)  |
| 089 P6c | decode at 2 and 4 sequences below stock's, gpt-oss                  | sign      | 0.778 / 0.778 / 0.773 / 0.784       | held (point)  |
| 089 P6d | decode at 2 and 4 sequences below stock's, Qwen3                    | sign      | 0.779 / 0.789 / 0.775 / 0.792       | held (point)  |
| 090 P1a | mean KL $\leq 0.01$ nats, gpt-oss                                   | threshold | 0.0019                              | held          |
| 090 P1b | mean KL $\leq 0.01$ nats, Qwen3                                     | threshold | 0.0005                              | held          |
| 090 P1c | 99.9th-percentile KL $\leq 0.5$ nats, gpt-oss                       | threshold | 0.121                               | held          |
| 090 P1d | 99.9th-percentile KL $\leq 0.5$ nats, Qwen3                         | threshold | 0.059                               | held          |
| 090 P2a | top-1 agreement $\geq 98\%$ , gpt-oss                               | threshold | 98.5                                | held          |
| 090 P2b | top-1 agreement $\geq 98\%$ , Qwen3                                 | threshold | 99.3                                | held          |
| 090 P2c | stock's -n-cpu-moe agreement within 1 point of ours, gpt-oss        | band      | 0.33                                | held          |
| 090 P2d | stock's -n-cpu-moe agreement within 1 point of ours, Qwen3          | band      | -0.07                               | held          |
| 090 P3a | mean NLL within $\pm 0.5\%$ of the reference, gpt-oss               | band      | 0.16                                | held          |
| 090 P3b | mean NLL within $\pm 0.5\%$ of the reference, Qwen3                 | band      | 0.19                                | held          |
| 091 P1a | ours leads FreeToken, gpt-oss 11% (CI > 1)                          | sign      | 1.371 [1.346, 1.395]                | held          |
| 091 P1b | ours leads FreeToken, gpt-oss 25% (CI > 1)                          | sign      | 1.296 [1.265, 1.326]                | held          |
| 091 P1c | ours leads FreeToken, Qwen3 12.5% (CI > 1)                          | sign      | 1.033 [1.021, 1.045]                | held          |
| 091 P1d | ours leads FreeToken, Qwen3 25% (CI > 1)                            | sign      | 1.045 [1.028, 1.063]                | held          |
| 091 P1e | lead in 10-35%, gpt-oss 11%                                         | band      | 37.1 [34.6, 39.5]                   | <b>failed</b> |
| 091 P1f | lead in 10-35%, gpt-oss 25%                                         | band      | 29.6 [26.5, 32.6]                   | held          |
| 091 P1g | lead in 0-15%, Qwen3 12.5%                                          | band      | 3.3 [2.1, 4.5]                      | held          |
| 091 P1h | lead in 0-15%, Qwen3 25%                                            | band      | 4.5 [2.8, 6.3]                      | held          |
| 091 P2a | ours $\geq 1.8\times$ llama.cpp, gpt-oss 11%                        | threshold | 1.810 [1.760, 1.860]                | held (point)  |
| 091 P2b | ours $\geq 1.8\times$ llama.cpp, gpt-oss 25%                        | threshold | 2.620 [2.520, 2.740]                | held          |
| 091 P2c | ours $\geq 1.8\times$ llama.cpp, Qwen3 12.5%                        | threshold | 1.780 [1.760, 1.810]                | <b>failed</b> |
| 091 P2d | ours $\geq 1.8\times$ llama.cpp, Qwen3 25%                          | threshold | 2.540 [2.490, 2.610]                | held          |
| 091 P3a | the law's gpt-oss table fetches no more than the headline machine's | equality  | 0,0,1,2,2                           | <b>failed</b> |
| 091 P3b | the law's Qwen3 table fetches no more than the headline machine's   | equality  | 0,0,1,1,2,3,4,4,5                   | <b>failed</b> |
| 091 P4a | ours at 25-45% of the limit, gpt-oss 11%                            | band      | 44.4 [43.3, 45.5]                   | held (point)  |
| 091 P4b | ours at 25-45% of the limit, gpt-oss 25%                            | band      | 29.9 [29.1, 30.8]                   | held          |
| 091 P4c | ours at 25-45% of the limit, Qwen3 12.5%                            | band      | 45.7 [45.1, 46.3]                   | <b>failed</b> |

continued

| Job     | Clause                                                              | Type      | Measured               | Status          |
|---------|---------------------------------------------------------------------|-----------|------------------------|-----------------|
| 091 P4d | ours at 25-45% of the limit, Qwen3 25%                              | band      | 34.3 [33.5, 35.2]      | held            |
| 091 P5  | ours-first and FreeToken-first orders agree within 3%, every cell   | band      | 1.6 / 2.0 / -0.6 / 0.1 | held (point)    |
| 092 P1a | ours leads FreeToken, gpt-oss 11% (CI > 1)                          | sign      | 4.460 [4.250, 4.690]   | held            |
| 092 P1b | ours leads FreeToken, Qwen3 12.5% (CI > 1)                          | sign      | 2.120 [2.020, 2.210]   | held            |
| 092 P1c | ours leads FreeToken, Qwen3 25% (CI > 1)                            | sign      | 2.680 [2.530, 2.830]   | held            |
| 092 P1d | lead in 10-35%, gpt-oss 11%                                         | band      | 346 [325, 369]         | <b>failed</b>   |
| 092 P1e | lead in 0-15%, Qwen3 12.5%                                          | band      | 112 [102, 121]         | <b>failed</b>   |
| 092 P1f | lead in 0-15%, Qwen3 25%                                            | band      | 168 [153, 183]         | <b>failed</b>   |
| 092 P2a | ours $\geq 1.8 \times$ llama.cpp, gpt-oss 11%                       | threshold | 1.640 [1.600, 1.670]   | <b>failed</b>   |
| 092 P2b | ours $\geq 1.8 \times$ llama.cpp, Qwen3 12.5%                       | threshold | 1.560 [1.480, 1.630]   | <b>failed</b>   |
| 092 P2c | ours $\geq 1.8 \times$ llama.cpp, Qwen3 25%                         | threshold | 2.180 [2.090, 2.270]   | held            |
| 092 P3a | ours $\geq 1.3 \times$ llama.cpp on Mixtral-8x7B, C = 2             | threshold | 1.015 [1.008, 1.024]   | <b>failed</b>   |
| 092 P3b | ours $\geq 1.3 \times$ llama.cpp on Mixtral-8x7B, C = 4             | threshold | –                      | <i>untested</i> |
| 092 P4a | the law’s gpt-oss table fetches no more than the headline machine’s | equality  | 0,0,0,0,0              | held            |
| 092 P4b | the law’s Qwen3 table fetches no more than the headline machine’s   | equality  | 0,0,0,0,0,0,0,0        | held            |
| 092 P5a | ours-first and FreeToken-first orders agree within 3%, Qwen3 12.5%  | band      | 7.0                    | <b>failed</b>   |
| 092 P5b | ours-first and FreeToken-first orders agree within 3%, Qwen3 25%    | band      | 2.6                    | held (point)    |
| 093 P1a | oracle (W = all) faster than base, gpt-oss 11%                      | sign      | 0.923 [0.914, 0.933]   | <b>failed</b>   |
| 093 P1b | oracle (W = all) faster than base, gpt-oss 25%                      | sign      | 1.398 [1.384, 1.412]   | held            |
| 093 P1c | oracle (W = all) faster than base, Qwen3 12.5%                      | sign      | 0.916 [0.911, 0.922]   | <b>failed</b>   |
| 093 P1d | oracle (W = all) faster than base, Qwen3 25%                        | sign      | 1.247 [1.235, 1.259]   | held            |
| 093 P2a | oracle gain 20-80%, gpt-oss 11%                                     | band      | -7.7 [-8.6, -6.7]      | <b>failed</b>   |
| 093 P2b | oracle gain 20-80%, gpt-oss 25%                                     | band      | 39.8 [38.4, 41.2]      | held            |
| 093 P2c | oracle gain 20-80%, Qwen3 12.5%                                     | band      | -8.4 [-8.9, -7.8]      | <b>failed</b>   |
| 093 P2d | oracle gain 20-80%, Qwen3 25%                                       | band      | 24.7 [23.5, 25.9]      | held            |
| 093 P3a | W = 16 captures $\geq$ half of the W = 0 gain, gpt-oss 11%          | threshold | –                      | <i>untested</i> |
| 093 P3b | W = 16 captures $\geq$ half of the W = 0 gain, gpt-oss 25%          | threshold | 87                     | held (point)    |
| 093 P3c | W = 16 captures $\geq$ half of the W = 0 gain, Qwen3 12.5%          | threshold | –                      | <i>untested</i> |
| 093 P3d | W = 16 captures $\geq$ half of the W = 0 gain, Qwen3 25%            | threshold | 100                    | held (point)    |
| 093 P3e | W = 4 captures $\geq$ half of the W = 0 gain, gpt-oss 11%           | threshold | –                      | <i>untested</i> |
| 093 P3f | W = 4 captures $\geq$ half of the W = 0 gain, Qwen3 12.5%           | threshold | –                      | <i>untested</i> |
| 093 P4a | ours with the oracle at 45-65% of the limit, gpt-oss 11%            | band      | 44.0 [43.2, 44.7]      | <b>failed</b>   |
| 093 P4b | ours with the oracle at 45-65% of the limit, gpt-oss 25%            | band      | 43.7 [42.5, 44.8]      | <b>failed</b>   |
| 093 P4c | ours with the oracle at 45-65% of the limit, Qwen3 12.5%            | band      | 44.8 [44.1, 45.5]      | <b>failed</b>   |
| 093 P4d | ours with the oracle at 45-65% of the limit, Qwen3 25%              | band      | 44.5 [43.3, 45.8]      | <b>failed</b>   |
| 093 P5a | oracle hit rate within 3 points of the optimum’s, gpt-oss 11%       | band      | -5.9                   | <b>failed</b>   |
| 093 P5b | oracle hit rate within 3 points of the optimum’s, gpt-oss 25%       | band      | 3.1                    | <b>failed</b>   |
| 093 P5c | oracle hit rate within 3 points of the optimum’s, Qwen3 12.5%       | band      | -9.2                   | <b>failed</b>   |
| 093 P5d | oracle hit rate within 3 points of the optimum’s, Qwen3 25%         | band      | 0.4                    | held            |
| 093 P5e | oracle admissions within 30% of the optimum’s reads, gpt-oss 11%    | band      | -9                     | held            |
| 093 P5f | oracle admissions within 30% of the optimum’s reads, gpt-oss 25%    | band      | 31                     | <b>failed</b>   |
| 093 P5g | oracle admissions within 30% of the optimum’s reads, Qwen3 12.5%    | band      | -35                    | <b>failed</b>   |
| 093 P5h | oracle admissions within 30% of the optimum’s reads, Qwen3 25%      | band      | 7                      | held            |
| 093 P6a | no-overlap 5-25% slower than base, gpt-oss 11%                      | band      | -0.5 [-1.0, -0.1]      | <b>failed</b>   |
| 093 P6b | no-overlap 5-25% slower than base, gpt-oss 25%                      | band      | 0.2 [0.1, 0.2]         | <b>failed</b>   |
| 093 P6c | no-overlap 5-25% slower than base, Qwen3 12.5%                      | band      | -0.0 [-0.1, 0.0]       | <b>failed</b>   |

*continued*

| Job     | Clause                                                          | Type      | Measured             | Status          |
|---------|-----------------------------------------------------------------|-----------|----------------------|-----------------|
| 093 P6d | no-overlap 5-25% slower than base, Qwen3 25%                    | band      | 0.1 [0.0, 0.2]       | <b>failed</b>   |
| 093 P6e | all-CPU 0-10% slower than base, gpt-oss 11%                     | band      | 7.8 [7.2, 8.4]       | held            |
| 093 P6f | all-CPU 0-10% slower than base, gpt-oss 25%                     | band      | 9.3 [8.7, 9.8]       | held            |
| 093 P6g | all-CPU 0-10% slower than base, Qwen3 12.5%                     | band      | 14.7 [14.1, 15.3]    | <b>failed</b>   |
| 093 P6h | all-CPU 0-10% slower than base, Qwen3 25%                       | band      | 14.6 [14.0, 15.2]    | <b>failed</b>   |
| 093 P7a | oracle gain below 15% at the GPU-bound cell, gpt-oss 40%        | threshold | 48.5 [46.2, 50.9]    | <b>failed</b>   |
| 093 P7b | oracle gain below 15% at the GPU-bound cell, Qwen3 43.75%       | threshold | 56.5 [53.3, 59.2]    | <b>failed</b>   |
| 094 P1a | bypass oracle's reads within 15% of the optimum's, gpt-oss 11%  | band      | 30.8                 | <b>failed</b>   |
| 094 P1b | bypass oracle's reads within 15% of the optimum's, gpt-oss 25%  | band      | 31.1                 | <b>failed</b>   |
| 094 P1c | bypass oracle's reads within 15% of the optimum's, gpt-oss 40%  | band      | 29.0                 | <b>failed</b>   |
| 094 P1d | bypass oracle's reads within 15% of the optimum's, Qwen3 12.5%  | band      | 43.4                 | <b>failed</b>   |
| 094 P1e | bypass oracle's reads within 15% of the optimum's, Qwen3 25%    | band      | 46.6                 | <b>failed</b>   |
| 094 P1f | bypass oracle's reads within 15% of the optimum's, Qwen3 43.75% | band      | 38.5                 | <b>failed</b>   |
| 094 P2a | bypass oracle faster than base, gpt-oss 11%                     | sign      | 1.033 [1.029, 1.038] | held            |
| 094 P2b | bypass oracle faster than base, gpt-oss 25%                     | sign      | 1.181 [1.175, 1.187] | held            |
| 094 P2c | bypass oracle faster than base, gpt-oss 40%                     | sign      | 1.232 [1.220, 1.243] | held            |
| 094 P2d | bypass oracle faster than base, Qwen3 12.5%                     | sign      | 1.002 [0.997, 1.006] | held (point)    |
| 094 P2e | bypass oracle faster than base, Qwen3 25%                       | sign      | 1.131 [1.122, 1.139] | held            |
| 094 P2f | bypass oracle faster than base, Qwen3 43.75%                    | sign      | 1.189 [1.178, 1.200] | held            |
| 094 P2g | bypass gain 10-45%, gpt-oss 11%                                 | band      | 3.3 [2.9, 3.8]       | <b>failed</b>   |
| 094 P2h | bypass gain 10-45%, gpt-oss 25%                                 | band      | 18.1 [17.5, 18.7]    | held            |
| 094 P2i | bypass gain 10-45%, Qwen3 12.5%                                 | band      | 0.2 [-0.3, 0.6]      | <b>failed</b>   |
| 094 P2j | bypass gain 10-45%, Qwen3 25%                                   | band      | 13.1 [12.2, 13.9]    | held            |
| 094 P3a | bypass hit rate within 4 points of the optimum's, gpt-oss 11%   | band      | -8.2                 | <b>failed</b>   |
| 094 P3b | bypass hit rate within 4 points of the optimum's, gpt-oss 25%   | band      | -3.2                 | held            |
| 094 P3c | bypass hit rate within 4 points of the optimum's, gpt-oss 40%   | band      | -1.4                 | held            |
| 094 P3d | bypass hit rate within 4 points of the optimum's, Qwen3 12.5%   | band      | -10.9                | <b>failed</b>   |
| 094 P3e | bypass hit rate within 4 points of the optimum's, Qwen3 25%     | band      | -5.3                 | <b>failed</b>   |
| 094 P3f | bypass hit rate within 4 points of the optimum's, Qwen3 43.75%  | band      | -1.3                 | held            |
| 094 P4a | law over-predicts the bypass oracle by 10-40%, gpt-oss 11%      | band      | 6.9                  | <b>failed</b>   |
| 094 P4b | law over-predicts the bypass oracle by 10-40%, gpt-oss 25%      | band      | 5.8                  | <b>failed</b>   |
| 094 P4c | law over-predicts the bypass oracle by 10-40%, Qwen3 12.5%      | band      | 4.7                  | <b>failed</b>   |
| 094 P4d | law over-predicts the bypass oracle by 10-40%, Qwen3 25%        | band      | 8.5                  | <b>failed</b>   |
| 094 P5a | bypass beats prefetch by $\geq 15\%$ , gpt-oss 11%              | threshold | 6.9                  | <b>failed</b>   |
| 094 P5b | bypass beats prefetch by $\geq 15\%$ , Qwen3 12.5%              | threshold | 5.9                  | <b>failed</b>   |
| 094 P5c | bypass and prefetch within 15%, gpt-oss 25%                     | band      | -27.6                | <b>failed</b>   |
| 094 P5d | bypass and prefetch within 15%, gpt-oss 40%                     | band      | -21.3                | <b>failed</b>   |
| 094 P5e | bypass and prefetch within 15%, Qwen3 25%                       | band      | -19.7                | <b>failed</b>   |
| 094 P5f | bypass and prefetch within 15%, Qwen3 43.75%                    | band      | -26.6                | <b>failed</b>   |
| 094 P6a | W = 16 keeps $\geq 70\%$ of the W = all gain, gpt-oss 11%       | threshold | 176                  | held            |
| 094 P6b | W = 16 keeps $\geq 70\%$ of the W = all gain, gpt-oss 25%       | threshold | 63                   | <b>failed</b>   |
| 094 P6c | W = 16 keeps $\geq 70\%$ of the W = all gain, gpt-oss 40%       | threshold | 34                   | <b>failed</b>   |
| 094 P6d | W = 16 keeps $\geq 70\%$ of the W = all gain, Qwen3 12.5%       | threshold | –                    | <i>untested</i> |
| 094 P6e | W = 16 keeps $\geq 70\%$ of the W = all gain, Qwen3 25%         | threshold | 94                   | held            |
| 094 P6f | W = 16 keeps $\geq 70\%$ of the W = all gain, Qwen3 43.75%      | threshold | 53                   | <b>failed</b>   |

Table 7: The grid: Table 1’s protocol on every card and host we measured, each cell against its own machine’s limit (the exact optimum’s reads on the same trace, the machine’s highest probed host rate, the card’s datasheet rate). *Bound*: what binds the limit; *host* when the optimum’s reads alone take longer than the GPU’s whole read, so the GPU has slack, *both* when the limit runs experts on the CPU until the two paths take equally long, which is where the all-in-VRAM speed can be exceeded. Ours and FreeToken are launch 2 (FreeToken’s backend carried over from the headline machine’s selection); the 40 and 43.75% budgets do not fit a 24 GB card. †On the RTX 3090 FreeToken’s carried-over hybrid backend ran below llama.cpp; its own calibration recommends its offload backend there, which we did not run, so those cells are reported, not counted. Mixtral-8x7B (26 GB of experts, top-2 of 8) is ours against llama.cpp only, on the RTX 3090, launch 1.

| Model                                                                                                                             | Experts on GPU | llama.cpp (tok/s) | FreeToken (tok/s) | Ours (tok/s) | Ours ÷ FreeToken (95% CI)      | Ours ÷ llama.cpp     | Limit (tok/s) | Bound | Ours, % of limit |
|-----------------------------------------------------------------------------------------------------------------------------------|----------------|-------------------|-------------------|--------------|--------------------------------|----------------------|---------------|-------|------------------|
| <b>RTX 5090, 9950X3D</b> (listing B): $B_{\text{host}}$ 88 GB/s, $B_{\text{gpu}}$ 1792 GB/s; tables 0,0,1,1,2 / 0,0,1,1,2,3,3,4,5 |                |                   |                   |              |                                |                      |               |       |                  |
| gpt-oss-120b                                                                                                                      | 11%            | 34.9              | 54.0              | 69.9         | 1.29 [1.28, 1.31]              | 2.00                 | 172           | host  | 41               |
| gpt-oss-120b                                                                                                                      | 25%            | 40.0              | 85.6              | 109.2        | 1.27 [1.25, 1.29]              | 2.73                 | 430           | host  | 25               |
| gpt-oss-120b                                                                                                                      | 40%            | 47.7              | 132.2             | 152.5        | 1.15 [1.13, 1.17]              | 3.20                 | 519           | both  | 29               |
| Qwen3-30B-A3B                                                                                                                     | 12.5%          | 19.5              | 38.7              | 40.0         | 1.03 [1.02, 1.04]              | 2.05                 | 96            | host  | 42               |
| Qwen3-30B-A3B                                                                                                                     | 25%            | 22.1              | 54.7              | 63.1         | 1.15 [1.13, 1.17]              | 2.86                 | 214           | host  | 30               |
| Qwen3-30B-A3B                                                                                                                     | 43.75%         | 28.4              | 103.1             | 108.2        | 1.05 [1.04, 1.06]              | 3.81                 | 308           | both  | 35               |
| <b>RTX 5090, 9950X</b> (job 089): $B_{\text{host}}$ 71 GB/s, $B_{\text{gpu}}$ 1792 GB/s; tables 0,0,1,2,3 / 0,0,1,2,2,3,4,5,6     |                |                   |                   |              |                                |                      |               |       |                  |
| gpt-oss-120b                                                                                                                      | 11%            | 28.9              | 48.0              | 57.9         | 1.21 [1.19, 1.22]              | 2.00                 | 140           | host  | 41               |
| gpt-oss-120b                                                                                                                      | 25%            | 33.6              | 78.7              | 94.1         | 1.20 [1.16, 1.22]              | 2.80                 | 351           | host  | 27               |
| gpt-oss-120b                                                                                                                      | 40%            | 40.1              | 122.2             | 133.6        | 1.09 [1.08, 1.11]              | 3.34                 | 514           | both  | 26               |
| Qwen3-30B-A3B                                                                                                                     | 12.5%          | 16.1              | 32.7              | 33.5         | 1.03 [1.02, 1.04]              | 2.08                 | 78            | host  | 43               |
| Qwen3-30B-A3B                                                                                                                     | 25%            | 18.5              | 51.4              | 53.9         | 1.05 [1.04, 1.06]              | 2.91                 | 174           | host  | 31               |
| Qwen3-30B-A3B                                                                                                                     | 43.75%         | 23.8              | 98.3              | 95.8         | 0.97 [0.96, 0.99]              | 4.02                 | 305           | both  | 31               |
| <b>RTX 4090, i5-12400</b> (job 091): $B_{\text{host}}$ 37 GB/s, $B_{\text{gpu}}$ 1008 GB/s; tables 0,0,1,2,2 / 0,0,1,1,2,3,4,4,5  |                |                   |                   |              |                                |                      |               |       |                  |
| gpt-oss-120b                                                                                                                      | 11%            | 18.1              | 23.8              | 32.7         | 1.37 [1.35, 1.39]              | 1.81                 | 74            | host  | 44               |
| gpt-oss-120b                                                                                                                      | 25%            | 21.0              | 42.5              | 55.0         | 1.30 [1.26, 1.33]              | 2.62                 | 184           | host  | 30               |
| Qwen3-30B-A3B                                                                                                                     | 12.5%          | 10.5              | 18.2              | 18.8         | 1.03 [1.02, 1.05]              | 1.78                 | 41            | host  | 46               |
| Qwen3-30B-A3B                                                                                                                     | 25%            | 12.3              | 30.0              | 31.3         | 1.05 [1.03, 1.06]              | 2.54                 | 91            | host  | 34               |
| <b>RTX 3090, i9-11900KF</b> (job 092): $B_{\text{host}}$ 45 GB/s, $B_{\text{gpu}}$ 936 GB/s; tables 0,0,0,0,0 / 0,0,0,0,0,0,0,0   |                |                   |                   |              |                                |                      |               |       |                  |
| gpt-oss-120b                                                                                                                      | 11%            | 22.8              | 8.4 <sup>†</sup>  | 37.3         | 4.46 [4.25, 4.69] <sup>†</sup> | 1.64                 | 88            | host  | 42               |
| Qwen3-30B-A3B                                                                                                                     | 12.5%          | 12.3              | 9.0 <sup>†</sup>  | 19.2         | 2.12 [2.02, 2.21] <sup>†</sup> | 1.56                 | 49            | host  | 39               |
| Qwen3-30B-A3B                                                                                                                     | 25%            | 13.8              | 11.2 <sup>†</sup> | 30.1         | 2.68 [2.53, 2.83] <sup>†</sup> | 2.18                 | 110           | host  | 27               |
| Mixtral-8x7B Q4_K_M $C=2$ (25%)                                                                                                   |                | 8.22              | —                 | 8.34         | —                              | 1.015 [1.008, 1.024] | —             | —     | —                |

## C The grid

Table 7 lists every cell of Table 1’s protocol on the four machines we measured, each against its own limit; Section 7 discusses it and Fig. 5 plots the fractions. The 24 GB cards run the cells that fit (the VRAM arithmetic, written on the machine from the GGUF header, skipped Qwen3 at 43.75%: 28.0 GB of slots, dense weights and KV against a 23 GB limit). The predictions for jobs 091 and 092 and their outcomes are in `prereg/grid_outcome_091.md` and `092.md`: on the 4090 our lead over FreeToken at gpt-oss 11% (37%) exceeded the predicted 10–35% band, ours ran  $1.78\times$  llama.cpp at Qwen3 12.5% against a predicted 1.8, the law’s tables copied one more expert than the headline machine’s at one entry each (the prediction reasoned from the link alone; the CPU is slower in the same proportion), and ours’ fraction of the limit at Qwen3 12.5% (45.7%) exceeded the predicted 25–45%; on the 3090 FreeToken’s carried-over backend ran below llama.cpp, ours ran  $1.56\text{--}1.64\times$  llama.cpp at two cells against a predicted 1.8, Mixtral tied at  $C=2$  against a predicted  $1.3\times$  and its  $C=4$  cell was lost to the job’s deadline, and one launch-order check exceeded 3% (7%).

## D Foresight in the engine

Table 8 lists every cell of the hit-optimal oracle’s measurement of Section 4 on the first host (a Ryzen 9 9950X whose probe reads 44 GB/s by the CPU at the helper count, 46 over the link, 49 with both and 50.5 at best; the law’s FETCH tables 0,0,1,2,3 for gpt-oss and 0,0,1,2,3,4,5,6,7 for Qwen3; every cell but gpt-oss 40% and Qwen3 43.75% host-bound at the limit), and Table 9 the bytes-optimal oracle beside the hit-optimal one on the second (a Ryzen 9 7950X: 44, 46, 50 and 50.6 GB/s; the same tables). The hit-optimal oracle’s window sweep on the first host



Table 8: The hit-optimal oracle in the engine (job 093; RTX 5090, Ryzen 9 9950X, host memory 50.5 GB/s at best). tok/s of the online policy and of the oracle with the whole sequence in view; their ratio, paired by sequence, with its 95% interval; hit rate and host reads per token (misses plus admissions) of each, with the optimum’s in parentheses; speed as a percentage of this host’s limit; the accounting’s foresight-only term  $v(F)$  for the cell in tok/s; the share of that term the oracle recovered and the share of the gap to the limit it closed.

| Cell         | online | oracle | ratio             | hits % (opt.) | reads/token (opt.) | % of limit  | $v(F)$ | recovered % | closed % |
|--------------|--------|--------|-------------------|---------------|--------------------|-------------|--------|-------------|----------|
| gpt-oss 11%  | 47.4   | 43.7   | 0.92 [0.91, 0.93] | 58 → 67 (73)  | 63 → 82 (38)       | 47.6 → 44.0 | 68.8   | −27         | −16      |
| gpt-oss 25%  | 77.6   | 108.5  | 1.40 [1.38, 1.41] | 79 → 92 (89)  | 34 → 31 (15)       | 31.2 → 43.7 | 124.4  | 76          | 41       |
| gpt-oss 40%  | 112.4  | 166.9  | 1.48 [1.46, 1.51] | 89 → 98 (95)  | 19 → 15 (7)        | 22.1 → 32.8 | 174.0  | 92          | 42       |
| Qwen3 12.5%  | 27.1   | 24.8   | 0.92 [0.91, 0.92] | 60 → 66 (75)  | 160 → 195 (97)     | 48.9 → 44.8 | 40.2   | −28         | −18      |
| Qwen3 25%    | 44.0   | 54.9   | 1.25 [1.24, 1.26] | 78 → 89 (89)  | 90 → 88 (43)       | 35.7 → 44.5 | 71.2   | 52          | 31       |
| Qwen3 43.75% | 78.7   | 123.1  | 1.56 [1.53, 1.59] | 92 → 99 (96)  | 38 → 30 (14)       | 26.1 → 40.8 | 121.3  | 103         | 49       |

(ratio to the online policy; Fig. 2 plots the same runs as fractions of the limit): gpt-oss 25% runs 0.81, 0.98, 1.35, 1.40 and  $1.40\times$  at  $W=2, 4, 16, 64$  and all, Qwen3 25% 0.99, 1.07, 1.25, 1.25 and 1.25; at gpt-oss 11% and Qwen3 12.5% every window loses ( $0.78\text{--}0.92\times$ ). An unpaced copy path (each copy published two steps after it is issued), run at the four host-bound cells, reaches 98–99% hits at three of them (75% at Qwen3 12.5%) and runs slower than the paced oracle at three ( $0.96, 1.18, 0.72$  and  $1.14\times$  the online policy); at gpt-oss 11% it is the faster of the two and still loses, so copies that arrive in time do not remove the loss there: the extra bytes do. With no FETCH table (every miss on the CPU) the online policy runs 8–15% slower at the four host-bound cells; with the engine’s within-layer overlap turned off it changes by under 1% at every cell ( $0.5\%$  faster to  $0.6\%$  slower). The law, with its GPU term calibrated on the online run, evaluated on each hit-optimal run’s own counters over-predicts its time by 24–39% at the four cells where that oracle gains on the first host (12.25 against a measured 9.22 ms at gpt-oss 25%, 7.89 against 5.99 at 40%, 22.5 against 18.2 at Qwen3 25% and 11.3 against 8.1 at 43.75%) and by 37–53% on the second; on the bytes-optimal oracle it is within 9% at every cell. The bytes-optimal oracle with a 16-token window keeps 34–176% of the whole-sequence gain at the five cells where that gain is positive (more than all of it at gpt-oss 11%, where the shorter window admits less).

**Why the bytes-optimal oracle reads more than the optimum.** Its misses are  $1.29\text{--}1.47\times$  the optimum’s reads. The engine’s paced copy path moves one 16 MB piece at a time from pageable host memory on a single copier thread, behind the demand copies, and publishes a landed copy at the start of the next step; at 16–30 admissions per token (210–280 MB) it lags the admissions by steps, and an expert still in flight is a miss the CPU reads again. The last four columns of Table 9 simulate MIN with bypass on the measured routing trace with the engine’s own semantics (the lookahead within the sequence, residents carried over) and a copy that lands  $d$  steps after it is issued, during which the expert misses and is not re-admitted: at  $d=1$ , the ideal, the misses come within 2–14% of the limit’s optimum at the host-bound cells (the optimum also sees across sequences); the engine’s misses meet the simulation at  $d=2.6\text{--}3.6$  at the host-bound cells and at or below 1.4 at the GPU-bound ones, where the admissions are few. The predictions and their outcomes are in `prereg/foresight_outcome_093.md` and `094.md`: of 36 clauses for the first job, 9 held, 2 held on the point estimate, 21 failed and 4 were untested; of 38 for the second, 12 held, 1 on the point estimate, 24 failed and 1 was untested (Appendix B). The second job’s failures are of size, not direction: the bypass oracle’s reads were predicted within 15% of the optimum’s and came 29–47% above (the copy latency); its gain was predicted at 10–45% at the host-bound cells and came below that at the two lowest budgets; the law was predicted to over-predict it by 10–40% and over-predicted by 5–9%; and the two oracles were predicted within 15% of each other at 25% and above, where the hit-optimal one was  $1.25\text{--}1.38\times$  faster.

## E The expert cache in more detail

This section holds the measurements of Section 7 that the main text summarises: the choice of the FETCH split, the hosts with half the link, long outputs, FreeToken’s own headline model, two other llama.cpp-based systems, output parity, and the cost of prompts and batches. Job numbers refer to the scripts on the `gpu` branch.

**Choosing the split per machine.** On listing B, the law’s table beats the fixed table we used before at every budget of both models, by 3.3–8.0% (launch 1). Table 10 shows why at one budget: a single missed expert should run on the CPU (about  $155\mu\text{s}$ , hidden behind the GPU’s other experts) rather than be copied (about  $225\mu\text{s}$ , on the critical path); that entry alone is worth 3.7%. Fetching everything, FreeToken’s offload strategy expressed in our

Table 9: The bytes-optimal oracle (MIN with bypass) beside the hit-optimal one on a fourth host (RTX 5090, Ryzen 9 7950X, host memory 50.6 GB/s at best). Per cell: the online policy’s speed; each oracle’s ratio to it with the whole sequence in view (paired, 95% interval), its hit rate, its host reads per token as the engine counts them (misses, each run on the CPU or fetched on demand; admissions in parentheses) and the law on its own counters against its measured time; the optimum’s reads and hit rate from the trace; and MIN with bypass simulated on the trace with the engine’s lookahead (within the sequence) under a copy latency of 1 (ideal), 2 and 3 steps, with the latency at which the simulation meets the engine’s misses.

| Cell         | online | ratio | bypass oracle, W=all |       |          |     | ratio | hit-optimal oracle, W=all |       |          |      | optimum |      | simulated misses/token |     |     |     |
|--------------|--------|-------|----------------------|-------|----------|-----|-------|---------------------------|-------|----------|------|---------|------|------------------------|-----|-----|-----|
|              |        |       | hits                 | reads | (adm.)   | law |       | hits                      | reads | (adm.)   | law  | reads   | hits | d=1                    | 2   | 3   | fit |
| gpt-oss 11%  | 48.0   | 1.03  | [1.03, 1.04]         | 65    | 50 (16)  | +7% | 0.97  | [0.96, 0.98]              | 71    | 42 (39)  | +19% | 38.3    | 73   | 39                     | 44  | 48  | 3.6 |
| gpt-oss 25%  | 78.6   | 1.18  | [1.18, 1.19]         | 86    | 20 (9)   | +6% | 1.63  | [1.61, 1.66]              | 95    | 7 (23)   | +53% | 15.3    | 89   | 18                     | 19  | 21  | 2.6 |
| gpt-oss 40%  | 113.8  | 1.23  | [1.22, 1.24]         | 94    | 9 (4)    | +2% | 1.56  | [1.54, 1.59]              | 98    | 3 (12)   | +39% | 6.8     | 95   | 9                      | 10  | 11  | <1  |
| Qwen3 12.5%  | 27.3   | 1.00  | [1.00, 1.01]         | 64    | 138 (30) | +5% | 0.95  | [0.94, 0.95]              | 67    | 128 (65) | +11% | 96.5    | 75   | 98                     | 124 | 145 | 2.7 |
| Qwen3 25%    | 45.0   | 1.13  | [1.12, 1.14]         | 83    | 64 (21)  | +9% | 1.41  | [1.39, 1.43]              | 91    | 34 (52)  | +37% | 43.4    | 89   | 47                     | 57  | 65  | 2.8 |
| Qwen3 43.75% | 83.2   | 1.19  | [1.18, 1.20]         | 95    | 19 (9)   | +1% | 1.62  | [1.58, 1.65]              | 99    | 3 (27)   | +47% | 13.7    | 96   | 18                     | 21  | 24  | 1.4 |

Table 10: FETCH tables at 43.75% on Qwen3-30B-A3B (launch 1): entry  $n$  is how many of a layer’s  $n$  missed experts are copied over PCIe; the rest run on the CPU. Ratio to the fixed table, paired, 95% CI.

| Table                    | $f(1..8)$         | tok/s | vs fixed             |
|--------------------------|-------------------|-------|----------------------|
| Lite                     | 0,0,1,1,1,2,2,3,3 | 109.0 | 1.057 [1.052, 1.062] |
| Law                      | 0,0,1,1,2,3,3,4,5 | 108.0 | 1.048 [1.041, 1.055] |
| Law, fetch a single miss | 0,1,1,1,2,3,3,4,5 | 104.1 | 1.010 [1.007, 1.012] |
| Fixed                    | 0,1,1,2,3,3,4,5,6 | 103.1 | 1                    |
| None (CPU only)          | 0, . . . , 0      | 100.7 | 0.977 [0.970, 0.984] |
| All (fetch every miss)   | 0,1, . . . , 8    | 98.7  | 0.957 [0.953, 0.962] |

engine, is the worst table, and running everything on the CPU the second worst. The law’s table was within 1% of the best of five alternatives on launch 1 and tied with it on launch 2. Given each run’s measured misses, Eq. (3) ranks the tables correctly except law against lite, which are within 1%, and overstates their differences about  $1.8\times$ .

**Machines with half the link.** On RTX 5090s whose links we measured at 27–30 GB/s instead of 53 (jobs 085, 085b and 088: gpt-oss next to a Core i9-14900K, Qwen3 next to a Ryzen 9 7900 in a PCIe 4.0 slot), the law’s tables fetch less (gpt-oss 0,0,0,1,1, Qwen3 0,0,1,1,2,2,3,3,4) and beat the fixed table by 21–34% on gpt-oss and 9–18% on Qwen3, about twice the gain on listing B for Qwen3 and four to six times for gpt-oss: the per-machine split matters most where the link is slow. On the Ryzen, ours leads FreeToken by  $1.04\text{--}1.08\times$ ; with the fixed table it would trail at 12.5%. On the Core i9 the first comparison (job 085) was confounded: FreeToken’s hybrid backend ran gpt-oss at 12–66 tok/s, below its own offload backend (25–86), because its CPU executor spread 23 threads across the 8 performance and 16 efficiency cores and its calibration then sent two thirds of the misses over the slow link, the same failure Section 3 reports for llama.cpp; the  $1.4\text{--}2.0\times$  lead that gave us measured a thread configuration. We reran it (job 088, a second rental of the same CPU model) with FreeToken’s executor on 8 threads pinned to the performance cores: its hybrid backend then runs  $1.4\text{--}3.1\times$  faster than before (37.5, 62.1 and 92.5 tok/s; its calibration now copies 51.7% of misses), beating its own offload backend at 11 and 25% and tying it at 40%, and ours leads its better configuration by 1.58, 1.52 and  $1.51\times$  (lowest interval bound 1.47). Restricted to the CPU alone, FreeToken’s 8-thread executor decodes 13.7 tok/s at every budget. llama.cpp on the same rental runs 40.0 tok/s with one thread per performance core (−9.4% from the law’s prediction) and 32.8 with all 24 threads (+25.9%): the hybrid-core failure is a thread-count choice, and the law holds once the choice is right. The law predicted llama.cpp on the Ryzen within +2.7%. Figure 6 puts every same-machine comparison on one axis.

**Long outputs.** Our lead does not come from the first 256 tokens. On gpt-oss, on ten held-out hard problems (AIME 2022, problems 11–15 of both exams) that neither system has finished reasoning about after 2,048 tokens (at most 1 of 10 reach an answer), ours runs 1.36 [1.34, 1.38], 1.32 [1.30, 1.34] and 1.19 [1.17, 1.21] times FreeToken at 11,

25 and 40% over all 2,048 tokens, and 1.37 [1.34, 1.39], 1.33 [1.31, 1.35] and 1.19 [1.17, 1.22] over tokens 257–2,048 (job 083, listing A); both engines stream one event per token. On Qwen3, over tokens 257–2,048 of ten held-out MATH-500 problems, ours runs 1.043 [1.020, 1.070], 1.056 [1.036, 1.077] and 1.054 [1.025, 1.086] times FreeToken at 12.5, 25 and 43.75%. Both leads are larger than over the same runs’ first 256 tokens. Turning off FreeToken’s prefill overlap, which lends it cache slots during each prompt, changes its overall speed by 1% or less.

**FreeToken’s own headline model.** On FreeToken’s own headline model, Qwen3.6-35B-A3B in BF16 (256 experts of 6.3 MB per layer, hybrid linear attention), the two systems are close: ours runs 1.030 [1.023, 1.038] times FreeToken at 12.5% and ties at 25 and 37.5% (1.010 [0.998, 1.021], 1.003 [0.994, 1.013]), at  $1.98\text{--}2.48\times$  llama.cpp (job 086, listing A). FreeToken as shipped, with its automatic backend and cache size, runs this model at 93.8 tok/s on that machine, above the 77–83 tok/s its paper reports on an RTX 5090, so our setup does not handicap it; FreeToken’s fetch cap and thread count were left at their defaults, while our FETCH table was swept. Our lead is therefore model-dependent: largest with few large experts (gpt-oss, top-4 of 128), smallest with many small ones.

**Other systems and outputs.** Two newer llama.cpp-based systems, Pipelined Sharding [37] and leloch’s expert cache [20], run gpt-oss-120b at  $1.28\text{--}1.44\times$  stock llama.cpp; ours ran  $1.42\text{--}2.34\times$  the better of them (job 077). Teacher-forced on fixed text (1536 steps per model) against stock llama.cpp with every weight in the VRAM of an RTX PRO 6000 (job 090), our cache at 25% has a mean KL divergence over the full vocabulary of 0.0019 nats on gpt-oss and 0.0005 on Qwen3 (99.9th percentile 0.121 and 0.059, maximum 0.21 and 0.13), picks the same next token on 98.5% and 99.3% of steps, and changes the loss by +0.16% and +0.19%. Stock llama.cpp’s own `--n-cpu-moe` placement diverges from the same reference by as much (mean KL 0.0017 and 0.0008, agreement 98.8% and 99.3%): the differences are those of running experts on the CPU (the helpers quantise activations to Q8\_0 where the GPU path uses Q8\_1), not of the cache. Task accuracy was not measured.

**Prompts and batches.** The cache serves single-token steps; prompts and batches above one run llama.cpp’s stock graph with every expert in host memory, and pay for the cache’s resident experts and its one-thread CPU side. In `llama-batched-bench` at 25% on the stock-clock host (job 089, both models) our prefill at 512 and 2,048 tokens runs 79–83% of stock llama.cpp’s and our decode at 2 and 4 parallel sequences 77–79% of it (the tool feeds both engines the same random tokens).

## F Artifact description

- **System.** A patch to llama.cpp commit 4da6337 (the expert cache, FETCH, the mailbox and the `llama-ec-bench` instrumentation: teacher-forced loss, logit dumps, routing lookahead records); built with CUDA 13.0 for compute capability 12.0.
- **Jobs.** One shell script per job; each installs FreeToken at commit 0d652e7, downloads the models from Hugging Face, runs the bandwidth probes, writes any on-machine prediction, runs the client, and stores raw rows (one JSON line per problem), server logs and engine counters. A small runner rents a Vast instance, streams a heartbeat, fetches the results and only then destroys the instance.
- **Analysis.** `scripts/*_stats.py` compute every interval from the raw rows; `scripts/perlayer_model.py`, `speed_limit_v2.py`, `shapley_gap.py`, `policy_study.py`, `foresight_exact.py`, `batchk_trace.py`, `fig_audit_sol.py` and `scorecard.py` produce the model comparison, the limit, the accounting, the trace studies, the audit figure and the scorecard; `scripts/wsg_tables.py` and `wsg_numbers2.py` generate every table, figure and number in this paper from their outputs. The audit’s 52 published rows and the physical speed-of-light model they are scored against are `data/published_measurements.json`, `data/audit/` and `scripts/audit.py`.
- **Cost.** The 55 rentals (34 distinct machines) of jobs 058–094 cost 44.7 US dollars at their hourly GPU prices, by the ledger in the repository; storage and download charges, which the ledger omits, add to this.